MaiAgent User Guide

Table of Contents

1. Getting Started

1.1 Apply for an Account

1.2 Log In to the Platform

1.3 Forgot Password

2. MaiAgent Work | Employee Productivity

2.1 MaiAgent Work Overview

2.2 MaiGPT

2.3 Create Documents and Spreadsheets with Templates

2.4 AI Chat

2.5 Conversation Handoff

2.6 AI Memory

2.7 Canvas

2.8 Deep Research

2.9 AI Image Generation

2.10 AI Video Generation

2.11 AI Presentations

2.12 Connectors

2.13 Use on Different Devices

2.14 Web Version

2.15 iOS App

2.16 Use on LINE

2.17 Use on Microsoft Teams

2.18 Frequently Asked Questions

2.19 Usage and Quotas

2.20 Choose the Right Model

2.21 Security and Privacy

2.22 Agent Marketplace

2.23 Use Official Built-in Agents

2.24 Enable and Manage Agents

2.25 AI Meeting Records

2.26 Create Meeting Records

2.27 Live Meeting Transcription

2.28 Meeting Records Quick Access

2.29 Meeting Backup Audio

2.30 Offline Meeting Transcription

2.31 Meeting Record Details

2.32 Meeting Access Permissions

2.33 AI Meeting Summaries and Q&A

2.34 Downloads and Knowledge Base Integration

2.35 Internal Q&A

2.36 Use Cases: Internal Knowledge Management

2.37 Professional Knowledge Management

2.38 Meeting Requirements Management

2.39 Financial Analysis Assistant

2.40 Billing Assistant

2.41 Credit Analysis Assistant

2.42 Use Cases: Image and Vision

2.43 Invoice Recognition

2.44 Official Document Scanning

2.45 Business Card Scanning

2.46 ID Document Recognition

3. MaiAgent Serve | Customer Service

3.1 MaiAgent Serve Overview

3.2 Connect Messaging Platforms

3.3 Web Chat Overview

3.4 Connect a Messaging Platform: Website

3.5 Chat Without Logging In Using a Contact Source ID

3.6 Citation Display

3.7 Markdown Rendering

3.8 Embed MaiGPT in Window Mode

3.9 Connect a Messaging Platform: LINE
3.10 Connect a Messaging Platform: FB Messenger
3.11 Connect a Messaging Platform: Instagram
3.12 Connect a Messaging Platform: Telegram
3.13 Connect a Messaging Platform: Microsoft Teams
3.14 Connect a Messaging Platform: WhatsApp
3.15 Connect a Messaging Platform: Email
3.16 Connect a Messaging Platform: Slack
3.17 Connect a Messaging Platform: LINE WORKS
3.18 Connect a Messaging Platform: Team+
3.19 Messaging Platform Settings Overview
3.20 Basic Settings
3.21 Appearance Settings
3.22 Response Loading Indicator
3.23 Feature Settings
3.24 Interaction Settings
3.25 Features for All Conversations
3.26 Conversation Viewer
3.27 Inbox Search
3.28 Messaging Platform Analytics
3.29 Human Agent Handoff Settings
3.30 Human Agent Takeover and Replies
3.31 Human Agent Handoff Notifications
3.32 Contacts
3.33 Webhook
3.34 AI Customer Service Quality Management
3.35 Voice Customer Service
3.36 Phone Customer Service (Call Center)
3.37 Voice Agent
3.38 IVR Customer Intent Recognition
3.39 Voice Call Summaries
3.40 Voice Call Quality Assessment
3.41 Use Cases: Text Customer Service
3.42 Regulatory Information Assistant
3.43 Product Information Assistant

4. MaiAgent Build｜Agent Development

4.1 MaiAgent Build Overview

4.2 Agent Builder Overview

4.3 I Want to Do X — What Should I Use?

4.4 AI Assistants

4.5 AI Assistant Features

4.6 Create an AI Assistant

4.7 Create an AI Assistant with the Agent Template Library

4.8 System Prompt Design Guide

4.9 Add a System Prompt to an AI Assistant

4.10 Multilingual Support

4.11 Cross-Conversation Search

4.12 Knowledge Base Label Filtering

4.13 Customize Maximum Output Tokens

4.14 Agent Scheduling

4.15 Set Up Agent Scheduling

4.16 Code Interpreter (Code Execution)

4.17 Browser Tool (Browser Operations)

4.18 Knowledge Base

4.19 Knowledge Base Overview

4.20 Create a Knowledge Base: Basic Settings

4.21 Crawler

4.22 Use the Crawler (Data Scraping) Feature

4.23 Create Frequently Asked Questions (FAQ)

4.24 Document Management: Tags and Metadata

4.25 Move Knowledge Base Files

4.26 PDF Document Translation

4.27 Search Testing

4.28 Database

4.29 Text to SQL Features

4.30 Create and Configure a Database

4.31 Text to SQL with the MaiAgent Database

4.32 Text to SQL with Supabase

4.33 Link AI Assistants and Group Permissions

4.34 Tools

4.35 Tool Features Overview
4.36 Create MCP Tools
4.37 Create API Tools
4.38 Configure Tools for an AI Assistant
4.39 Customize Tool Message Display
4.40 Configure Contact MCP Credentials
4.41 Built-in AI Image Generation Tool
4.42 Gemini 3 Image Generation Tool
4.43 Quick Menu
4.44 Agent UI
4.45 Skills
4.46 Skill Features Overview
4.47 Skill Management
4.48 Agent Teams
4.49 Create and Orchestrate Teams
4.50 Team Execution Records
4.51 Hook Message Interception
4.52 Add and Configure Hooks
4.53 Hook Execution Logs
4.54 AgentOps
4.55 Test Dataset Management
4.56 Automated Evaluation and AI Assistant Monitoring
4.57 Evaluation Insights Report
4.58 Tool Execution Records
4.59 Add Conversation Records to FAQ
4.60 AI Gateway
4.61 API Call Logs
4.62 MaiAgent MCP
4.63 Use Cases: External Data Integration
4.64 Agricultural Data Queries
4.65 Use Cases: AI Assistant Optimization Consultant

5. Organization Management and Governance

5.1 Organization Overview
5.2 Organization and Member Management

5.3 Organization Management
5.4 Delete an Organization
5.5 Credits Billing System
5.6 Member Credits Quota Management
5.7 Messaging Platform Credits Quota Management
5.8 Roles and Permissions
5.9 Role Permission Management
5.10 MaiGPT Role Usage Permissions
5.11 Role Resource Access Permissions
5.12 Filter Role Members
5.13 Differences Between Roles and Contacts
5.14 Role Permission Planning Guide
5.15 Member Management
5.16 Model Access Permissions
5.17 Activity Log
5.18 Knowledge Management Permissions: Query Metadata
5.19 Start Building with the Query Builder
5.20 Third-Party Login (SSO)
5.21 SAML Integration
5.22 ADFS Integration
5.23 EIP Integration
5.24 Keycloak Integration
5.25 Identity Synchronization (LDAP)
5.26 Usage Analytics
5.27 Personal Profile: Custom Attributes

6. Plans and Usage

6.1 Trial and Subscription Plans
6.2 Usage Calculation

7. Need Help?

7.1 Scan a QR Code for Real-Time Interaction

Getting Started

Apply for an Account

Please contact us via email and provide the following information:

```
Company Information:
Tax ID Number:
Name:
Phone Number:
LINE ID (for easier communication, not required):
```

Send the information to **sales@maiagent.ai** 😊

Log In to the Platform

If you already have an account and password, visit the MaiAgent AI Assistant Platform login page:

https://admin.maiagent.ai/login

Enter your account and password to log in and start using the platform.

Forgot Password

If you forgot your login password, click "Forgot Password?" to reset it.

1. Click the Forgot Password Button

2. Enter Your Registered Email

3. Check Your Email for the Password Reset Email and Click the "Reset Password" Button

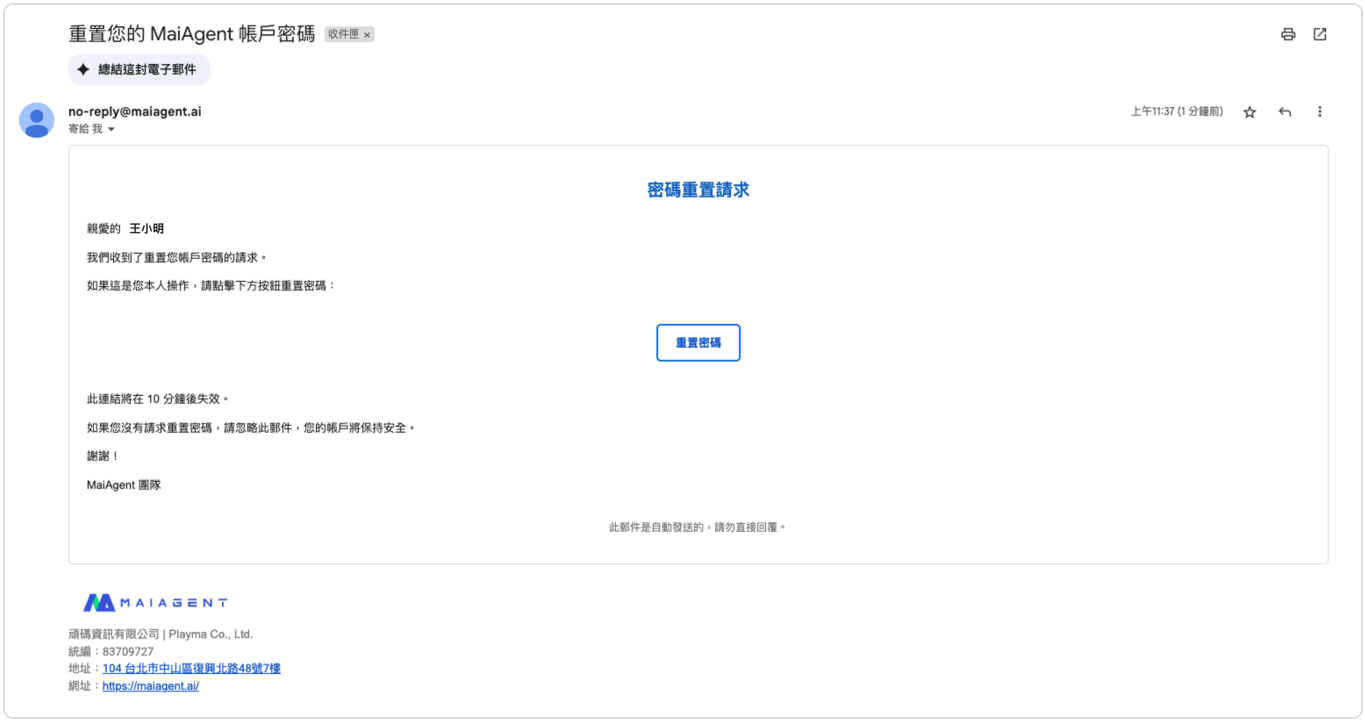

4. Set Your New Password

MaiAgent Work｜Employee Productivity

MaiAgent Work Overview

Help everyone focus on more valuable work.

MaiAgent Work is an AI workspace for employees. Search company knowledge, organize meetings, and use Agents from the marketplace for everyday tasks, all in one place. You do not need to build an AI assistant yourself; activate your account to get started. Administrators grant access to features and models based on job roles, while the company centrally manages usage and conversation records.

Products

MaiGPT	Enterprise AI chat connected to company knowledge. Choose from multiple models, get answers with knowledge base citations, and generate Word documents, Excel spreadsheets, and presentations directly. Supports Canvas, Deep Research, image and video generation, connectors, and AI memory.	MaiGPT
Agent Marketplace	Choose official built-in Agents by task without building your own. Once your organization enables them, members can use them directly in conversations.	Agent Marketplace
AI Meeting Records	Transcribe in real time or upload recordings to automatically generate summaries, decisions, and action items, and ask questions about the meeting content.	AI Meeting Records

Get Started by Role

Your Role⋯	Recommended Starting Point
Employee	MaiGPT: Start with AI Chat, then try Create Documents and Spreadsheets from Templates
Frequent Meeting Participant	Create Meeting Records, then use AI Meeting Summaries and Q&A to organize key points after the meeting
Organization Administrator	First configure MaiGPT Role Access Permissions and Member Credits Quota Management, then enable official Agents in the Agent Marketplace
Department Knowledge Owner	First Create a Knowledge Base so colleagues can look up department information through MaiGPT or Internal Q&A

A Common Starting Point: Internal Knowledge Q&A

Most companies start with “one department’ s manuals and SOPs”：

Organize One Department’ s Documents

Select the manuals, SOPs, and frequently asked questions most relevant to your work and upload them to a Knowledge Base.

Let Colleagues Ask Questions Directly

Colleagues ask questions in MaiGPT or Internal Q&A. MaiGPT answers include citations to knowledge base sources; click a citation to view the original document.

Expand to the Next Department

After verifying answer quality, apply the same approach to other departments. Control access through Roles and Permissions.

Use Cases

Internal Knowledge Management: professional knowledge, meeting needs, financial analysis, billing costs, and credit analysis

Image Vision: recognition of invoices, official documents, business cards, and identity documents

> To design your own AI assistants and integrate enterprise systems, see MaiAgent Build Overview. To bring AI to customer touchpoints, see MaiAgent Serve Overview.

MaiGPT

MaiGPT is your company's dedicated AI chat assistant. You can talk to it just like ChatGPT, but it can also query your company's internal data, connect to your work tools, and all conversations stay within the company — nothing leaks out.

Overview: MaiGPT Feature Modules

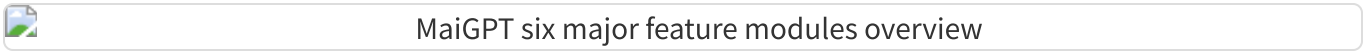

MaiGPT six major feature modules overview

AI chat is the core experience — choose a model, upload files, and ask questions directly. Other features extend the chat capabilities: Canvas lets you edit AI-generated documents in real time, Deep Research has AI automatically search the web to produce comprehensive reports, AI Image Generation lets you describe visuals with text, and Connectors let AI access your work tools like Notion and Slack.

Six Features at a Glance

Feature	What It Does	Typical Use Cases
AI Chat	Multi-model conversation with file uploads, knowledge base Q&A, and web search	Daily Q&A, document analysis, translation
Canvas	Edit AI-generated long-form content in real time	Writing reports, emails, proposals, creating charts
Deep Research	AI automatically searches the web to produce structured research reports	Market research, competitive analysis, topic research
AI Image Generation	Generate images from text descriptions with iterative editing	Marketing materials, presentation illustrations, concept visualization
Connectors	Connect external work tools so AI can query them directly	Search Notion notes, find Slack messages, locate Google Drive files
Multi-Device Access	Web, iOS App, LINE, Microsoft Teams	Use MaiGPT anytime, anywhere

How Is It Different from Agent Builder?

MaiGPT vs Agent Builder: positioning, audience, use cases, and operation at a glance

	MaiGPT	Agent Builder
Audience	All company employees	IT admins Developers
Purpose	Chat directly with AI to get daily work done	Build AI assistants for customers or internal use
How to Use	Just type, like ChatGPT	Configure role instructions, knowledge bases, tools, and other modules

Quick rule of thumb: Want to use AI yourself → MaiGPT; Want to build AI assistants for others → Agent Builder.

Recommended Getting Started Order

If this is your first time using MaiGPT, follow this order to get started:

```
Step 1 → Start a conversation — pick an AI model and ask a question
Step 2 → Try uploading a file — let AI analyze the document for you
Step 3 → Use Canvas — ask AI to write a report or email for you
Step 4 → Expand as needed: connect Notion
 Slack, use Deep Research
```

You don't need to learn everything at once. Start with chat + file uploads, then gradually explore advanced features.

Still Have Questions?

FAQ — Answers about usage limits, model selection, security, and privacy

Create Documents and Spreadsheets with Templates

The Document and Spreadsheet shortcuts on the MaiGPT home page let you select a template and data sources before returning to the conversation and asking AI to generate content. Documents work well for reports, summaries, and proposals, while spreadsheets are ideal for organizing data into fields and rows.

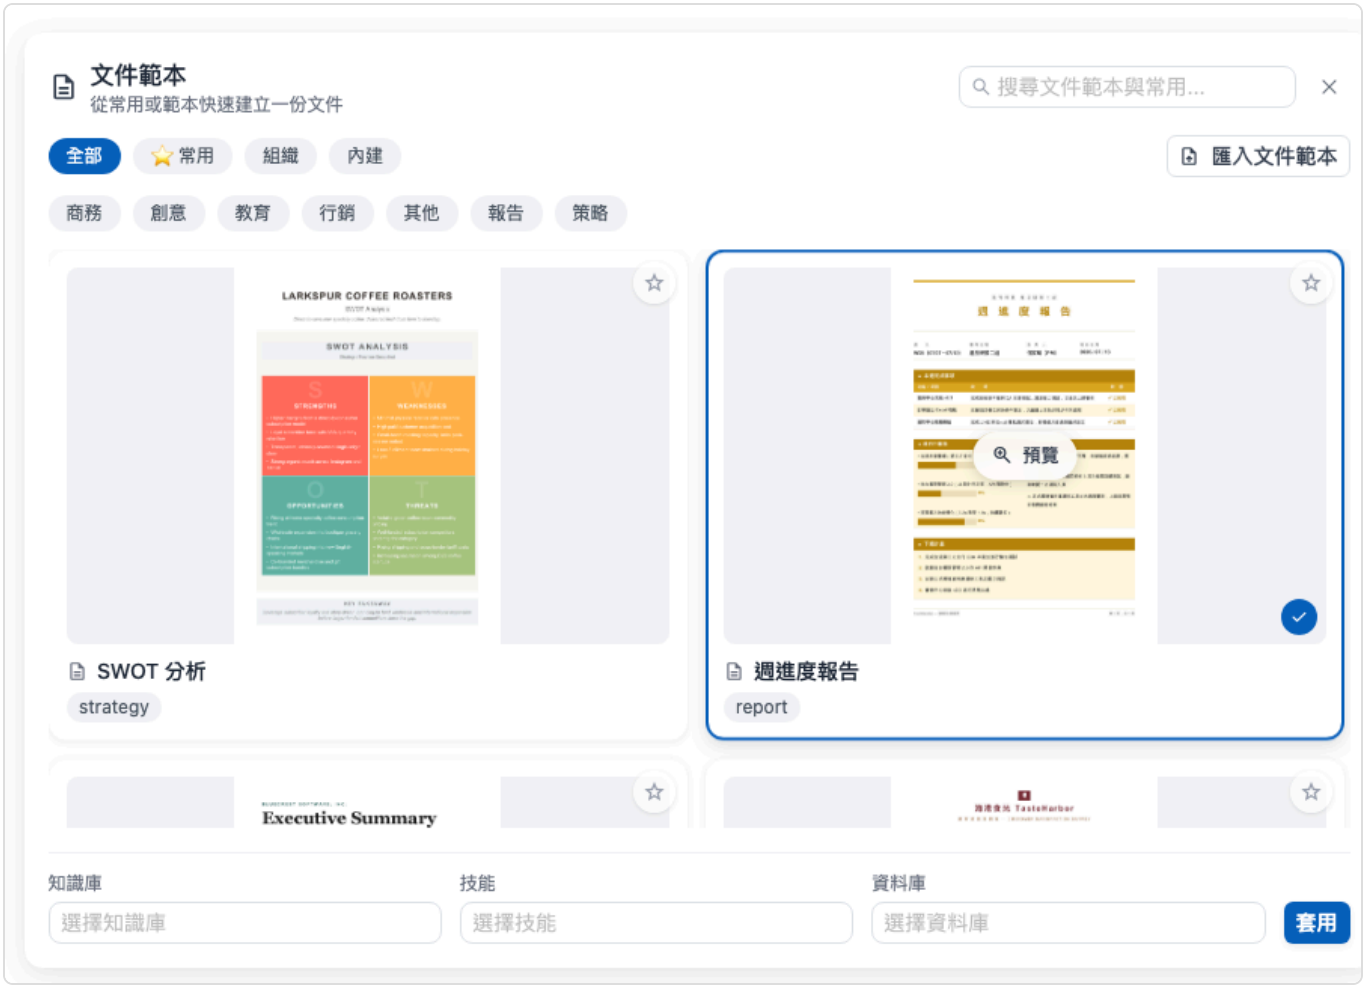

Filter document templates by source and category, then add a knowledge base, skill, or database below

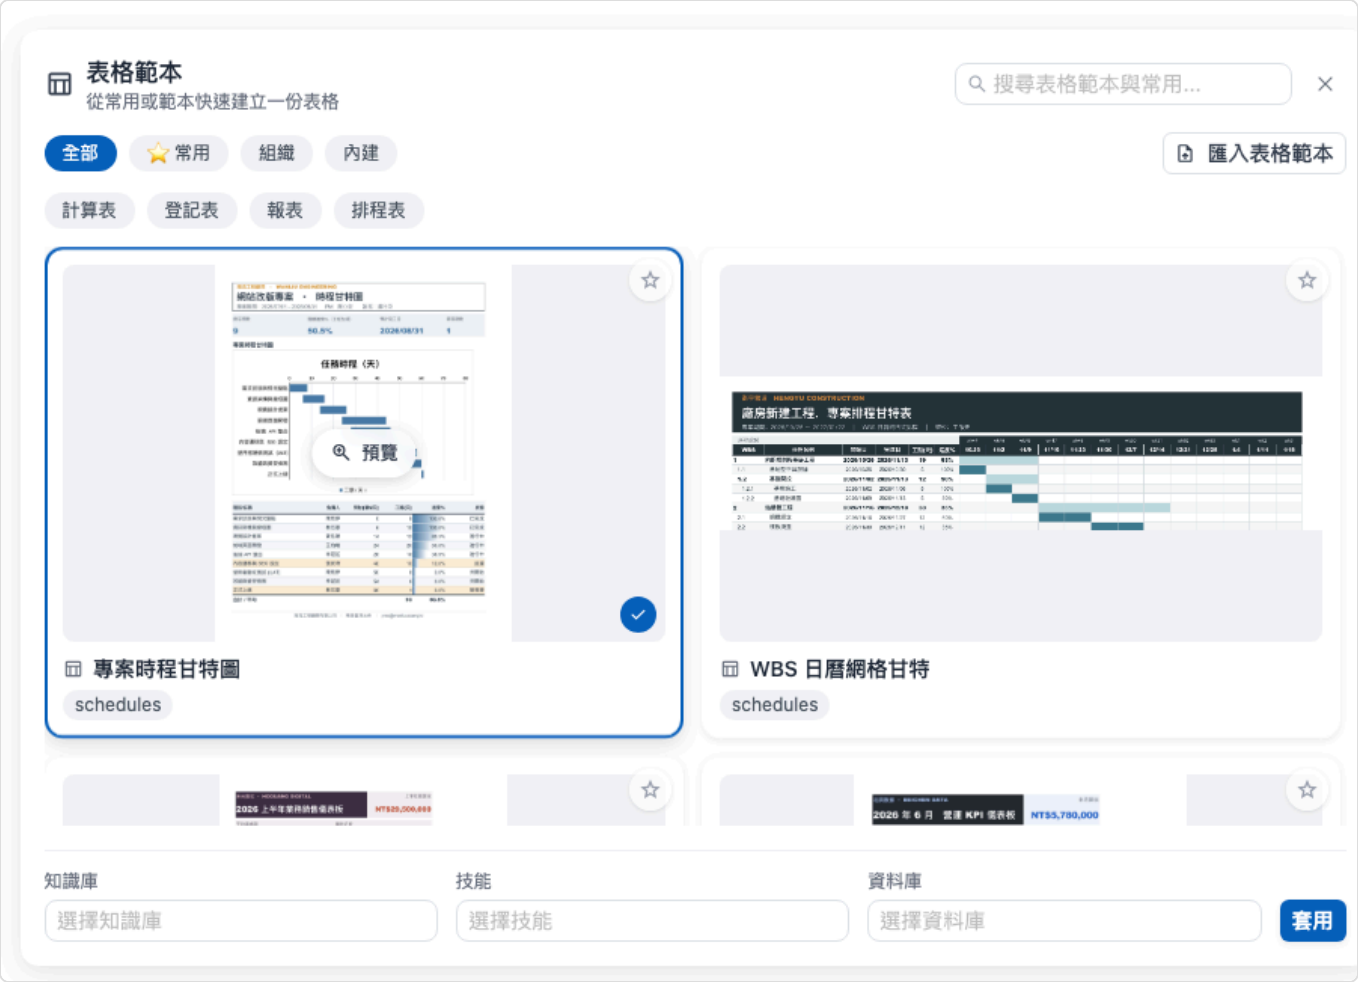

The spreadsheet shortcut provides built-in templates such as Gantt charts, dashboards, registers, and spreadsheets

Use Case

Project manager Mai needs to prepare a weekly progress report. She selects the weekly report template under Document, adds the project knowledge base, and applies the template. MaiGPT brings the template and selected data sources into the conversation, where she can specify the scope for the week and generate a draft.

Create a Document or Spreadsheet

Choose a Document or Spreadsheet

In the Productivity section of the MaiGPT home page, select Document or Spreadsheet.

Find a Template

Use search, source, or category filters to find a suitable template. You can also select the star icon to add frequently used templates to Favorites.

> Organization templates are provided by your company, while Built-in templates are included with MaiGPT. The templates displayed depend on your organization's settings.

Add Data Sources

At the bottom of the panel, select the knowledge bases, skills, or databases to reference. Add only the sources needed for the content to keep the subsequent conversation focused.

Apply the Template and Describe Your Request

Select a template, then click Apply. The selected template appears above the input box, which is also prefilled with prompt text. Close the template panel, add the time period, audience, output tone, or required fields to your message, and submit the generation request.

When generation is complete, MaiGPT lists a summary of the results and the generated files in its response. Review the content, then click Download to get the document or spreadsheet.

MaiGPT displays the weekly report summary, generated file name, and download button

The workflow is the same for spreadsheets. The following example uses the Project Schedule Gantt Chart template to generate an Excel file from three fictitious work items. When it is ready, review the item summary, file name, and file size in the response, then download the `.xlsx` file.

The spreadsheet result lists three work items and provides an Excel file for download

Manage Your Own Templates

Select Import Document Template or Import Spreadsheet Template in the upper-right corner of the panel to add custom templates permitted by your organization. Imported templates appear in the template list and can be reused later.

> If the data source list is empty, first confirm that you have access to the relevant knowledge base, skill, or database, then contact your organization administrator for assistance.

AI Chat

AI Chat is the core experience of MaiGPT. Select an AI model, type your question, and get an answer. You can also upload files for AI analysis, query your company knowledge base, search the web in real time, and enable deep thinking to tackle complex problems.

Chat Capabilities Overview

Capability	What It Does	Best For
Select AI Model	Switch between different models (Claude, GPT, Gemini)	Choosing the best AI for your task
Upload Files & Images	Upload documents for AI to read, summarize, and analyze	Contract review, report analysis, image recognition
Knowledge Base Q&A	Query your company's internal knowledge base	Company policies, product manuals, SOPs
Web Search	Search the web in real time for the latest information	News, market trends, competitor intelligence
Deep Thinking	AI reasons before answering, ideal for complex questions	Multi-step analysis, logical reasoning, decision-making
Chat Memory	Remembers conversation context and your preferences	Continuing discussions, long-term collaboration
Share Conversations	Share an entire conversation with members of your organization so they can view it or copy it and continue asking questions	Sharing results, continuing work
Conversation Handoff	Summarize a long conversation for handoff and continue in a new conversation	Long-running discussions, high context usage

---

Quick Start

Type your question — Enter your question in the input box and press Enter to send

Upload a file — Click the paperclip icon or drag and drop a file directly

Switch models — Click the model name at the top and select a different AI

Not sure where to start? Just type your question. The AI will automatically respond based on your input — no setup required.

Select AI Model

MaiGPT includes multiple built-in AI models that you can freely switch between based on your task. Each model has its own strengths, and choosing the right one helps you get better answers.

How to Switch Models

Click the model name at the top of the chat area, then select your desired model from the dropdown menu.

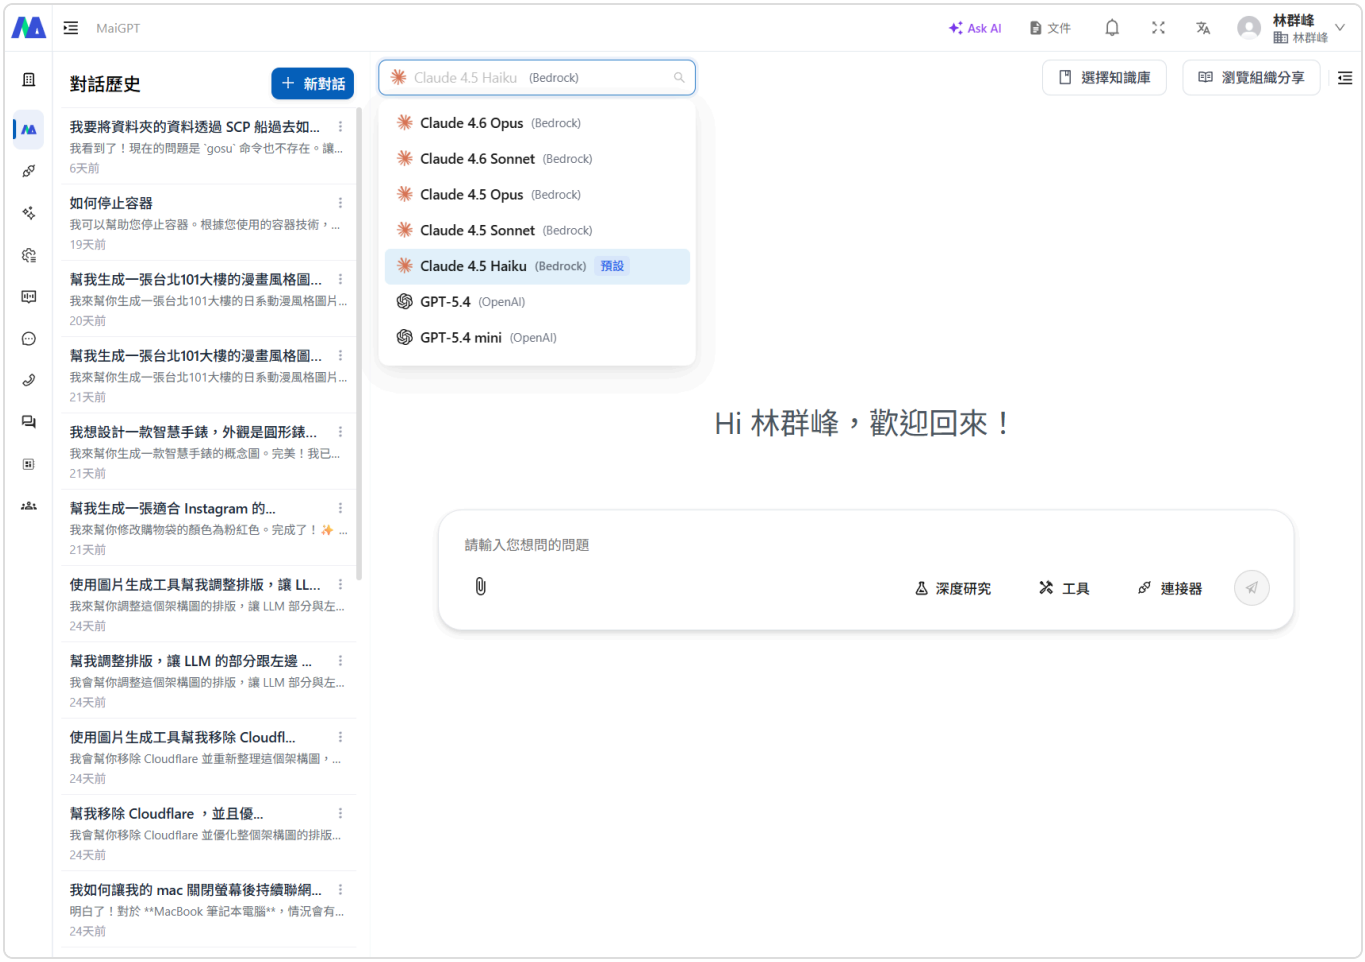

Click the model name at the top to open the dropdown and switch AI models

After switching, all subsequent responses will use the new model. Your previous chat history remains unaffected.

How to Choose a Model

My Task	Recommended Model	Reason
Writing reports, analyzing documents	Claude	Strong Chinese comprehension, excels at long-form writing
Code, technical questions	Claude GPT	Both excel here — try both and compare
Quick Q&A, translation	Gemini	Faster response times
Creative ideation, brainstorming	Claude GPT	Better creative quality
Need the latest information	Pair with web search	Any model works

> Not sure which to pick? Start with the default model. If you're not satisfied with the response quality, try switching to another model and asking the same question to compare which answer better meets your needs.

FAQ

Will my previous conversation disappear after switching models?

No. Your chat history is fully preserved — only subsequent responses will be generated by the new model.

Why can't I see certain models?

Available models are configured by your company administrator. If you need access to a specific model, contact your IT admin.

Do different models cost the same?

Different models have different usage calculations, but this is typically managed by your company, so you don't need to worry about costs.

Upload Files & Images

You can send files or images to MaiGPT and let the AI read, summarize, analyze, or translate them for you.

How to Upload

There are two ways:

Click the paperclip icon — Located next to the input box, click it to select files
Drag and drop — Drag files from your desktop into the chat area

After uploading, the file will appear above the input box. Then type your question and press Enter to send.

Supported Formats

Type	Formats
Documents	PDF, Word (.docx), Excel (.xlsx), PowerPoint (.pptx), TXT
Images	PNG, JPG, GIF
Other	CSV, JSON, Markdown

What You Can Do

Uploaded Content	What You Can Ask AI
A contract PDF	"What clauses in this contract are unfavorable to us?"
An Excel report	"Summarize the key trends in this report"
A product image	"What are the product specifications in this image?"
A meeting presentation	"List the three key takeaways from this presentation"
An English document	"Translate this into Traditional Chinese"

Tips

> Upload multiple files at once You can select and upload multiple files at the same time, and the AI will analyze them together. For example, upload three quotes and ask "Compare the differences between these three quotes."

File size limits The size limits for documents and audio files in a chat depend on your company's settings; the current default is 10 MB. Images have a separate limit. If a file exceeds the limit, MaiGPT displays the next steps instead of uploading it directly.

When File Content Cannot Be Read

For example, a project manager wants to summarize a PDF sent by a partner, but the file was corrupted during transmission. She clicks Add → Upload File, selects the file, and enters a request for a summary. If the system cannot read the content, the response may explain that it cannot provide a summary and ask her to upload the file again. When she sees this type of message, she can first confirm that the original file opens normally, then obtain the complete file and upload it again.

In this verification, MaiGPT explained that it could not summarize the file and asked the user to upload it again.

If parsing fails, first open the original file on your computer to confirm that it is not corrupted and that its content is complete. Then download it again or ask the sender to resend it. Do not use the file name as the basis for a summary. After uploading the file again, verify that the response is actually based on the file's content.

What to Do If a File Exceeds the Limit

For example, project manager Mai wants MaiGPT to summarize an 11 MB fictional company report. After selecting the file, a message indicates that it exceeds the 10 MB limit. Mai can then upload and attach the file through a knowledge base before returning to the chat to continue asking questions.

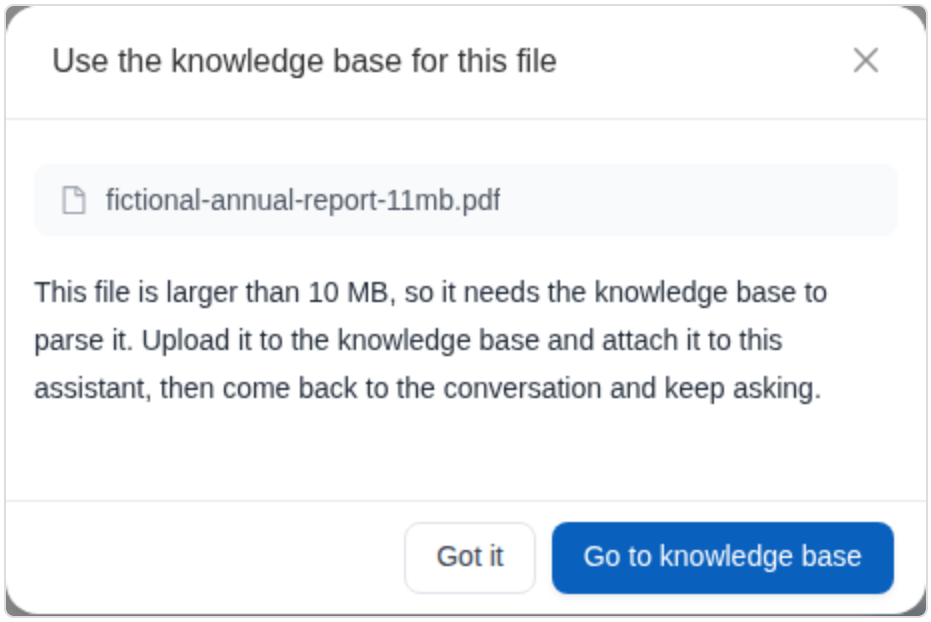

When a file exceeds the chat limit, follow the on-screen instructions to process it through a knowledge base

Confirm the file name and current size limit in the notification window.

If you have permission to use knowledge bases, click Go to Knowledge Base. Create or open a knowledge base, upload the file, and wait for processing to finish.

In the knowledge base's Linked AI Assistant settings, link the knowledge base to the MaiGPT assistant you want to use. Return to AI Chat, click Select Knowledge Base, select the knowledge base you just used, and then enter the question you want MaiGPT to process.

> The instructions above apply to the MaiGPT admin web interface. Web Chat end users do not have access to knowledge bases. If Go to Knowledge Base is not displayed, contact your organization administrator and ask them to add the file to the assistant's knowledge base. Alternatively, paste the sections you need analyzed, or split the file into smaller files and upload them again. Audio files are more likely to reach the limit, so try compressing them first.

Knowledge Base Q&A

MaiGPT can connect to your company's knowledge base, allowing you to query internal company data through natural conversation. AI responses include source references that you can click to view the original documents.

How to Use

Just ask your question. If your MaiGPT is connected to a company knowledge base, the AI will automatically search for relevant information to answer your question.

"What is the company's leave request process?"

"What are the differences between Product A and Product B?"

"What should new employees do on their first day?"

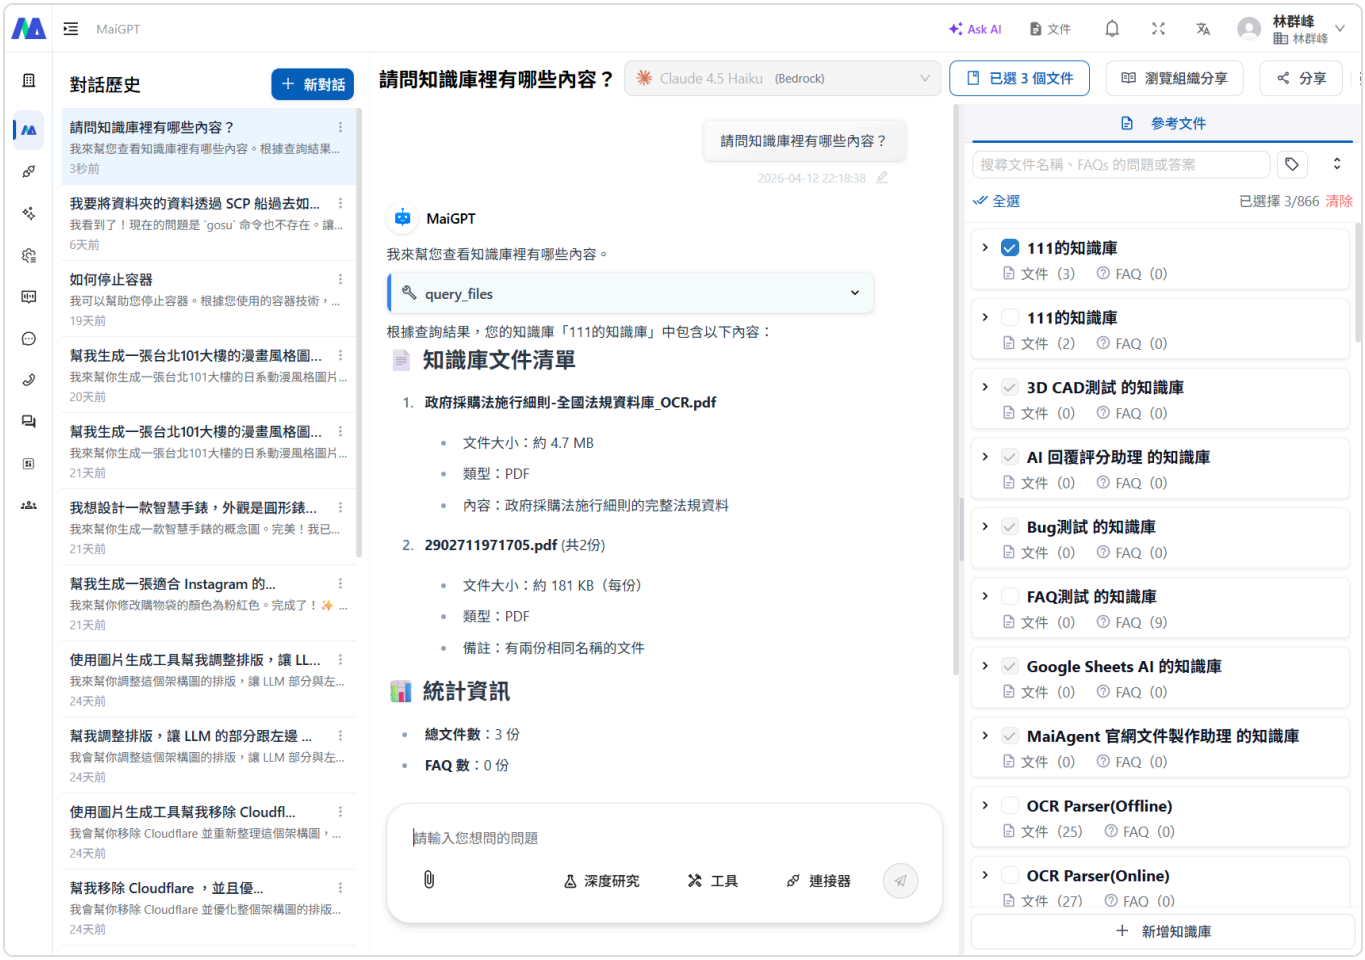

After selecting a knowledge base and asking a question, the AI uses the query_files tool to search the knowledge base and return results

Answers Include Sources

AI responses indicate data sources, typically displayed as citation markers or links. Click them to view the relevant sections of the original document.

This allows you to:
- **Verify answer accuracy** — See which document the AI based its answer on
- **Dive deeper** — If the summary isn't enough, click through to read the full original text

Can't Find an Answer?

If the AI responds with "I couldn't find relevant information in the knowledge base," it may be because:

Situation	What to Do
The data hasn't been uploaded to the knowledge base yet	Contact your admin to upload the relevant documents
Your question is too vague	Try asking with more specific keywords
Knowledge base access restrictions	You may not have access permissions for that data — contact your admin

> "Not found" doesn't mean the company doesn't have the data Knowledge base content is uploaded and maintained by administrators. If you're sure certain data exists but the AI can't find it, notify your admin to confirm whether it has been imported into the knowledge base.

Web Search

MaiGPT can search the web in real time to get the latest information for answering your questions. It's ideal for queries about news, market trends, competitor intelligence, and other scenarios that require up-to-date data.

How to Use

Click the web search button next to the input box, or mention that you need the latest information in your question:

"Search for today's Taiwan stock market closing price"

"What are the recent AI-related regulatory updates?"

"Look up what new products competitor X has recently launched"

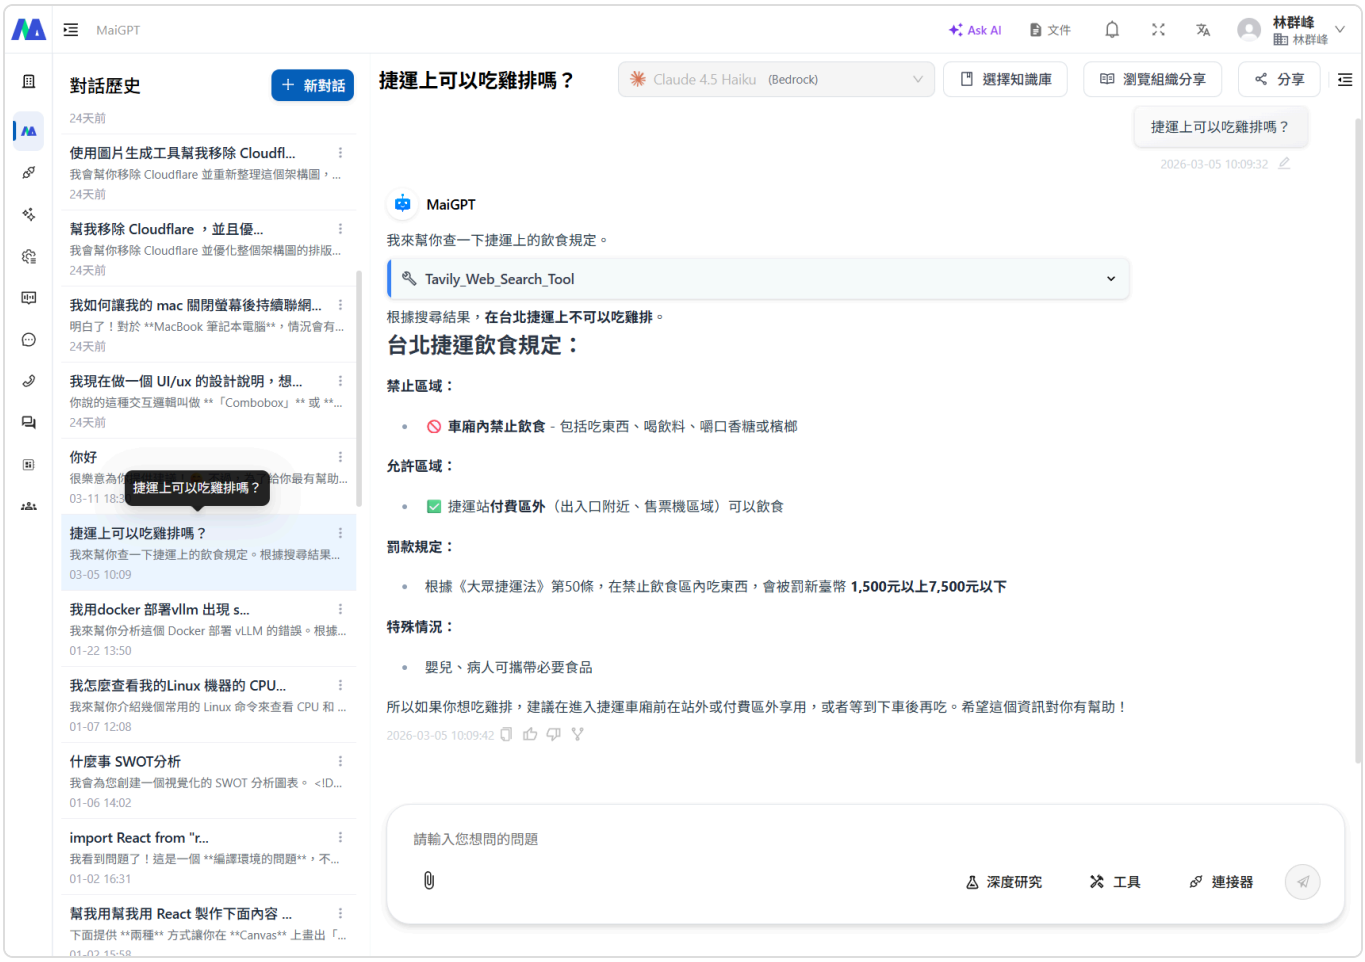

After using the Web Search Tool to search web sources, the AI compiles and responds with results and source links

The AI automatically searches multiple websites, compiles the findings into a comprehensive answer, and includes source links.

When to Use Web Search

Good For	Not Needed
Latest news and current events	Company internal data (use the knowledge base)
Product prices and specifications	Writing reports, translation (no external data needed)
Regulatory and policy updates	Analyzing uploaded files
Market trend research	Casual conversation, brainstorming

> Web Search vs Deep Research Web search is great for quickly looking up a single question — you get answers in seconds. If you need to deeply research a topic and produce a comprehensive report, use Deep Research.

Deep Thinking

Deep Thinking lets the AI "think it through" before answering, making it ideal for complex questions that require reasoning, analysis, or multi-step judgment. When enabled, you can see the AI's thinking process.

How to Use

Click the Deep Thinking button next to the input box to enable it, then ask your question as usual.

The AI response will have two parts:

Thinking process (expandable/collapsible) — Shows how the AI analyzed your question
Final answer — The complete response after reasoning

When to Enable

Good to Enable	Not Needed
Complex analysis and comparisons	Simple Q&A ("What's the weather today")
Logical reasoning required	Translation
Math calculations, data analysis	Casual conversation
Multi-faceted decision-making	Text polishing, rewriting

> Deep Thinking takes longer Because the AI needs to think before answering, response times are slower than normal mode. If your question isn't complex, stick with the standard mode.

Chat Memory

MaiGPT remembers your conversation content, allowing you to continue previous discussions without re-explaining the background every time.

Two Types of Memory

Conversation Memory (Short-Term)

Within the same conversation, the AI remembers everything discussed earlier. You can continue the discussion seamlessly:

You: "Summarize the key Q1 performance highlights for me"

AI: (responds with the summary)

You: "Now add a Q2 comparison" ← The AI knows you're talking about performance

Starting a new conversation clears the short-term memory, and the AI starts fresh.

Long-Term Memory

MaiGPT can remember important information across conversations, such as your role, preferences, and frequently used formats.

You: "Remember that I'm in the marketing department, and I always want reports in bullet-point format"

AI: Got it, I'll remember that.

Even in new conversations afterward, the AI will remember these preferences.

Tips

> Make the most of a single conversation If you're working on a topic (such as writing a report), try to complete it within the same conversation. This way the AI retains all the context, resulting in better responses.

> Need a fresh start? If you want to switch to a completely different topic, start a new conversation to prevent content from the previous topic from influencing responses on the new one.

Organize Conversations: Pin Frequently Used Chats

As your conversations pile up, important or frequently used ones get pushed down by newer chats and become hard to find. You can "pin" them to keep them fixed at the top of your "Chat History" for one-click access anytime.

How to Pin

In the "Chat History" panel on the left, hover your mouse over the conversation you want to pin
Click the ⋮ menu that appears on the right

Select Pin

Hover over a conversation and click the ⋮ menu to see the "Pin" option

Once pinned, the conversation moves to the Pinned section at the top of "Chat History," with a 📌 icon next to its title.

Pinned conversations are grouped at the top in the "Pinned" section, while regular conversations are sorted by time below

Unpin

Click the ⋮ menu on a pinned conversation and select Unpin. The conversation will return to the regular time-sorted list below.

Sorting & Rules

Item	Description
Sort order	Both the Pinned section and the regular list are sorted by the time of the last message, from newest to oldest — regardless of when pinning occurred
Pin limit	There is no limit — you can pin multiple conversations at the same time
Persistence	Pin status is saved and persists after refreshing or logging in again

Item	Description
Collapse section	Click the "Pinned" header to expand or collapse the entire section
Scope	You can only pin your own conversations

Good Scenarios for Pinning

Finance
Banking: Pin reconciliation chats or financial plan comparison conversations you regularly track

Tech
Electronics Manufacturing: Pin frequently referenced product spec or technical troubleshooting conversations

Educational Institutions: Pin recurring admissions Q&A or curriculum planning discussions

Service
Retail: Pin ongoing campaign planning or membership benefits inquiry conversations

Or any workflow you return to and continue every day

> Pinning only affects your own view Pinning is a personal organization method — it only changes the sort order you see in "Chat History." It does not affect other team members, nor does it alter the conversation content itself.

Share Conversations

You can share an entire conversation with members of your organization. Anyone with access can open the link to view the full conversation. If you allow copying, they can also copy it to their own conversation list with one click and continue asking questions.

Use Case

A knowledge base administrator uses MaiGPT to compile a "Frequently Asked Questions for New Employees" conversation and wants everyone in the organization to use it. At the top of the conversation, they click Share and select Quick Share with Entire Organization. The link is immediately copied to the clipboard and can be pasted into an announcement. Members can open the link to view the entire conversation. To ask follow-up questions about their own situation, they can click Continue Conversation to copy it to their conversation list and keep chatting without affecting the original conversation.

How to Share

The share option appears at the top of the conversation page only after the current conversation contains at least one message.

1. Click "Share"

Click Share at the top of the conversation page to display two options:

Quick Share with Entire Organization: Immediately generates a link and copies it to the clipboard. Everyone in the organization can view and copy the conversation, with no additional settings required.

Customize Sharing Scope: Choose the recipients and permissions yourself, as described in the next step.

The Share button provides two options: Quick Share with Entire Organization and Customize Sharing Scope

2. Set the Sharing Scope and Permissions

After selecting Customize Sharing Scope, configure the following in the dialog:

Field	Description
Sharing Scope	Choose Everyone in Organization, Specific Groups, or Specific Members. For groups or members, use the selection field below to select one or more recipients
Share Title	Optional. The title recipients see on the shared page and in the Browse Organization Shares list. If left blank, the conversation title is used
Description	Optional. Briefly describe the conversation's key points or purpose

Field	Description
Sharing Permission	Can Copy (default) lets recipients copy the conversation to their own list and continue asking questions. Read Only lets recipients view it only

分享

分享範圍

組織內全員
一鍵複製分享連結（預設全員可見）

特定群組
選擇特定群組成員可查看

特定成員
僅限指定成員查看

分享標題

新進人員常見問題

描述

報到、設備、報帳、會議室、通訊工具五題問答，可複製後依自己的情況追問。

35 / 500

分享權限

可複製

只讀

Set the sharing scope, title, description, and permissions in the same dialog

3. Generate and Send the Link

Click Generate Share Link. The link is automatically copied to the clipboard, and you can click Copy Link to copy it again. Send the link to the recipients. They must sign in to MaiAgent to open it.

The link is copied automatically when generated, and you can click Copy Link to copy it again

What Recipients See

After signing in and opening the link, recipients arrive at the shared page. The title, sharer, sharing time, and message count appear at the top, with the entire conversation below. They can view the content but cannot type on this page.

When permission is set to Can Copy, a Continue Conversation button appears at the bottom. Clicking it copies the conversation into a new conversation owned by the recipient, displays "Copied to your conversation list!", and automatically opens it so they can continue asking questions.

When permission is set to Read Only, the page displays "This conversation is shared as read-only" and does not show a copy button.

Members outside the sharing scope, or accounts from another organization, see "Unable to load this shared conversation. The link may no longer be valid."

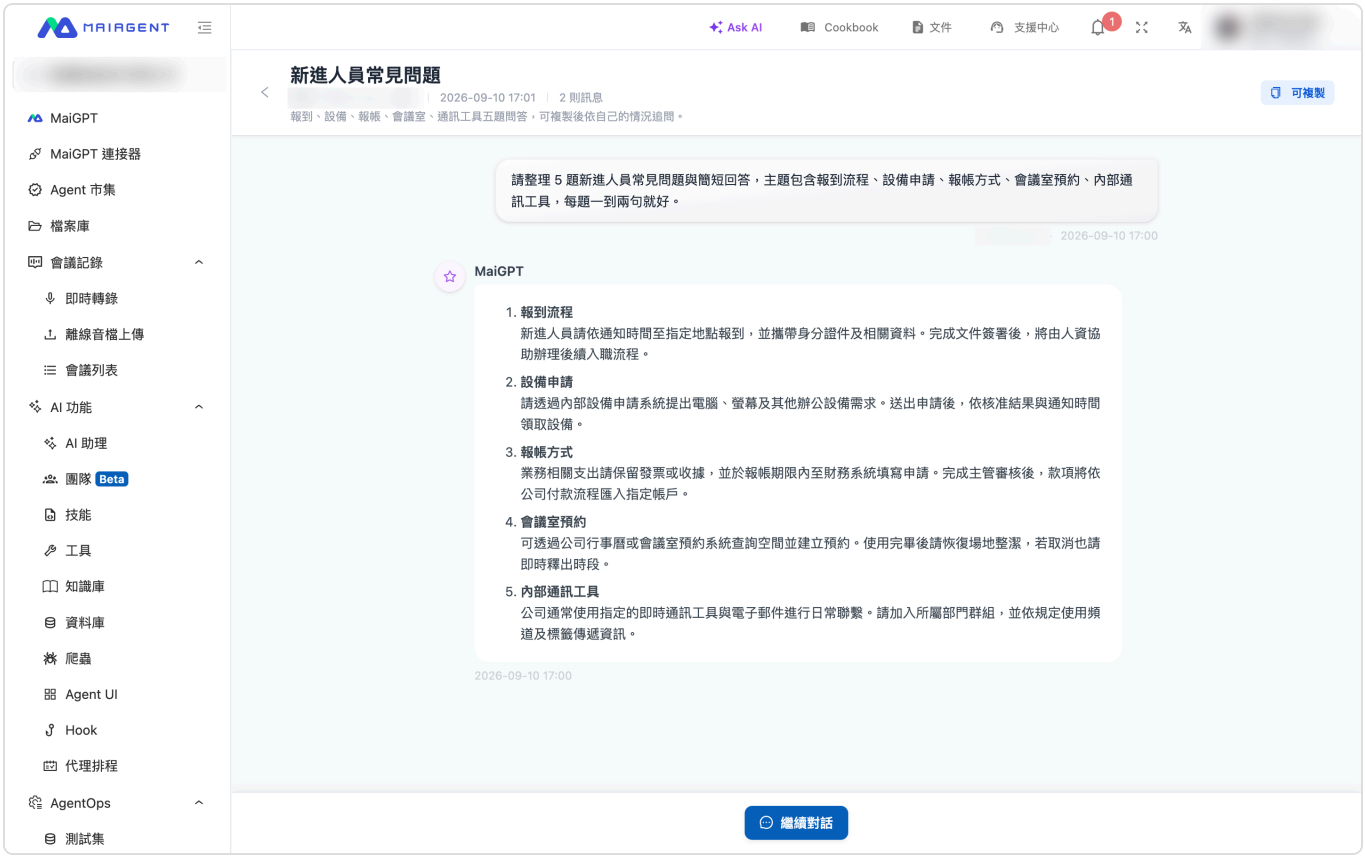

A shared page with Can Copy permission has a Continue Conversation button at the bottom; click it to copy the conversation to your own list

> The copy is a new conversation After clicking Continue Conversation, the recipient continues in a new conversation under their own account. Your original conversation is unchanged, and you cannot see their subsequent messages. The two conversations remain independent.

Browse Organization Shares and Manage My Shares

Browse Organization Shares at the top of the conversation page lists all shares in the organization that you are allowed to view. Search by title or description, then click View to open a shared page.

On this page, click My Shares to manage every conversation you have shared:

Action	Description
Edit	Change the share title, description, and sharing permission. The sharing scope and recipients cannot be changed after creation. To change recipients, cancel the share and create a new one
Copy Link	Copy the same share link again

Action	Description
Cancel Sharing	After cancellation, others who open the link see that it cannot be loaded. Conversations already copied to their own lists are unaffected

Use My Shares to edit a share, copy its link, or cancel sharing

Sharing Rules

Item	Description
Shared content	The share captures the content at the moment you click Share. Messages subsequently added to the original conversation do not appear on the shared page. Create a new share to share newer content
Excluded content	The AI's thinking process and tool execution records do not appear on the shared page. Only the conversation text, code, and Canvas content are included
Who can open it	Only signed-in members of the same organization. When the scope is limited to groups or members, anyone not selected cannot open it
Expiration	There is no expiration date. The share remains available until you cancel it
After the original conversation is deleted	The shared page remains available to view and copy, preserving the version captured when it was shared
Who can manage it	Only the person who created the share can edit or cancel it

---

Review Earlier Content While a Response Is Being Generated

While MaiGPT is generating a long response, you can scroll up to review earlier content. The page stays where you scrolled instead of being pulled back down as new text appears. When you scroll **above the**

latest question, a white Back to Latest pill appears above the input box. Click it to return to the latest question and resume automatic scrolling.

Use Case

A customer service manager asks MaiGPT to compile a very long training checklist. While the AI is still writing, she wants to review the requirements she provided in the previous turn, so she scrolls up. The screen stays at the section she wants to read without being pushed down by newly generated text, and Back to latest appears above the input box. After checking the requirements, she clicks the pill to return to the position of her latest prompt and continues watching MaiGPT finish the rest of the response.

How to use it

1. Scroll up

While MaiGPT is generating a response, scroll up using the mouse wheel, a touch gesture, or the scroll bar.

2. Read in place

The screen stays at the position you scrolled to, and newly generated content does not move your current position. When you scroll above the latest prompt, Back to latest appears above the input box.

After you scroll above the latest prompt, “Back to latest” appears above the input box. Click it to return to the position of that prompt.

3. Return to the latest content

Click Back to latest to return to the position of the latest prompt and resume automatically following new content. You can also scroll back down to the bottom to resume following the response.

> Back to latest appears only when you are above the latest prompt. If you scroll down within the same long response to read content that has already been generated, the screen also stays in place, but the pill does not appear. Scroll to the bottom to resume following the latest progress.

Next steps

Write documents with Canvas

Meeting records quick access

Return to the MaiGPT overview

Conversation Handoff

The longer you chat in the same conversation, the more the assistant needs to remember. When the content approaches its limit, the assistant starts to lose track of what you discussed at the beginning.

With “Conversation Handoff,” you can ask the assistant to organize your current progress into a **handoff summary**, then start a new conversation with one click to continue your work. The assistant in the new conversation reads this summary first, so you do not need to explain the background again. The original conversation is also preserved for you to revisit at any time.

> Conversation handoff is available only after MaiAgent enables the system's global conversation handoff feature. This global switch is not available in the admin console, and customer administrators cannot configure it themselves. Contact MaiAgent to enable it. Once enabled globally, conversation handoff is available in MaiGPT and Web Chat. For Web Chat, an administrator must also turn on “Enable Conversation Handoff” for the conversation platform (see Enable in Web Chat). Other conversation platforms, such as LINE, Facebook Messenger, and Microsoft Teams, are not supported.

Use Case

A marketing specialist is planning an autumn product launch in MaiGPT. She has spent the afternoon discussing the venue, guest count, budget allocation, and introduction scripts for two featured products with the assistant. She notices that “Context Usage” below the input box has reached 82%, and a prompt bar above the input box suggests preparing a handoff.

She clicks “Prepare Handoff.” The assistant creates a handoff summary listing the event goals, the confirmed venue and budget, the introduction scripts in progress, and the remaining tasks. After reviewing the summary, she clicks “Continue in New Conversation.” The screen switches to a new conversation, and the assistant first says, “We've picked up where you left off. The current tasks are…” She then asks the assistant to write the invitation without pasting the earlier discussion again.

Where You Can Use It

Entry Point	Users	Description
MaiGPT	Organization members	The MaiGPT conversation page works the same way in every layout. Once the global conversation handoff feature is enabled, you can use it without configuring a separate platform switch
Web Chat	Website visitors, organization members	The Web Chat conversation window embedded on a website is also available to visitors who are not logged in. After the global conversation handoff feature is enabled, an administrator must also turn on “Enable Conversation Handoff” for that conversation platform (off by default)

> The Web Chat assistant must use Agent mode The assistant prepares the handoff summary. If the assistant connected to Web Chat does not use “Agent mode” (or Agent Teams), it cannot generate a handoff summary. MaiGPT is not affected.

Enable in Web Chat

Conversation handoff and context usage in Web Chat have **two independent switches**, configured separately for each conversation platform and off by default. Visitors to a customer's embedded Web Chat see neither handoff prompts nor context usage until an administrator enables them. “Show Context Usage” takes effect as soon as it is enabled. “Enable Conversation Handoff” takes effect only after MaiAgent enables the global conversation handoff feature.

The customer service manager at an online retailer notices that website visitors often ask about several orders in a single conversation. As the conversation gets longer, the assistant loses track of order numbers mentioned at the start. She wants visitors to first see how much conversation space they have used, so she opens Feature Settings for the official website's Web Chat conversation platform, enables only Show Context Usage, and saves. When visitors reopen the conversation window and receive the assistant's first reply, “Context Usage” and a percentage appear below the input box.

1. Open the Conversation Platform's Feature Settings

In the admin console, go to Customer Service Conversations → Conversation Platforms. Edit the Web Chat platform you want to enable, then click Feature Settings in the left menu.

2. Enable the Switches You Need and Save

The “Basic Features” section has two switches. You can enable either one or both:

Switch	Effect in Web Chat When Enabled
Enable Conversation	When the global conversation handoff feature is enabled, a handoff prompt bar appears starting at 80% context usage, the assistant can generate a handoff summary card, and visitors

Switch	Effect in Web Chat When Enabled
Handoff	can continue in a new conversation. If the global feature is disabled, turning on this switch produces no visible change
Show Context Usage	Displays the context usage percentage below the input box

After configuring the switches, click Save.

“Enable Conversation Handoff” and “Show Context Usage” under “Feature Settings” → “Basic Features”

> The two switches are independent: If you enable only “Show Context Usage,” only the percentage appears, without a handoff prompt bar. If you enable only “Enable Conversation Handoff,” usage is not displayed below the input box, but the handoff prompt bar still appears at the 80%, 90%, and 100% usage thresholds (the global conversation handoff feature must be enabled). Conversation handoff also requires the global switch: MaiAgent controls the global conversation handoff switch. It is not available in the admin console, so customer administrators must contact MaiAgent to enable it. When the global switch is off, handoff is unavailable on all platforms, including MaiGPT, and enabling “Enable Conversation Handoff” produces no visible change. “Show Context Usage” is unaffected by the global switch. The settings page for internal platforms does not display these two switches. For their locations on the settings page, see Feature Settings.

View Context Usage

A “Context Usage” row below the conversation input box contains a short progress bar and a percentage. It shows how much of the assistant's available memory space this conversation has used.

MaiGPT: Always displayed; no configuration is needed.

Web Chat: Displayed only after an administrator enables Show Context Usage for that conversation platform (see Enable in Web Chat). This switch is independent of “Enable Conversation Handoff.”

Display rules:

In a new conversation, the row appears only after the assistant's first reply.

It updates after each completed assistant reply.

The maximum displayed value is 100%.

The closer usage gets to 100%, the more advisable it is to continue in a new conversation. Once it reaches 100%, the assistant starts losing access to the earliest conversation content.

“Context Usage” below the input box updates after each assistant reply

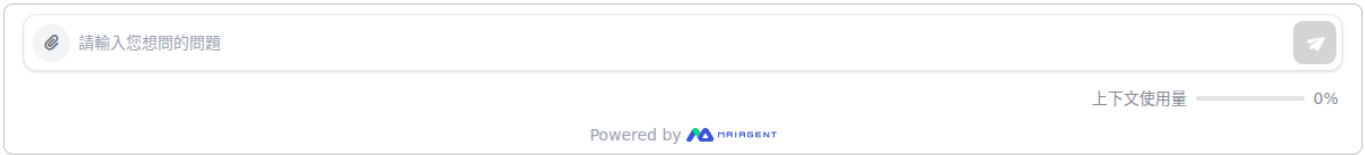

After “Show Context Usage” is enabled in Web Chat, “Context Usage” and a percentage appear below the input box

---

Three Ways to Start a Handoff

1. Ask the Assistant Directly

At any time, you can enter and send a message like this:

Hand off this conversation to a new conversation

The assistant prepares a handoff summary (see “Handoff Summary” below). You do not need to wait for usage to increase before handing off.

2. Use the Prompt Bar

When context usage reaches **80%**, a prompt appears above the input box:

This conversation is approaching its context limit. We recommend preparing a summary and handing off to a new conversation to continue

The prompt bar has two buttons on the right:

Prepare Handoff: Automatically sends “Hand off this conversation to a new conversation” for you, and the assistant prepares the handoff summary.

Close: Hides the prompt bar so you can continue in the original conversation.

After you close it, the prompt bar appears again when usage crosses the next threshold (**90%**, **100%**). At 100%, the prompt changes to:

This conversation has reached its context limit. To preserve the completeness of the conversation content, we recommend preparing a summary and handing off to a new conversation to continue

> Closing the prompt bar applies only to the currently open page. After refreshing the page, the prompt bar appears again if usage is still above the threshold.

3. Respond to the Assistant's Proactive Question

When context usage reaches 80%, the assistant adds a question at the end of that reply, such as:

This conversation is getting quite long. Would you like me to summarize the key points and hand off to a new conversation to continue?

If you agree with a response such as “Yes” or “Sure,” the assistant prepares a handoff summary.

If you say “No” or ignore the question, the assistant continues as usual without repeating the question every time.

The assistant reminds you once more at each of the 90% and 100% thresholds.

The assistant's question and the prompt bar appear at the same time. You can use either method to start the handoff.

Handoff Summary

The assistant's summary remains in the original conversation as a "Handoff Summary" card with five sections:

Section	Content
Goals	What this conversation aims to accomplish
Decided	Confirmed decisions and conclusions
In Progress	Items still being worked on and their current status
To Do	What needs to be done next
Key References	Important numbers, names, and names of previously uploaded files

The summary uses the language of your conversation. Sections with no content are hidden.

At the bottom of the card is a "Continue in New Conversation" button, with a note beside it indicating which new conversation will be created.

After you enter a handoff request, the assistant leaves a handoff summary card in the original conversation

> The system does not switch conversations automatically After the handoff summary is generated, you remain in the original conversation. Review the summary and confirm it is correct before clicking “Continue in New Conversation.” If important information is missing, ask the assistant to add it in the original conversation, or request another handoff.

You can hand off the same conversation more than once. Each handoff creates a new summary, and only the **latest summary** displays the button.

Continue in a New Conversation

Click “Continue in New Conversation”

Click Continue in New Conversation at the bottom of the handoff summary card. While the conversation is being created, the button displays “Creating New Conversation…”.

Enter the New Conversation

The screen switches to a new conversation whose title is the original title followed by “(Continued)”. The system sends the first message on your behalf, with the text “Continuing from the handoff for 〈original conversation title〉”. The message is labeled “Sent by System”.

The Assistant Confirms the Continuation Status

After reading the handoff summary, the assistant first confirms the continuation status in one sentence, such as “We've picked up where you left off. The current tasks are…”. You can then chat as usual.

The system sends a continuation message in the new conversation, and the assistant reads the summary and confirms the continuation status

After the handoff, the conversation list shows the relationship between the two conversations. The original is labeled “Handed Off To,” and the new one is labeled “Continued from Original Conversation.” The button in the original conversation changes to “Go to New Conversation”. Clicking it again returns to the same new conversation without creating another one.

The conversation list shows the handoff relationship between the two conversations

If creating the new conversation fails, the screen displays “Failed to create a new conversation. Please try again,” and the button becomes available to click again.

What You Can Do After Continuing in a New Conversation

Continue your work directly: The assistant already knows the goals, decisions, and tasks in the handoff summary, so you can give it the next instruction immediately.

Use the same assistant: The new conversation uses the same assistant as the original, with the same model, knowledge base, and tool settings.

Revisit the original conversation: The complete original conversation remains in the conversation list for you to open and review whenever needed.

Hand off again: When the new conversation gets long, you can hand off to another new conversation. Usage is calculated afresh in each new conversation.

The assistant in the new conversation can see only the handoff summary, not the other messages in the original conversation. The following content is not carried over:

Files and images uploaded in the original conversation (the summary lists only their filenames)

Files generated in the original conversation and Canvas content

Conversation details outside the summary

If the new conversation needs this content, upload the files again or paste the relevant passages into the new conversation.

> If the original conversation or handoff summary was deleted before continuing, the assistant says, “I couldn't find the handoff summary from the original conversation. Please tell me directly what you'd like to continue.” Simply explain what you want to continue.

Notes for Web Chat Visitors

Web Chat visitors who are not logged in can also hand off conversations using the steps above. Keep in mind:

The new and original conversations are saved only in the visitor's current browser. They are not available in another browser or on another device.

After the visitor's conversation history expires, the original conversation can no longer be opened, but an already created new conversation remains usable.

Unsupported Scenarios

Scenario	Description
Other conversation platforms such as LINE, Facebook Messenger, Microsoft Teams, and Slack	Usage and prompt bars are not displayed, and the assistant does not proactively ask about handoff
Automatically switching to a new conversation	Not supported. You must click “Continue in New Conversation” yourself
Handing off to a different assistant	Not supported. The new conversation always uses the original assistant
Carrying the original conversation's attachments and generated files into the new conversation	Not supported. Upload them again in the new conversation
Compressing content and continuing within the same conversation	Not supported. A handoff always opens a new conversation

---

Frequently Asked Questions

Can I hand off while usage is still low?

Yes. Enter “Hand off this conversation to a new conversation” directly. You do not need to wait for the prompt bar.

Can I still hand off if I accidentally closed the prompt bar?

Yes. Ask the assistant to hand off directly in the input box.

How is context usage calculated?

The percentage represents how much of the assistant's available memory space the conversation currently occupies, estimated from the conversation content and the assistant's model. Different models can remember different amounts, so conversations of the same length may have different usage percentages.

Can I keep chatting in the original conversation after a handoff?

Yes. The original conversation is not locked, but its usage is already high, so we recommend continuing in the new conversation.

Related Features

AI Chat

AI Memory

The AI Memory page lets you view and manage the preferences and facts currently saved to your account. You can add, edit, delete, or reactivate your memories here.

> This item appears on the account page only when your organization has enabled AI Memory.

When to use it

For example, a project manager wants to organize the communication preferences saved to her account. She can open AI Memory to review existing entries, then manually add a preference to "reply in Traditional Chinese." If her needs change, she can edit or delete it at any time.

Open the memory list

Open the account menu in the upper-right corner and select Profile. In the organization section, select AI Memory. The page lists your memories. If there are none, it displays "No memories yet."

Empty state of the AI Memory list

Add a memory manually

On the AI Memory page, select Add in the upper-right corner.

Choose a type and enter a topic, summary, and details. Use English `snake_case` for the topic, such as `communication_style`. You can put each detail on a separate line.

Fields required when adding a memory manually

Select Save. The new memory appears in the list.

Ask the assistant to remember in a conversation

Besides adding a memory manually on the AI Memory page, you can tell the assistant “Remember…” in a conversation. The assistant turns the information into a memory and saves it to your memory list for use in later conversations.

The assistant may also save facts it considers useful over the long term, such as your job role, preferred response language, or frequently used skills.

Requirements

The assistant can write a memory during a conversation only when all of these conditions are met:

The organization has enabled AI Memory.

Vector Memory is enabled in the AI assistant's settings.

The person in the conversation is an organization member. Conversations with anonymous visitors do not create memories.

> The assistant does not save sensitive personal information such as credit card numbers, national ID numbers, passport numbers, bank account numbers, mobile phone numbers, email addresses, passwords, or API keys. If content contains such information, it is not saved as a memory, and the assistant explains that it cannot save it.

Example

For example, a marketing planner is preparing a presentation for a new product in MaiGPT and wants to avoid repeating its code name in every conversation. She types, “Remember: our demo product's code name is Aurora. Use that code name whenever you write product descriptions.” After using Save Memory, the assistant confirms that it has remembered. She can then see the memory on the AI Memory page, labeled User Request. When she starts another conversation and asks for a product description, the assistant uses the code name Aurora.

Ask the assistant to remember during a conversation; the steps show Save Memory

> If a memory written by the assistant has the same topic as an existing memory, it updates that entry instead of creating a duplicate.

Memory source labels

Every memory in the list shows its source. Memories written by the assistant during a conversation display one of these labels:

Label	Meaning
User Request (green)	Content you explicitly asked the assistant to remember during a conversation
Assistant Initiated (purple)	Content the assistant decided to save during a conversation

The list may also show other source labels, such as Added Manually for memories you added on this page, or AI Extraction for information automatically extracted from conversations by the system. If you do not want to keep a memory, edit or delete it as described below.

Source labels in the memory list: AI Extraction and User Request

Manage existing memories

In the memory list, you can:

Edit the topic, summary, and details.

Delete memories that are no longer accurate or that you do not want the AI to use. Deleted memories cannot be restored.

Reactivate expired entries in the “Expired Memories” section.

Select Refresh to show the latest list.

> Memories may affect answers in later conversations. Correct or delete information that is outdated, inaccurate, or should not be kept long term.

Canvas

Canvas is the area MaiGPT uses to display documents and charts. When AI generates content in Canvas, you can view the result alongside the conversation and use preview, source code, and download features depending on the content type. Markdown documents currently support previewing. To edit text manually, download the file first, then open it in a text editor.

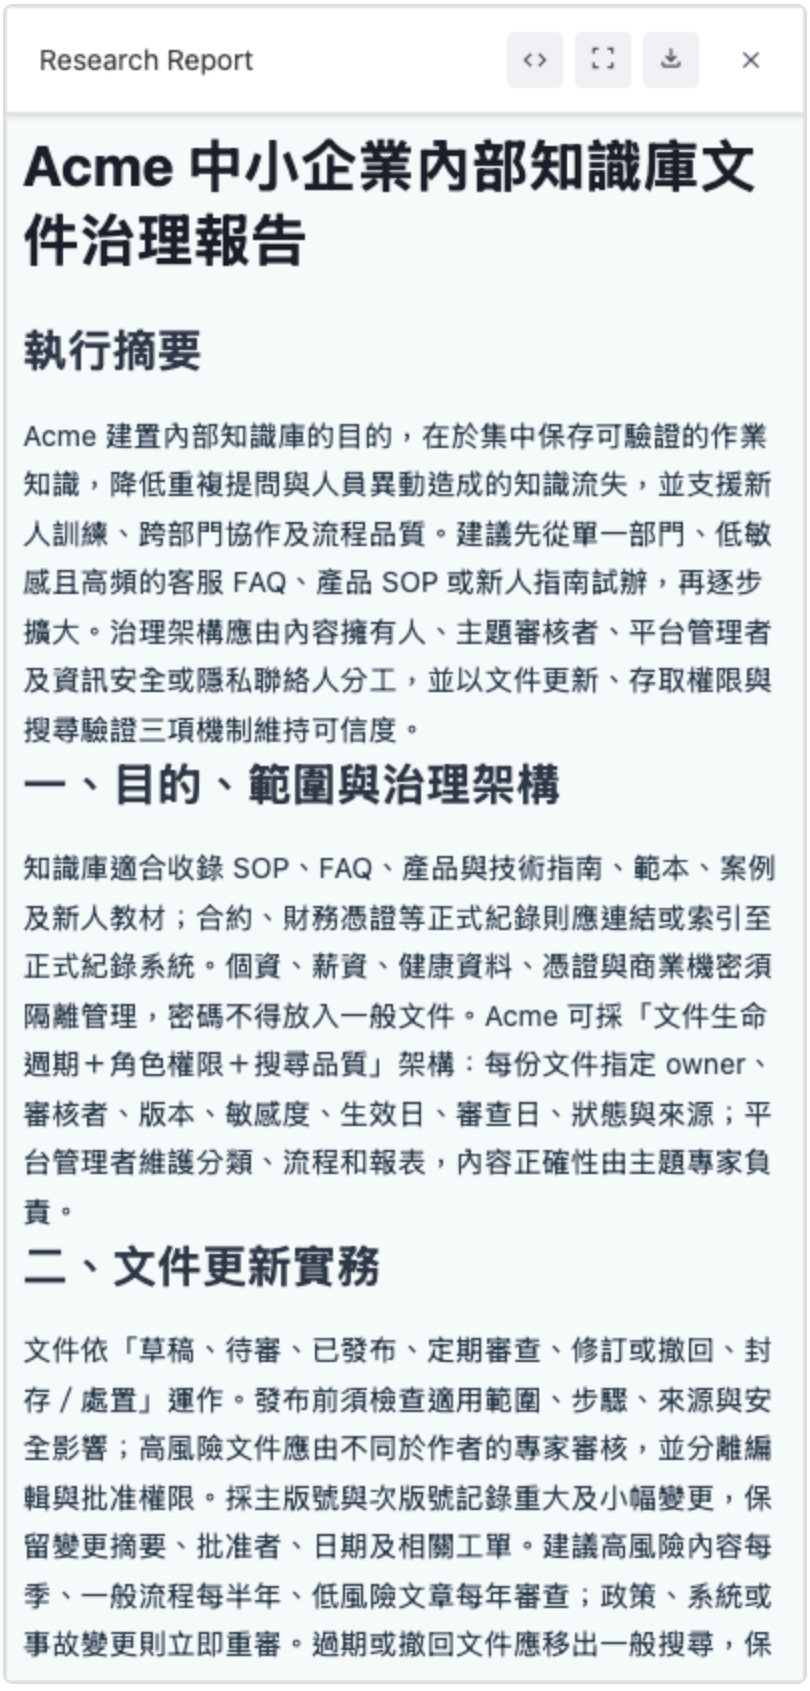

Markdown Canvas displays the document content, with source code, full-screen, and download controls in the upper-right corner

When to use Canvas

Scenario	Example prompt
Write a report	"Write a Q1 performance summary for me"
Write an email	"Help me reply to this customer's inquiry"
Organize meeting notes	"Format the following content as meeting notes"
Draw a flowchart	"Draw a leave request flowchart"
Create a table	"Create a table comparing Plan A and Plan B"

How to use it

Open Canvas

AI automatically opens Canvas when generating longer documents. You can also ask it directly:

"Use Canvas to write a..."

Edit content

After Canvas opens, you can:

View the document — Read the content in Markdown Canvas. Switch to source code mode to view the Markdown text; this mode is not a text-editing field.

Download and edit — Download the Markdown file, then edit it in a text editor. Changes to the local file are not automatically written back to Canvas.

Request changes — Clearly specify the paragraph you want to adjust in the conversation, such as "Make the second paragraph more concise."

Export

The upper-right corner of Markdown Canvas provides a Download menu:

Original format (.md) — Save the Markdown file so you can edit it in a text editor

PDF document — Download a PDF for reading and sharing

To copy the Markdown text, switch to code mode, then use the copy button in the content area.

What Canvas can create

Type	Suitable for	Example
Documents	Reports, emails, proposals, announcements	"Write a new product launch announcement"
Tables	Comparative analysis, data organization	"Compare quotes from three vendors"
Charts	Flowcharts, organization charts, sequence diagrams	"Draw the customer complaint handling process"

Tips

> Let AI help you iterate You do not need to get everything right on the first try. Ask AI to create an initial version, then refine it step by step: 1. "Write an initial draft" 2. "Make the tone more formal" 3. "Add supporting data" 4. "Turn the last paragraph into a call to action"

> Keep the versions you need Download the file first when you need to preserve the current content. Markdown Canvas preview and source code modes do not support direct text editing; edit the downloaded local file instead.

Write documents

You can ask AI to write reports, emails, proposals, announcements, or meeting notes, then view the result in Canvas. You can download Markdown documents and edit them manually.

Quick start

Tell AI directly what you want to write:

"Write a reply to a customer's request for a quote"

"Use Canvas to organize today's meeting notes"

"Write an internal announcement for a new feature launch"

After AI generates the document in Canvas, you can open the preview and download the result.

Edit and revise

Edit manually after downloading

Unlike Word, Markdown Canvas does not let you click text to edit it directly. Use the download menu in the upper-right corner to save the Markdown file, then open it in a text editor. Source code mode is only for viewing and copying.

Ask AI to make changes

Enter clear revision instructions in the conversation, such as:

You say	AI does
"Make the tone more formal"	Adjusts the tone throughout the document
"The second paragraph is too long. Make it more concise"	Shortens the specified paragraph
"Add a paragraph about the schedule"	Adds a paragraph
"Convert it to a bulleted list"	Reformats the content
"Translate it into English"	Translates the entire document

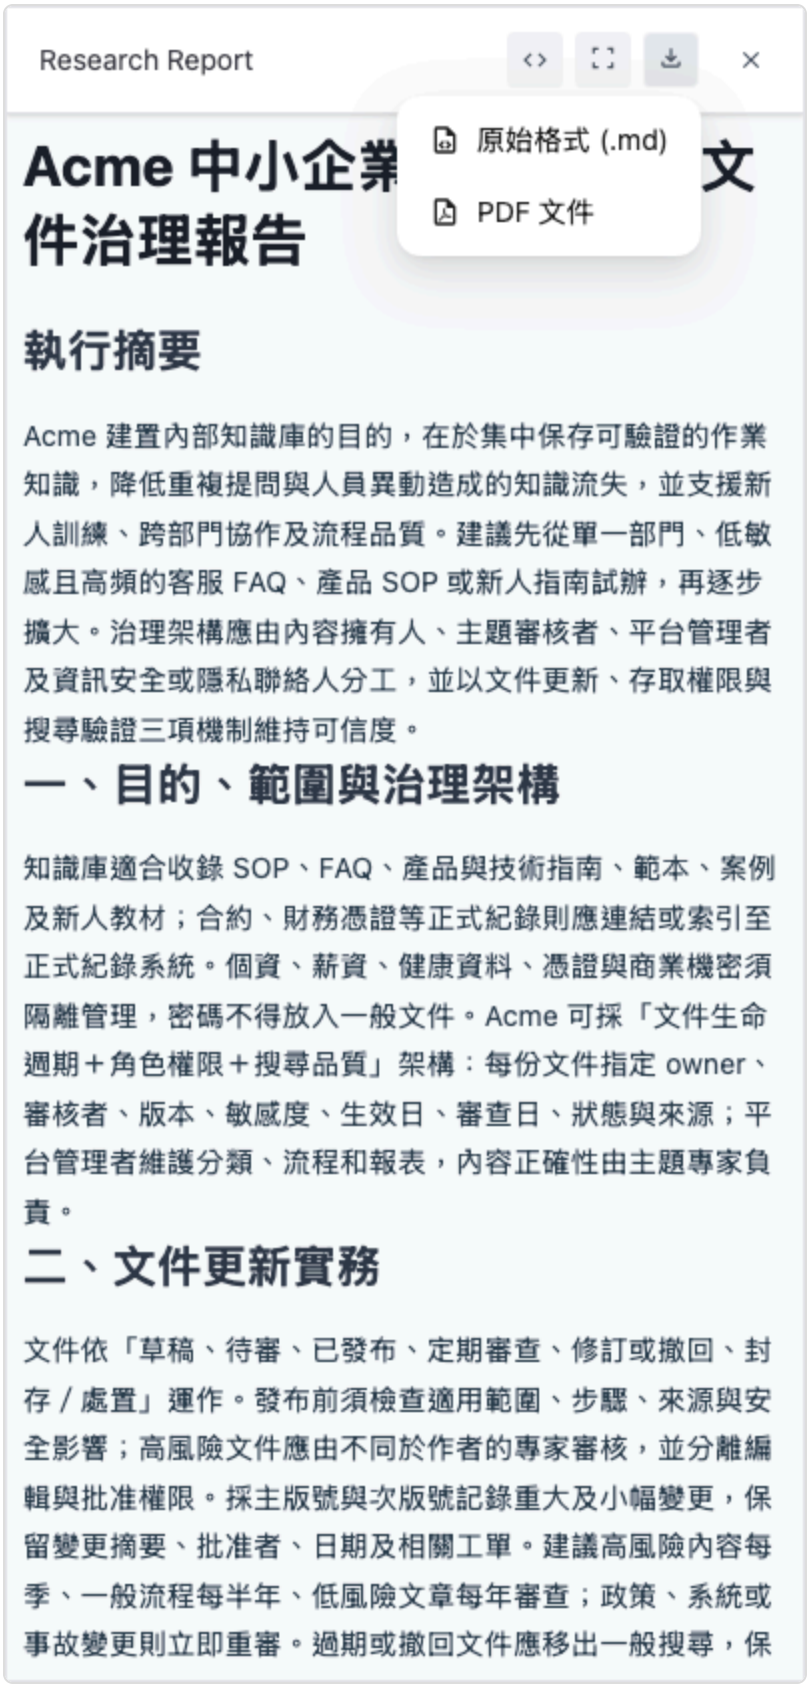

The download menu provides the original Markdown format and PDF

Common document templates

Try these prompts to quickly create common documents:

Document type	Example prompt
Weekly report	"Write this week's report. Highlights: completed Project A, started Requirement B, and will work on C next week"
Email	"Write an email politely but clearly declining a vendor's quote"
Meeting notes	"Format the following content as meeting notes: (paste notes)"
Proposal	"Write an outline for a proposal to introduce AI customer service, including a benefits analysis"
Announcement	"Write a system maintenance announcement for Saturday from 2:00 to 6:00 a.m."

Export documents

See Export above for how to download and copy Markdown documents. You can edit downloaded local files with the appropriate tools.

Create charts

Describe what you need in natural language, and AI can create flowcharts, organization charts, sequence diagrams, and other charts for you in Canvas. You do not need to learn any syntax—simply say what you want to draw.

Quick start

Tell AI directly what you want to draw:

"Draw a leave request flowchart"

"Draw our department's organization chart"

"Draw the process from a user placing an order to shipment"

Supported chart types

Chart type	Example prompt	Purpose
Flowchart	"Draw the reimbursement process"	Steps and decision branches

Chart type	Example prompt	Purpose
Organization chart	"Draw the department structure"	Reporting relationships
Sequence diagram	"Draw a sequence diagram for API calls"	Sequence of system interactions
Gantt chart	"Draw a Gantt chart for the project schedule"	Project scheduling
Mind map	"Organize the marketing strategy as a mind map"	Brainstorming and categorization

Edit charts

Want to make changes after the chart is generated? Just say what you need:

You say	AI does
"Add an approval step after the second step"	Adds a node
"Change the flowchart to run from left to right"	Adjusts the direction
"Remove the last step"	Deletes the node
"Change the colors to shades of blue"	Adjusts the style

Export charts

When finished, you can:

Copy — Paste into a presentation or document
Take a screenshot — Capture the chart directly in Canvas

Next steps

Deep Research
Back to MaiGPT overview

Deep Research

Deep Research is MaiGPT's in-depth research feature. Enter a research topic, and AI proposes a research plan. It then searches for and organizes information based on the direction you confirm and produces a report. It is suitable for market research, competitive analysis, and technical research.

How It Differs from Regular Chat

	Regular Chat	Deep Research
Best suited for	Quick questions and everyday tasks	Research that requires searching and organizing multiple sources
How it works	Enter a question and receive a response	Review the research plan and cost estimate first, then choose an output format
Research results	Chat response	Canvas document, chat response, or downloadable files
Waiting time	Usually shorter	Varies with the topic's scope and search conditions

When to Use It

A corporate planning manager is preparing to evaluate how to adopt generative AI. He enables Deep Research in MaiGPT and enters "Research the current state, challenges, and success stories of generative AI adoption among Taiwanese businesses." He first reviews the research plan and estimated Credits, then selects Reply directly in the chat. Once the research is complete, he can read the organized findings and cited sources in the original conversation as a basis for further discussion.

Start a Research Task

The following example covers topic research without attached files.

1. Enable Deep Research

Open a New Chat in MaiGPT, click Add (＋) next to the input box, and select Deep Research. Once enabled, a label with the same name appears below the input box.

Select "Deep Research" in the Add menu

> If you cannot see Deep Research, contact your organization administrator to check your MaiGPT role permissions.

2. Enter a Research Topic

Describe the topic, scope, and purpose of the report, then click the send button at the bottom right of the input box. If you want to confirm the direction first, explicitly request, "Please propose a research plan first and wait for my confirmation before proceeding."

More general	More specific
"Research AI"	"Research the current state, challenges, and success stories of generative AI adoption among Taiwanese businesses"
"Analyze competitors"	"Compare the enterprise knowledge base features of three specified products and summarize their use cases"

Confirm that "Deep Research" is enabled before submitting your research topic

3. Review the Plan and Cost Estimate

Read the research direction proposed by AI and the estimated Credits. If the scope is unsuitable, request adjustments in the same conversation first. Proceed to the next step once both the direction and the estimate are acceptable.

1. **整體概況與導入現況**：分析台灣企業導入生成式 AI 的普及程度、主要應用領域、企業規模與產業差異，以及近年的政策、投資與市場趨勢。
2. **主要應用場景與效益**：整理企業在客服、製造、研發、行銷、知識管理、軟體開發與行政流程等場景的導入方式，並評估生產力、成本、營收、品質或決策效率等成效。
3. **關鍵企業與成功案例**：盤點具代表性的台灣企業及其生成式 AI 專案，說明導入動機、技術架構、治理方式、執行成果與可複製的成功因素。
4. **技術與組織挑戰**：探討資料品質與安全、模型準確性、幻覺、系統整合、算力與成本、人才短缺，以及員工採用和流程改造等問題。
5. **法規、資安與 AI 治理**：研究個資、機密資料、智慧財產權、著作權、供應商風險與台灣及國際 AI 法規對企業導入的影響，並整理常見治理做法。
6. **導入策略與衡量方法**：比較自建、採購與混合式方案，歸納從試點、概念驗證到規模化部署的實務路徑，以及 ROI、使用率、品質與風險等評估指標。
7. **未來展望與建議**：研判代理型 AI、多模態模型、主權 / 企業級 AI、產業專用模型等趨勢，並提出台灣企業未來 1 至 3 年的優先行動建議。

預估一次完整研究的成本約為 18,905 credits；實際消耗將依實際搜尋與 token 使用量計費，可能高於或低於此估算。

The research plan includes a cost estimate; the amount shown is only an estimate for this demonstration

4. Choose an Output Format and Confirm the Start

Select a format in the Choose the research report's output format before confirming the start card:

Option	How results appear after research is complete
Canvas document	Preview the report in Canvas, view its source, read it in full screen, or download it
Reply directly in the chat	Display the report in the original conversation for convenient reading within the chat
Download files (docx／PDF)	Generate downloadable Word and PDF files for saving or sharing

Choose Canvas, a chat response, or downloadable files based on the report's purpose

> Selecting a format also sends a confirmation message. The system then checks the available quota and starts the research if the check passes. Review the plan and cost estimate before selecting a format. Once selected, the card retains your chosen format, and you cannot select another format on the card.

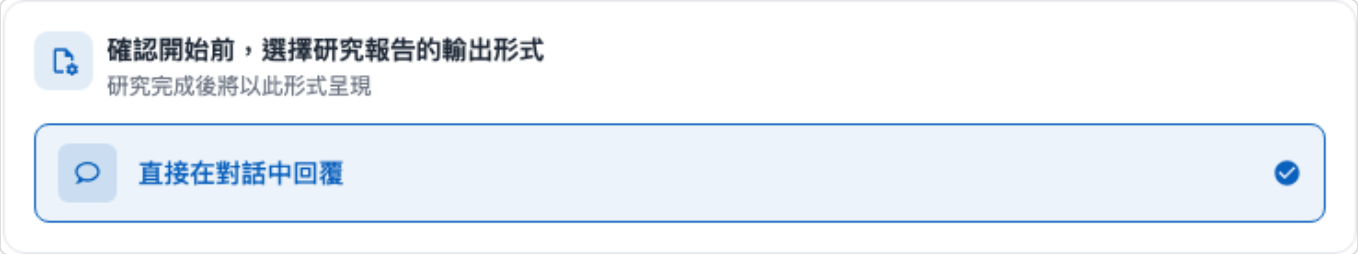

After selection, the card retains only the output format used for this research

5. Wait for the Research Results

Once research starts, the conversation displays Research progress. AI searches for and organizes information, then generates a report. The time required varies with the research scope, sources, and generation conditions. Refer to the actual progress shown.

Credits Usage Estimates and Quota Prechecks

Deep Research may perform multiple searches and analyses. The cost estimate in the plan helps you decide whether to proceed. When you confirm execution, the system checks the available quota against the estimate.

How the Estimate Is Calculated

The estimate first uses the median actual usage from up to 20 of your organization's most recent research tasks.

If your organization has no research history, it uses the median of recent research tasks across the platform. If there is still no history, it uses the default estimate configured in the system.

This is a reference based on historical usage, rather than a fixed price quoted for the current topic. Actual charges depend on usage such as searches and Tokens, and may be higher or lower than the estimate.

> The estimate is not a spending limit and does not guarantee that you will have sufficient quota throughout the research. The check before confirmation uses the estimate and balance at that time. Subsequent actual usage may still exceed the estimate.

The Research Plan Includes a Cost Estimate

When AI proposes a plan, it includes estimated Credits and a note that actual usage may be higher or lower than the estimate. AI generates this text, so the wording and amount may vary. Read the notice in your current conversation.

Quota Check Before Execution

After you confirm execution, the system checks the organization wallet and the quotas applicable to this task. If any applicable quota is insufficient, the research will not start:

Check	What is checked	Details
Organization wallet	Whether available Credits can cover the estimated usage	About Credits
Individual quota	The member's available quota, if configured	Member Credits quotas
Shared role quota	The available shared quota of the member's role, if configured	Role quotas
Conversation platform quota	The available quota for the platform used by this conversation, if configured	Conversation platform Credits quotas

What Happens If the Check Blocks Research

AI displays a notice about insufficient quota. Follow the notice and contact your organization administrator for assistance.

Research does not start. The original conversation and research plan remain at the stage awaiting confirmation.

Once the quota has been replenished, ask AI to proceed in the same conversation. For example, "The quota has been replenished. Please start the research according to the plan you just proposed." You do not need to enter the topic again.

What to Do When Quota Is Insufficient

Your role	What to do
Regular member	Share the insufficient quota notice with your organization administrator and explain the research purpose and estimated usage
Organization administrator	Follow the notice to check the organization wallet, individual quota, role quota, or conversation platform quota, then take action according to the organization's budget

Frequently Asked Questions

Q: Does seeing a cost estimate mean that amount has already been deducted?

The estimated amount is not deducted directly. **The estimate and quota check themselves are free, but AI responses such as generating the research plan and confirming details in the conversation may still consume Credits.** Therefore, research that has not started or is blocked by the quota check does not mean the entire conversation is free.

Q: Why is actual usage higher or lower than the estimate?

The estimate uses historical research costs. Differences in the number of searches, source lengths, and Token usage for this task lead to differences in actual usage.

Q: Do I need to start a new research task after replenishing the quota?

No. Return to the original conversation and ask AI to proceed according to the confirmed plan.

Q: Why do I not see a cost estimate?

Deployments without Credits billing enabled do not provide cost estimates. If an estimate is temporarily unavailable, the system still checks the existing available quota rules. The absence of an estimate does not mean the task is free.

Read and Export the Report

Read the Report

Once research is complete, open the results in the format you selected earlier. As you read the conclusions, also check the cited sources and the scope of applicability provided in the report.

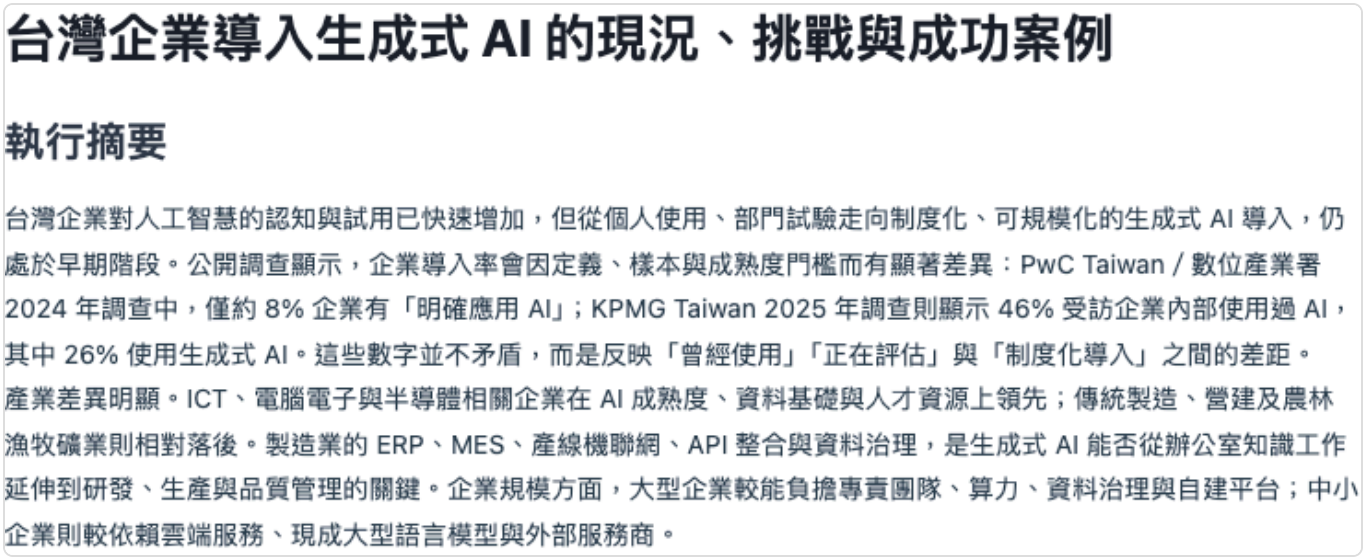

台灣企業導入生成式 AI 的現況、挑戰與成功案例

執行摘要

台灣企業對人工智慧的認知與試用已快速增加，但從個人使用、部門試驗走向制度化、可規模化的生成式 AI 導入，仍處於早期階段。公開調查顯示，企業導入率會因定義、樣本與成熟度門檻而有顯著差異：PwC Taiwan / 數位產業署 2024 年調查中，僅約 8% 企業有「明確應用 AI」；KPMG Taiwan 2025 年調查則顯示 46% 受訪企業內部使用過 AI，其中 26% 使用生成式 AI。這些數字並不矛盾，而是反映「曾經使用」「正在評估」與「制度化導入」之間的差距。產業差異明顯。ICT、電腦電子與半導體相關企業在 AI 成熟度、資料基礎與人才資源上領先；傳統製造、營建及農林漁牧礦業則相對落後。製造業的 ERP、MES、產線機聯網、API 整合與資料治理，是生成式 AI 能否從辦公室知識工作延伸到研發、生產與品質管理的關鍵。企業規模方面，大型企業較能負擔專責團隊、算力、資料治理與自建平台；中小企業則較依賴雲端服務、現成大型語言模型與外部服務商。

When a chat response is selected, the completed report can be read in the original conversation; the image shows an excerpt from the beginning

When you select Canvas document, the report appears in the Canvas panel. Click Full screen view at the top right to enlarge it for reading. Switch to code mode displays the Markdown source. Click Switch to preview mode to return to the formatted view.

Canvas full screen preview preserves the report title, body, and action buttons at the top

Edit the Report

The Markdown Canvas used for research output on this page provides preview, source viewing, and download features. **It does not support editing by typing directly.** To edit the report manually, download the original Markdown file, or select Download files (docx／PDF) before starting the research and edit the downloaded Word file with an appropriate editor.

For additional searches or analysis, ask follow-up questions in the conversation. New AI responses are still charged based on actual usage.

Export

If you select Download files (docx／PDF), Generating downloadable files… appears after the report is complete. Wait until the files have been generated, then click Download for the corresponding file. If you select Canvas, click Download at the top right of the panel and choose Original format (.md) or PDF document.

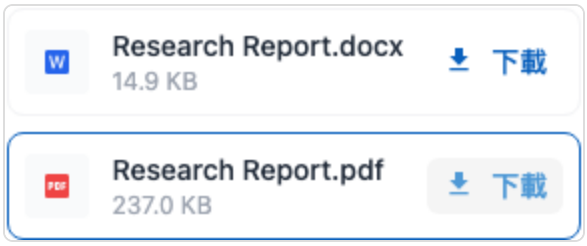

Once the files have been generated, click "Download" next to the Word or PDF file

The Canvas download menu offers the original Markdown and a PDF document

Next Steps

Learn about Canvas

View Credits and usage history

Return to the MaiGPT overview

AI Image Generation

MaiGPT has built-in AI image generation. Describe the visual you want in text, and AI will generate an image for you. Chinese descriptions are supported — no need to write in English. After generation, you can iteratively edit the image until you're satisfied.

Generate Images from Text Descriptions

Describe the image you want directly in the conversation:

"Draw an illustration of a modern office with bright tones"

"Generate a product branding image for marketing with a gradient blue background"

"Draw a cute mascot — a cat wearing a business suit"

Description Tips

The more specific your description, the closer the image matches your expectations:

Generic Description	Better Description
"Draw a cat"	"An orange cat sitting on a windowsill, sunset in the background, watercolor style"
"Make a poster"	"An A4 portrait poster with the title '2025 Annual Conference,' tech style, blue-purple tones"
"Product image"	"A cup of coffee on a white background, overhead angle, a few coffee beans beside it, commercial photography style"

You can specify:
- **Style** — Watercolor, flat design, realistic, cartoon, commercial photography
- **Color tone** — Bright, warm, cool, monochrome
- **Composition** — Overhead, front-facing, close-up, panoramic
- **Purpose** — Social media post, presentation graphic, marketing material

Downloading Images

After an image is generated, click it to view in full size. Right-click or long-press to save it to your device.

Editing and Iterating on Images

Not satisfied with the AI-generated image? No need to start over — you can keep refining it in the same conversation until you're happy with the result.

How to Edit

After an image is generated, tell AI what you want to change directly in the conversation:

"Change the background to white"

"Change the person's expression to a smile"

"Add the company logo in the bottom right corner"

"Make the overall tone warmer"

AI will generate a new version based on your instructions.

Generating Based on an Existing Image

You can also upload an image and have AI use it as a reference to generate a new one:

(Upload a product photo)

"Using this image's style as a reference, generate similar marketing material for me"

"Convert this image to cartoon style"

Tips

> Edit Step by Step Don't request too many changes at once. Modify one thing at a time for the best results: 1. First confirm the composition and subject are OK 2. Then adjust the color tone and style 3. Finally add details (text, logo, etc.)

> Not Happy? Start Over If modifications keep drifting off course, just re-describe from scratch. AI generates independently each time, so starting fresh is sometimes faster than continuing to edit.

Next Steps

Connectors

Back to MaiGPT Overview

AI Video Generation

MaiGPT can generate videos from text descriptions. The conversation guides you through selecting the video quality and aspect ratio, then displays a summary of your settings and the estimated Credits before submission. You do not need to remember the technical names of video models.

> Video generation consumes Credits. Before submitting, review the estimated usage on the confirmation screen, then decide whether to generate the video.

Use Case

A marketer wants to create a short social media video for a new product. In MaiGPT, they describe the appearance, setting, and camera movement for a fictional smart speaker, select the video quality and a landscape aspect ratio based on the campaign requirements, then review the duration, settings, and estimated Credits on the confirmation card. Generation starts only after they confirm.

Generate a Video from Text

1. Describe the Video

In a new MaiGPT conversation, describe the subject, setting, action, and camera effects. For example:

Generate a 4-second product showcase video: a white smart speaker slowly rotates on a table in a bright living room as the camera smoothly moves closer. Do not show any people or text.

2. Choose Generation Settings

Use the options in the conversation to select the video quality, then choose the aspect ratio. Available options may vary; refer to the options displayed at the time.

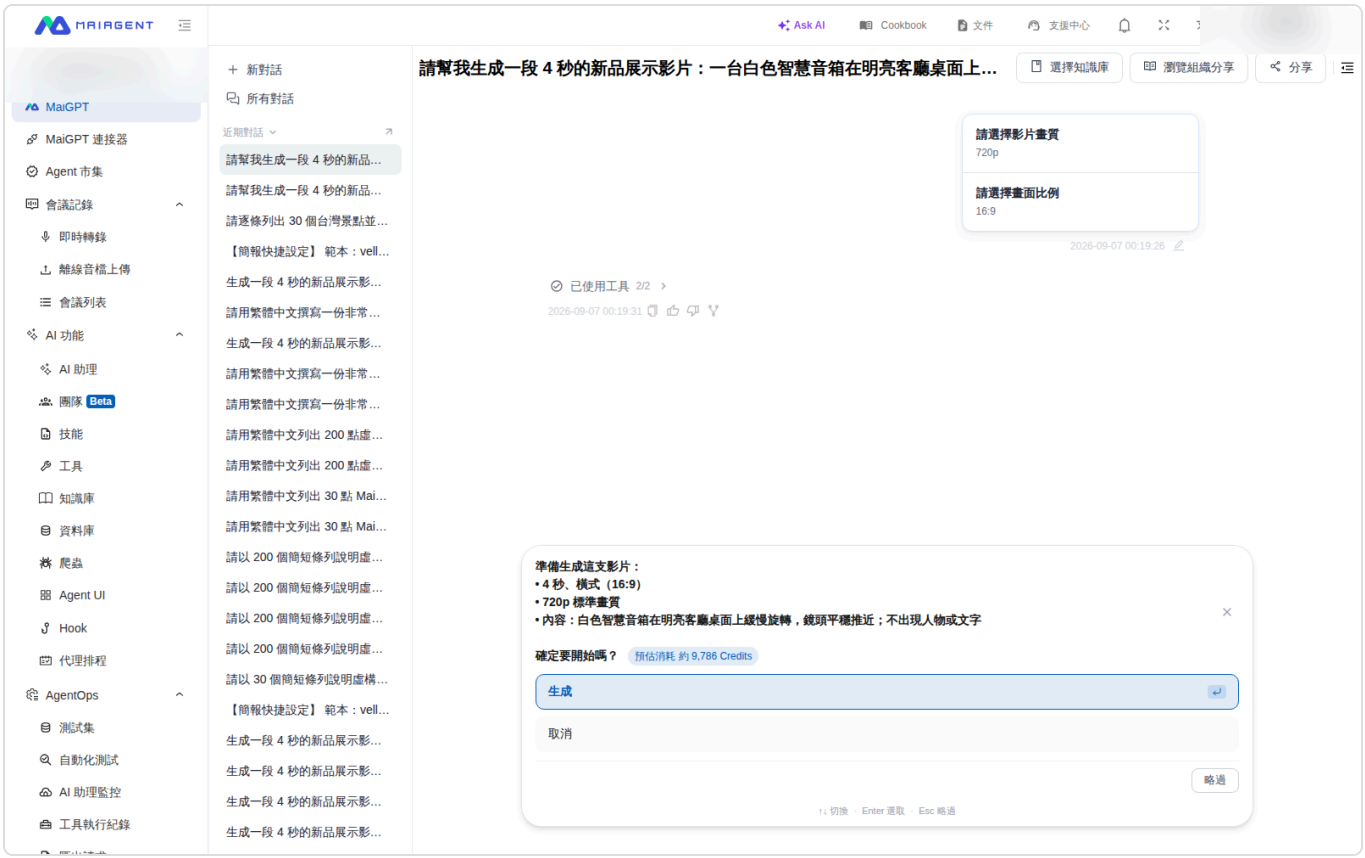

After selecting the quality and aspect ratio, the confirmation card displays the duration, aspect ratio, quality, content, and estimated Credits

3. Confirm and Generate

Review the content description, duration, aspect ratio, quality, and estimated Credits on the confirmation card. If the settings are correct, click Generate. To stop, click Cancel.

4. Preview the Result

After submitting, wait for the generated result in the same conversation.

Next Steps

AI Image Generation

Usage and Credits

Back to MaiGPT Overview

AI Presentations

After you ask AI to create a presentation in a MaiGPT conversation, the presentation panel opens on the right. You can preview each slide, edit content with AI, and export to PDF or PPTX. If the panel is collapsed, click the Open Presentation card in the conversation to reopen it.

Connectors

Connectors let MaiGPT access your third-party tools such as Notion, Slack, Google Drive, and GitHub. Once connected, you can query and interact with data in these tools directly from the conversation.

Available Tools

Tool	What You Can Do After Connecting
Notion	Search pages, query databases, create new pages
Slack	Search messages, view channel content
Google Drive	Search and read documents
GitHub	View repositories, issues, pull requests

> Available connectors are configured by your company administrator. If the tool you need is not listed, contact your admin.

How to Connect

1. Open Connector Settings

Click the tools icon in the chat interface to find the connectors section.

2. Authorize Your Account

Select the tool you want to connect, click "Connect," and you'll be redirected to that tool's authorization page. Follow the prompts to complete authorization.

3. Start Using

Once connected, ask questions directly in the conversation:

"Search Notion for pages about the Q1 plan"

"Look up the quote document in Google Drive"

Managing Connected Tools

You can do the following at any time in the connector settings:

View — Currently connected tools
Disconnect — Remove a tool's connection
Reauthorize — Reauthorize if a connection has expired

Notion

After connecting Notion, you can search Notion pages, query databases, and even create new pages directly in MaiGPT.

Connection Steps

Open connector settings and find Notion
Click "Connect" to go to the Notion authorization page
Select the pages or workspace you want to grant MaiGPT access to
Click "Allow" to complete authorization

What You Can Do

Once connected, say the following in the conversation:

You Say	MaiGPT Does
"Search Notion for pages about product specifications"	Searches and lists relevant pages
"Check the Notion project database for incomplete tasks"	Queries the database and filters results
"Create a new page in Notion titled Meeting Minutes"	Creates a new page
"Save this summary to Notion"	Writes content to Notion

Important Notes

Authorization Scope When authorizing Notion, you can choose which pages to share with MaiGPT. Only pages you've authorized can be searched. If a page can't be found, verify that you've authorized access to it.

Slack

After connecting Slack, you can search Slack messages and channel content in MaiGPT to quickly find previous discussions.

Connection Steps

Open connector settings and find Slack
Click "Connect" to go to the Slack authorization page
Confirm the authorization scope and click "Allow"

What You Can Do

You Say	MaiGPT Does
"Search Slack for discussions about server migration"	Searches relevant messages
"Check what's been discussed recently in #general"	Summarizes channel content
"Find the link John mentioned last week"	Searches a specific person's messages

---

Google Drive

After connecting Google Drive, you can search and read your Google Drive documents in MaiGPT.

Connection Steps

Open connector settings and find Google Drive
Click "Connect" to go to the Google authorization page
Select your Google account and allow access

What You Can Do

You Say	MaiGPT Does
"Search Google Drive for the quote document"	Searches and lists relevant files
"Open the last meeting minutes"	Finds and reads the document content
"Summarize that market analysis report in Drive"	Reads the document and produces a summary

Important Notes

> Only Google Native Formats Are Supported Currently, only Google native formats such as Google Docs, Sheets, and Slides can be read. Office formats like .pptx, .docx, and .xlsx cannot be read directly. Please convert them to Google format in Google Drive before referencing them.

> Shared Drive and External Sharing Limitations Files stored in a Shared Drive or existing as shortcuts may still fail to be read even after authorization. Ensure the authorized account has access permissions to the Shared Drive, and use the original file link instead of shortcuts.

GitHub

After connecting GitHub, you can query repositories, issues, pull requests, and other development-related information in MaiGPT.

Connection Steps

Open connector settings and find GitHub

Click "Connect" to go to the GitHub authorization page

Select the repositories to authorize and confirm

What You Can Do

You Say	MaiGPT Does
"Check what open issues main-repo has recently"	Lists issues
"Show me a summary of changes in PR #123"	Reads the PR and summarizes
"Search the repo for code related to authentication"	Searches code

---

Next Steps

Manage access permissions

Back to MaiGPT Overview

Use on Different Devices

MaiGPT offers multiple ways to chat with AI on any device. Conversation history syncs automatically, so you can continue a discussion from whichever device you log in on.

Supported Devices and Platforms

Platform	Description	Best For
Web Version	Use directly in your browser with no installation	Office computers and laptops
iOS App	Native iPhone iPad app	Travel, commuting, and mobile work
LINE	Chat through the LINE official account	Users who regularly use LINE
Microsoft Teams	Use directly in Teams	Teams whose company already uses Teams

---

Which Should You Choose?

The most complete experience → Web version (supports all features)
The most convenient access on the go → iOS App or LINE
Integration with your workflow → Microsoft Teams

Conversation history is shared across all platforms. You can start a conversation on your computer and continue on your phone when you're out.

Next Steps

Web Version
iOS App
Use on LINE
Use on Microsoft Teams
Back to MaiGPT Overview

Web Version

The MaiGPT web version works in any browser with no software installation required.

How to Use

Open your browser and go to the MaiGPT URL provided by your company
Log in with your account and password (or use SSO)
Start chatting

Recommended Browsers

Browser	Support
Google Chrome	Recommended
Microsoft Edge	Supported
Safari	Supported
Firefox	Supported

> Use the latest version of your browser for the best experience.

Next Steps

iOS App
Using on LINE
Back to MaiGPT Overview

iOS App

MaiGPT offers an iOS App so you can use it anytime on iPhone and iPad.

Installation

Open the App Store

Search for "MaiGPT" or scan the QR code provided by your company

Download and install

Login

After opening the app, log in with the same account and password as the web version. Conversation history syncs across devices.

Features

The app supports the same core features as the web version:

AI Chat

Upload photos (take a photo directly or select from your photo library)

Knowledge Base Q&A

Web Search

> Some advanced features (Canvas, Deep Research) are recommended to be used on the web version for a better experience.

Next Steps

Using on LINE

Using on Microsoft Teams

Back to MaiGPT Overview

Use on LINE

If your company has connected MaiGPT to a LINE official account, you can chat with AI directly in LINE without opening a separate web page.

Getting Started

Scan the LINE QR code provided by your company, or search for the official account name

Add as a friend

Send a message to start chatting

What You Can Do on LINE

Feature	Supported
Text conversation	Yes
Send images for AI recognition	Yes
Knowledge Base Q&A	Yes
Upload files	Depends on configuration
Canvas	Not supported (use the web version)
Deep Research	Not supported (use the web version)

> LINE is suited for quick Q&A and simple tasks. For advanced features like Canvas or Deep Research, switch to the web version.

Next Steps

Using on Microsoft Teams

Web

Back to MaiGPT Overview

Use on Microsoft Teams

If your company has integrated MaiGPT into Microsoft Teams, you can chat with AI directly within Teams.

Getting Started

Find the MaiGPT app in the left sidebar of Teams (or have your IT admin install it)

Click to open

Send a message to start chatting

What You Can Do on Teams

Feature	Supported
Text conversation	Yes
Send images	Yes
Knowledge Base Q&A	Yes
Upload files	Depends on configuration
Canvas	Not supported (use the web version)
Deep Research	Not supported (use the web version)

> The Teams integration lets you use AI without leaving your work environment. It's ideal for quick Q&A and daily work assistance.

Next Steps

Web

Using on LINE

Back to MaiGPT Overview

Frequently Asked Questions

The most common questions when using MaiGPT.

Question Categories

Category	What It Covers
Usage and Quotas	What operations consume usage, how to save usage
How to Choose Between Models	Which scenarios Claude, GPT, and Gemini are best for
Security and Privacy	How data is protected, whether conversations can be leaked

---

Quick Answers

How is MaiGPT different from ChatGPT?

MaiGPT is your company's dedicated AI assistant. Data stays within the company, it can query the company knowledge base, it supports multiple AI models, and it's managed centrally by your company's IT team.

Do I need to pay?

No. Usage is managed and paid for by your company.

Are my conversations secure?

All conversations are processed within the company environment and are not used for AI model training. See Security and Privacy for details.

Next Steps

Usage and Quotas

How to Choose Between Models

Security and Privacy

Back to MaiGPT Overview

Usage and Quotas

MaiGPT uses pay-as-you-go billing — you pay for what you use. Below is an overview of usage basics.

What Is Usage

Each time you chat with AI, the system calculates usage based on the length and complexity of the messages. Usage is typically managed and paid for by your company, so you don't need to pay anything yourself.

What Operations Consume Usage

Operation	Usage
Regular chat	Low
Uploading files for analysis	Medium (depends on file size)
Knowledge Base Q&A	Medium
Deep Research	High (performs multiple searches and analyses)
AI Image Generation	Medium
Deep Thinking	Higher (more than regular chat)

How to Check Usage

> How to check usage varies by company configuration. If your interface shows usage information, you can view it directly. Otherwise, contact your admin to learn about your usage status.

Tips for Saving Usage

More precise questions use less usage — Specific questions produce shorter responses and consume less usage

Make good use of the same conversation — Continuing a conversation is more efficient than starting a new one, since you don't need to repeat background information

Choose the right feature — Simple questions don't need Deep Research or Deep Thinking

---

Next Steps

How to Choose Between Models

Security and Privacy

Back to MaiGPT Overview

Choose the Right Model

MaiGPT offers multiple AI models. Here's a quick guide to help you understand the differences.

Quick Comparison

Model	Strengths	Speed	Best For
Claude	Chinese comprehension, long-form writing, analytical reasoning	Medium	Writing reports, analyzing documents, complex reasoning
GPT	Well-rounded general ability, code	Medium	Code, translation, general Q&A
Gemini	Fast response speed, multimodal	Fast	Quick Q&A, image understanding
DeepSeek	Strong reasoning ability	Medium	Math, logical reasoning
Gemma	Google open-source, lightweight and efficient	Fast	Lightweight deployment, edge computing
Grok	Real-time information, humorous style	Medium	Current events Q&A, creative writing
Llama	Open source, privately deployable	Varies	On-premises deployment needs

---

Not Sure How to Choose?

Use this decision flow:

General work (writing documents, research, translation) → Use the default model

Not satisfied with quality → Try switching to Claude or GPT

Need maximum speed → Use Gemini

Complex reasoning or calculations → Use Claude or DeepSeek

> There is no "best" model Every model has strengths and weaknesses. The most practical approach is to try the same question with different models and find the one you prefer.

Next Steps

Select an AI Model (How-To)

Security and Privacy

Back to MaiGPT Overview

Security and Privacy

MaiGPT is designed from the ground up to protect enterprise data security. Below are answers to your most frequently asked security questions.

FAQ

Can my conversation content be seen by external parties?

No. Your conversation data is restricted to your company's environment and cannot be accessed by other companies or external parties.

Will AI use my conversations to train models?

No. The AI models used by MaiGPT do not use your conversation data for training.

Can company administrators see my conversations?

This depends on your company's permission settings. Consult your IT admin for the specific privacy policy.

Are uploaded files secure?

Uploaded files are stored in your company's secure environment, and all transfers are encrypted end-to-end.

Are connectors (Notion, Slack, etc.) secure?

Connectors use standard OAuth authorization. MaiGPT does not store your account passwords. You can revoke authorization at any time.

Usage Recommendations

> Avoid entering highly confidential information Although MaiGPT has comprehensive security protections, avoid entering passwords, credit card numbers, national ID numbers, or other highly sensitive information directly in conversations.

Next Steps

Usage and Quotas

Back to MaiGPT Overview

Agent Marketplace

Official built-in Agents are AI assistants that MaiAgent has designed, tested, and listed in the Agent Marketplace. Their role instructions, models, and tools are already configured. Once an organization administrator enables one, members can start chatting immediately. There is no need to build an Agent, adjust its role instructions, or prepare a knowledge base first.

In the Agent Marketplace, official built-in Agents and Agents your organization creates appear on separate tabs: Official Built-in and Created by Organization.

How They Differ from Organization-created Agents

	Official built-in Agent	**Organization-created Agent**
Creator	Designed, tested, and maintained by MaiAgent	Created with Agent Builder
Setup before use	Ready to use	Must be built manually from scratch
Role instructions tools model	**Cannot be changed**	Fully customizable
Knowledge base	You can import your own documents	Fully customizable
Version updates	Automatically follows the latest official version	Maintained by your organization
Quantity limit	Plan-specific **limit on simultaneously enabled Agents**	Depends on the plan (see Trial and Subscription Plans)
Who can use it	An administrator chooses the available users	Determined by role permissions

In short: **Use an official built-in Agent to get started quickly; build your own for customization.** You can use both. Many organizations use official built-in Agents for common tasks and create their own Agents for company-specific processes.

What You Can and Cannot Change After Enabling One

When you enable an official built-in Agent, the system creates a **locked assistant** in your organization from the official blueprint.

You can:

Import your own documents into this Agent's dedicated knowledge base.

Choose which roles and groups can use it.

Complete any connector authorization it needs.

You cannot:

Change its role instructions (System Prompt).

Add or remove tools or skills.

Switch its language model.

> It is normal if there is no edit button for an official built-in Agent in the AI Assistants list. MaiAgent maintains its settings centrally to provide consistent quality across organizations and automatically apply official updates. If you need custom role instructions or a different set of tools, build an assistant with Agent Builder, or apply an official template from the Agent Template Library to create an assistant you can edit freely.

Versions Update Automatically

Each official built-in Agent card displays a version. When MaiAgent releases a new version to improve prompts, correct behavior, or extend capabilities, **organizations that have enabled the Agent automatically receive the latest version**. You do not need to, and cannot, upgrade manually or pin a version.

> If an Agent's response style or output format changes, MaiAgent has usually released a new version. Documents you imported into its knowledge base are unaffected.

The Number You Can Enable Depends on Your Plan

Official built-in Agents have a **limit on the number enabled at the same time** and a **cooldown period** after activation. Allowances differ by plan. See Trial and Subscription Plans • Feature Differences.

> If your organization's plan does not include official built-in Agents, the Official Built-in tab will not appear in the Agent Marketplace. Contact the sales team to upgrade your plan.

For details on enabling Agents, calculating the cooldown period, and what happens when you reach the limit, see Enable and Manage.

Available Scenarios

Official built-in Agents fall into two groups, **job functions** and **industries**, with 14 scenario categories in total:

Job Functions

Scenario category	Example Agents
Human resources	Interview question generation, resume evaluation, job description writing
Legal	Contract risk review, nondisclosure agreement review, regulatory compliance checks
General affairs	Administrative announcement writing, purchase price comparisons, document classification and filing
Accounting	Expense receipt organization, financial report summaries and analysis, business tax Form 401 analysis
Project management	PRD writing, requirements organization and analysis, competitor analysis
Personal productivity	Meeting summaries, presentation design, email reply writing

Industries

Scenario category	Example Agents
Healthcare	ICD-10 classification code suggestions, discharge record summaries, nursing shift handoff reports
Manufacturing	Quality issue 8D reports, SPC process control analysis, equipment knowledge Q&A
Banking	Credit assessment reports, KYC due diligence summaries, financial regulation interpretation lookup
Securities	Individual stock research, internal compliance policy Q&A, premeeting preparation for financial advisors
Retail	Daily sales summaries, product listing copy, customer complaint and review analysis

Scenario category	Example Agents
General administration	Citizen petition classification and assignment, press release drafting, preliminary grant application review
Government accounting	Budget execution and final accounts analysis
Legal affairs and procurement	Summaries of regulatory interpretations, tender and procurement document checks

> The table lists only representative examples for each category. MaiAgent continues to add official built-in Agents. For the current list, check the Official Built-in tab in the Agent Marketplace.

Next Steps

Want to get started → Use an Official Built-in Agent

Are you an organization administrator enabling Agents for your team → Enable and Manage

Use Official Built-in Agents

This page is for **all members**. It explains how to find the right official built-in Agent in the Agent Marketplace, how to work with it, and who to contact if you run into an issue.

If you are an organization administrator deciding which Agents to make available to your team, see Enable and Manage.

How Official Built-in Agents Work

Every official built-in Agent follows the same workflow: **you provide source material and requirements, it asks follow-up questions until the details are clear, and then it produces the result.**

The four-step workflow for official built-in Agents

1. Follow the Greeting to Enter Your Task or Upload Files

When you start a conversation, the Agent first sends a **greeting** explaining what it can help with and what you need to provide. Follow its instructions: type a description of your request or upload files, such as contracts, résumés, reports, or meeting recordings, depending on the Agent.

2. The Agent Offers Options and Recommendations to Clarify the Output

The Agent does not immediately start working as soon as it receives your material. It asks about key decisions—for example, “How many years of experience should these interview questions target?” or

“Should the report use bullets or prose?” —and **provides options directly**, so you can click one or add a short response.

At this stage, some Agents also **retrieve external information through connectors**, such as data from your calendar, cloud drive, or a third-party system. The first time a connector is used, you will be asked to authorize access.

> Do not skip this step. The extra 30 seconds you spend answering clarification questions directly improves the accuracy of the output.

3. Generate a Draft

Once the details are confirmed, the Agent generates the result. You can continue asking for changes in the conversation—adjust the tone, add sections, or change the structure—without starting over.

4. Choose Whether to Export to a Specific File Format

When you are satisfied with the draft, ask the Agent to export it as a file. Supported formats vary by Agent and commonly include Word, Excel, PDF, and presentation files. The generated file appears in the conversation for you to download.

Choose an Agent from the Agent Marketplace

Open the Agent Marketplace

After signing in, select Agent Marketplace from the left sidebar.

Switch to the “Official Built-in” Tab

At the top of the marketplace, switch from the Organization-built tab to Official Built-in.

Search or Filter by Scenario Category

Search: Use the search bar at the top to search by Agent name, description, or scenario keyword

Scenario categories: Narrow the list using category tags such as Human Resources, Legal, Healthcare, or Manufacturing, or select All Scenarios to see everything

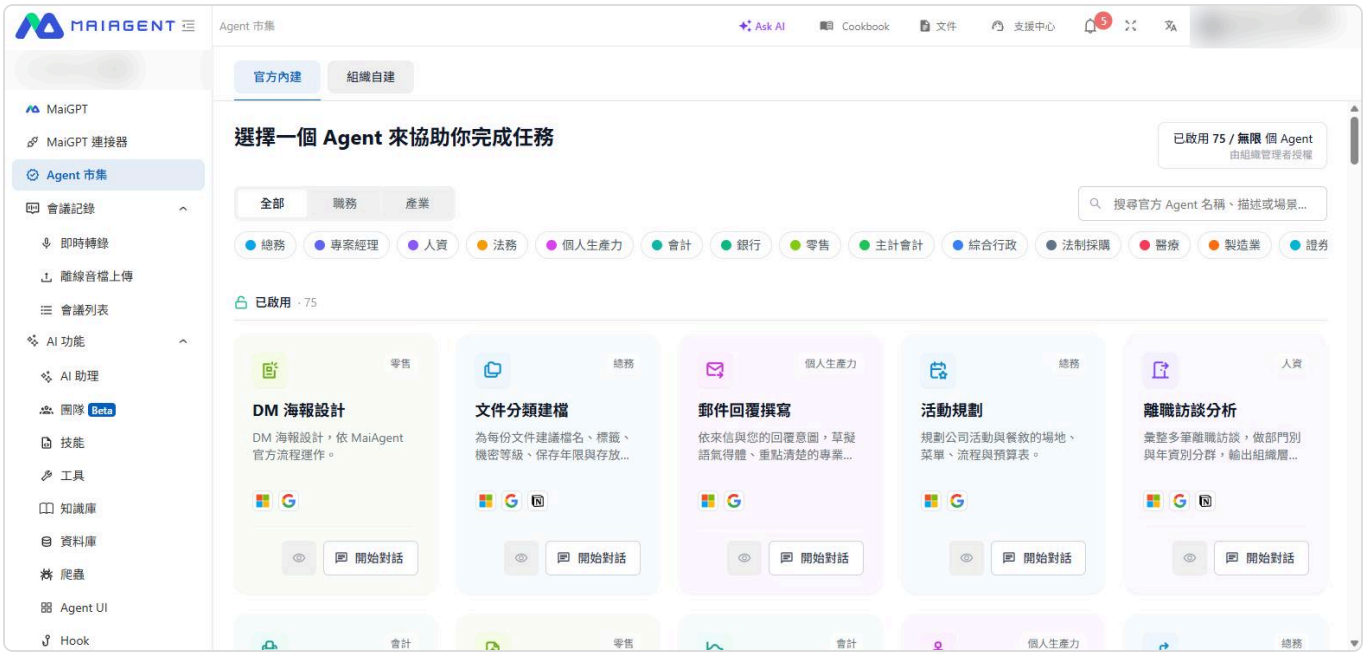

Agent Marketplace → Official Built-in: tabs and scenario categories appear at the top, with the enabled count in the upper-right corner

Cards are divided into two sections:

Section	Meaning	Available on the Card
Enabled	Your organization has enabled the Agent, so you can use it immediately	Preview and Start Conversation
Not Enabled	Your organization has not enabled the Agent	Preview only; no Start Conversation option

The upper-right corner of the page shows the number of Agents currently Enabled by your organization and its limit. This allowance is managed by the organization administrator.

Preview Before You Decide

Not sure what an Agent actually does? Select Preview (the eye icon) in the upper-right corner of its card. **You can preview it without enabling it.**

The window is titled Conversation × Guide Sync and plays two synchronized columns:

Left column: Replays one Example Scenario, from the greeting and the user's request through each step completed by the Agent

Right column—Guide ・ What This Step Does: As the conversation advances, the corresponding STEP 01, STEP 02, and so on is highlighted, explaining what happens and which capabilities are used

After playback finishes, the right column changes to Guide ・ Select Each Step for Details. Select any step to expand it, or choose Replay to watch again.

Two more details appear at the bottom of the window:

Connectors: The third-party services used by this Agent, such as Microsoft 365 or Google Workspace

Enablement status: Agents that are not enabled display Must Be Enabled by an Organization Administrator; enabled Agents offer Start Conversation

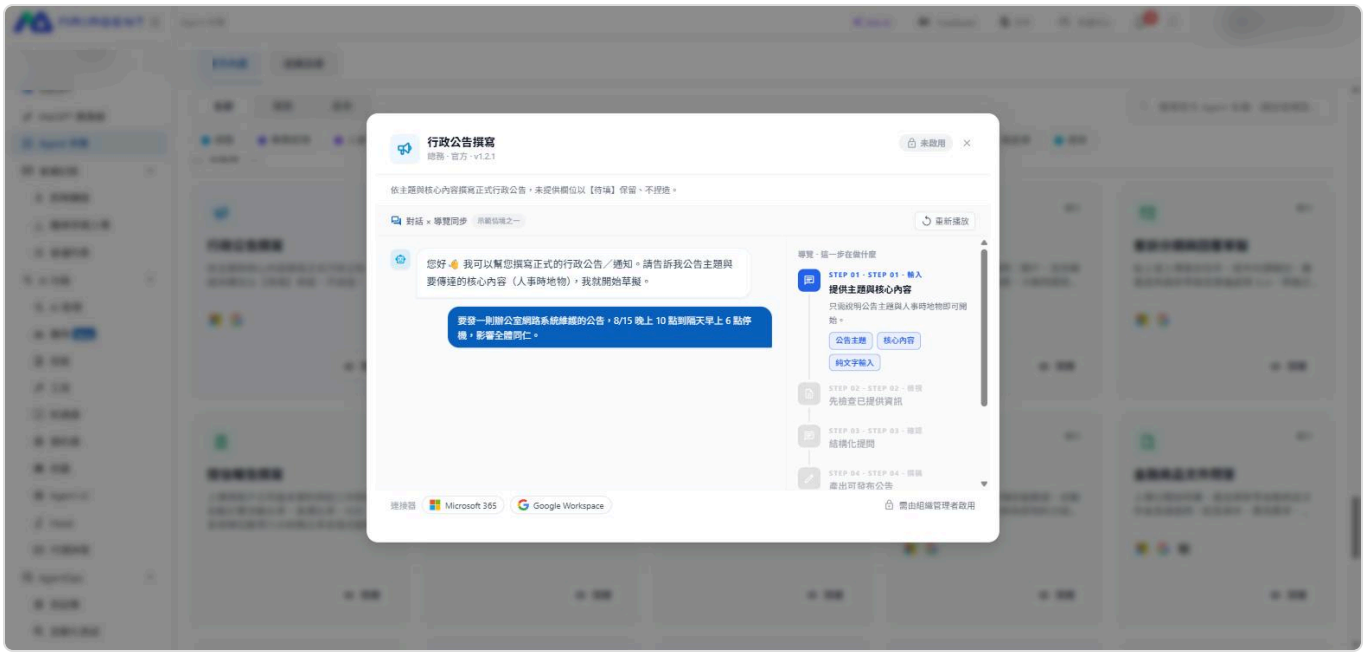

Preview window: the left column replays an example conversation while the right column highlights the corresponding guide step

Start a Conversation

Select Start Conversation on a card in the Enabled section to open a new conversation, then follow the four steps under How Official Built-in Agents Work.

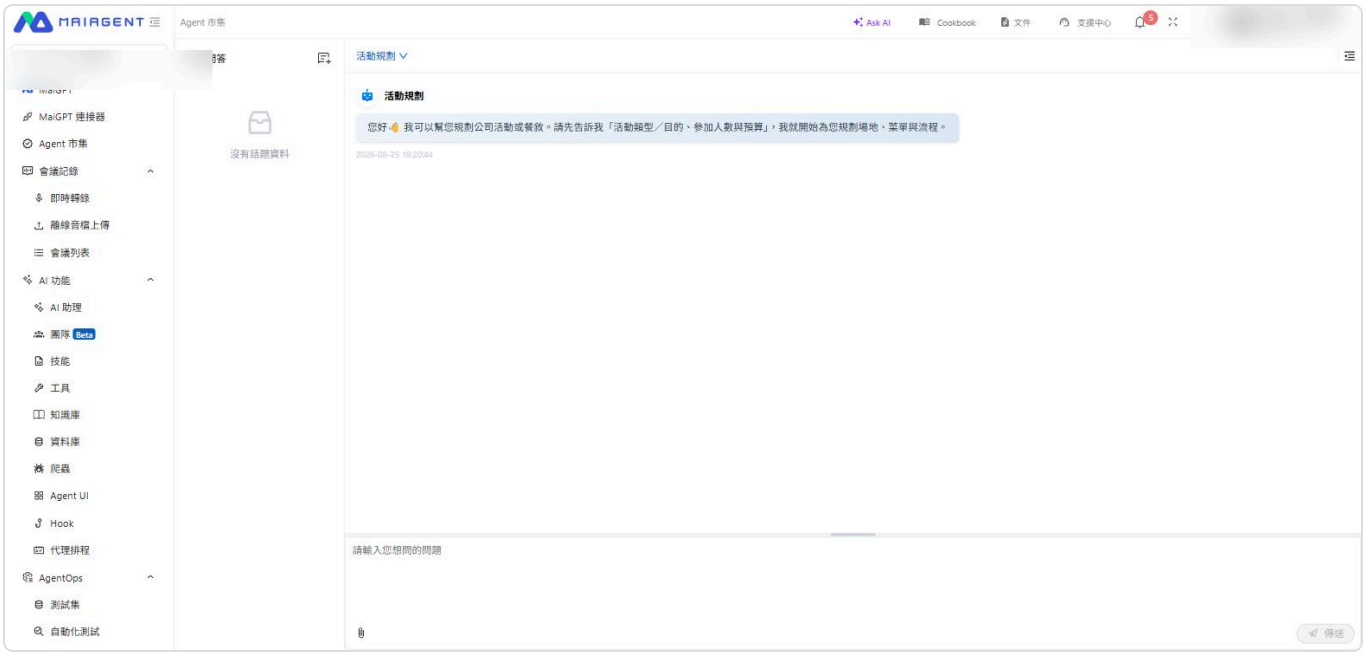

Step 1: The greeting tells you exactly what to provide; use the paper clip in the lower-left corner to upload files

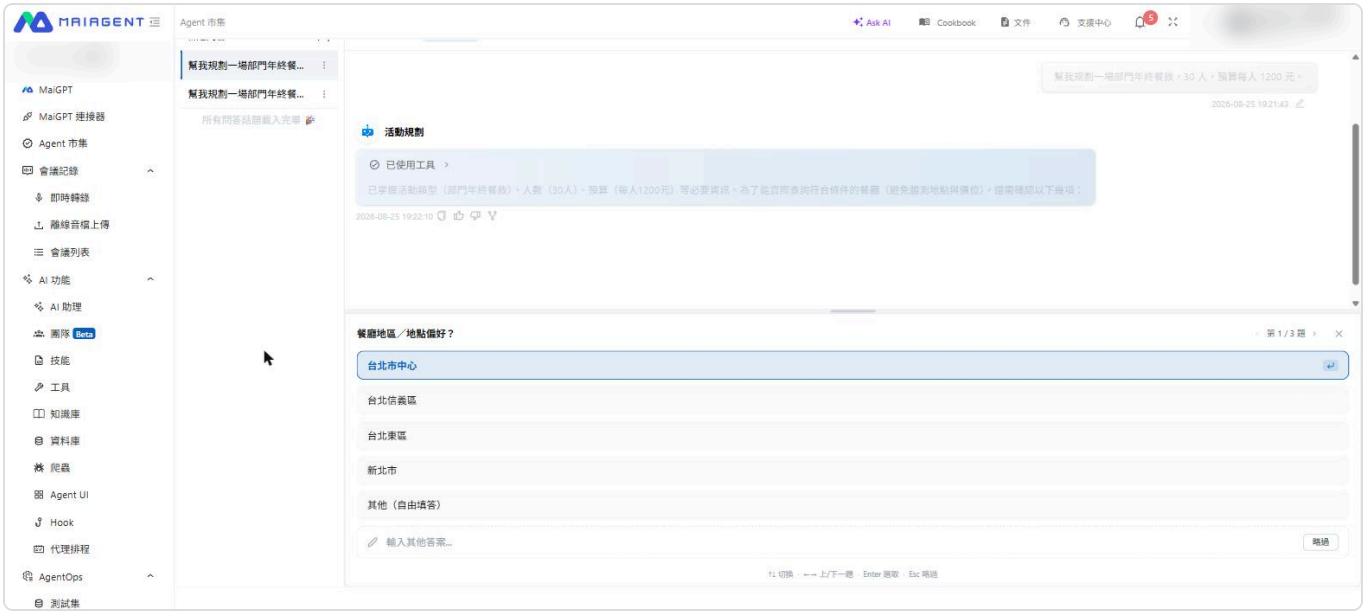

Step 2: The Agent asks follow-up questions using option cards; the upper-right corner shows how many questions remain

Step 3: After confirming the details, the Agent generates a structured draft

Return to Recently Used Agents

After you have **had a conversation with at least one official built-in Agent**, a Recently Used row appears at the top of the marketplace, ordered from most to least recent conversation. Select Continue Conversation to resume your previous work without browsing through the categories again.

> This row includes only Agents that are enabled and that you have personally used. It is normal for the row not to appear when you have just joined an organization and have not started using any Agents.

When an Official Built-in Agent Is Unavailable

These three situations look similar but have different causes:

What You See	What It Means	What to Do
The Agent is in the Not Enabled section and has Preview but no Start Conversation option	Your organization has not enabled this Agent	Ask your organization administrator to enable it
The Agent is in the Enabled section but displays No Permission to Use	Your organization enabled it, but you are not in this Agent's Authorized Users list	Ask your organization administrator to add your role/group to the authorized users

What You See	What It Means	What to Do
The Official Built-in tab does not appear in the marketplace	Your plan does not include official built-in Agents	Contact the sales team to upgrade your plan. See Trial and Subscription Plans

> If a search returns no results, select Clear Search to view every Agent in the category, or select Clear Selected to remove the category filter.

If a category has no enabled Agents, it displays No Enabled Agents in This Category and a Request Enablement button. This button takes you to Built-in Agent Management, but an organization administrator with the required permission must still enable the Agent.

Enable and Manage Agents

> This page is for organization administrators. For instructions for general members, see Use Official Built-in Agents.

Open Built-in Agent Management

Go to Organization Settings → Built-in Agent Management.

You can also select Go to Built-in Agent Management in the Agent Marketplace to open it directly.

> If Built-in Agent Management is not in the menu, your plan does not include official built-in Agents. The screen will display Built-in Agents cannot be enabled under the current plan. Please upgrade your plan. Contact the sales team to upgrade. If you can view the page but the switches are locked, you do not have permission to manage enablement. Contact your organization administrator.

Built-in Agent Management: enablement count and filters appear at the top; action icons and an enablement switch appear on the right of each row

Each row shows the Agent's name, scenario category, version, and current Authorized: {Role Name} setting. Deactivated Agents display Deactivated, Knowledge Base Retained, while Agents in cooldown display Cooldown: N Days Remaining.

Three action icons, followed by the enablement switch, appear from left to right on each row:

Icon	Purpose
Configure Authorized Users	Specify which roles/groups can use the Agent
Clear Knowledge Base	Delete all imported documents (cannot be undone)
Import to Knowledge Base	Open the Agent's dedicated knowledge base to upload documents

Use the search box at the top to search by keyword. You can also filter by status—All
Enabled
Not Enabled
Knowledge Base Retained—or by scenario category.

> Select the row itself, rather than an icon, to open the Agent's Preview window. You can see how it works before deciding whether to enable it.

Enable or Deactivate: Select “Apply Changes” to Take Effect

> Changing a switch does not take effect immediately. All changes remain pending until you select Apply Changes at the bottom to submit them together.

Switch the Agents You Want to Enable or Deactivate

In the list, set the switch to Enabled or Deactivated. The row is then labeled Pending Enablement or Pending Deactivation.

Review the Change Summary at the Bottom

A fixed action bar appears at the bottom of the screen. When changes are pending, it shows the number Pending Enablement, the number Pending Deactivation, and the resulting number of Enabled Agents versus the limit. It also reminds you that an Agent cannot be deactivated within 7 days of enablement. The Apply Changes button becomes available at this point.

Select “Apply Changes”

Select Apply Changes to submit everything at once. To discard the changes, select Cancel to restore all switches.

After a switch is turned off, the row is labeled “Pending Deactivation,” and the counts and cooldown reminder update in the action bar below

If an individual item fails after the changes are applied, the screen identifies the reason for each failure, such as reaching the limit, an active cooldown period, or insufficient permissions. Successful changes are unaffected.

Enablement Limit and Cooldown Period

Concurrent Enablement Limit

The Enabled Count at the top of the page shows how many slots are currently used and the maximum allowed. The limit depends on your plan. See Trial and Subscription Plans ・ Feature Differences.

When the limit is reached, switches for Agents that are not enabled are locked and display Enablement limit reached. Deactivate an Agent or upgrade your plan.

7-day Cooldown Period

Each time an Agent is enabled, it cannot be deactivated for 7 days. During this period, the list displays Cooldown: N Days Remaining, and the switch is locked with the message In cooldown. Can be deactivated in N days.

> The cooldown period restarts each time the Agent is enabled. If you deactivate and then re-enable it, you must wait another 7 days before deactivating it again. We recommend using Preview to confirm that an Agent is suitable before enabling it, because you cannot swap it out during the cooldown period.

When You Want to Downgrade Your Plan

> If the number of Agents currently enabled exceeds the new plan's limit, the system blocks the plan change. You must first deactivate enough Agents to meet the new plan's limit. Because deactivation is also subject to the 7-day cooldown period, allow time for the cooldown when planning a downgrade.

What Happens When You Deactivate an Agent

> Deactivation does more than hide the Agent. The system removes the assistant along with its conversation channels: The assistant disappears from the AI Assistants list, and members can no longer chat with it Its conversation channels, such as Web Chat, and group permissions are removed Its Agent schedules stop running Your imported knowledge base is retained, with the status Deactivated, Knowledge Base Retained

If you re-enable it later, the system creates **a new assistant** based on the latest official blueprint and **automatically reconnects the original knowledge base**, so you do not need to upload the documents again. However, conversations from before deactivation do not appear under the new assistant.

Configure Authorized Users

Enabling an Agent does not make it available to the entire organization. You must also specify which roles/groups can use it.

Select Configure Authorized Users on the Agent's row. In the drawer, select the roles/groups that can use it, then select Save.

Configured Agents display Authorized: {Role Name} in the list

Agents without authorized users display Not Yet Available

Select the roles/groups allowed to use this Agent in the drawer

If a member sees No Permission to Use on an enabled Agent in the Agent Marketplace, they have not been added to its authorized users. Return here and add them.

Import to the Dedicated Knowledge Base

You cannot change the role instructions or tools of an official built-in Agent, but **you can give it your organization's documents**. Every enabled Agent has a dedicated knowledge base.

Select Import to Knowledge Base on the Agent's row to open its knowledge base settings. From there, use it like any other knowledge base: upload documents, create FAQs, and configure tags. See Knowledge Base Overview.

For example, you could import your company's contract templates and review guidelines into a "Contract Risk Review" Agent, or internal administrative policies into an "Administrative Process Q&A" Agent, so its answers reflect your organization's actual rules.

Clear the Knowledge Base

Select Clear Knowledge Base on the Agent's row to delete its imported documents.

This permanently deletes all imported documents and cannot be undone. Before proceeding, make sure the documents are backed up elsewhere.

The outcome depends on the Agent's current status:

Agent Status	After Clearing
Enabled	Only the knowledge base **content** is deleted. The knowledge base itself remains connected to the Agent, and you can continue importing new documents
Deactivated, Knowledge Base Retained	**The entire knowledge base is deleted and disconnected**. Re-enabling the Agent later creates a new, empty knowledge base

Connectors: Which Ones You Need to Configure First

Some official built-in Agents must connect to third-party services, such as Google Workspace, Microsoft 365, or Notion, to retrieve data. Cards indicate one of two authorization types:

Label	Meaning	Who Takes Action
Administrator Must Configure Connection First	The connector uses organization-level credentials	An **organization administrator** completes the connection setup before members can use it
No label	Personal user authorization	**Each member** authorizes their own account in the conversation the first time the connector is used

> Administrator Must Configure Connection First does not mean that you must authorize every user. It is a one-time, organization-level setup. For connectors without this label, each member authorizes their own account, and the administrator does not need to be involved.

For more information, see Connectors.

AI Meeting Records

Meeting Records is the voice transcription and AI summary feature of the MaiAgent platform, offering both **live** and **offline** modes. Live mode provides low-latency transcription and translation during meetings; offline mode lets you upload audio recordings for high-accuracy transcription and speaker identification. Both modes support AI auto-summary, Q&A, download/export, and knowledge base integration.

Two Modes

	Live Meeting Records	Offline Meeting Records
Input Method	Real-time microphone recording	Upload audio files (mp3, mp4, m4a, aac, wav)
Transcription Features	Low-latency real-time transcription	Multi-language high-accuracy transcription
Translation	Low-latency real-time translation (up to 3 languages)	Post-transcription translation
Speaker Identification	Depends on STT engine support	High-accuracy speaker identification and management
Timeline	Real-time marking	High-accuracy timeline recognition
Subtitle Playback	—	Dynamic subtitle playback, synced with audio
Backup Audio	Synchronized local backup recording in the browser prevents audio loss during network outages (see Meeting Record Backup Audio)	—

Common Features

Feature	Description
Multi-language Transcription	Supports speech recognition in multiple languages

Feature	Description
Multi-language Translation	Translates transcribed content into multiple target languages
Auto Summary	Automatically generates meeting highlights using a custom AI assistant
AI Q&A, Agent Ticket Creation	Ask questions about the transcript; if the selected AI assistant is configured with ticket creation tools, you can ask it to turn meeting action items directly into tasks or tickets
Download & Export	Download MP3 recordings, plain text transcripts, and SRT subtitle files
Knowledge Base Integration	Upload transcripts to the knowledge base for Agent reference

Workflow

Live Mode:

```
Create Meeting → Configure STT Engine & Translation → Start Recording → Low-
latency Real-time Transcription
                                                    ↓
                                        Meeting Ends → Review Transcript
                                                    ↓
                                        AI Summary
 AI Q&A
 Download & Export
```

Offline Mode:

```
Upload Audio (mp3/mp4/m4a/aac/wav) → System Transcription & Speaker
Identification → Review Transcript
                                                          ↓
                                    Dynamic Subtitle Playback
 AI Summary
 AI Q&A
 Download & Export
```

Meeting Record Statuses

Status	Description
Pending (Pending)	Created or uploaded, processing has not started
Recording (Recording)	Transcription in progress in live mode
Processing (Processing)	After recording ends or audio upload, the system is processing
Completed (Completed)	Processing complete, transcript can be viewed and summaries generated
Failed (Failed)	An error occurred during processing

Deployment Options

Meeting Records supports multiple deployment options:

SaaS: Use the MaiAgent cloud service directly

Private Cloud: Deploy in the customer's cloud environment

On-Premises: Install on the customer's local servers

Permissions

Meeting Records requires organization-level activation before use, with three permission levels:

View Permission: Browse the meeting list, view transcripts, and download files

Transcription Permission: Create meetings, start recordings, upload audio files, edit and delete meeting records

Management Permission: Determines whether a user can see all meetings in the organization, or only meetings they participate in

The meeting list **only shows meetings you created or were added to** by default. For details on who can see which meetings and how to share meetings with colleagues, see Meeting Access Control.

Create Meeting Records

Create a New Meeting

Go to the "Meeting Records" page

Click "New Meeting"

Fill in the following information:

Field	Description	Required
Meeting Name	The meeting title for easy searching and identification	Yes
Participants	Select organization members; added members automatically receive viewing access to this meeting	No
External Participants	For customers, vendors, and other non-organization members. Enter a name and press Enter. Names are recorded only in the meeting information and do not receive viewing access. See Meeting Access Control for details	No
Speech Recognition Engine (STT)	Select the speech-to-text engine	Yes
Recognition Language	Languages used in the meeting (up to 3)	Yes

Select a Speech Recognition Engine

The system provides multiple STT (Speech-to-Text) engines, each with different language support and recognition quality:

Engine	Features
Deepgram	High-quality real-time transcription, supports multiple languages
Azure	Microsoft Speech Services, excellent Chinese and English recognition
AssemblyAI	High accuracy, ideal for English scenarios
WhisperLive	Based on OpenAI Whisper, supports a wide range of languages
Bronci	Customized engine

Different engines support different languages and quantities. After selecting an engine, the system automatically displays the available language list.

Configure Real-time Translation (Optional)

If participants speak different languages, enable real-time translation:

Turn on the "Enable Translation" toggle

Select target translation languages (up to 3)

Select the LLM model for translation

Translation occurs simultaneously with transcription, and the transcript displays both the original text and translation results.

Link Knowledge Base (Optional)

You can specify a knowledge base to link when creating the meeting. After the meeting ends, you can upload the transcript to the knowledge base with one click, allowing the Agent to access meeting content.

Manage Meeting Records

On the meeting records list page, you can:

Search: Search by meeting name

Filter: Filter by status (Pending, Recording, Processing, Completed, Failed)

Date Range: View records within a specific time range

Edit: Modify meeting name, language settings, translation settings, and knowledge base links

Delete: Remove unwanted meeting records

Live Meeting Transcription

Live mode provides low-latency voice transcription and translation, allowing you to see the real-time transcript while the meeting is in progress.

Start Recording

After creating a meeting record, click "Start Recording" to begin live transcription.

The system will:

Create a real-time communication room (LiveKit)
Start the speech recognition engine
Begin receiving microphone audio and transcribing in real time

Transcription View

During recording, the screen displays in real time:

Recording Timer: Shows the elapsed recording time
Live Transcript: Speech recognition results appear on screen in real time
Speaker Labels: Indicate the speaker for each text segment
Language Tags: Display the original language of each text segment
Translation Results: If translation is enabled, translated text appears below the original text
Transcription Statistics: Number of completed segments, in-progress segments, and total word count
Local Backup Status: "Recording Local Backup" appears below the meeting name, indicating that the browser is simultaneously recording a backup audio file so no audio is lost if the network disconnects. See Meeting Record Backup Audio for details

Low-latency Real-time Transcription

The core feature of live mode is low latency — transcribed text appears on screen within seconds of speaking, allowing participants to verify the record in real time.

Low-latency Real-time Translation

When translation is enabled, the system translates content into the specified target languages simultaneously with transcription, also with low latency:

Supports up to **3 target languages** for simultaneous translation

Translation is processed in real time by the selected LLM model

Translation results are attached below each transcript segment

Ideal for international meetings and multilingual team communication

Controls

Control	Description
Microphone Toggle	Temporarily mute or unmute the microphone
Stop Recording	End the meeting recording; the system begins post-processing
Search Transcript	Search for keywords in the live transcript
Add Remove External Participants	Record or remove customers, vendors, and other non-organization members who join unexpectedly at any time in the "External Participants" section of the Participants card on the right, without interrupting the recording. See Meeting Access Control for details

After Recording Ends

After clicking "Stop Recording," the meeting status changes to "Processing" and the system performs the following in the background:

Audio Conversion: Convert the raw audio to MP3 format

Transcript Organization: Organize live transcription results into a structured transcript file

Cloud Storage: Upload the audio and transcript to S3 storage

Once processing is complete, the status automatically updates to "Completed," and you can access the meeting detail page to view the full content. The local backup audio uploads and merges in the background. When complete, the details page displays "Backup Complete." See Meeting Record Backup Audio for details.

Meeting Records Quick Access

In the **Office Apps** section on the MaiGPT home page, select Meeting Records to start recording immediately or continue viewing an existing meeting under **Recent Records**.

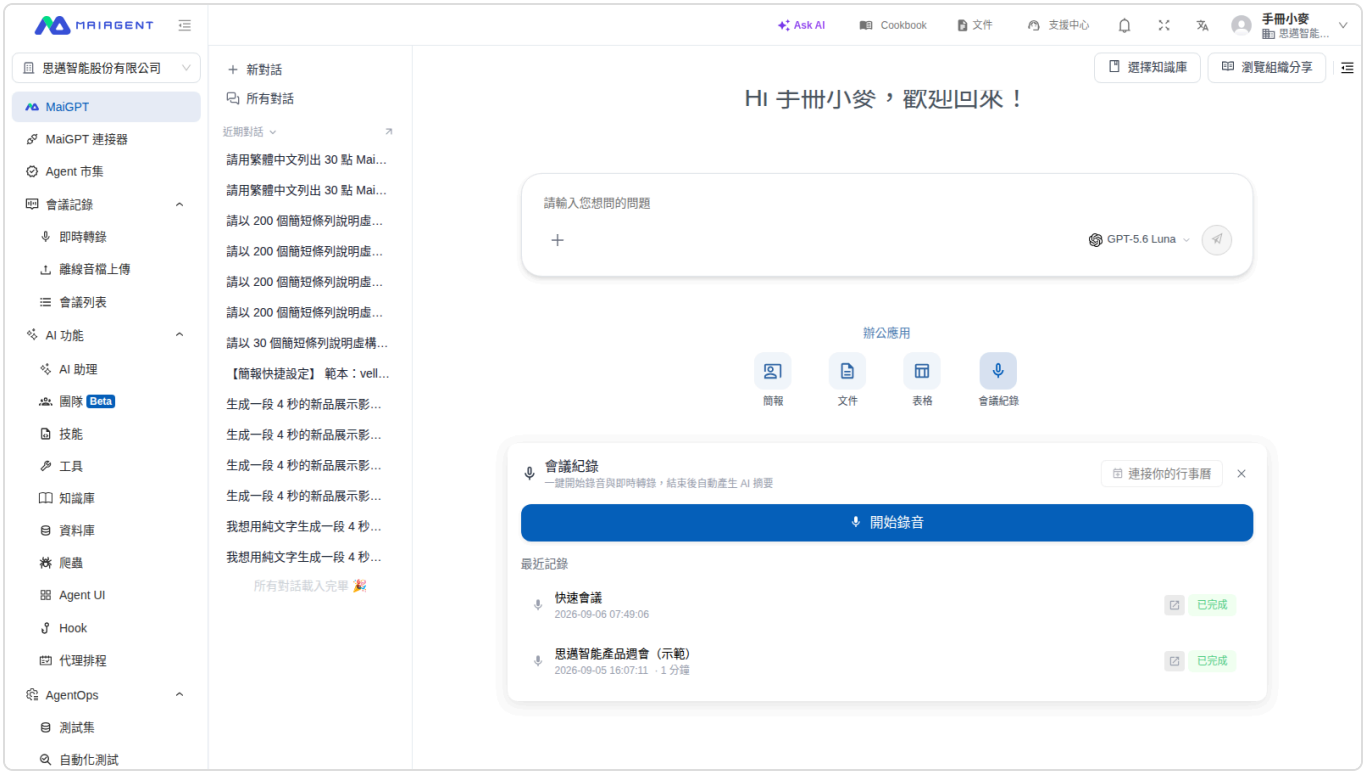

The Meeting Records quick access panel lets you start recording, connect your calendar, and view recent records

Use Case

After a regular project meeting, a project manager wants to check in MaiGPT whether a summary has been generated. She opens Meeting Records, selects the meeting under **Recent Records**, and checks the summary status in the panel. When finished, she selects Back to Entry to return to the start recording screen. If she needs the transcript or other complete details, she can select the button on the right side of the record to open the full record page. This leaves the MaiGPT chat page; after reviewing the record, she can return to MaiGPT from the sidebar.

View Recent Records

1. Open the Quick Access Panel

On the MaiGPT home page, select Meeting Records.

2. Select a Meeting

Under **Recent Records**, select a meeting to view its summary or summary status directly in the panel.

3. Return to the Entry Screen

On the summary screen, select Back to Entry to return to the calendar and start recording screen.

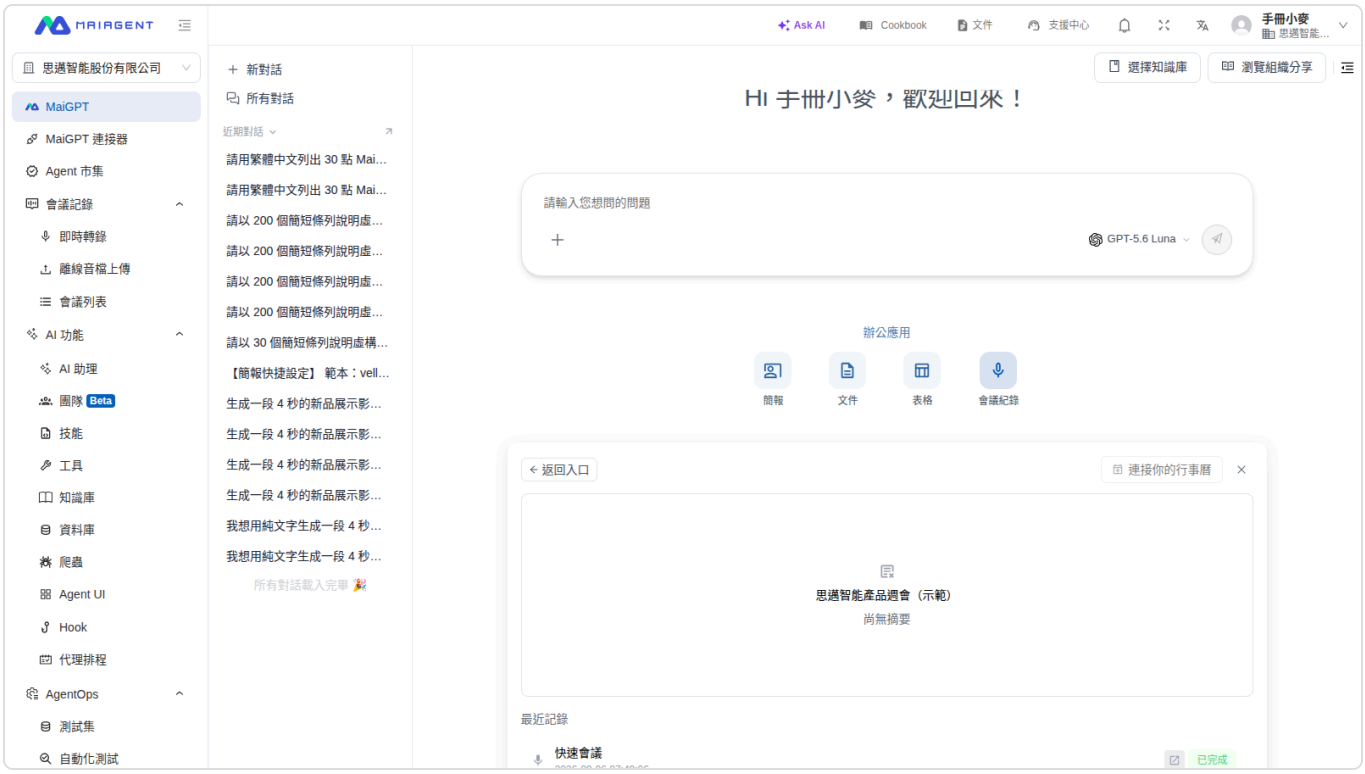

After selecting a recent record, you can view its summary status and return to the entry screen

4. Open the Full Record

When you need the transcript or other complete details, hover over the button on the right side of the record. When the **View Full Record** tooltip appears, select the button to open the full record page for that meeting. This leaves the MaiGPT chat page; after reviewing the record, select MaiGPT from the sidebar to return.

Recent Record Statuses

Screen Status	Meaning	Recommended Action
Generating	The meeting is still being transcribed, processed, or summarized	Reopen the quick access panel later to check the latest status
Generation Failed	Meeting transcription or summary generation failed	Open the full meeting record to review the failure status

Screen Status	Meaning	Recommended Action
No Summary Yet	The meeting is complete, but no summary is available to display	Open the full meeting record to generate or view the summary

> If Recent Records displays No Meeting Records Yet, complete a meeting recording first, then return to this panel to view it.

For complete instructions on recording, transcription, and summaries, see Create a Meeting Record, Live Meeting Transcription, and AI Meeting Summary & Q&A.

Meeting Backup Audio

Live meeting records send microphone audio over the network for cloud transcription. Previously, if a laptop lost its connection during a meeting, audio from the disconnected period could not be sent, leaving a gap in both the cloud recording and transcript. **Backup audio** fills this gap: when recording starts, the browser also records a local backup and uploads it in segments. Even if the network disconnects, local recording continues, and the segments upload automatically when the connection is restored. After the meeting, the system combines all segments into a complete backup audio file—the most complete recording in the system.

> Backup audio is created only for live meeting records (recorded with a microphone). Meetings created by uploading an offline audio file already have a complete file and do not generate backup audio.

Use Case

A customer service manager uses a laptop in a conference room to record a product meeting with Live Transcription. Halfway through, the Wi-Fi becomes unstable. The transcript stops updating, but "Recording Local Backup" remains visible below the title, so she knows no audio is being lost and continues as usual. After clicking "Stop Recording," she sees "Backup Complete" on the meeting details page and switches the player to "Backup Audio" to confirm that the discussion during the outage is included. She then clicks "Retranscribe," changes the source to "Backup Audio," reviews the estimated cost, and confirms. A few minutes later, a backup version appears in the Transcript tab, where she can view the complete meeting without asking colleagues to reconstruct the missing section from memory.

Backup Status During Recording

After you click Start Recording, the backup status appears below the meeting name:

Display Text	Meaning
Recording Local Backup	Cloud recording and local backup are running simultaneously; everything is working normally
Syncing Backup, N Segments Remaining	The network is unstable or disconnected. Recorded segments are queued locally and upload automatically when the connection is restored; no action is required
Local Backup Unavailable (Reason); Meeting Continues	Backup cannot be enabled (for example, the browser is unsupported, browser storage is unavailable, or another tab is already recording). The meeting and

Display Text	Meaning
Normally	cloud recording are unaffected

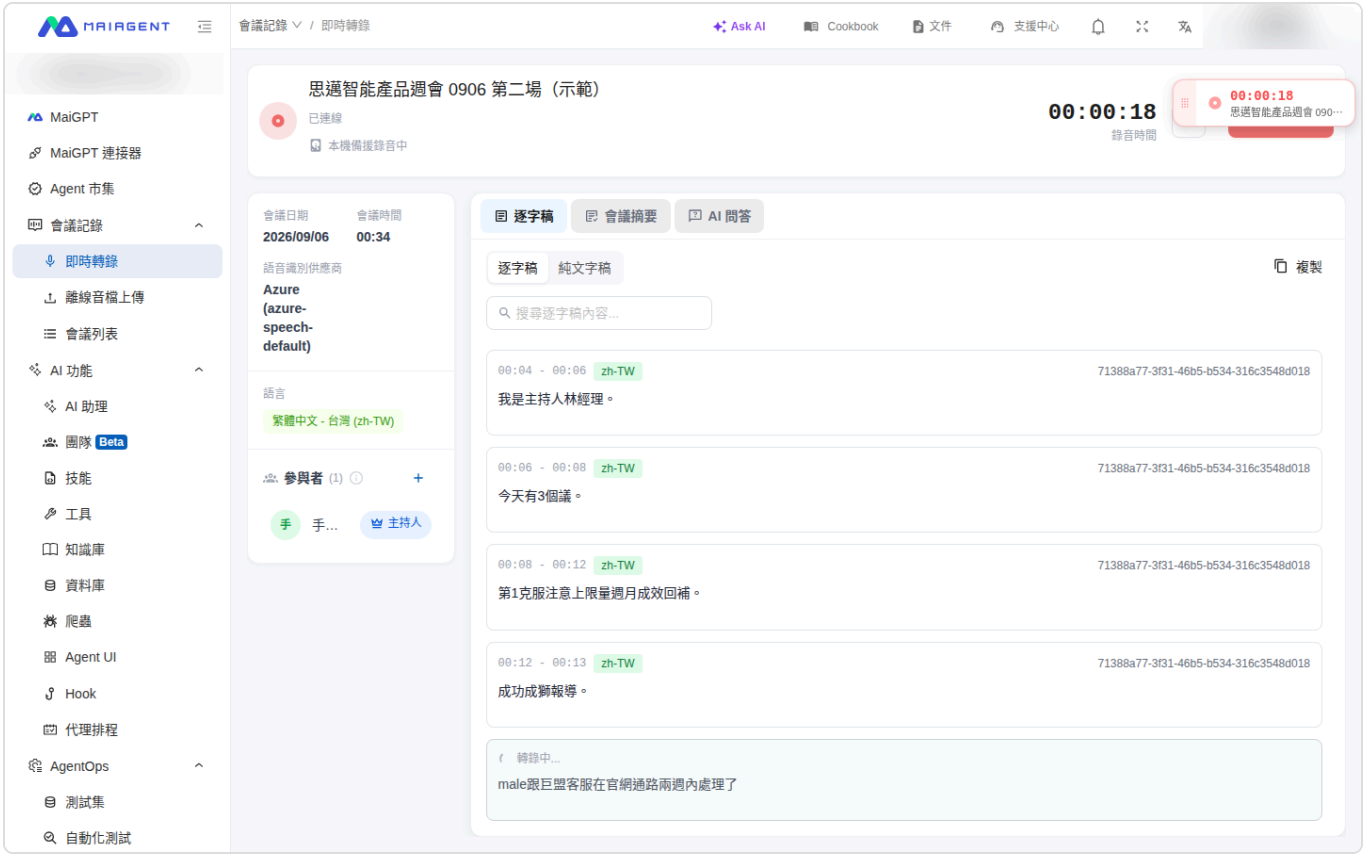

The local backup status appears below the meeting name during recording

- Backup recording runs in the browser. Keep the entire meeting open in the same browser tab. Only one tab can record a meeting; a second tab displays "Recording in Another Tab." - When you click the microphone icon to the left of "Stop Recording" to mute, the corresponding section of the backup audio is also muted, so content you do not want recorded is not captured. - Local backup data awaiting upload is limited to 1 GB. Once the limit is reached, you cannot start another backup recording.

Backup Status After the Meeting

After you click Stop Recording, the meeting status first changes to "Completed." Backup audio has its own sync status, displayed next to the title on the meeting details page:

Status	Meaning
Backup Complete	The meeting is complete and the backup audio is fully synchronized—this backup is the most complete recording in the system
Syncing Backup	The meeting is complete, but the local backup audio is still uploading and merging. The status automatically changes to "Backup Complete" when finished
Backup Incomplete	The backup has not synchronized after more than one day. The system combines the uploaded segments into a backup audio file with the final section missing

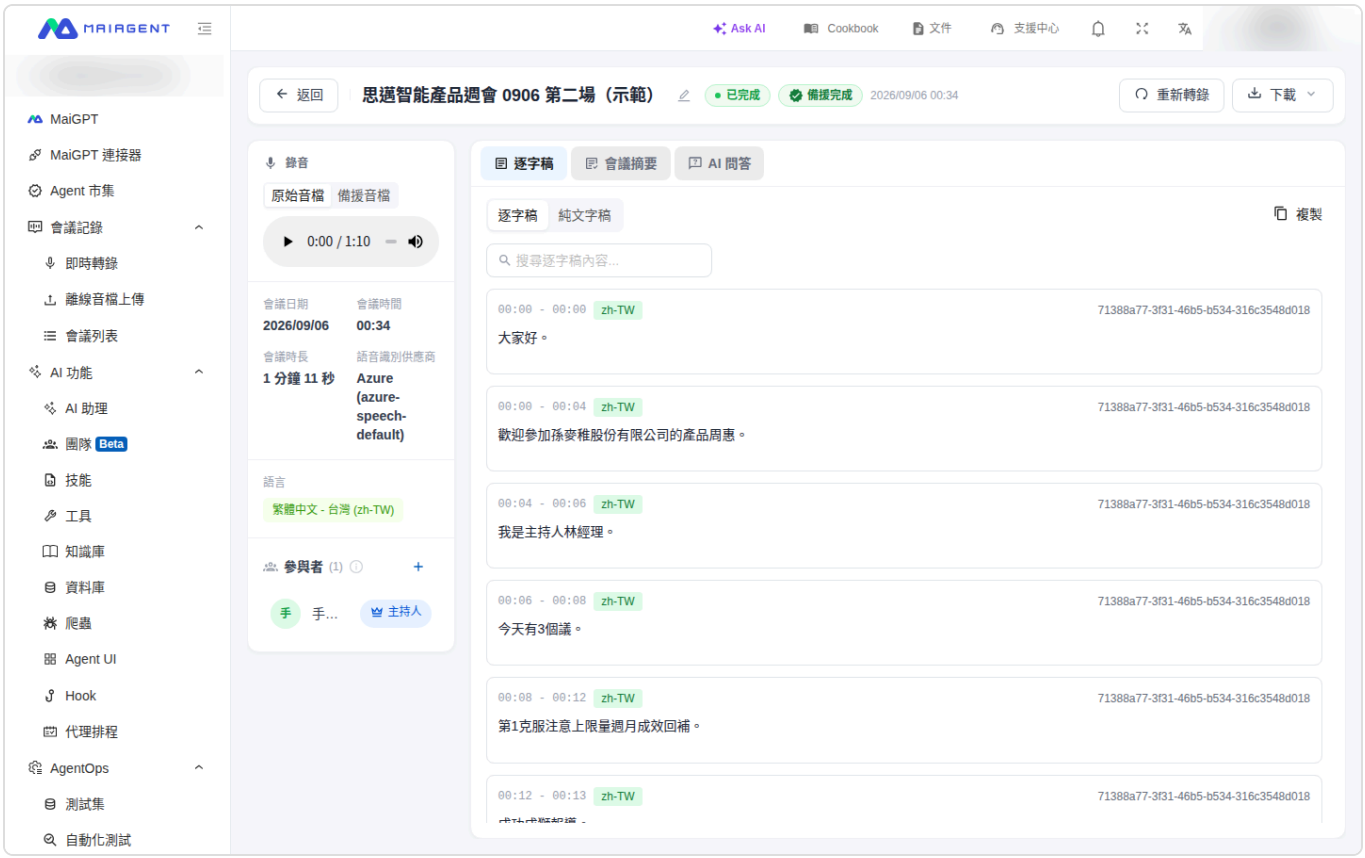

Meeting details page: the "Backup Complete" label and audio source switch

In the meeting list, meetings whose backups are not yet finalized display an additional icon beside the status column. Hover over it to view details. Meetings with completed backups have no additional indicator.

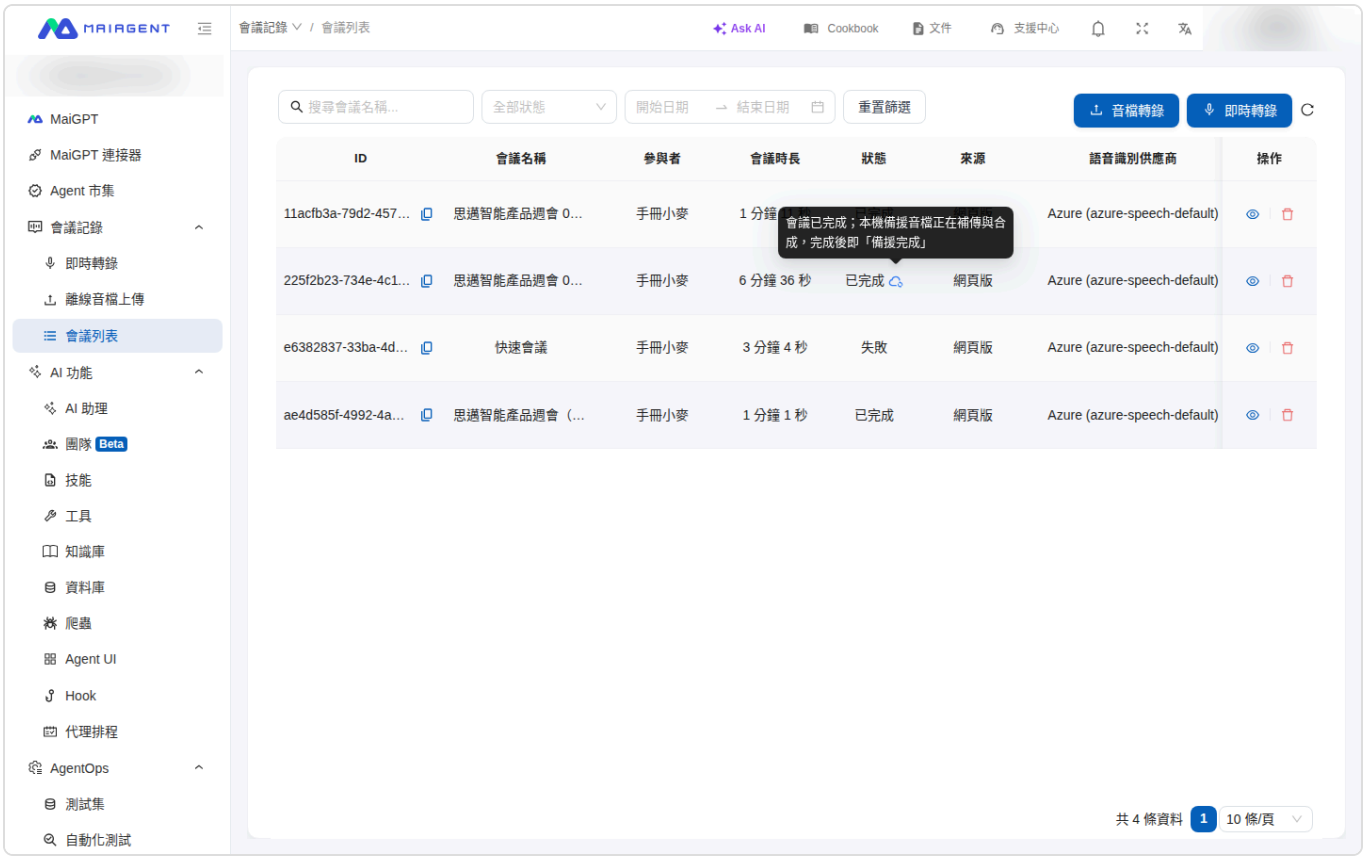

Meeting list: the "Syncing Backup" icon beside the status column

> What happens if I close the browser immediately after losing the network connection? Segments that have not uploaded remain on the local device. Open the MaiAgent admin console again within one day, and the system automatically resumes uploading while the meeting displays "Syncing Backup." If you do not return within one day, the system combines the uploaded portions and labels the backup "Backup Incomplete."

Switch to Backup Audio for Playback

For a meeting with backup audio, an Original Audio
Backup Audio switch appears above the player in the "Recording" section on the left side of the details page. This switch affects **playback only**: meeting summaries, AI Q&A, knowledge base uploads, and MP3 downloads continue to use the original recording (unless you select another transcript source when generating a summary, as described below).

Retranscribe Using Backup Audio

If the original transcript is missing a section due to a network outage, you can retranscribe the backup audio to create a new transcript version. The existing transcript is retained.

Open Retranscription

Open the meeting details page and click Retranscribe in the upper-right corner.

Select "Backup Audio" and Confirm the Cost

Under Transcription Source, select Backup Audio. Billing information appears at the bottom of the dialog, including the billable item (audio transcription), backup audio duration, and estimated cost. It also explains that "you are charged only if transcription succeeds; failed attempts are free" and "each retranscription is billed separately."

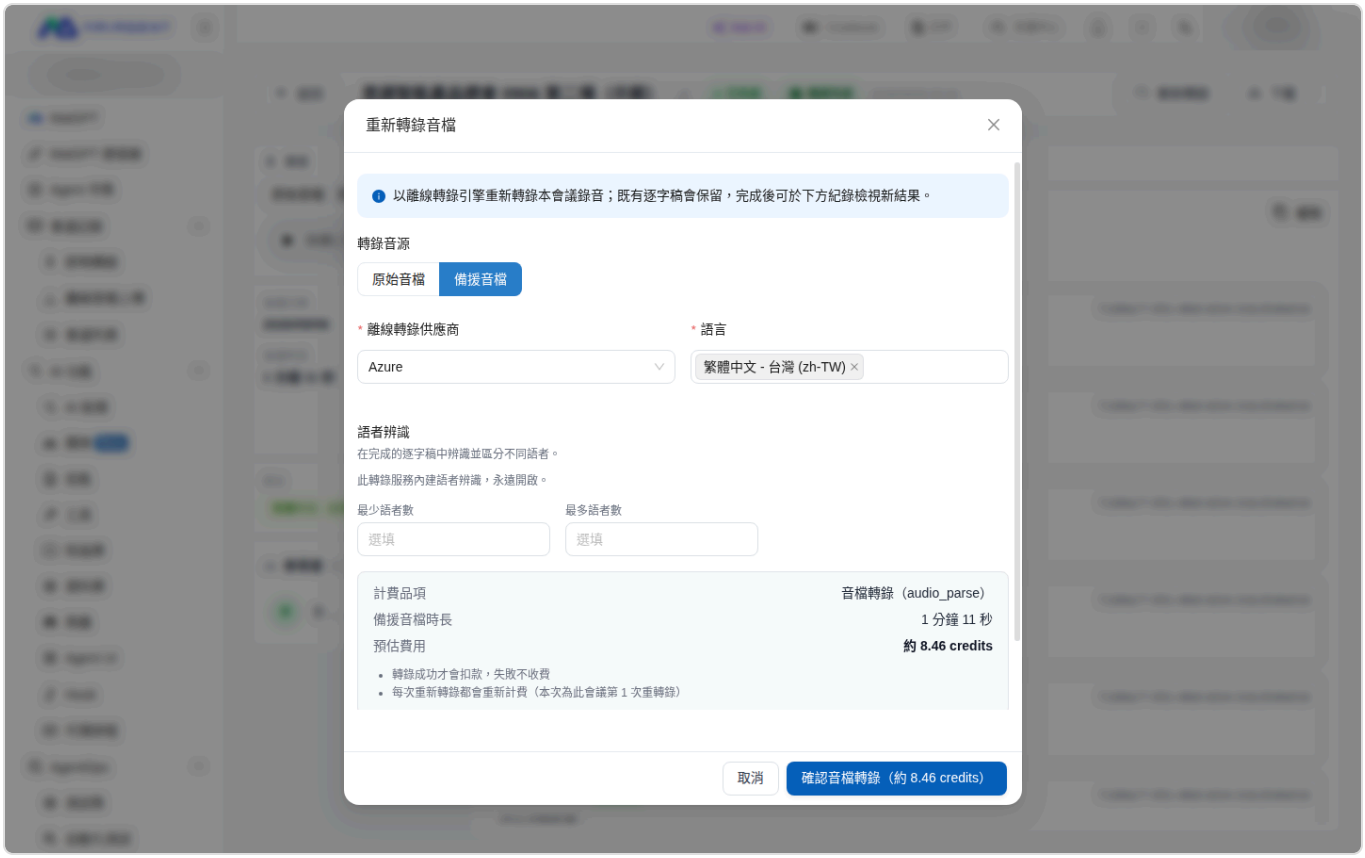

The dialog discloses the estimated cost and billing rules

Submit the Transcription

Click Confirm Audio Transcription (Approx. X credits)—the button displays the estimated charge. After submission, the screen indicates that transcription has started and that you will be charged only if it succeeds. A "Processing" entry appears under "Retranscription History."

View the Backup Transcript Version

After transcription is complete, a version labeled "Backup" appears in the version menu on the Transcript tab (for example, "Azure • 09/06 00:42 • Backup"). When selected, the player **automatically switches to the backup audio**, allowing you to click the transcript and jump accurately to the corresponding timestamp.

Version menu on the Transcript tab: original transcript and backup version

- The backup transcript aligns only with the backup audio. If you manually switch to "Original Audio" while viewing the backup version, the screen warns: "The transcript and the audio being played come from different recordings. Clicking the transcript cannot jump precisely to that sentence." - If the organization wallet has insufficient funds, "Backup Audio" is disabled and displays "Insufficient wallet balance. Please add funds first." No transcription job is created, and no charge is made. - If the system detects that the end of the original transcript may be missing, a "The end of the transcript may be missing approximately N minutes" notice appears at the top of the Transcript tab, with a Retranscribe Using Backup Audio shortcut.

Generate a Summary from a Backup Transcript

When you click Generate Summary on the Meeting Summary tab, Transcript Source defaults to "Original Transcript." If a backup version is available, you can select it so the summary includes the recovered sections. The dialog then states: "The backup transcript timeline corresponds to the backup audio and differs from the original audio."

Select a transcript source when generating a summary

Further Reading

Live Meeting Transcription: Start recording and view the recording screen

Meeting Record Details: Transcript, player, and meeting information

AI Meeting Summary & Q&A: Generate summaries and manage summary versions

Offline Meeting Transcription

Offline mode lets you upload pre-recorded audio files. The system performs high-accuracy transcription and speaker identification, generating a transcript with timeline markers.

Supported Audio Formats

Format	Extension
MP3	`.mp3`
MP4	`.mp4`
M4A	`.m4a`
AAC	`.aac`
WAV	`.wav`

Upload Audio

Go to the "Meeting Records" page

Select upload audio

Choose a local audio file (supports mp3, mp4, m4a, aac, wav)

The system begins background processing

After uploading, the meeting record status changes to "Processing" and the system automatically starts transcription.

Multi-language High-accuracy Transcription

Offline mode uses a high-accuracy speech recognition engine (such as WhisperX), which provides higher recognition accuracy compared to live mode:

Supports multi-language recognition

Higher recognition accuracy than live transcription

Ideal for important meetings that require precise records

Speaker Identification and Management

Offline mode supports automatic Speaker Diarization, where the system automatically distinguishes between different speakers:

Auto Identification: The system automatically labels different speakers (e.g., SPEAKER_00, SPEAKER_01)

Custom Names: Replace system labels with real names (e.g., "Manager Zhang", "Team Lead Li")

Management Interface: Manage all speaker label names in one place

After modifying speaker labels, all related segments in the transcript are updated simultaneously.

High-accuracy Timeline Recognition

Each transcribed text segment includes precise start and end times, allowing you to:

Pinpoint the exact time of each utterance

Sync with the audio player for jump navigation

Export as SRT subtitle files with timestamps

Dynamic Subtitle Playback

Offline mode provides dynamic subtitle playback that syncs the audio player with the transcript:

Sync Highlighting: During audio playback, the current transcript segment is automatically highlighted

Click to Jump: Click any transcript segment to jump the audio to the corresponding time point

Search and Navigate: After searching for a keyword, jump directly to that segment's audio position

This feature lets you quickly find specific discussion content in a meeting without listening to the entire recording from the beginning.

Meeting Record Details

After meeting recording processing is complete, you can view the full meeting content on the detail page.

Page Structure

The detail page contains the following sections:

Audio Player: Play the meeting recording MP3. For live meetings with backup audio, you can switch between "Original Audio" and "Backup Audio" above the player. See Meeting Record Backup Audio for details

Meeting Info Card: Meeting name, participants, recording duration, creation time

Content Tabs: Transcript, Plain Text, Summary, AI Q&A, Download

Transcript Tab

Displays the complete transcript in timeline format:

Timestamps: Start time of each text segment

Speaker Names: Indicate who is speaking

Language Tags: Original language labels

Translation Content: If translation is enabled, displays translation results

Keyword Search: Search for specific content in the transcript

Plain Text Tab

Presents the transcript in plain text format for easy reading and copying of the complete meeting content. Also supports keyword search.

Meeting Info Sidebar

The right sidebar displays:

Basic Information: Meeting date, meeting time, meeting duration, speech recognition provider

Participants: The access roster for this meeting (the creator is labeled as host); you can add or remove members — see Meeting Access Control for details

External Participants: Names of attendees who are not organization members. You can add or remove names; this is a text-only roster and does not represent access permissions. See External Participants for details

Meeting Access Permissions

Meeting content often involves sensitive information, so the meeting list **does not show all meetings in the organization**. By default, each member can only see meetings they are involved in; to share a meeting with colleagues, grant access through the "Participants" roster within the meeting.

Who Can See Which Meetings

Identity	Visible Scope in Meeting List
Regular Member	Only meetings they created + meetings they were added to as a participant
Organization Owner	All meetings in the organization
Member with "Meeting Record Management Permission"	All meetings in the organization

Meetings without access permission will display a "meeting not found" message even when accessed directly via URL — the system will not reveal its existence.

Meeting list: The "Participants" column shows the access roster for each meeting

> The person who created the meeting automatically becomes the host and can always see the meetings they created.

Participants: The Meeting Access Roster

The "Participants" card on the right side of the meeting detail page is the **access roster** for that meeting — only people on the roster can see the meeting.

會議日期

2026/08/17

會議時間

18:29

會議時長

54 分鐘 5 秒

語音識別供應商

-

參與者 (3)

陳 陳文彬 主持人

王 王志明

林 林佳蓉

逐字稿講者 (3)

陳文彬 林佳蓉 王志明

Participants card: Host label, member list, add participant button

Card contents:

Host: The meeting creator, with a host label next to their name

Members: Organization members who have been added

External Participants (bottom of card): Names of attendees who are not organization members; this is **a text-only roster and does not represent access permissions**. See External Participants for details

The "Participants" column in the meeting list also displays this roster, showing `+N` when there are more than 2 people.

> "Participants" and "External Participants" are two different things. External participants are recorded by name only and cannot access the meeting; to grant viewing access, you must add an organization member to the participants roster.

Who Can View and Who Can Edit the Roster

Action	Condition
View participants roster	Anyone who can see the meeting
Add Remove participants	Organization Owner, member with "Meeting Record Management Permission", or **any participant already on the roster**

Add a Participant

1. Open the meeting detail page

Click on the meeting you want to share in the meeting list.

2. Click the "+" in the upper right corner of the participants card

The card will expand a member selection dropdown.

3. Select an organization member and click "Add"

會議日期
2026/08/17
會議時間
18:29
會議時長
54 分鐘 5 秒
語音識別供應商
-
參與者 (3)
選擇成員
黃郁婷（huang.yut...
許珮琪（xu.peiqi@...
陳
陳文彬
主持人
王
王志明
林
林佳蓉

Select the person to add from organization members

The dropdown only lists members from the **same organization**, with keyword search available
Members already on the roster will not appear in the dropdown
After successfully adding, the person can **immediately** see the meeting in their meeting list

Remove a Participant

Click the "X" on the right side of the member in the participants roster and confirm to remove them. The person immediately loses viewing access to the meeting.

> The host cannot be removed. The meeting creator's row does not have a remove button, ensuring the meeting always has at least one person who can manage it.

External Participants: Record Non-Organization Members

Attendees such as customers, vendors, and consultants who are **not organization members** cannot be added to the "Participants" roster, but you can record them under "External Participants." External participants are **a text-only roster**: their names are recorded in the meeting information to provide a complete attendee list, but they **do not receive viewing access or notifications and do not need an account**.

Use Case

A sales manager is holding a product onboarding meeting with two colleagues from her team and two customer contacts. She opens "Live Transcription" and selects the two colleagues under "Participants" so they can view the transcript after the meeting. She then enters the customer contacts' names and company under "External Participants." Halfway through the meeting, the customer's technical manager

joins unexpectedly. She clicks "+" in the "External Participants" section on the right and adds the manager's name without interrupting the recording. After the meeting, the Participants card on the meeting detail page clearly separates the two groups: colleagues appear under "Participants," while customers appear under "External Participants." The full attendee list is easy to review, while the customers still cannot access the meeting content.

Record Participants When Creating a Meeting

In the "Live Transcription" setup form, the "External Participants" field appears below "Participants." Enter a name and press Enter to turn it into a tag. You can add multiple names in succession. When recording starts, the roster is created along with the meeting.

"Participants" only lets you select organization members and grants them viewing access; "External Participants" lets you freely enter names that are recorded only in the meeting information

> The "External Participants" field is currently available in the Live Transcription setup form. It is not available in the offline audio upload form or the edit dialog in the meeting list. For offline meetings or names added afterward, use the sidebar on the meeting detail page (see the next section).

Add or Remove Participants in the Sidebar During or After Recording

The Participants card on the right side of both the Live Transcription page (while recording) and the meeting detail page (after the meeting) includes an "External Participants" section. You can add or remove names at any time:

1. Add

Click "+" next to "External Participants," enter a name, and press Enter or click "Add." The roster updates immediately after it is saved and remains current after you leave, return to, or refresh the page.

2. Remove

Click the "X" next to a name to remove it. No confirmation dialog appears; if you remove someone by mistake, you can add the name again immediately.

The sidebar has two sections: "Participants" are organization members with viewing access, while "External Participants" are a text-only roster

Anyone who can add or remove external participants has the same permissions as someone who can edit the meeting: members with "Transcription Permission" who can view the meeting. Other users can only view the roster

If saving fails, the page displays "Failed to update external participants." The roster remains unchanged, and recording is not affected

Roster Rules

Rule	Description
Maximum number	Up to 20 people per meeting. Once the limit is reached, no more names can be added and a notification appears
Name length	Up to 50 characters per person. If the name is too long, a notification appears and the name is not automatically truncated
Duplicate names	Identical names are treated as the same person and merged automatically
Whitespace	Empty input or input containing only spaces is not added to the roster

These rules apply consistently to the setup form, the sidebar during recording, and the sidebar after the meeting.

> Meetings started from a calendar invitation (for example, starting a recording after connecting a calendar in the MaiGPT meeting record shortcut panel) automatically import the invitee list. Invitees who are organization members become "Participants" and receive viewing access, while everyone else is recorded as an "External Participant." You can still adjust the roster afterward in the sidebar.

Grant "See All Meetings" Access

For scenarios like auditing, handover, or department managers who need an overview of all meetings in the organization, grant access through role permissions:

1. Go to the Role Permission page in Organization Settings

See Role Permission Management.

2. Edit role permissions and check "Meeting Record Management Permission"

Check "Meeting Record Management Permission" to allow role members to see all meetings in the organization

"Meeting Record Management Permission" is located at the bottom of the permission list and is an independent main permission:

Unchecked: Role members can only see meetings they participate in (created or added to)

Checked: Role members can see all meetings in the organization

When checking "Meeting Record Permission", the system will **automatically check** this permission as well; if you only want the role to see their own meetings, manually uncheck it

> Members with "Meeting Record Management Permission" who only want to view their own meetings can filter using the API parameter ?scope=mine.

FAQ

A colleague says they can't see a meeting I created?

Add them to the meeting's participants roster — it takes effect immediately.

How do I record customers or vendors who attend a meeting?

Add their names under "External Participants." This is only a roster and does not let them access the meeting content. If the person is an organization member, add them to the "Participants" roster instead to grant viewing access.

After unchecking "Meeting Record Management Permission", can I still see old meetings?

You can only see meetings you created and meetings where you were added to the participants roster. Meetings created by others that you were not added to will disappear from the list — if needed, ask a participant of that meeting to add you.

AI Meeting Summaries and Q&A

Meeting summaries are processed in the background, so you do not need to remain on the summary page while a long meeting is being processed. You can leave the page during generation and return later to check the progress. When processing is complete, the summary is created as a new version.

Use Case

A project manager has just finished a long quarterly meeting and wants to organize the decisions before leaving work. She opens the Meeting Summary tab for the meeting and starts generating a summary with the default template. After seeing that the system is analyzing the transcript, she leaves the page to work on other tasks. When she returns to the same meeting later, the page continues to display the progress. Once processing is complete, she can review the decisions and action items in the newly created summary version.

Custom Summary Agent

The meeting records summary feature lets you select a custom AI assistant (Agent) to generate summaries. Different Agents can have different role instructions to produce summaries in different styles:

General Summary Agent: Organize meeting highlights, decisions, and action items

Technical Meeting Agent: Focus on technical decisions, architecture discussions, and technical debt items

Business Meeting Agent: Focus on client requirements, sales progress, and opportunity tracking

Generate a Summary

Go to the meeting detail page

Switch to the Meeting Summary tab

To customize the summary format, select the AI assistant to generate the summary; if none is selected, the default template is used

Click Generate Summary

The AI uses the complete transcript to organize the meeting highlights, decisions, and action items. When you use the default template, the page displays the Analyzing transcript N/M status followed by the summary compilation status. N represents the number of completed segments, and M represents the

total number of segments. When you select a custom AI assistant, the assistant processes the transcript directly and the segmented N/M progress is not displayed.

Whether you use the default template or a custom AI assistant, you can leave the page while processing continues in the background. Refresh the page or return to the meeting later to see the current status.

Leave the page and return later to continue viewing the default template's segmented progress

When the summary is complete, the new version appears in the version list

> When the summary is complete, a new version is created automatically; you do not need to save it separately. If generation fails, the page prompts you to try again. Confirm that the transcript is available, then click Regenerate.

Summary Version Management

Each time a summary is generated or edited, the system retains a complete version history:

Add Summary: Each summary generation creates a new version

Edit Summary: Modify existing summary content

Version History: View all summary versions and their creation times

Delete Summary: Remove unwanted summary versions

Copy a Summary

When viewing a saved summary, a Copy Summary button appears in the toolbar above it. Use the button to copy the entire summary to the clipboard without manually selecting all the text, making it easy to paste into email, documents, or messaging apps.

Go to the meeting detail page and switch to the "Meeting Summary" tab

Select a saved summary

Click Copy Summary

After the "Copied successfully" notification appears, paste it into another application

The toolbar in the "Meeting Summary" tab

The summary's original text is copied with its paragraph structure preserved.

> Copy Summary is available only when viewing a saved summary. It does not appear while editing a summary or previewing an unsaved draft. The button is disabled when the summary is empty.

If "Copy failed" appears, the browser has usually blocked clipboard access. Select and copy the text manually instead.

Read a Summary Aloud

When reading is inconvenient, have the system read the summary aloud. When viewing a saved summary, a Read Aloud button appears in the toolbar. Click it to hear the summary.

Go to the meeting detail page and switch to the "Meeting Summary" tab

Select a saved summary

Click Read Aloud. The button displays "Preparing⋯" before playback begins

During playback, click Pause to pause, Resume to continue, or Stop to end playback

During playback, "Read Aloud" changes to "Pause" and "Stop"

Before playback, the system automatically removes Markdown symbols such as `#`, `*`, and `-` from the summary so they are not read aloud.

Voice and Supported Languages

The voice is determined by **the meeting's spoken language or target translation language**, not by the actual language of the summary. The following seven languages are currently supported:

Language	Language Code
Traditional Chinese (Taiwan)	`zh-TW`
English (United States)	`en-US`
Japanese	`ja-JP`
Korean	`ko-KR`
Thai	`th-TH`
Indonesian	`id-ID`
Vietnamese	`vi-VN`

> If the meeting language is not on the list above, such as Simplified Chinese zh-CN or Cantonese zh-HK, the toolbar displays a message saying that the language is not currently supported and the default language will be used. You can still use Read Aloud, but it uses a Traditional Chinese voice.

Limitations

Read Aloud is available only when viewing a saved summary; it does not appear while editing a summary or previewing an unsaved draft

The button is disabled when the summary is empty

Playback stops automatically when you switch summary versions or leave the "Meeting Summary" tab, so audio does not continue in the background

The button is not displayed if the browser does not support audio playback

If playback fails, "Text-to-speech failed. Please try again later" appears; try again later

AI Q&A

In addition to the summary feature, you can ask questions directly about the transcript content:

Switch to the "AI Q&A" tab

Select the AI assistant to use

Enter your question (e.g., "What decisions were made in this meeting?" "What are the next action items?")

The AI finds relevant content from the transcript and provides an answer

This feature is especially useful for quickly reviewing specific content from long meetings without re-reading the entire transcript.

Agent Ticket Creation

After a meeting ends, action items and decisions usually need to be manually entered into a project management or ticketing system. You can let the AI assistant handle this step: in the AI Q&A tab, ask it to turn the items discussed in the meeting directly into tasks or tickets.

> Meeting Records does not have a dedicated "Create Ticket" button. AI Q&A uses the full AI assistant conversation, so as long as the selected AI assistant is configured with tools that can create tickets, you can ask it to create tickets in the conversation — the ticket creation capability comes from the AI assistant, not from Meeting Records.

How It Works

Step	Details
① **Meeting Transcript**	Meeting transcription is complete, and the transcript has been uploaded to the knowledge base and parsed
② **Select AI Assistant**	In the AI Q&A tab, select an AI assistant configured with ticket creation tools
③ **Issue Ticket Creation Command**	For example, "Compile the action items from this meeting and create a ticket for each one"
④ **AI Assistant Calls Tools**	The assistant calls API tools or MCP tools based on the field mappings and ticket creation rules in its role instructions
⑤ **Create Tickets**	Tickets are written to your company's ticketing or project management system
⑥ **Report Results**	The assistant reports each ticket's number and link in the conversation

Prerequisites

Condition	Description
Meeting is linked to a knowledge base, and transcript parsing is complete	AI Q&A requires the transcript to be in the knowledge base before it can be used. You can link a knowledge base when creating a meeting; if not set at that time, you can also select a knowledge base in the AI Q&A tab and click Upload Transcript. See Download and Knowledge Base Integration for details
An AI assistant configured with ticket creation tools	The tool, not Meeting Records, is what actually calls the ticketing system. See the setup steps below
Ticketing system provides a callable API	MaiAgent communicates with external systems via API tools or MCP tools, without being limited to any specific vendor; as long as the other party provides an API for creating tickets, it can be integrated

Setup Steps

Create a ticket creation tool: On the Tools page, click Add API Tool and configure your company's ticketing system's "Create Ticket" API (you can paste a cURL command directly and the system will automatically parse the URL, headers, and HTTP method). See Create API Tool for details. If the other party provides an MCP server, you can also use MCP Tool instead.

Assign the tool to the AI assistant: Add the newly created tool to the AI assistant's available tools list. See Configure Tools for AI Assistant.

Define ticket creation rules in the role instructions: Whether the AI calls the tool at the right time and how fields are filled depends on the role instructions. Refer to the example below, and see the Role Instructions Design Guide for writing principles.

Use in meetings: Go to the AI Q&A tab on the meeting detail page, select this AI assistant, and issue a command directly, such as "Compile the action items from this meeting and create a ticket for each one."

> Ticket creation writes to an external system. It is recommended to verify field mappings with a test project or test environment before switching to your production ticketing system.

Role Instructions Example

```
# Role
You are a meeting action item ticketing assistant.

# Task
Users will ask you to create tickets based on meeting transcripts. Follow this
process:

1. Identify clear action items and decisions from the transcript; ignore pure
discussions or items that haven't been finalized
2. Organize each action item into a ticket with the following field mappings:
   - Title: One sentence describing what needs to be done, under 30 characters
   - Content: Add the background and rationale mentioned in the meeting
   - Assignee: The person explicitly assigned in the transcript
   - Due date: The date explicitly mentioned in the transcript
3. Before calling the tool, present the compiled ticket list to the user for
confirmation; only create after getting approval
4. After creation, report each ticket's number and link

# Constraints
- Only create tickets based on transcript content; do not add tasks that were not
discussed
- If the assignee or due date cannot be determined from the transcript, leave
blank and note it in the response; do not guess
- When the same item is mentioned multiple times in the meeting, merge into a
single ticket
```

Verification and Troubleshooting

The parameters actually received by the tool and the returned results can be viewed in AgentOps' Tool Execution Records — this is the fastest way to troubleshoot ticket creation issues.

Issue	Cause and Resolution
AI Q&A is unavailable	The meeting is not linked to a knowledge base, or the transcript is still being parsed. Select a knowledge base in the tab and upload the transcript, then try again after parsing is complete
AI listed the action items but did not actually create tickets	The tool was not added to the assistant's tool list, or the role instructions do not explicitly require calling the tool. Check steps 2 and 3
Ticket creation failed	Check the error message in the tool execution records. Common causes include missing required fields, expired authentication, or insufficient permissions
Ticket field content is incorrect	Make the field mapping rules more specific in the role instructions, and explicitly instruct it to leave fields blank rather than guess when information cannot be determined

Issue	Cause and Resolution
AI's response is unrelated to this meeting	The currently selected AI assistant may not be linked to this meeting's knowledge base; the interface will show a reminder. Switch to a linked assistant

Downloads and Knowledge Base Integration

Download Files

After meeting processing is complete, you can download the following files:

Format	Description	Use Case
MP3 Audio	Meeting recording in MP3 format	Archiving, offline playback
Plain Text Transcript	Complete transcript in plain text format	Copy and paste, document archiving
SRT Subtitle File	Transcript in subtitle format with timestamps	Video playback, subtitle import

On the meeting detail page, go to the "Download" tab and click the corresponding download button to get the file. The plain text transcript can also be copied with one click from the "Plain Text" tab.

Upload to Knowledge Base

After uploading the meeting transcript to the knowledge base, the Agent can answer questions based on meeting content:

On the meeting detail page or list page, select "Upload to Knowledge Base"

Select the target knowledge base (if a knowledge base was linked when creating the meeting, it will be pre-filled)

Confirm the upload

After uploading, the transcript becomes a document in the knowledge base, and the Agent can retrieve meeting discussion content through RAG.

Use Cases

Project Meetings: Import each project meeting transcript into the knowledge base so the Agent can answer "What was decided in the last meeting?"

Client Interviews: Import client interview records for easy reference of client requirements

Training Records: Import internal training content to build an organizational knowledge base

Internal Q&A

Introduction to Internal Q&A

Internal Q&A uses the built-in **Internal** channel that is automatically created when you create an AI assistant. No integration setup is required. Organization members can select an assistant from the **Agent Marketplace** in the left menu to open this page. Under **All Conversations**, the source of these conversations is shown as **Internal**.

💬 Conversation area

Chat with an AI assistant in real time

📋 Conversation history

View and manage all past conversations

🤖 AI assistant selection

Switch between AI assistants for conversations

⚙️ AI assistant settings

Adjust the AI assistant's permissions, tool configuration, and other settings for the current conversation in real time

Displayed only when you have permission to edit the assistant

📚 Knowledge base document management

Precisely control which documents the AI assistant can reference in each conversation

How do I use Internal Q&A?

1. Open the Internal Q&A page from the Agent Marketplace

Click **Agent Marketplace** in the left menu, then select the AI assistant you want to chat with to open this page.

---

2. Select an AI assistant

Use the blue drop-down menu in the upper-left corner to select the AI assistant you want to chat with.

3. Start a conversation

Continue a previous conversation: Select a topic from the previous Q&A topics on the left.
Start a new conversation: Click **New Q&A** in the left panel to start a new conversation.

> All conversations appear in the Q&A history in the upper-left corner.

After selecting a conversation, enter your question in the text box below to start chatting.

Continue a previous conversation (select a conversation)

Start a new conversation (new conversation)

4. Customize the document reference scope (optional)

Use the knowledge base management feature on the right to select the specific documents the AI assistant can reference in this conversation.

You can select the reference documents and FAQs separately for each message

You can change the selection before sending each question

Document access scope

When sending the first message, allow access to all documents, including `氣氛燈-系列.md`.

After referencing the document, the AI assistant responds: "Document content: The official website has been closed, and no products are sold externally anymore."

When sending the second message, clear the selection for `氣氛燈-系列.md`.

The AI assistant cannot reference it and responds: "I can't find the specific content of the **`氣氛燈-系列.md`** file."

FAQ access scope

When sending the first message, allow access to all FAQs, including "How do LLMs work?"

The AI assistant references the FAQs and provides a relevant answer.

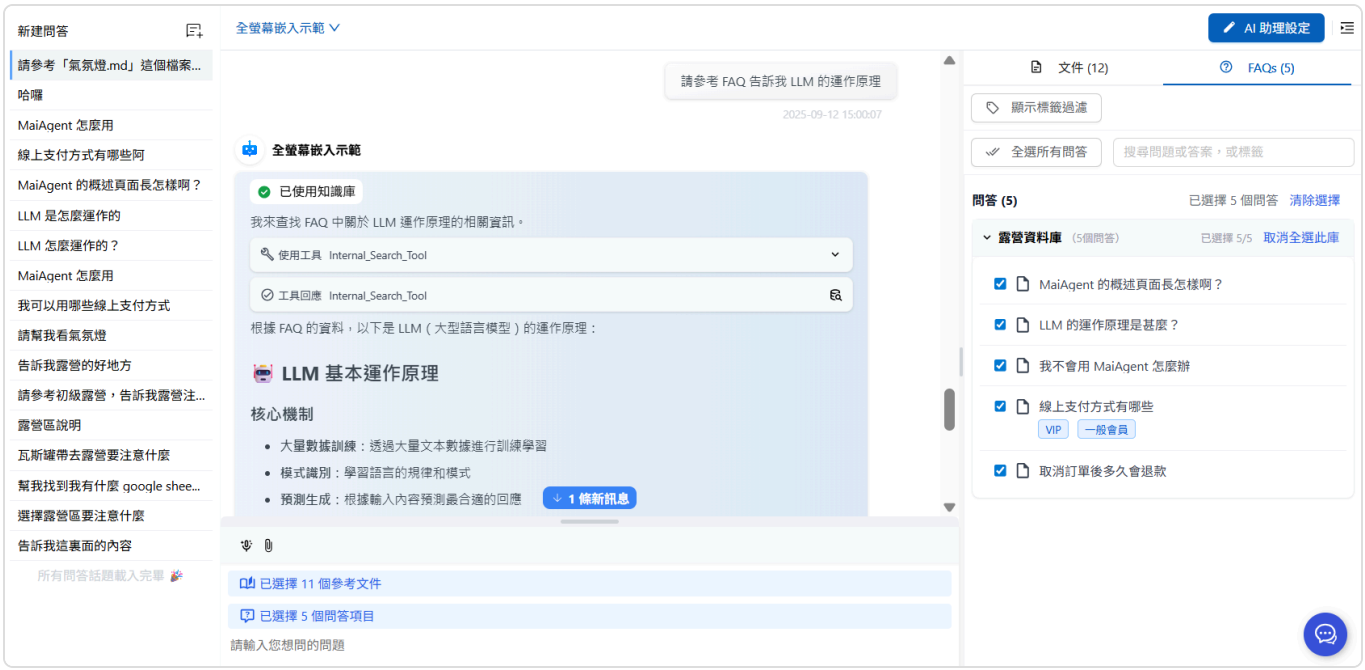

When sending the second message, clear the selection for "How do LLMs work?"

The AI assistant cannot reference "How do LLMs work?" and responds, "I can't find any information about how LLMs work in the FAQs."

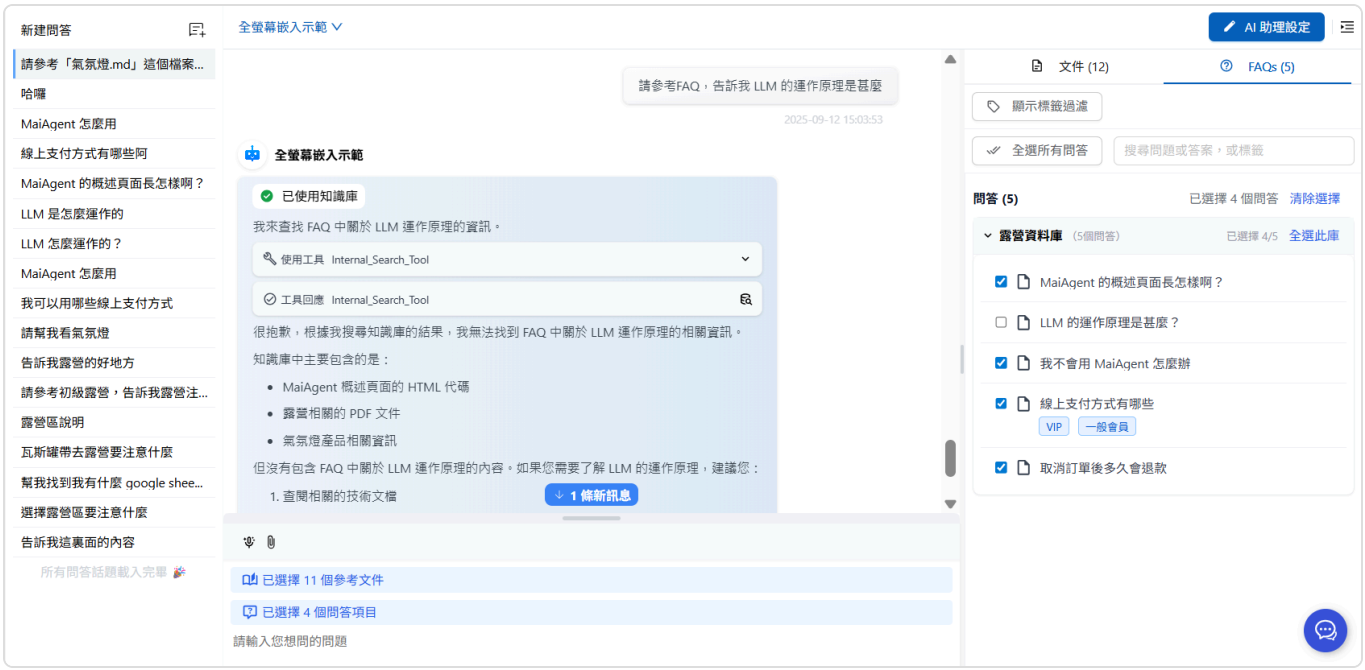

Filter by label

Use labels configured in the knowledge base to filter the documents that the AI assistant can reference:

> For label configuration instructions, see Document management: Labels and metadata.

a. Click **Show Label Filter** in the upper-right corner.

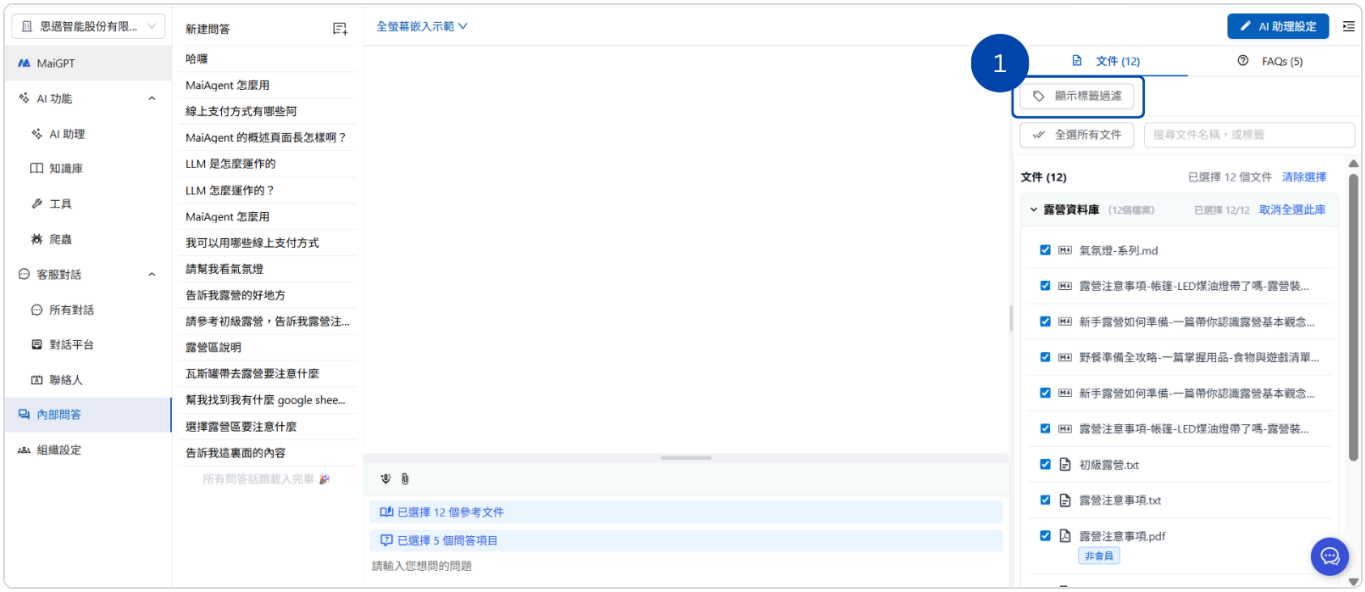

b. Select the matching method: Click **Include all labels (AND)** or **Include any label (OR)** to switch between them. (The image below shows the previous version's AND/OR buttons in the same location.)

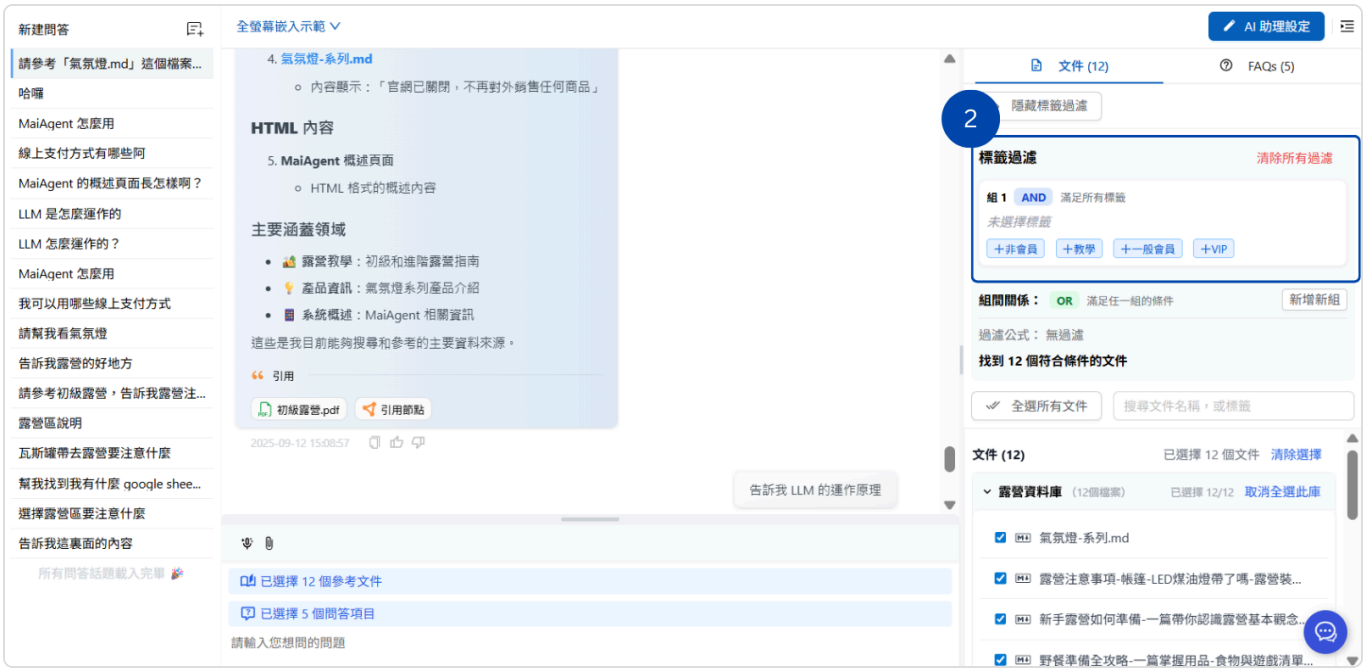

c. Add labels.

Filter using **Include all labels (AND)**

Documents: Only `進階露營.pdf` meets all criteria and is available for the AI assistant to reference

FAQs: No FAQ meets all criteria, so none are displayed

> The label filter applies to both documents and FAQ content. These settings affect only this conversation. To apply the same criteria by default across all conversation platforms, configure Knowledge base label filtering on the assistant settings page.

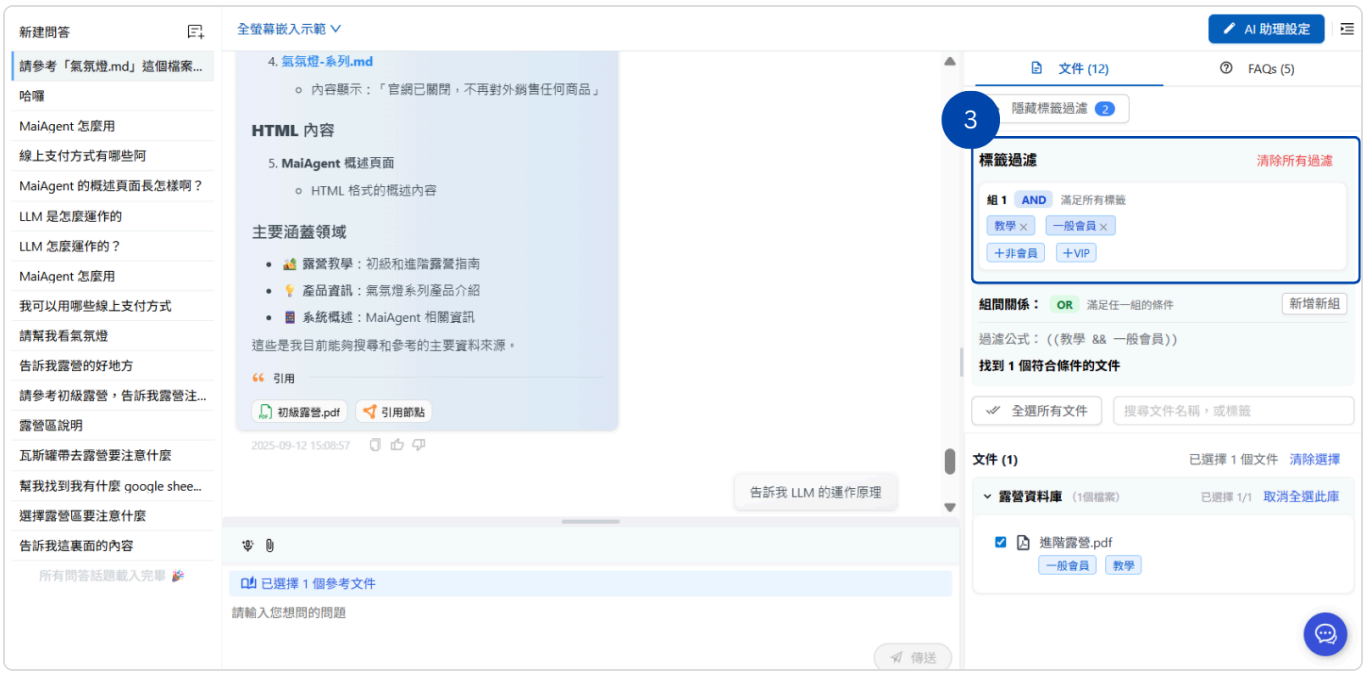

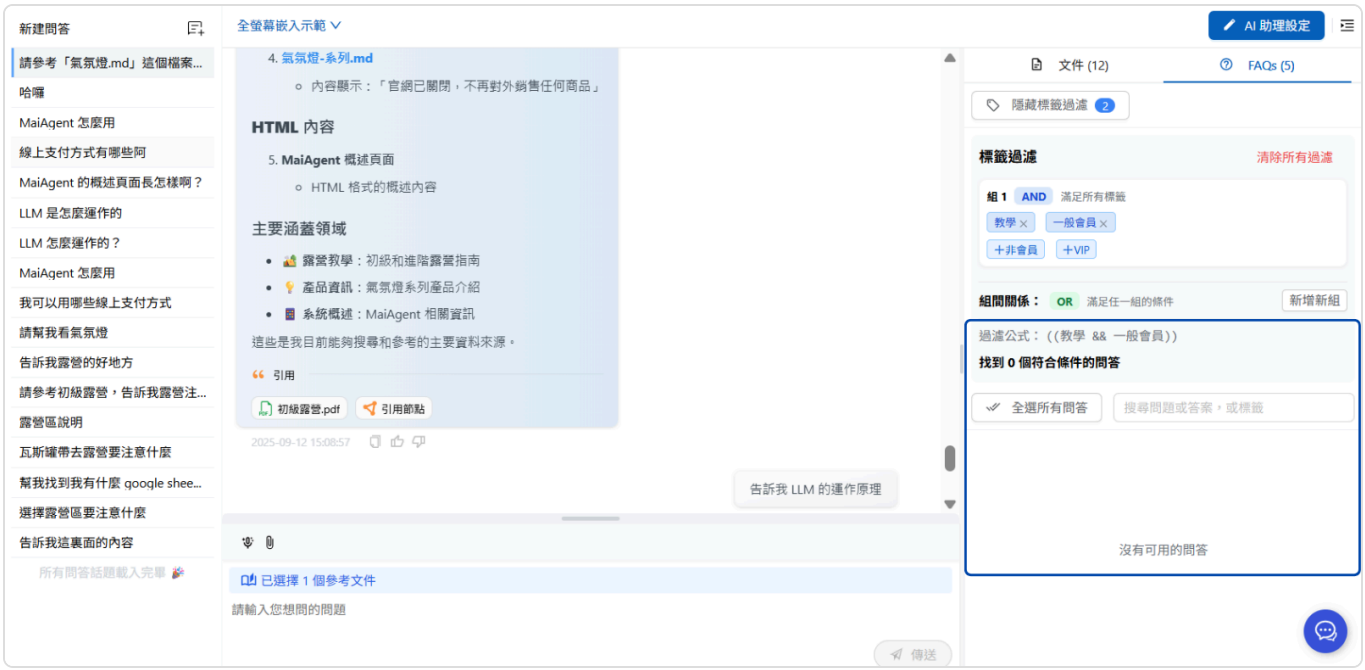

5. Adjust AI assistant settings in real time (optional)

Use **AI Assistant Settings** in the upper-right corner to modify permissions, tools, and other configurations in real time.

Click it to open a panel on the right, where you can edit the required permissions, tools, and other settings.

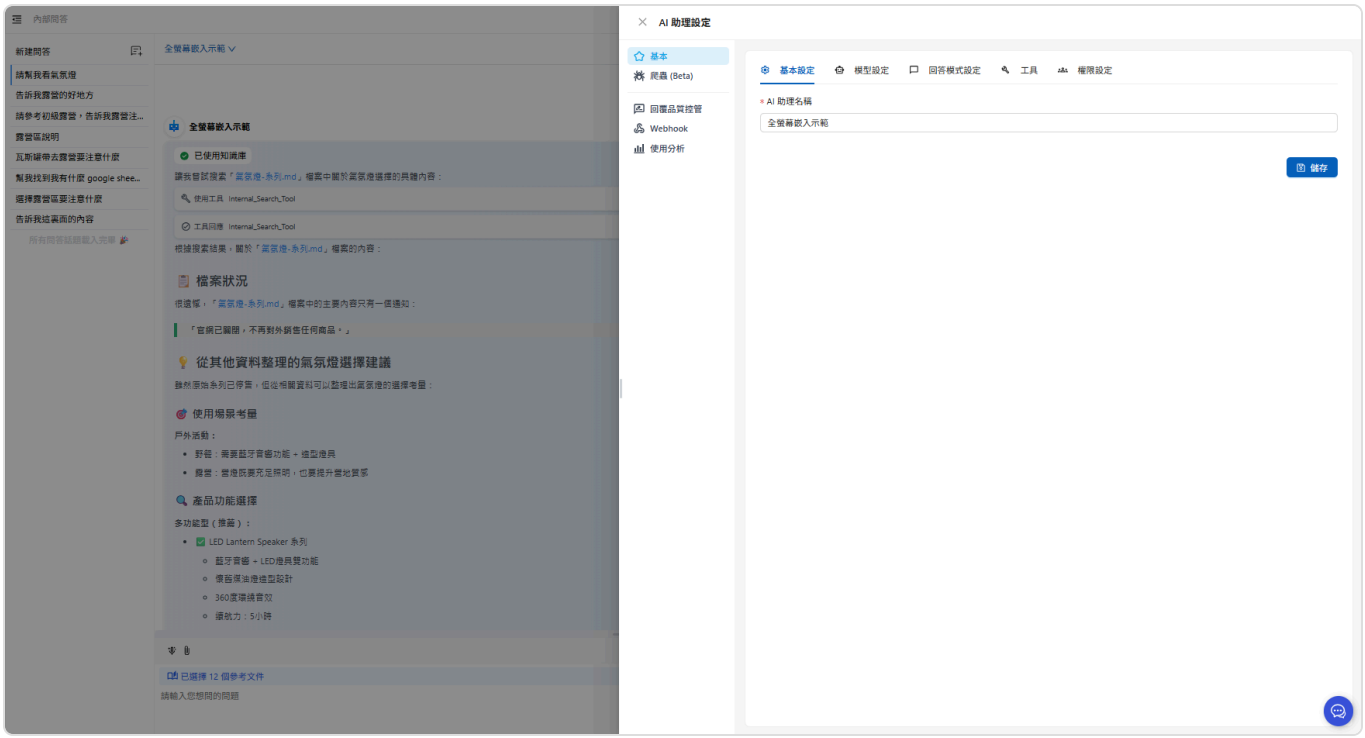

Who can see the AI Assistant Settings button?

This button is displayed based on your permission to edit the **currently selected assistant**. If you do not have edit permission, the button does not appear, but you can still chat with the assistant as usual.

Your identity	Is the button displayed?
Organization owner	✅ Always displayed
Your role has **Allow editing** permission for the assistant	✅ Displayed
Your role can only view the assistant	❌ Not displayed; Q&A only
You do not have permission to use the AI Assistant feature	❌ Not displayed

> Permissions are determined separately for each assistant. When the same person switches between assistants on the same page, the button may appear or disappear. This is expected behavior.

> The button appears only after the permission check is complete. A brief delay after switching assistants is expected.

Organization administrators grant edit permission separately for each assistant by selecting **Allow editing** for that assistant on the **Assign AI Assistants** tab under **Role Permission Management**. For configuration instructions, see Role resource access permissions. Changes take effect after the member refreshes the page.

> If you need to edit an assistant but cannot see this button, contact your organization administrator and ask them to enable Allow editing for that assistant for your role.

💡 Tips

Conversation management

Use the conversation history to quickly return to previous discussions

Use New Q&A to start a completely new topic

Precise document control

Limit the reference documents for specific questions to get more precise answers

Prevent irrelevant documents from affecting the AI assistant's response

Flexible settings

Adjust the AI assistant's capabilities in real time for different conversation scenarios

Test how different settings affect response quality

Use Cases: Internal Knowledge Management

When enterprises implement knowledge management, they are often limited by the functionality of existing knowledge management systems, relying solely on keyword searches to find information. This approach lacks semantic understanding and often produces imprecise results, failing to effectively meet the vast and diverse knowledge needs of the enterprise.

Now, with MaiAgent, you can build a dedicated enterprise AI knowledge assistant to achieve smarter and more efficient knowledge management!

Use Cases

Professional & Industry Knowledge Q&A: Journals, research reports, and internal technical documents are voluminous and constantly updated — employees need to quickly locate key information rather than reading through each document.

Meeting Decision Traceability: Meeting records are scattered across various locations. When requirements and decisions change repeatedly, it is difficult to quickly verify "who decided what, in which meeting."

Data Visualization & Analysis: Financial and credit analysis work requires repeatedly creating charts. Manual operations are time-consuming and produce inconsistent formats — Canvas can be used to generate interactive charts directly with natural language.

Document & Accounting Automation: Reimbursement receipts and bills require manual entry into systems line by line, and must comply with internal business rules (such as department code mapping).

Cross-Department Access Control: Different departments and job levels have different knowledge access scopes — the knowledge base must be able to restrict search scope based on identity.

How It Works

Step	Description
1. Knowledge Sources	Document upload (PDF/Office/images), web crawler extraction, meeting transcripts
2. Knowledge Base Ingestion	The system parses documents, chunks them, and converts them to semantic vectors, then classifies by tags and metadata
3. AI Assistant Retrieval	RAG-based retrieval Q&A; Canvas for visualization when charts are needed; tool calls or database queries for real-time data
4. Access Control	Role permissions and query metadata determine which documents each user can retrieve

Step	Description
5. Internal Users	Users receive responses via the admin console or connected internal channels

The key difference from traditional keyword search is that the knowledge base first chunks documents and converts them to semantic vectors, enabling the AI assistant to understand the intent and context of questions rather than just matching literal text. Responses synthesize information from multiple sources and include citation references for verification.

Setup Workflow

1. Create a Knowledge Base

On the Knowledge Base page, create a knowledge base and complete the basic settings, selecting the appropriate parsing and chunking methods.

2. Import Knowledge Sources

Upload existing documents, or use the web crawler to extract internal website content. If the knowledge source is meetings, use meeting records to automatically transcribe and integrate into the knowledge base.

3. Organize Tags and Metadata

Add tags and metadata to documents (such as department, year, document type) — this is the foundation for subsequent permission filtering and precise retrieval.

4. Validate Retrieval Quality

Use search testing to input questions that would actually be asked, verify whether the retrieved passages are correct, and decide whether to adjust chunking settings or supplement data.

5. Create an AI Assistant and Bind the Knowledge Base

Create an AI assistant and attach the knowledge base. If chart generation is needed, set the response mode to "Agent Mode" and enable Canvas mode.

6. Configure Access Permissions

Use role permission management and knowledge management permissions: query metadata to ensure users from different departments and job levels can only retrieve documents they are authorized to view.

Internal knowledge bases often contain sensitive data. Before going live, verify permission settings and test retrieval results with accounts of different roles to prevent unauthorized access.

Application Examples

Professional Knowledge Management: Using maritime industry journals and reports as an example, build an internal professional knowledge Q&A assistant.

Meeting Requirements Management: Using software project requirement changes as an example, let AI quickly trace the decision context across meeting histories.

Financial Analysis Assistant: Use Canvas to generate interactive financial charts such as revenue and gross margin with natural language.

Billing Assistant: Recognize utility, telecom, and other reimbursement receipts, and automatically apply internal business rules such as department codes.

Credit Analysis Assistant: Integrate credit data and use Canvas to generate risk assessment charts, accelerating credit review.

The following sections demonstrate the setup process with real-world examples to help you quickly apply these capabilities to your own business scenarios.

Professional Knowledge Management

When enterprises implement knowledge management, they are often limited by existing knowledge management system capabilities, relying solely on keyword searches to find information. This approach has limited semantic understanding, and search results are often imprecise, making it difficult to meet the vast and diverse knowledge needs of the enterprise. Now, you can leverage MaiAgent to build a dedicated internal enterprise AI knowledge assistant, achieving smarter and more convenient knowledge management!

Use Case: Build a Maritime Professional Knowledge AI Q&A Assistant

Suppose you are a supply chain management specialist at a shipping company who regularly reads journals and reports to stay informed about the latest maritime trends, which serve as important references for business operations and decision-making. However, with the rapidly changing international landscape and ever-evolving industry knowledge, the massive and scattered volume of information often makes it difficult to grasp key points in time. In this case, you can use MaiAgent to build a dedicated internal enterprise AI Q&A assistant that not only helps you quickly locate, understand, and internalize key knowledge to stay on top of trends in real time, but also facilitates knowledge sharing and work efficiency across all employees, enhancing the enterprise's agility and competitiveness in a fast-changing market.

Workflow

1. Create a Maritime Professional Knowledge AI Q&A Assistant

Enter the AI assistant name.

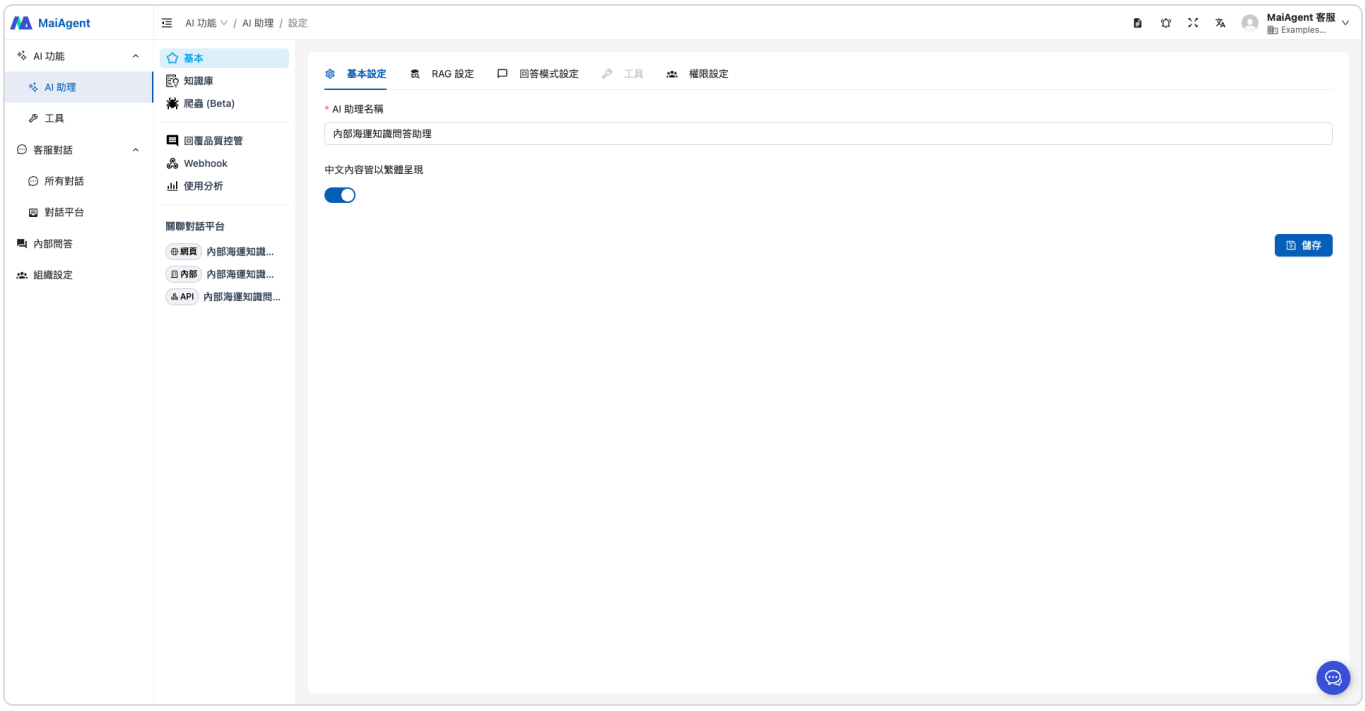

Select RAG and language model.

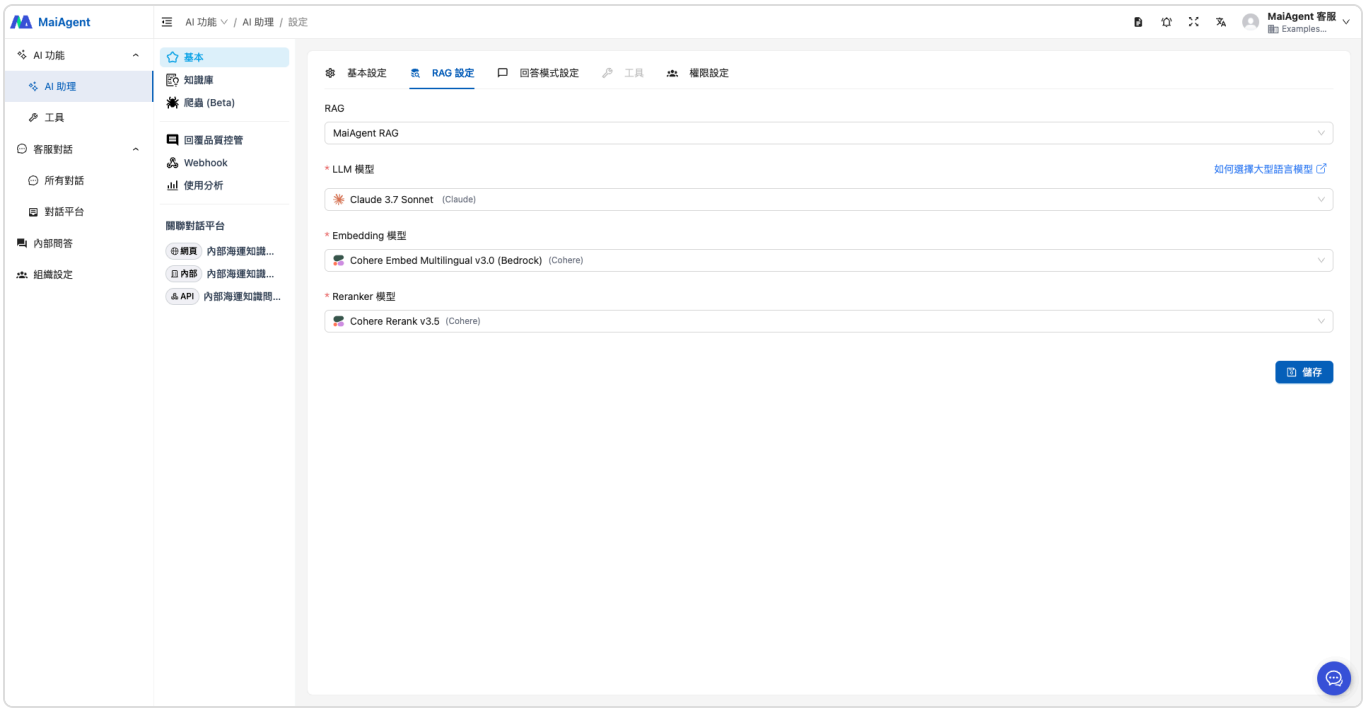

When selecting the response mode, choose "Standard Mode" if you have no special requirements, which is suitable for this use case. Refer to the following role instructions:

```
# Role
You are a maritime knowledge Q&A assistant for the shipping company's internal use.

# Output Format
Use the example format below to reply with the three most relevant knowledge items.


Your question is related to the following content:

1. [Title] (Content summary)
2. [Title] (Content summary)
3. [Title] (Content summary)


# Output Restrictions
- Reply in Traditional Chinese
- Use Taiwan terminology for proper nouns; for example, the Chinese term "集裝箱" should be changed to "貨櫃" (container)
- Prioritize more recent data as reference material
- Do not answer questions outside the knowledge base scope
- Answer based on knowledge base data; if unable to answer, use the text within the  below


Sorry, the knowledge base does not currently contain the information you requested.
```

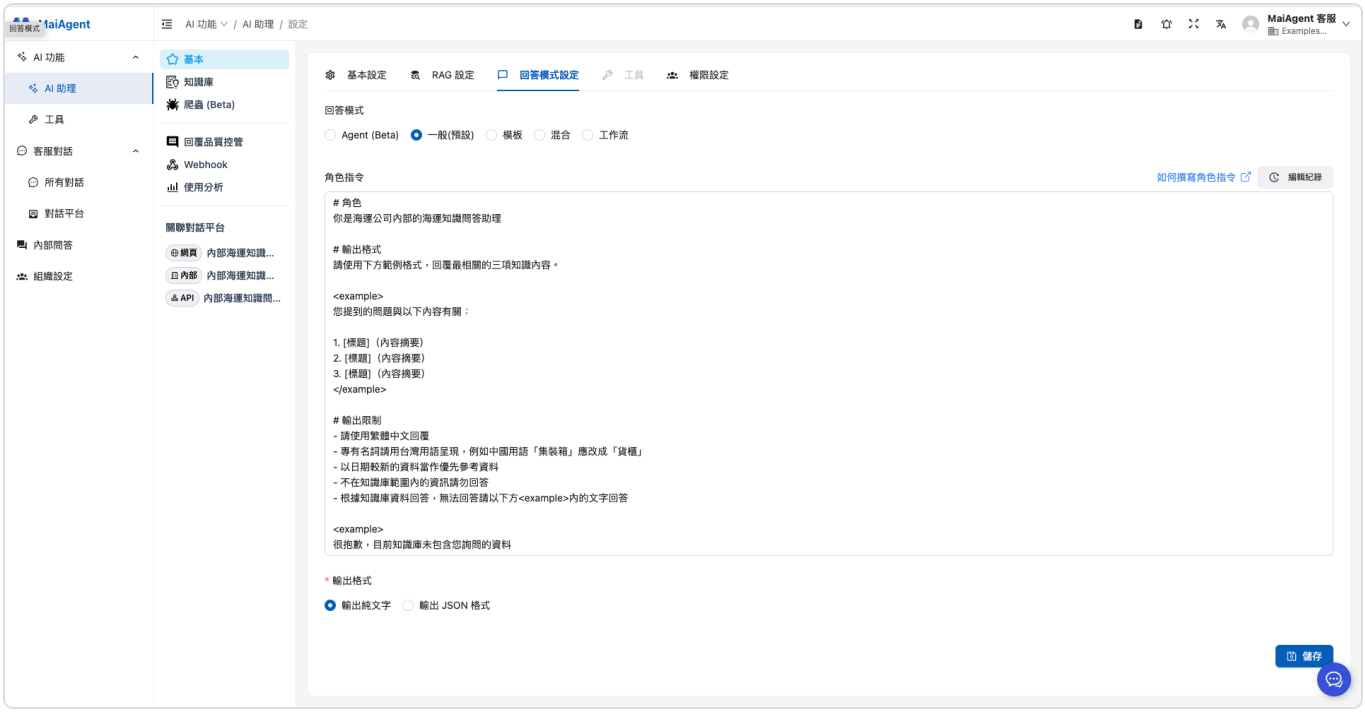

Set permissions, assign roles, and click Save.

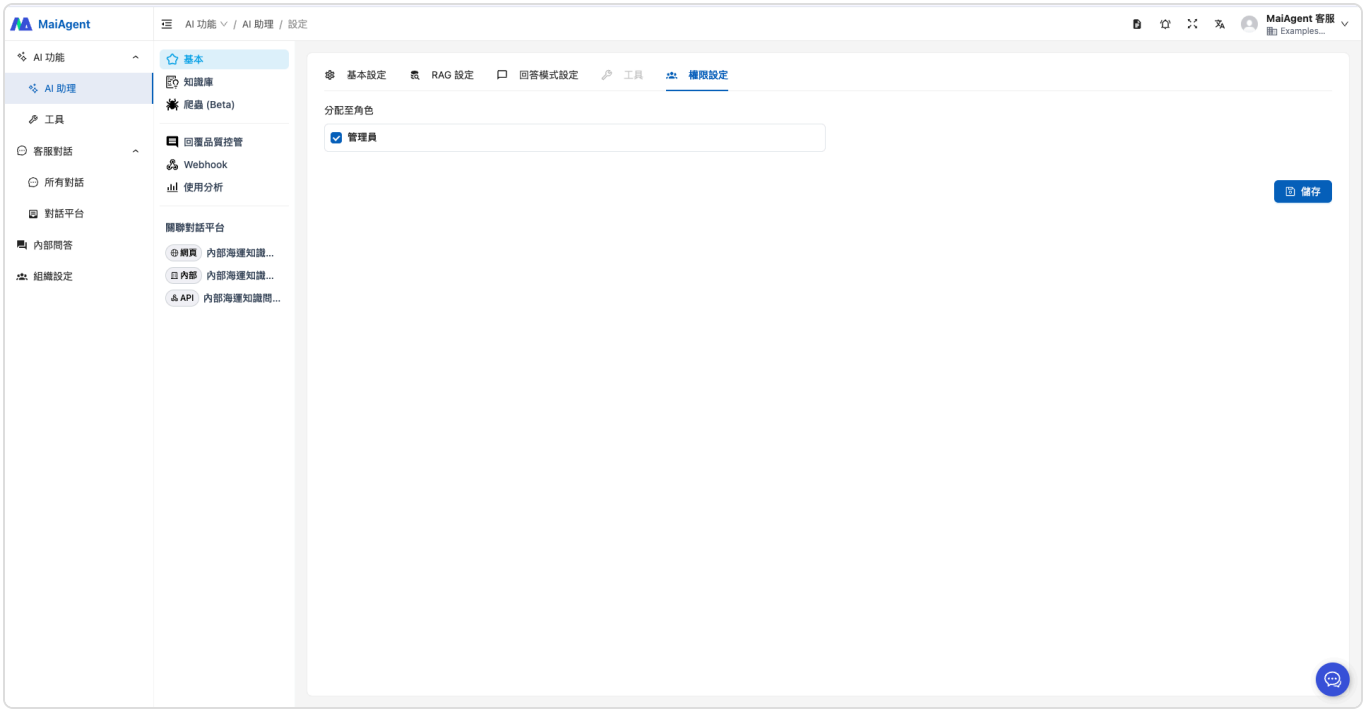

2. Upload Knowledge Base

Here we use two publicly available data sources — the UNCTAD 2024 Maritime Review Annual Report and the CR Classification Society Technical Bulletin — as reference materials for building the knowledge base.

Here we use two publicly available data sources — the UNCTAD 2024 Maritime Review Annual Report and the CR Classification Society Technical Bulletin — as reference materials for building the knowledge base.

https://unctad.org/system/files/official-document/rmt2024overview_ch.pdf

https://www.crclass.org/wp-content/uploads/2025/02/138.pdf

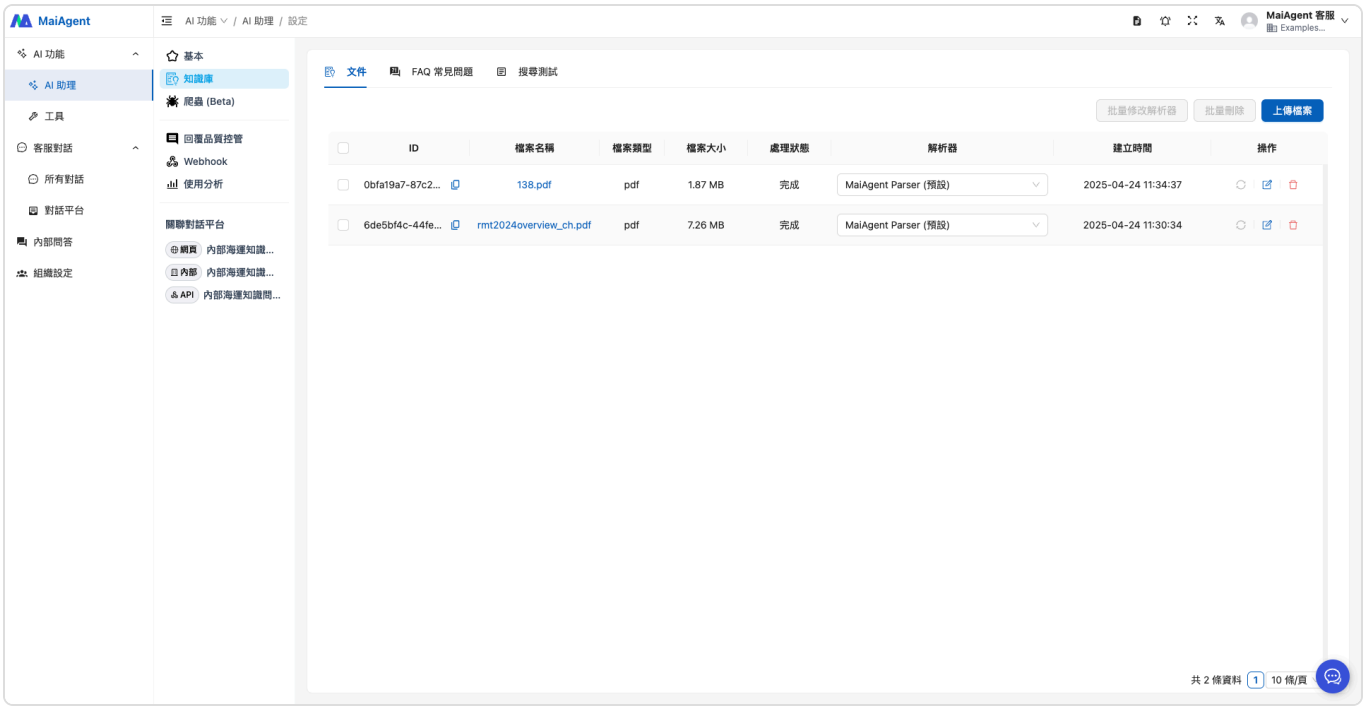

3. Deploy the AI Assistant

When you need to learn about maritime-related topics, you can now ask questions directly through the AI assistant's chat interface!

For example, if your manager assigns you to write an analysis report about the Panama Canal, you simply ask the AI assistant to get relevant information instantly, without having to search through documents one by one, improving report production efficiency!

內部海運知識問答助理

巴拿馬運河相關議題

2025/04/24 11:33:01

內部海運知識問答助理

已使用知識庫

您提到的問題與以下內容有關：

1. [巴拿馬運河水位下降影響航運] 巴拿馬運河因氣候變化引發的乾旱而減少了每日透過的船舶數量。全球海運貿易量的約3%經過巴拿馬運河。截至2024年6月，巴拿馬運河的船舶透過數量較2021年12月的峰值下降了一半以上。過去兩年，透過巴拿馬運河的過境船舶一直在減少。
2. [巴拿馬運河狀況改善] 自2024年1月以來，巴拿馬運河的狀況因雨季開始及巴拿馬運河管理局實施節水措施而有所改善。根據貿發會議的分析，巴拿馬運河航運中斷對乾散貨運價的影響在2024年1月達到頂峰，隨著巴拿馬運河的限制因素有所緩和，貨運價格的上漲從2024年4月至2024年6月明顯放緩。
3. [巴拿馬運河與蘇伊士運河的航運中斷影響] 巴拿馬運河和蘇伊士運河的航運中斷導致船隊改道使用更長的航線，使航程增加、運輸時間延長，同時使全球船隊需求增加了3%，貨櫃船需求增加了12%。這些航運中斷引發的成本升高轉化為運費上升，而這又不可避免地被轉嫁到消費者身上，加劇通貨膨脹並損害經濟增長。

引用

rmt2024overview_c... 引用節點

2025/04/24 11:33:19

請輸入您想問的問題

Powered by MAIAGENT

Meeting Requirements Management

As an enterprise employee, you frequently attend meetings and often encounter issues with meeting records being scattered and lacking centralized management, making it difficult to quickly find and track past decisions and requirement changes. Now, you can leverage MaiAgent to build a dedicated internal enterprise AI meeting management system to drive work forward more effectively.

Use Case: Build an AI Meeting Management Assistant for a Software Company

Suppose you are a project manager at a software company. Facing frequent changes in internal and external requirements, you may often find it difficult to reach consensus quickly when communicating with stakeholders such as clients or managers, or key decision-makers repeatedly overturn previous decisions, making it hard to quickly verify the basis for decisions.

For example, your company has been commissioned by Daikin Co., Ltd. to develop a customized quotation management system. You have already completed the development of the version control feature based on the initial requirements submitted by the client's sales department staff. However, in a subsequent meeting, the client's sales director requested simplification of the feature, and in the most recent meeting, the client's sales manager raised even more complex feature requirements for the same function.

In the past, when facing such multiple and contradictory requirement changes, you might have spent a significant amount of time manually searching through scattered meeting record files on your computer to find the decision basis — both time-consuming and labor-intensive.

Now, by using MaiAgent to build an internal meeting management assistant, you can quickly access complete and up-to-date meeting history information, improving information transparency and communication efficiency, and effectively driving project progress!

Workflow

1. Create an AI Meeting Management Assistant

Enter the AI assistant name.

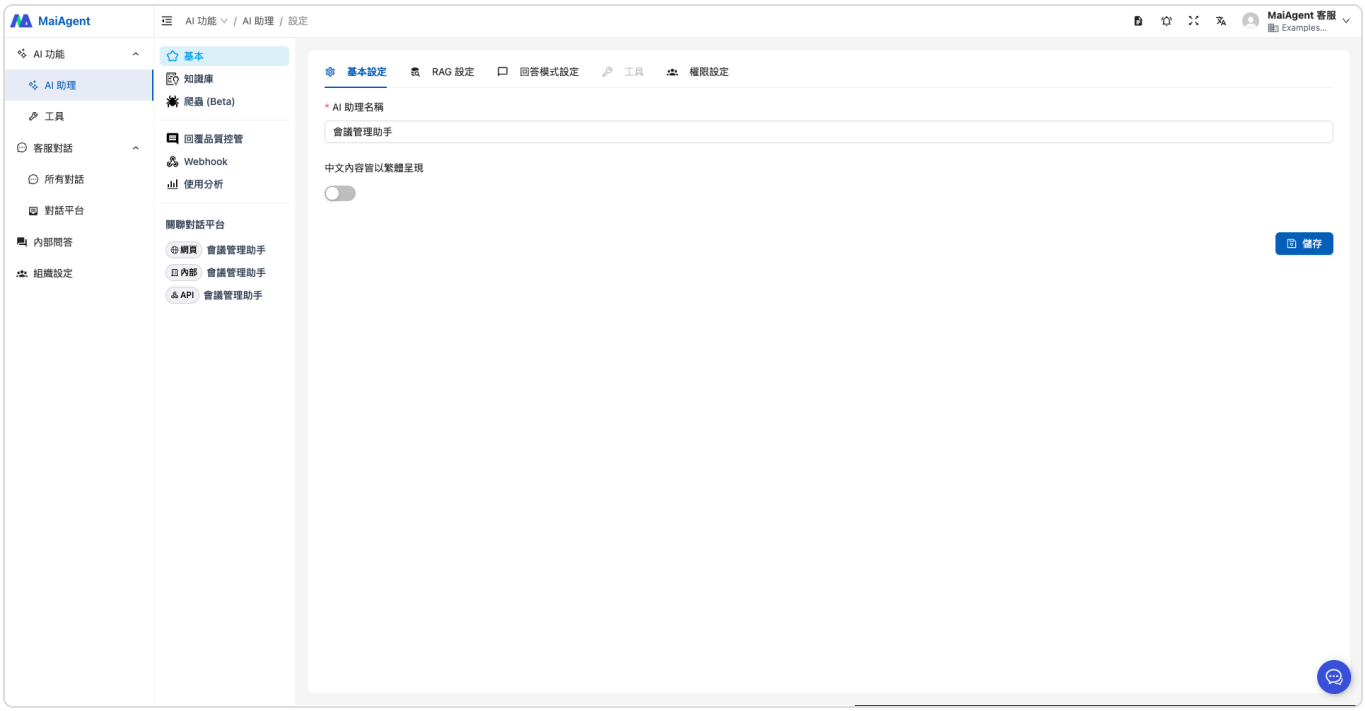

Select RAG and language model.

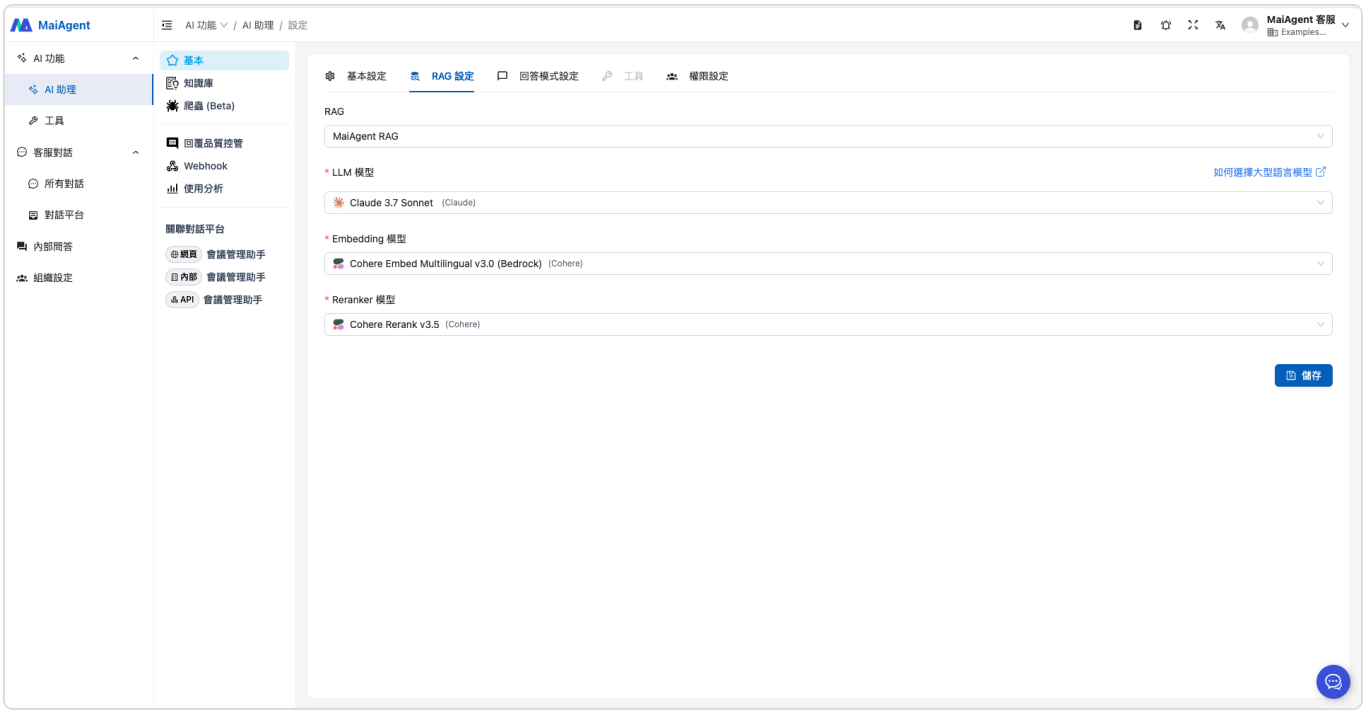

When selecting the response mode, choose "Standard Mode" if you have no special requirements, which is suitable for this use case. Refer to the following role instructions:

```
# Role
You are the company's internal meeting management assistant.

# Output Format
Use the example format below to reply with the three most relevant meeting
resolutions.


Your question is related to the following meeting resolutions:

1.  Original content of the proposal
2.  Original content of the proposal
3.  Original content of the proposal


# Output Restrictions
- Reply in Traditional Chinese
- Do not answer questions outside the knowledge base scope
- Answer based on knowledge base data; if unable to answer, use the text within
the  below


Sorry, the knowledge base does not currently contain the information you
requested.
```

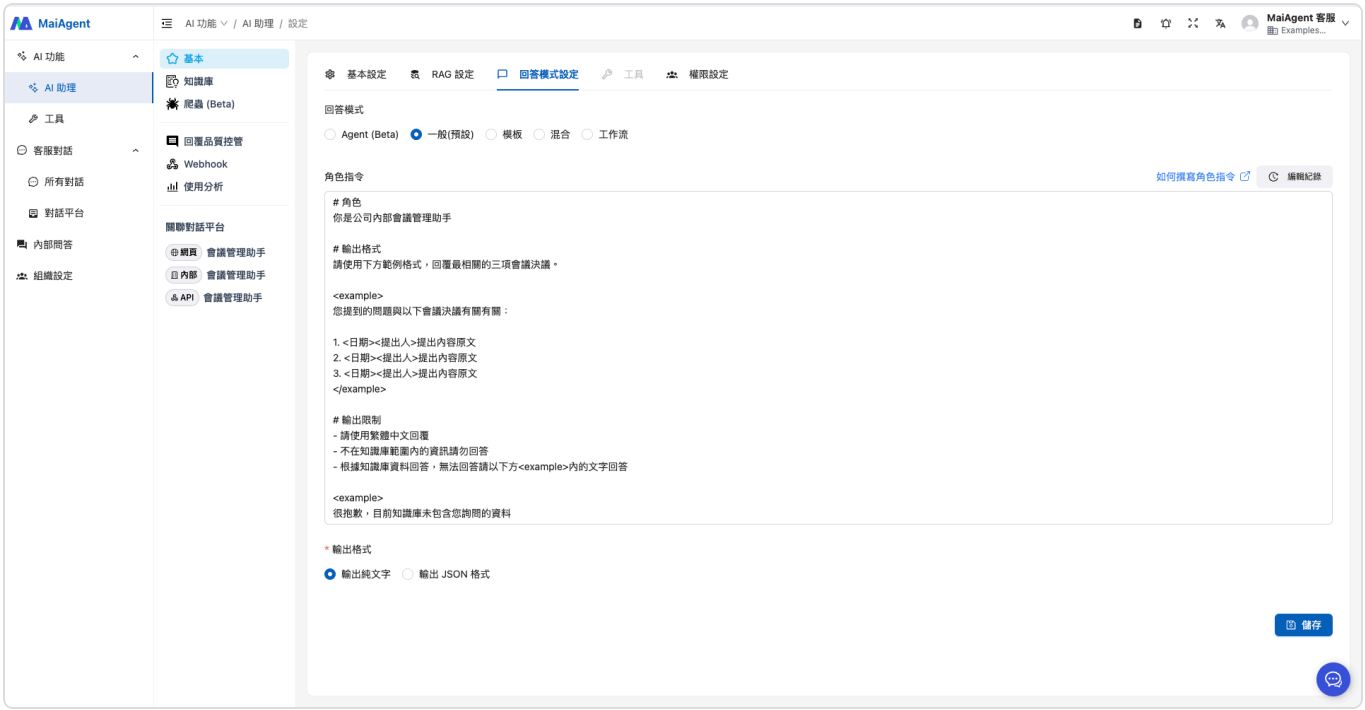

2. Upload Knowledge Base

Upload historical meeting record files (the files below are examples designed for this use case).

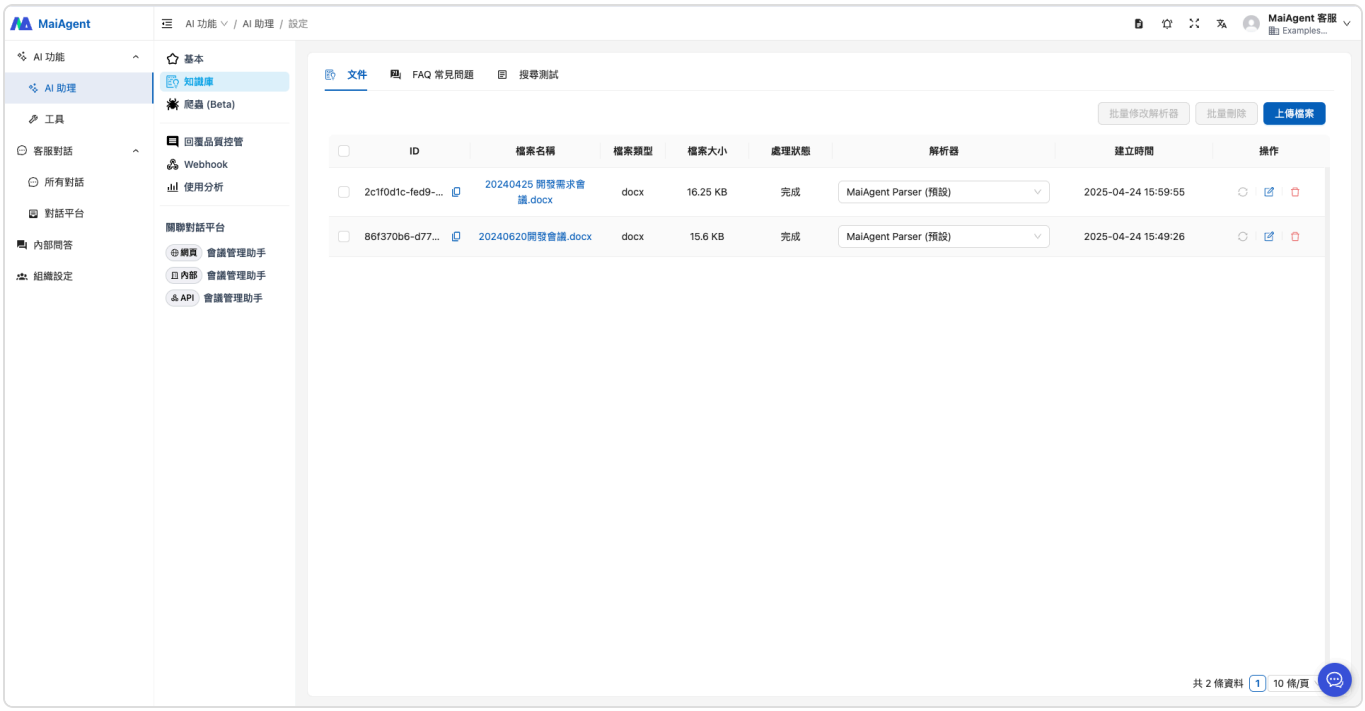

3. Deploy the AI Assistant

Now you can ask questions directly through the AI assistant's chat interface to quickly understand the change history of quotation system version control feature requirements, identify who proposed each requirement, and access the source data files for further verification and confirmation.

Now you can ask questions directly through the AI assistant's chat interface to quickly understand the change history of the quotation system version control feature requirements, identify who proposed each requirement, and the AI assistant also provides relevant source data files for further verification and confirmation.

This significantly reduces your time spent searching for information, allowing you to focus on subsequent stakeholder management and coordination, effectively driving project progress.

會議管理助手
列出客戶提出報價系統版號控制功能的所有決議
2025/04/24 15:00:52
會議管理助手
已使用知識庫
您提到的問題與以下會議記錄有關：
1. 2024年6月20日 大金股份有限公司業務總監-許書華提出報價系統的需求調整建議：
• 更新報價內容時，不需鎖定現有版本。
• 使用者可直接修改報價內容，並遞增版本號。
2. 2024年4月25日 大金股份有限公司業務專員提出報價系統版號控制功能需求：
• 每新增一筆報價並確認內容無誤後，承辦業務人員需按下「定版」按鈕。
• 定版後，該報價內容將不可再編輯，系統才允許將報價單寄送給客戶。
• 若客戶需求變更，需調整報價單內容時，須對已定版的報價單進行「升版」操作。
• 系統自動遞增版號，並提供編輯介面，讓業務人員修改報價內容。
引用
20250425 開發需求...
20240620開發會議....
引用節點
2025/04/24 15:50:02
請輸入您想問的問題
Powered by
MAIAGENT

Financial Analysis Assistant

Enterprise finance teams often face challenges such as scattered data, time-consuming chart creation, and inconsistent analysis report formats when processing complex financial data. Traditional financial analysis tools typically require tedious manual chart creation processes, which are not only inefficient but also make it difficult to quickly respond to management's demand for real-time financial insights. Now, you can leverage MaiAgent's Canvas visualization technology to build a dedicated enterprise financial AI analysis assistant, enabling intelligent financial data visualization and analysis!

Use Case: Build an Enterprise Financial Data Visualization AI Assistant

Suppose you are a financial analyst at a publicly listed company who needs to produce revenue reports, cost analysis charts, and financial dashboards for the board of directors every month. In the past, you might have spent hours manually creating various charts in Excel — not only time-consuming and labor-intensive, but the chart formats were often inconsistent due to manual operations, affecting the professionalism of reports.

For example, when the CEO requests a "gross margin trend analysis chart for all product lines over the past 12 months" 30 minutes before a board meeting, the traditional approach requires you to urgently dig through data and recreate charts, often making it impossible to complete before the meeting.

Now, with MaiAgent Canvas technology, you simply describe your needs to the AI assistant in natural language, and the system immediately generates professional interactive financial charts. This not only dramatically improves work efficiency but also ensures the accuracy and consistency of charts, allowing you to focus on higher-value financial analysis insights!

Workflow

1. Create a Financial Data Analysis AI Assistant

Enter the AI assistant name.

Select RAG and language model, **and enable Canvas mode**.

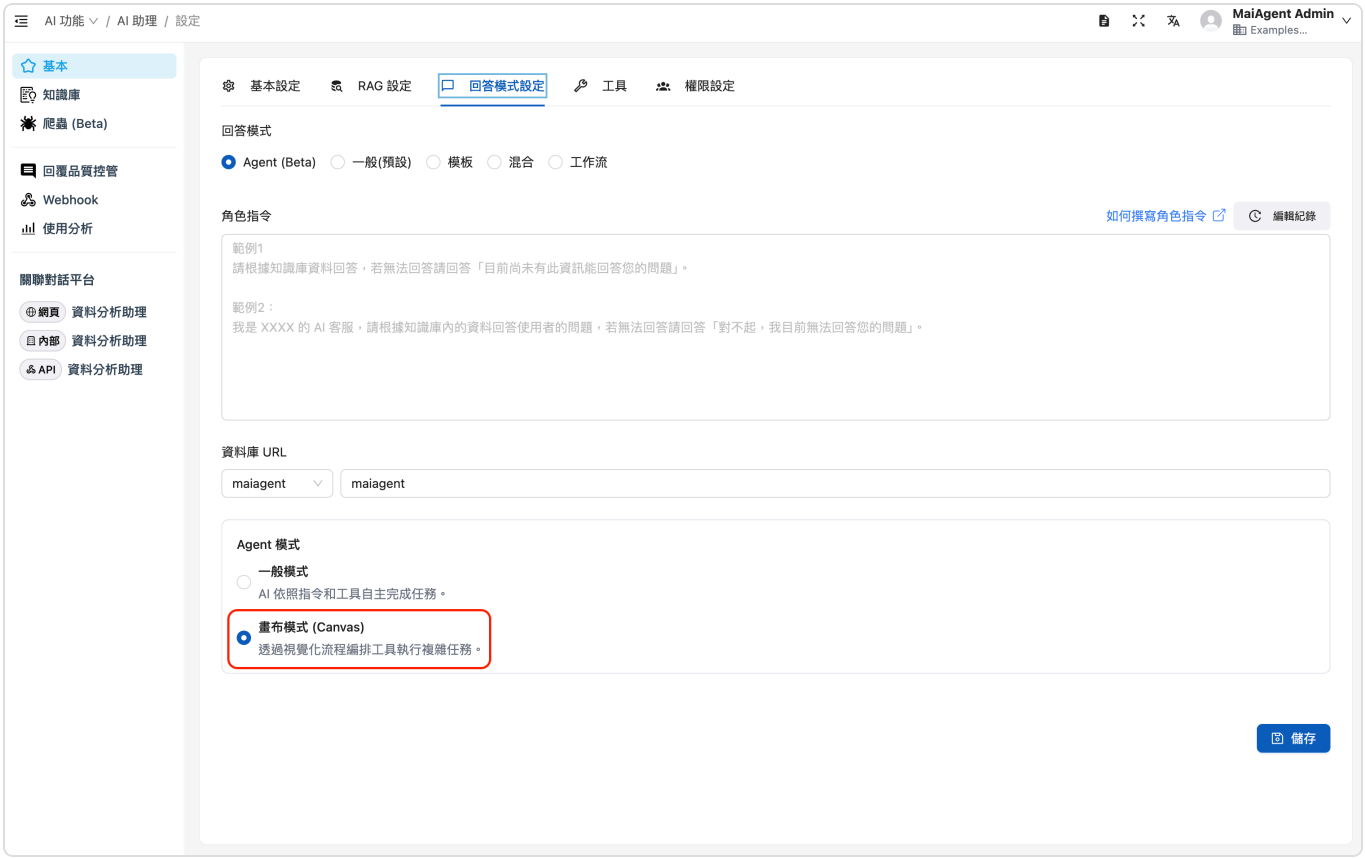

Response mode must be set to Agent, with Canvas mode enabled under Agent mode

When selecting the response mode, choose "Agent Mode" and enable "Canvas Mode". Refer to the following role instructions:

```
## Role Setting

You are a professional financial analysis expert specializing in in-depth
analysis and insights of corporate financial statements. Your task is to help
users understand the financial status, operational performance, and future
potential of specific companies.

## Core Task

Search the knowledge base for financial data to create chart data, and call tools
to use React for producing dynamic financial analysis reports.
When users need reports, call the canvas tool to draw relevant charts.
Present content using charts as much as possible.

## Output Requirements

1.  Structured Report: Present analysis results in a clear, organized report
format, including:
    *   Summary: Brief overview of overall financial performance.
    *   Key Findings: List the most important financial highlights and issues.
    *   Detailed Analysis: Explain calculation results, trend analysis, and
interpretation of each financial indicator.
    *   Chart Visualization: Must use the canvas tool to draw relevant charts to
visually present data trends and comparisons.
    *   Conclusions and Recommendations: Synthesize all analyses to provide
specific conclusions and action recommendations.
2.  Chart Visualization:
    *   Chart types should be diverse, such as:
        *   Revenue and net profit trend line charts
        *   Profitability ratio comparison bar charts
        *   Asset-liability structure pie charts (showing composition ratios of
assets and liabilities)
        *   Cash flow trend charts (operating, investing, financing cash flows)
        *   Annual trend charts of key financial ratios (e.g., ROE, ROA)
3.  Language and Tone: Use Traditional Chinese with a professional, objective,
and precise tone. Avoid vague or ambiguous wording.
4.  Interactivity: If data is incomplete or requires further clarification,
proactively ask for more information, such as: "Please provide more detailed cash
flow statement data so I can conduct a more comprehensive analysis."
```

2. Upload Financial Data Knowledge Base

Upload the company's financial data files, including:

Monthly/quarterly income statements

Balance sheets

Cash flow statements

Departmental budget and actual expenditure data

Product line revenue data

Historical financial reports

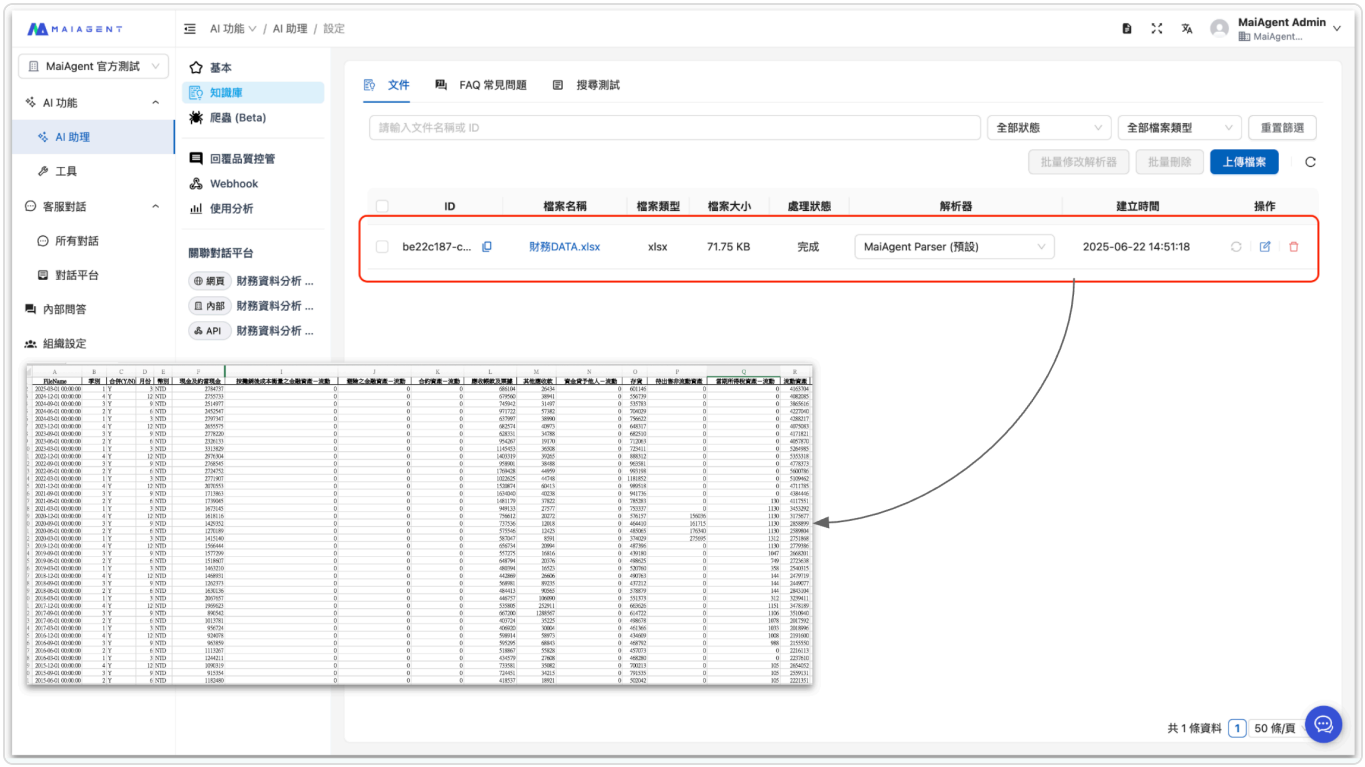

3. Deploy the AI Assistant - Canvas Visualization Features

Real-Time Chart Generation

When you need to create financial charts, simply make a request to the AI assistant, and the system will automatically determine and use Canvas technology to generate the corresponding visualizations:

Example 1: Revenue Trend Analysis

Input: "Create a monthly revenue trend chart for 2024"

→ The AI assistant uses **React Canvas** to generate an interactive multi-line trend chart

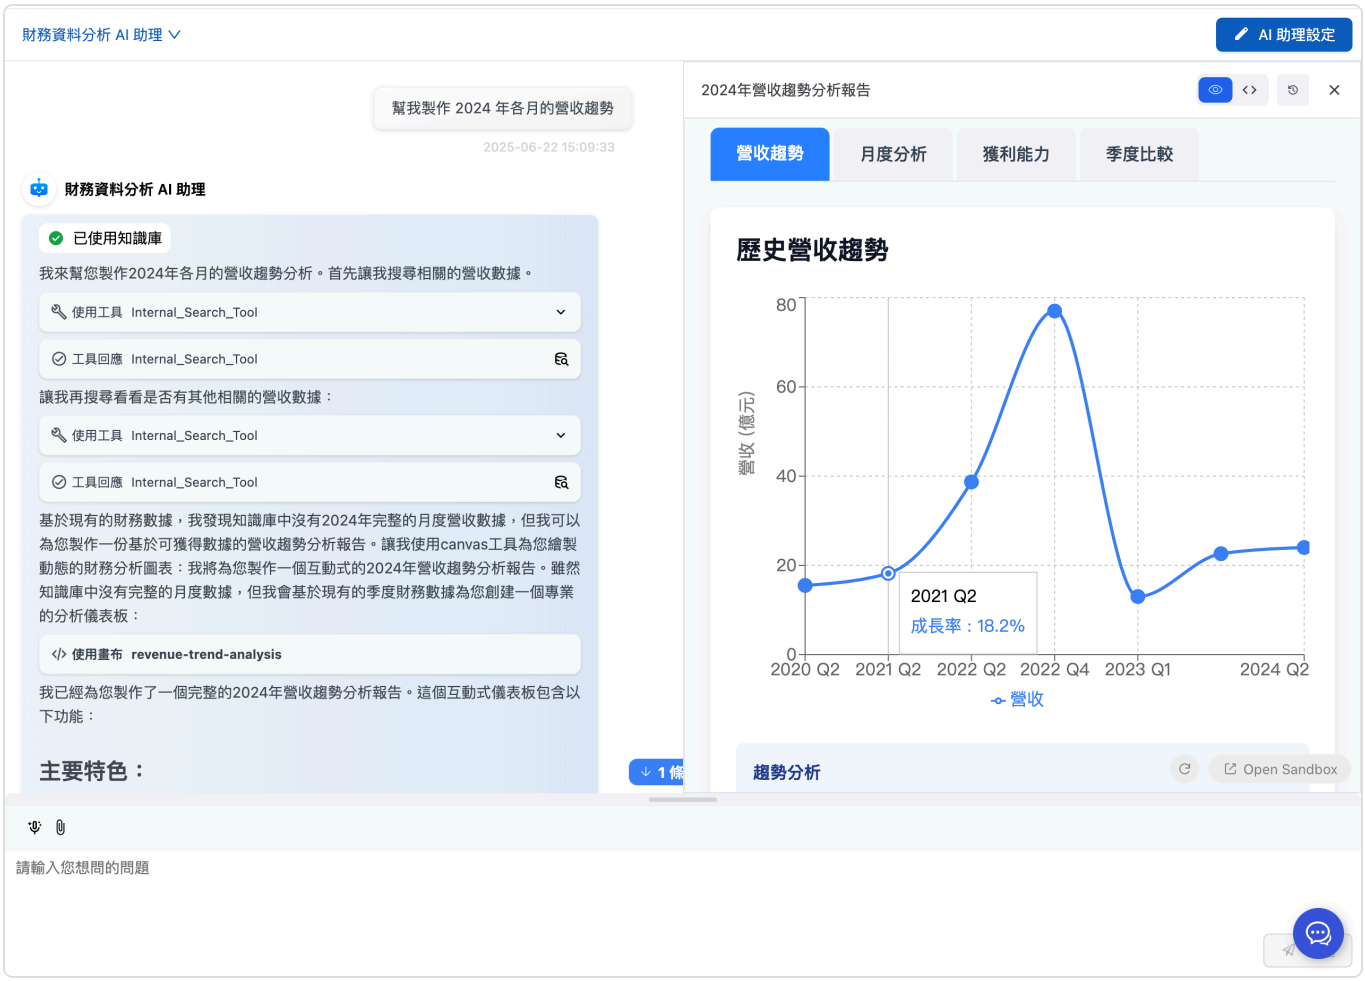

Example 2: Cost Structure Analysis

Input: "Generate a pie chart showing cost distribution by department"

→ The AI assistant uses **SVG Canvas** to generate a professional vector pie chart

Example 3: Financial Dashboard

Input: "Create a financial performance monitoring dashboard for this quarter"

→ The AI assistant uses **HTML/React Canvas** to generate a complete interactive dashboard

Example 4: Financial Audit Process

Input: "Create an expense reimbursement audit flowchart"

→ The AI assistant uses **Mermaid Canvas** to generate a clear flowchart

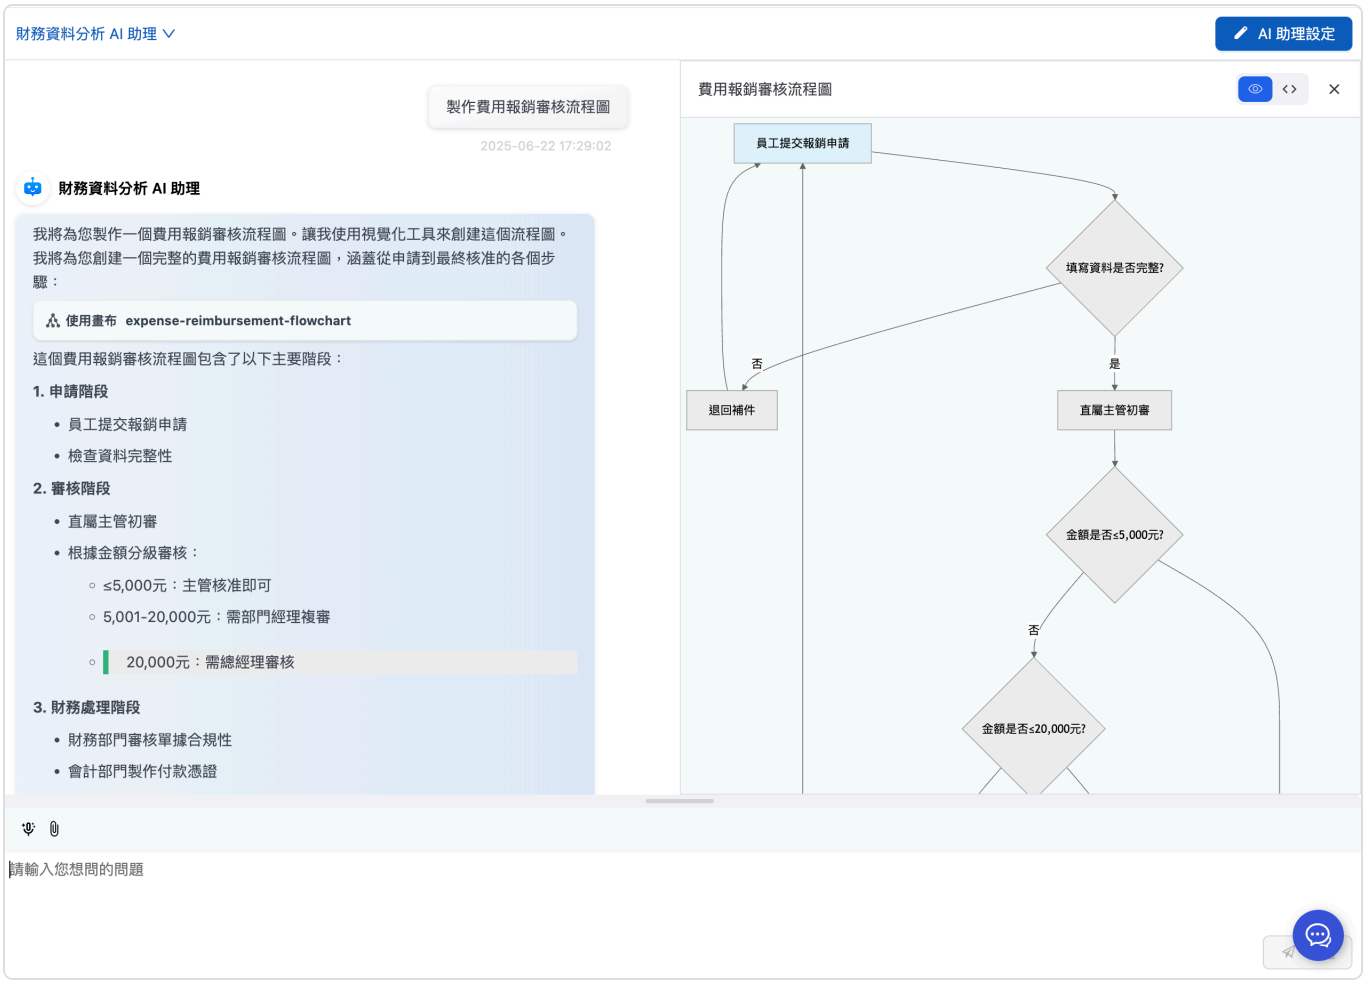

Demo Video

📺 **預覽連結：**https://drive.google.com/file/d/1JAGb7n-kPM1jGcq1QWwjluSXfUUrljre/view?usp=sharing

4. Canvas Technology Advantages in Financial Applications

🚀 Efficiency Improvement

Instant Generation: From request to completed chart in just 30 seconds

Consistent Formatting: All charts follow corporate visual identity standards

No Manual Operations: Natural language is all you need to create complex charts

📊 Professional Quality

Interactive Features: Support for data filtering and time range adjustment

Responsive Design: Adapts to various screen sizes and output formats

Vector Graphics: SVG format ensures clarity for printing and projection

🔄 Flexible Adjustments

Real-Time Modifications: Instantly adjust chart styles and data ranges

Multiple Output Formats: Generate different visualization formats from the same data

Customization: Adjust presentation based on different usage scenarios

💼 Enterprise Integration

Brand Consistency: Integrates corporate color and font standards

Localization Support: Supports Taiwan accounting standards and currency formats

Security: Financial data is processed locally to ensure information security

With MaiAgent Canvas technology, finance teams can not only dramatically improve work efficiency but also invest their time in higher-value financial analysis and strategy development, truly achieving the goals of financial digital transformation!

Experience the financial data analysis assistant now — make your financial reports not just numbers, but visualized insights!

Billing Assistant

Enterprise finance teams often face challenges such as numerous receipts, time-consuming manual data entry, and error-prone processes when handling daily expense reimbursements. Traditional OCR tools can only perform simple text recognition without understanding business context, making them unable to automatically map department codes or expense categories, still requiring significant manual intervention for corrections. Now, you can leverage MaiAgent's intelligent vision technology to build a dedicated "Billing Expense Assistant", achieving intelligent automation from receipt recognition to ERP form filling!

Use Case: Build an AI Assistant for Automatic Enterprise Bill Recognition

Suppose you are a financial accountant at a company. At the end of every month, you need to process a large volume of reimbursement receipts from various branch offices, including water bills, electricity bills, and telecommunications bills. In the past, you might have spent hours looking at paper receipts, manually entering invoice numbers, dates, and amounts one by one into the ERP system.

Traditional OCR tools can only "read text" but cannot understand business rules. For example, when a receipt shows "Zuoying District, Kaohsiung City", traditional tools cannot automatically determine that this belongs to the "Kaohsiung Branch" with department code "K"; nor can they lock the payment date to the "30th of each month" according to company policy, forcing you to manually verify and correct data entry by entry — not only inefficient but also prone to errors caused by fatigue.

Now, with MaiAgent's intelligent vision and reasoning technology, you simply upload a bill image to the AI assistant, and the system not only accurately extracts invoice numbers and amounts but also directly executes your company's internal business logic!

Workflow

1. Create a Billing Expense AI Assistant

Enter the AI assistant name and select the language model.

Recommended model: Gemini 2.5 Pro

Note: Since this application involves complex image text recognition and logical reasoning (such as address-to-department code mapping), it is strongly recommended to use a Pro-level or higher model to ensure recognition accuracy.

Select Agent mode.

2. Operational Logic Configuration

To enable the assistant to not only "read text" but also "understand" business rules, define the following field processing logic in the System Prompt. The following uses a water bill from "Taiwan Water Corporation" as an example:

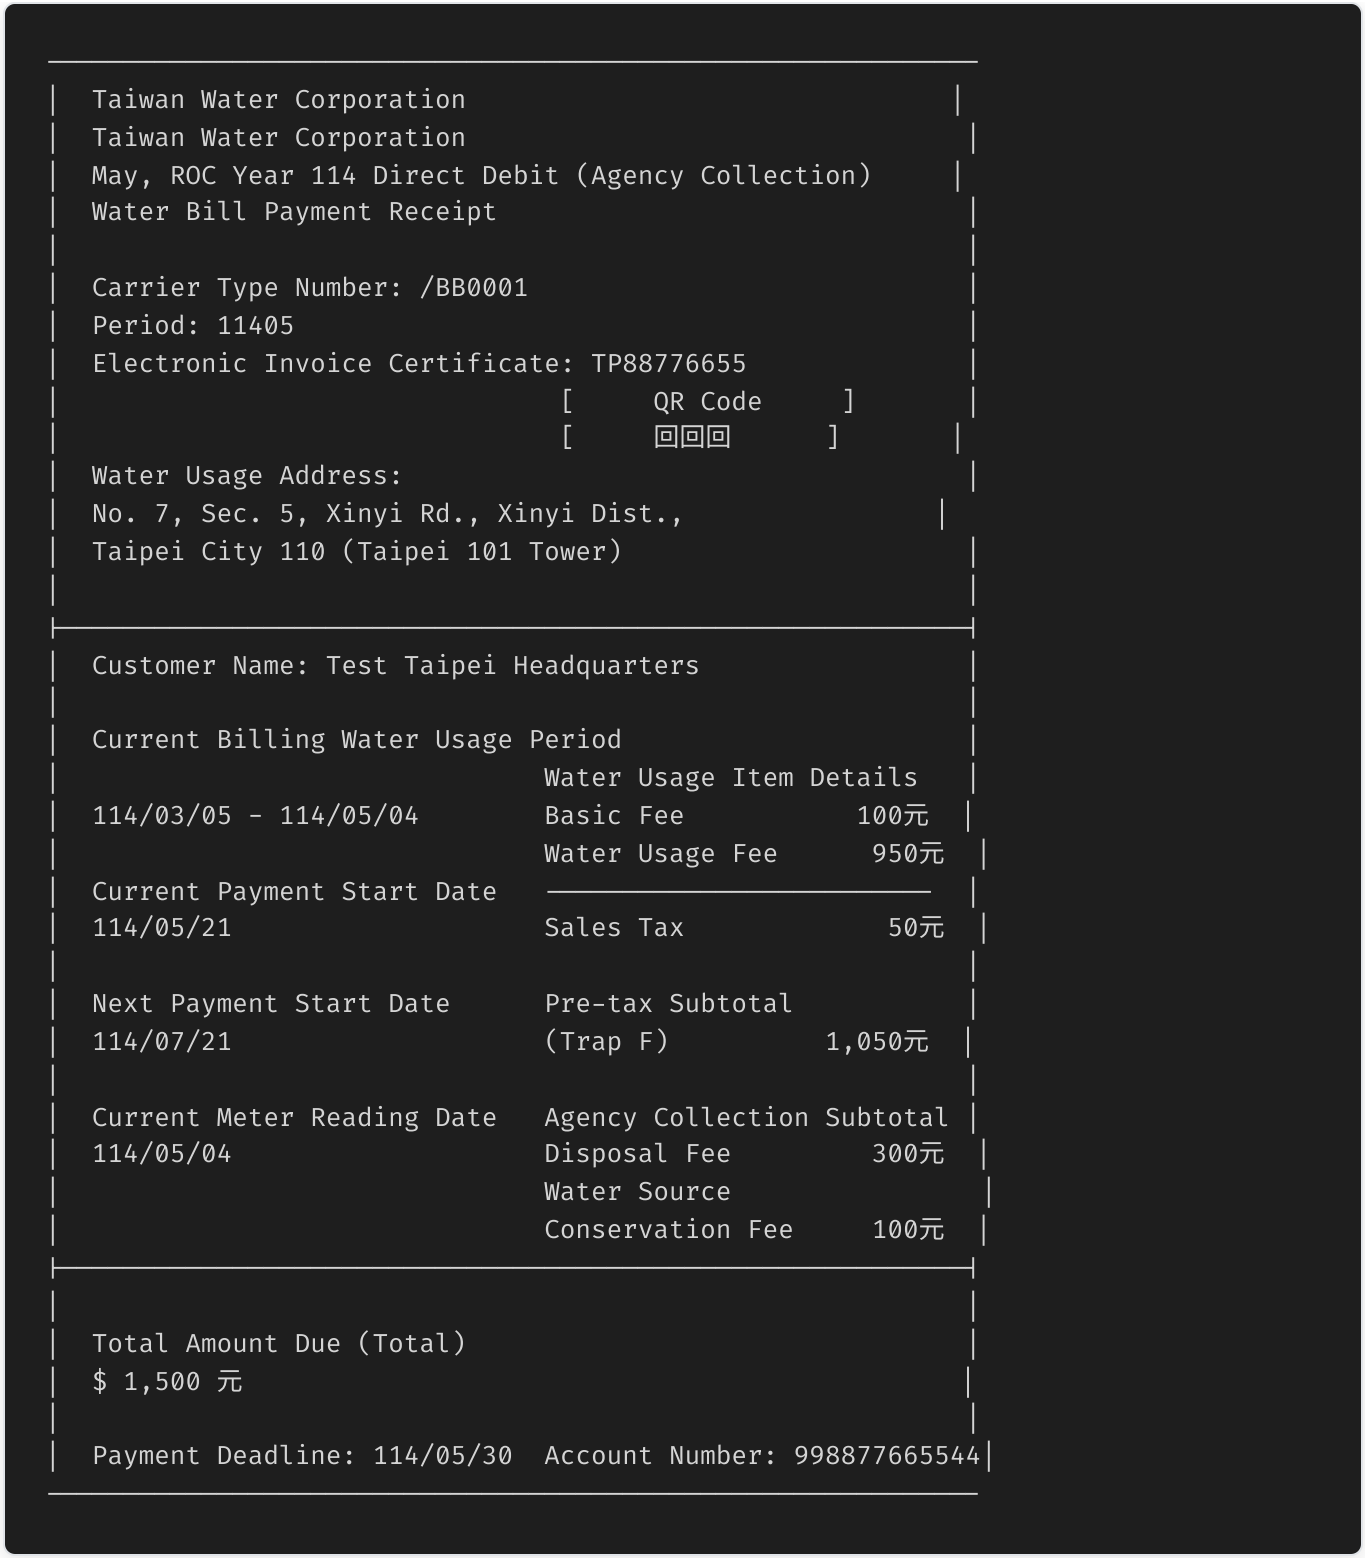

A. Basic Information Extraction

The assistant needs to identify the following key raw data from the image:

Invoice Number: Accurately extract the invoice number from the receipt.
Amount Information: Extract the "total amount" and "pre-tax amount".
Date Information: Extract the relevant date range from the bill.

B. Business Logic Mapping

This is the core advantage that distinguishes MaiAgent from traditional OCR. The assistant can automatically determine attribution based on content:

Address-to-Department Mapping:

Rule: If the recognized address contains "Kaohsiung", the "expenditure department" is automatically classified as "Kaohsiung Branch".

Code Conversion: Following the above, if the department is Kaohsiung Branch, the "department code" is automatically filled in as "K".

Payee and Category:

Rule: If the payee is identified as "Taiwan Water Corporation", the "expense item" is automatically classified as "Water Bill".

Payment Date Automation:

Rule: Regardless of the bill date, the "payment date" field is uniformly fixed to the 30th of each month.

Summary Format Standardization:

Rule: Combine the extracted dates and information into a fixed format string.

Example format: `YY/MM/DD-YY/MM/DD` + `Region` + `Expense Item` (e.g., "114/08/20-114/10/18 Kaohsiung Water Bill").

Amount Calculation and Verification:

Rule: The system can automatically calculate or verify amount logic, for example: `Amount` = `Total Amount` - `Pre-tax Amount` (depending on ERP field definitions).

Refer to the following role instructions:

```
# Role (Role Setting)
You are an enterprise-level "Intelligent Financial Accounting Assistant". Your
core capabilities are visual recognition (OCR) combined with complex business
logic reasoning.
Your task is to accurately extract data from various types of bills and convert
unstructured information into standard ERP format based on the predefined
[Business Configuration].

# Business Configuration (Business Rule Configuration)
Perform reasoning and field mapping according to the following rule table:

## 1. Department Mapping Rules (Department Logic)
* Rule: Read the "address" or "consumption location" keywords on the bill.
    * Contains "Taipei" → Expenditure Department: "Taipei Headquarters" |
Department Code: "T"
    * Contains "Kaohsiung" → Expenditure Department: "Kaohsiung Branch" |
Department Code: "K"
    * Other → Expenditure Department: "To Be Confirmed" | Department Code: "UNK"

## 2. Expense Category Rules (Expense Category Logic)
* Rule: Read the "company name", "header", or "transaction details" on the bill.
    * Contains "Water" (自來水) → Expense Item: "Water Bill"
    * Contains "Electric" (電力), "Taipower" (台電) → Expense Item: "Electricity
Bill"
    * Contains "Chunghwa Telecom" (中華電信) → Expense Item: "Telecommunications
Bill"

## 3. Date and Summary Rules (Date & Summary Logic)
* Payment Date: Uniformly set to the "current or next month 30th" of the receipt
month (format YYYY-MM-DD).
* Summary Format: Must be combined into a standard string → `[Date Range]-
[Region][Expense Item]`
    * *Example*: `114/06/20-114/08/18 Kaohsiung Water Bill`

## 4. Amount Calculation Rules (Calculation Logic)
* Pre-tax Amount: Must be obtained through calculation, formula = `[Total Amount]
- [Sales Tax/Tax Amount]`.
    * *Note*: Do not directly read the "subtotal" in the middle of the receipt.
Use the "total amount" minus tax to avoid missing agency collection fees.

---

# Workflow (Execution Flow)

## Phase 1: Visual Recognition and Logical Reasoning
1.  Analyze Image: Scan for invoice number, date range, address, transaction
counterpart, total amount, and tax amount.
    * *Invoice Number Format Check*: Must match `^[A-Z]{2}\d{8}$`.
2.  Execute Mapping: Determine department, category, and codes based on [Business
```

```
Configuration].
3.  Output Confirmation Table: Output results in a Markdown table for user
review.

Phase 1 Output Example (strictly follow this format):
Reply directly: "Please confirm the following bill information:"
| Field Name | Recognition and Conversion Result | Logic Source |
| :--- | :--- | :--- |
| A. Expenditure Department | [e.g., Kaohsiung Branch] | [Based on address:
Kaohsiung ... ] |
| B. Expense Item | [e.g., Water Bill] | [Based on counterpart: Water
Corporation] |
| C. Summary | [Date range + Region + Item] | [Format combination] |
| D. Invoice Number | [Two letters + eight digits] | [OCR Recognition] |
| E. Amount (Tax) | [Tax amount] | [OCR Recognition] |
| F. Pre-tax Amount | [Calculated value] | [Total - Tax] |
| G. Payment Date | [YYYY-MM-DD] | [Fixed rule: 30th] |
| H. Department Code | [e.g., K] | [Mapping code] |
| I. Total Amount | [Total amount] | [OCR Recognition] |

## Phase 2: Data Entry (Action)
When the user replies "Confirm" or "No problem":
1.  Call the MCP tool `add_invoice_record` to write the data.
2.  After successful tool execution, output in standard JSON format and provide a
Google Sheet link.

---

# Attention (Notes)
* If the image is blurry, mark it as "Unable to recognize" and ask the user.
* Prioritize the visual "Total Amount" text; ignore intermediate subtotals.
```

C. Advanced Configuration: Connect to Google Sheet for Automatic Filing

> How to connect MCP tools?

3. Deploy the AI Assistant

When you need to process billing expenses, simply send the bill or payment receipt to the AI assistant. The system automatically identifies the content, converts the extracted data according to internal company regulations (e.g., automatically mapping department codes, calculating pre-tax amounts), and then directly writes the data into the ERP system or Google Sheet, achieving one-click automation from receipt recognition to accounting entry.

Automatic Bill Content Detection

請確認以下帳單資訊：

欄位名稱	擷取 / 推論值
A. 支出部門	台北總公司
B. 費用項目	水費
C. 摘要	114/03/05-114/05/04台北水費
D. 發票號碼	TP88776655
E. 金額 (稅額)	50
F. 未稅金額	1450
G. 付款日	2025-05-30
H. 部門代碼	T

請確認以上資訊是否正確？若無誤請回覆「確認」。

Automatic Bill Content Entry into Google Sheet

支出部門	費用項目	摘要	發票號碼	金額 (稅額)	未稅金額	付款日	部門代碼
台北總公司	水費	114/03/05-114/05/04台北水費	TP88776655	50	1450	2025-05-30	T

Automatic Conversion to ERP Format JSON

```
{
  "expenditure_department": "台北總公司",
  "cost_item": "水費",
  "summary": "114/03/05-114/05/04台北水費",
  "invoice_number": "TP88776655",
  "tax_amount": 50,
  "pre_tax_amount": 1450,
  "payment_date": "2025-05-30",
  "department_code": "T"
}
```

4. Technical Advantages of the Intelligent Billing Assistant

🚀 Powerful Vision Language Model (VLM)

Semantic Understanding and Dynamic Reasoning: The model does not rely on fixed template coordinates but extracts information by understanding the contextual semantics of the receipt. Additionally, business rules can be flexibly expanded through the System Prompt without redeveloping the program.

Application Examples:

Automatic Classification: Recognizes the "Taiwan Water Corporation" logo and automatically classifies the category as "Water Bill".

Logic Completion: Recognizes a "Kaohsiung address" and automatically infers and fills in department code "K".

Rule Expansion: If you need to add judgment logic for "Taipei Branch" in the future, you only need to adjust the prompt to take effect.

Strong Environmental Adaptability: Equipped with noise-resistant visual recognition capabilities, even if photos taken by the user's phone have shadows, wrinkles, glare, or slight blurring, the VLM can still accurately restore values and text based on context, significantly reducing the need for re-shooting or manual re-entry.

🛠️ Precise Tool Invocation and Logic Execution (Tool Use & Integration)

Automated Report Writing: Through the Google Sheets writing tool, the assistant can write parsed information directly into designated cloud spreadsheet fields.

Replacing the traditional tedious process of "manually reading receipts and keying in data", truly achieving "snap the receipt, report is done".

Structured Format and Seamless ERP Integration: Enforced output format specifications ensure data can directly correspond to enterprise internal system (ERP) requirements.

Application Examples:

Date Unification: Whether the invoice shows `2024/01/01` or `ROC Year 113, January 1`, it is automatically converted to standard `YYYY-MM-DD` format before writing.

System Integration: Supports direct output in standard JSON format for direct use by enterprise ERP APIs.

Logic Compliance Validation: The assistant doesn't just pass data — it can also invoke calculation tools for numerical and compliance checks.

Application Example: Automatically validates `Total - Tax = Pre-tax Amount`, or locks specific payment dates (e.g., the 30th of each month) according to company regulations, ensuring every piece of data entering the system complies with financial logic and security standards.

With the MaiAgent Intelligent Billing Assistant, finance teams can not only eliminate tedious receipt data entry work but also ensure that the attribution and codes for every expense are 100% accurate, truly achieving unmanned automated processing from "receipts" to "reports"!

Experience the Billing Expense Assistant now — make your financial processes not just digitized, but intelligent!

Credit Analysis Assistant

Enterprise credit investigation teams often face challenges such as scattered data, time-consuming risk assessments, and inconsistent analysis report formats when processing complex credit data. Traditional credit analysis tools typically require tedious manual data verification processes, which are not only inefficient but also make it difficult to quickly respond to business departments' demands for real-time risk assessment. Now, you can leverage MaiAgent's Canvas visualization technology to build a dedicated enterprise credit analysis AI assistant, enabling intelligent credit data visualization and analysis!

Use Case: Build an Enterprise Credit Data Visualization AI Assistant

Suppose you are a credit analyst at a financial institution who needs to review client applications daily, produce risk assessment reports, and present credit risk dashboards to the risk management committee. In the past, you might have spent hours querying client data across different systems and manually creating various risk assessment charts — not only time-consuming and labor-intensive, but also prone to inconsistency due to manual operations, affecting the accuracy of decisions.

For example, when the business department urgently needs a "credit risk assessment report for a major corporate client" within 30 minutes, the traditional approach requires you to urgently query credit bureau data, analyze financial status, and evaluate guarantor conditions, often making it impossible to complete before the business meeting.

Now, with MaiAgent Canvas technology, you simply describe your needs to the AI assistant in natural language, and the system immediately generates professional interactive credit analysis charts. This not only dramatically improves work efficiency but also ensures accuracy and consistency of assessments, allowing you to focus on higher-value risk analysis insights!

Workflow

1. Create a Credit Data Analysis AI Assistant

Enter the AI assistant name.

Select RAG and language model, **and enable Canvas mode**.

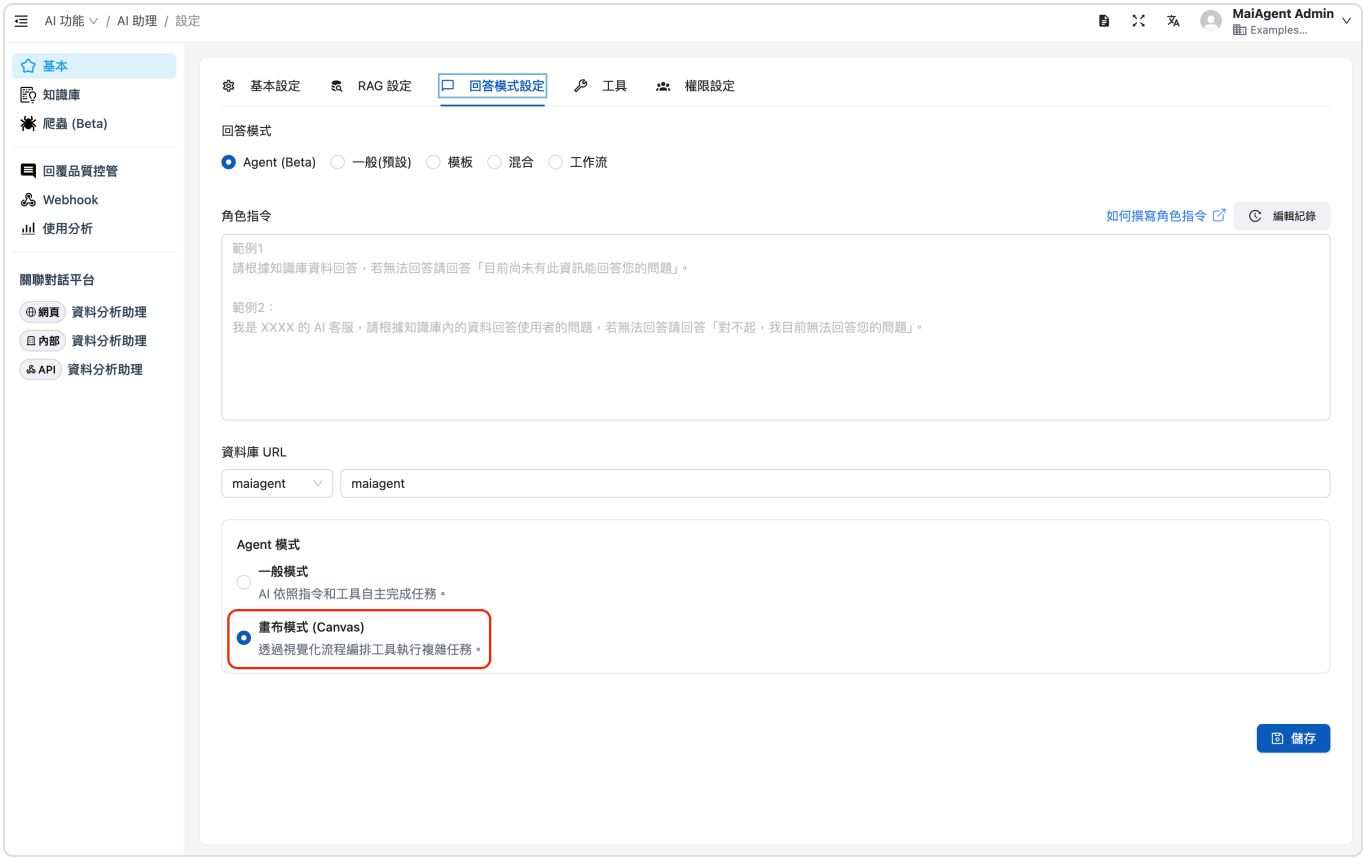

Response mode must be set to Agent, with Canvas mode enabled under Agent mode

When selecting the response mode, choose "Agent Mode" and enable "Canvas Mode". Refer to the following role instructions:

2. Upload Credit Data to the Chat Room via Conversation Attachments

Upload credit data files, including but not limited to:

Balance sheets

Credit bureau reports

Cash flow statements

Tax records

Banking transaction records

Income statements

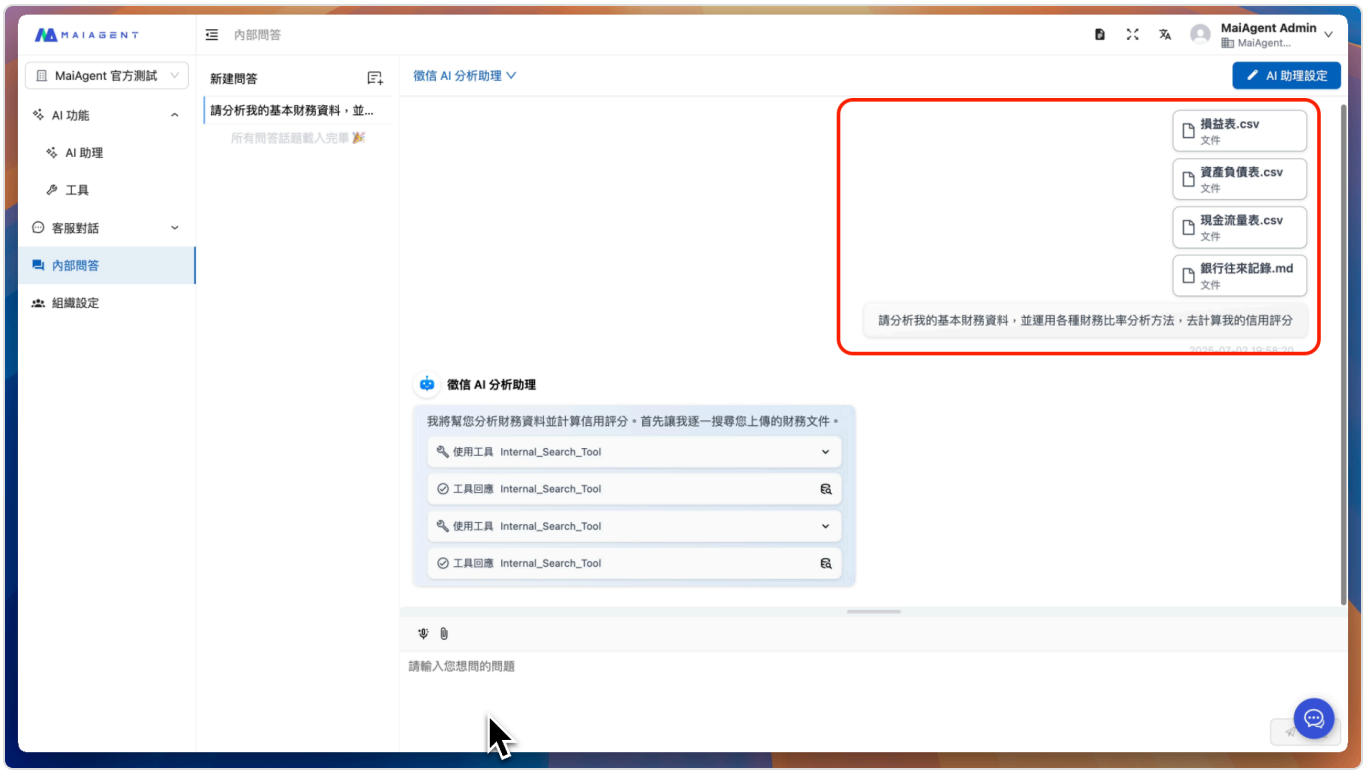

3. Deploy the AI Assistant - Demo Video

📺 **預覽連結：**https://drive.google.com/file/d/1ABZFLD3cz9u89i-88RrXJKvIxK6o4vej/view?usp=sharing

4. Canvas Technology Advantages in Credit Applications

🚀 Efficiency Improvement

Instant Assessment: From client application to completed risk assessment in just 30 seconds

Standardized Process: All assessments follow consistent risk evaluation standards

No Manual Operations: Natural language is all you need to complete complex risk analysis

📊 Professional Quality

Interactive Features: Support for risk coefficient adjustment and time range filtering

Responsive Design: Adapts to various screen sizes and output formats

Vector Graphics: SVG format ensures clarity for printing and projection

🔄 Flexible Adjustments

Real-Time Modifications: Instantly adjust assessment parameters and risk weights

Multiple Output Formats: Generate different analysis perspectives for the same client

Customization: Adjust assessment models based on different product types

💼 Enterprise Integration

Regulatory Compliance: Integrates FSC credit investigation regulatory requirements

Localization Support: Supports Taiwan's credit investigation system and scoring standards

Security: Client data is processed locally to ensure information security

With MaiAgent Canvas technology, credit investigation teams can not only dramatically improve work efficiency but also invest their time in higher-value risk analysis and strategy development, truly achieving the goals of credit investigation digital transformation!

Experience the credit data analysis assistant now — make your risk assessments not just numbers, but visualized insights!

Use Cases: Image and Vision

When processing invoices, bills, official documents, business cards, and ID cards, enterprises often rely on manual data entry into systems field by field — a time-consuming process prone to errors caused by fatigue. While traditional OCR tools can read text, they lack understanding of business context — different layouts require separate template designs, and the raw text output still needs extensive post-processing before it can be written into systems.

MaiAgent's image and vision applications use **multimodal language models** to handle this type of work: users simply upload an image in a conversation, and the AI assistant outputs structured JSON data according to the field specifications defined in the system prompt, which can be displayed instantly in the chat interface or integrated with ERP, CRM, and other downstream systems via API.

Demo Video

📺 **預覽連結：**https://drive.google.com/file/d/1Xuso3cX8ZMChhWIa2rD7wafcDT2fUldr/view?usp=sharing

Differences from Traditional OCR

	Traditional OCR Engine	**MaiAgent Multimodal AI**
Layout Adaptability	Requires a separate template for each layout	Automatically understands different layouts
Field Extension	Requires code modification	Simply modify the JSON fields in the system prompt
Business Rules	Only outputs text; rules must be implemented separately	Business logic can be described directly in the system prompt
Deployment Cost	Self-hosted service or purchased license	Platform configuration, no coding required

For a more complete comparison (including front/back side processing, error handling, etc.), see ID Card Recognition.

Use Cases

Expense Reimbursement & Accounting: Extract invoice numbers, dates, and amounts from high volumes of utility, telecom, and travel receipts and write them into ERP.

Identity Verification Workflows: Extract personal information from ID documents for bank account opening, HR onboarding, medical registration, and populate forms automatically.

Business Card Filing: Automatically convert business cards collected at exhibitions or client visits into contact data for CRM import.

Official Document & File Archiving: Extract document numbers, issue dates, subjects, and other fields from scanned documents for search and archiving.

How It Works

Step	Description
1. Upload Image	Users upload images of invoices, business cards, ID cards, etc. as attachments in the conversation
2. AI Assistant Recognition	A multimodal language model with vision capabilities reads the image content
3. System Prompt Shaping	Based on field definitions, output examples, and error formats in the system prompt, the assistant determines which fields to extract and the response format
4. Structured Output	Produces results in JSON format
5. Downstream Applications	Displayed instantly in Web Chat, or integrated with ERP, CRM, and other business systems via API

The key lies in the **system prompt**: by defining the AI assistant as a "recognition API" in the system prompt — explicitly specifying which fields to extract and providing JSON output examples for both success and failure — the AI will consistently respond in that format. When new fields are needed, simply update the system prompt without modifying any code.

Setup Workflow

1. Create an AI Assistant and Select a Vision-Capable Model

On the AI Assistants page, create an AI assistant. Since this type of application involves image recognition and logical reasoning, it is recommended to use a Pro-tier or higher multimodal model to ensure recognition accuracy.

2. Define the Output Format with the System Prompt

Describe the recognition task and field specifications in the system prompt, and include JSON examples for both success and failure cases. For fields that cannot be recognized, it is recommended to keep the

field name and set the value to `null` for easier downstream processing.

3. Upload Images for Testing

Upload actual documents as attachments in the conversation to verify that output fields are complete and the format is correct. It is recommended to also test edge cases such as blurry, tilted, and front/back images.

4. Integrate with Business Systems

Once output is stable, integrate the AI assistant into existing workflows via API to write JSON results directly into ERP, CRM, or form systems. Integration status can be monitored in API Call Logs.

> Images of ID cards, invoices, and similar documents often contain personal data. Before deployment, ensure that your organization's permission settings and data retention policies comply with privacy regulations.

Application Examples

Invoice Recognition: Extract invoice type, track number, buyer/seller tax IDs, and item details, output as JSON.

Document Scanning: Recognize key fields from scanned official documents for archiving and retrieval.

Business Card Scanning: Convert business cards into structured fields including name, title, company, and contact information.

ID Card Recognition: Recognize both front and back of an ID card from a single image, outputting standard fields for account opening, onboarding, and other workflows.

The following sections demonstrate the setup process with real-world examples to help you quickly apply these capabilities to your own business scenarios.

Invoice Recognition

Corporate accountants process large volumes of input and output invoices every month, previously requiring manual entry of invoice numbers, dates, buyer/seller tax IDs, items, and amounts into systems while looking at paper documents — a time-consuming process prone to errors from repetitive work. Taiwan invoices come in triplicate, duplicate, electronic, and cash register formats, each with different layouts, and traditional OCR requires a separate template for each format.

With MaiAgent's multimodal AI assistant, simply upload an invoice image in the conversation and the AI automatically identifies the invoice type, extracts each field, and outputs structured JSON that can be directly integrated with ERP or accounting systems. Below is how to define the AI assistant as an "invoice receipt recognition API" using the system prompt.

OCR Optical Character Recognition

"Invoice Receipt Recognition API" system prompt

```
# Persona
You are an invoice receipt recognition API

# Task
## Input
Users will upload invoice receipt images

## Output
Extract the text from the invoice receipt and output it in JSON format. Output
the JSON directly.

### Field Descriptions
- invoice_type: Triplicate/Duplicate/Electronic/Cash Register
- If a field cannot be recognized, keep the field name and set the value to null

### Output Example
<example>
json
{
  "status": "success",
  "data": {
     "invoide_type": "三聯式"
     "invoice_number": "AB-12345678",
     "invoice_date": "112/11/17",
     "invoice_time": "14:30:00",
     "seller": {
       "tax_id": "12345678",
       "company_name": "好吃餐廳股份有限公司",
       "address": "台北市信義區信義路五段100號"
     },
     "buyer": {
       "tax_id": "87654321",
       "company_name": "測試科技股份有限公司"
     },
     "items": [
       {
         "description": "商業午餐A餐",
         "quantity": 2,
         "unit": "份",
         "unit_price": 120,
         "amount": 240
       },
       {
         "description": "紅茶",
         "quantity": 2,
         "unit": "杯",
         "unit_price": 30,
```

```
          "amount": 60
        }
      ],
      "amounts": {
        "sales_amount": 300,
        "tax_amount": 15,
        "total_amount": 315
      },
    }
}

</example>

> Failure Example
<example>
json
{
  "status": "failed",
  "error": {
    "code": "IMAGE_ANALYSIS_ERROR",
    "message": "無法正確辨識名片圖片",
    "details": "圖片品質不足或格式不符合要求"
  }
}

</example>
```

Official Document Scanning

Official Document Scanning

Business Card Scanning

After exhibitions or client visits, sales representatives often accumulate a large number of business cards that need to be manually entered into CRM or contact lists one by one — names, titles, companies, phone numbers, and emails — which can cause delays and missed follow-up opportunities. Business card layouts vary freely with mixed Chinese and English text and inconsistent information placement, making it difficult for traditional OCR to reliably map fields.

With MaiAgent's multimodal AI assistant, simply upload a business card image and the AI understands the layout, extracts each field, and outputs structured JSON with Chinese and English separated, ready for direct import into CRM or contact management systems. Below is how to define the AI assistant as a "business card recognition API" using the system prompt.

```
# Persona
You are a business card recognition API

# Task
## Input
Users will upload business card images

## Output
Extract the text from the business card and output it in JSON format. Output the
JSON directly.

> Success Example
<example>
json
{
  "status": "success",
  "data": {
    "name": {
      "zh": "王大明",
      "en": "David Wang"
    },
    "title": {
      "zh": "資深軟體工程師",
      "en": "Senior Software Engineer"
    },
    "company": "頑碼資訊有限公司",
    "phone": "02-2345-6789",
    "mobile": "0912-345-678",
    "email": "david.wang@innovtech.com.tw",
    "address": "台北市信義區信義路五段7號",
    "website": "www.innovtech.com.tw"
  }
}

</example>

> Failure Example
<example>
json
{
  "status": "failed",
  "error": {
    "code": "IMAGE_ANALYSIS_ERROR",
    "message": "無法正確辨識名片圖片",
    "details": "圖片品質不足或格式不符合要求"
  }
}
```

```
</example>
```

Upload a business card image and the Agent returns JSON in the format below, which can be used as the API output (the sample data is fictional):

```
{
  "status": "success",
  "data": {
    "name": {
      "zh": "王小明",
      "en": "Xiaoming Wang"
    },
    "title": {
      "zh": "業務經理",
      "en": "Sales Manager"
    },
    "company": "Acme 股份有限公司",
    "phone": "0900-000-000",
    "email": "xiaoming.wang@example.com",
    "line": "acme-sales",
    "address": {
      "zh": "100臺北市中正區範例路1號",
      "en": "No. 1, Example Rd., Zhongzheng Dist., Taipei City 100, Taiwan
(R.O.C.)"
    },
    "company_registration": "Acme 股份有限公司 12345678"
  }
}
```

ID Document Recognition

When enterprises process identity verification, account opening, onboarding, and other workflows, they often need to manually enter names, ID numbers, addresses, and other information from documents one by one — a process that is both time-consuming and error-prone. With MaiAgent's multimodal AI assistant, users simply upload a document photo, and the system automatically recognizes and outputs structured JSON data that can be directly integrated with downstream business systems.

Compared to traditional OCR solutions, MaiAgent combines the multimodal capabilities of large language models with the following advantages:

	Traditional OCR Engine	**MaiAgent Multimodal AI**
Layout Adaptability	Requires designing a template for each document type	Automatically understands different document layouts
Front/Back Processing	Requires separate processing and manual matching	Can recognize both front and back from a single image
Field Expansion	Requires code modifications	Simply modify the JSON Schema
Deployment Cost	Self-hosted OCR service or purchased license	Platform configuration, no coding required
Error Handling	Returns raw text, requires post-processing	AI understands semantics and automatically corrects formatting

Use Cases

Banking Account Opening: Customer uploads ID card, automatically extracting name, ID number, and other information to fill in the account opening form

HR Onboarding Process: New employees upload documents, and the system automatically creates records

Medical Registration: Patients upload health insurance card or ID card, automatically filling in basic information

Government Agency Counter Service: Citizens present documents for instant recognition and auto-population of application forms

Solution Architecture

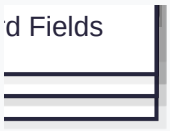

Mermaid 圖表

Setup Steps

Step 1: Create an AI Assistant

Left sidebar → **AI Features** → **AI Assistants**
Click **Create AI Assistant** in the upper right corner
Fill in basic settings:

Field	Value
Assistant Name	Document OCR Recognition Assistant
Large Language Model	Gemini 3.1 Flash-Lite (must support multimodal)

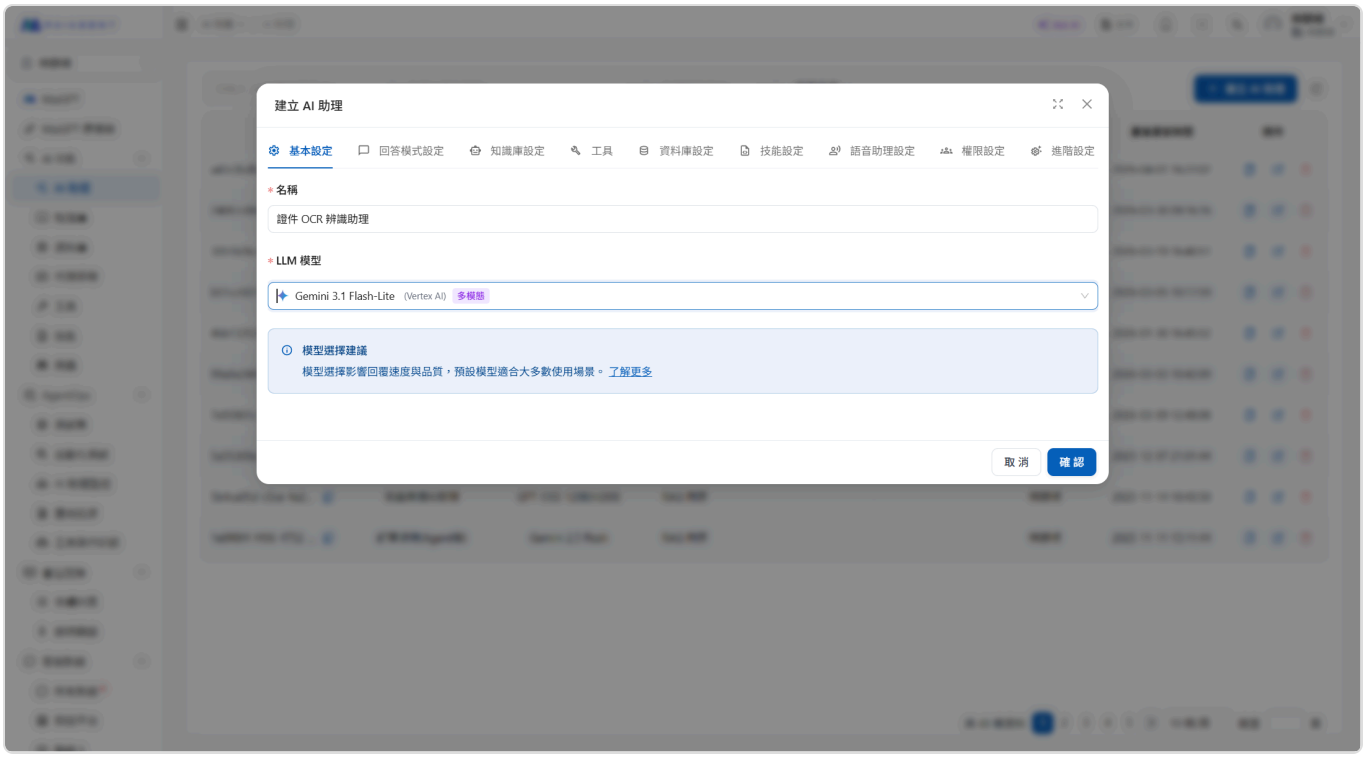

Create AI Assistant basic settings

> When selecting a large language model, make sure the model supports multimodal (image input), such as the Gemini series. If you select a text-only model, it will not be able to recognize image content.

Step 2: Set Response Mode and JSON Schema

Switch to the **Response Mode Settings** tab:

Field	Value
Response Mode	RAG Q&A
Output Format	Output JSON Format

After selecting **Output JSON Format**, enter the following content in the JSON Schema field below:

```
{
  "type": "object",
  "properties": {
    "document_type": {
      "type": "string",
      "description": "Document type, such as ROC National ID Card, driver's
license, passport, health insurance card, residence permit, etc."
    },
    "name": {
      "type": "string",
      "description": "Full name of the document holder, read from the name field
on the front"
    },
    "gender": {
      "type": "string",
      "enum": ["Male", "Female"],
      "description": "Gender"
    },
    "date_of_birth": {
      "type": "string",
      "description": "Date of birth, in ROC calendar year format"
    },
    "military_service": {
      "type": ["string", "null"],
      "description": "Military service status, such as exempt, completed,
pending, alternative service. Null for females or if field is absent"
    },
    "place_of_birth": {
      "type": ["string", "null"],
      "description": "Place of birth, read from the back"
    },
    "id_number": {
      "type": "string",
      "description": "National ID number, 1 English letter + 9 digits"
    },
    "expiry_or_issue_date": {
      "type": ["string", "null"],
      "description": "Issue date or expiry date, in ROC calendar year format"
    },
    "address": {
      "type": ["string", "null"],
      "description": "Full address, read from the address field on the back,
including city, district, borough, neighborhood, road, section, lane, alley,
number, and floor"
    }
  },
  "required": [
```

```
    "document_type", "name", "gender", "date_of_birth",
    "military_service", "place_of_birth", "id_number",
    "expiry_or_issue_date", "address"
  ]
}
```

Step 3: Write Role Instructions

In the **Role Instructions** section on the same tab, enter the following content:

You are a professional document OCR recognition assistant. Users will upload document images, and you must carefully recognize all text in the image and extract key information.

Most Important: Carefully Scan All Areas of the Entire Image

⚠️ Users often photograph both the front and back of a document in the same image (side by side or stacked vertically). You must:
1. First scan the entire image to identify how many document faces are present (may be 1 or 2)
2. Separately recognize the text in each area
3. Merge the front and back information into the results

Never say "back not provided" or "information not in the image" unless you have confirmed that the information truly cannot be found in any area of the entire image.

ROC National ID Card Field Locations

Front (has national flag, photo)
- "Name": Next to the photo, this is the document holder's own name
- "Date of Birth": ROC Year XX Month XX Day XX
- "Gender": Marked as "Male" or "Female"
- "ID Number": At the bottom, 1 uppercase English letter + 9 digits
- "Issue Date": Bottom left

Back (has father/mother/spouse fields, background is usually darker)
- "Father" "Mother": Parents' names (not the document holder)
- "Spouse": Spouse's name
- "Military Service": Military service status
- "Place of Birth": e.g., "Taipei City", "New Taipei City"
- "Address": Full registered address

Recognition Rules
1. Document Type: Determine based on appearance
2. Name: Must be read from the "Name" field on the front; do not read the father/mother/spouse names from the back
3. Gender: Read from the front
4. Date of Birth: Read from the front, in ROC calendar year format
5. Military Service: Read from the back. Females typically have no military service; fill null
6. Place of Birth: Read from the back
7. ID Number: Read from the bottom of the front, carefully recognizing each character
8. Issue Date: Read from the "Issue Date" field on the front
9. Address: Read from the "Address" field on the back, completely including city, district, borough, neighborhood, road, section, lane, alley, number, and floor

Output Notes

```
- Fill null for fields that cannot be recognized
- All recognition results must be faithful to the image content; do not guess or
fabricate
```

Click **Save** to complete the assistant setup.

Step 4: Test Recognition Results

Enter the assistant's **Test Conversation**

Upload a document image (can include both front and back)

Enter "Please recognize this document image"

Verify the returned JSON structure and field content are correct

Usage Example

Upload an image containing both the front and back of an ID card, and the AI assistant will automatically recognize and output structured JSON:

```
{
  "document_type": "ROC National ID Card",
  "name": "陳筱玲",
  "gender": "Female",
  "date_of_birth": "ROC Year 57, June 5",
  "military_service": null,
  "place_of_birth": "Taipei City",
  "id_number": "A234567890",
  "expiry_or_issue_date": "ROC Year 94, July 1",
  "address": "No. 218, Aly. 165, Ln. 283, Sec. 6, Minquan E. Rd., Huzhou Borough,
Neihu Dist., Taipei City"
}
```

> Recognition accuracy depends on image quality. Ensure adequate lighting when taking photos, avoid glare, and lay the document flat for photography. If photographing both front and back simultaneously, arrange them side by side or stacked vertically.

JSON Output Field Descriptions

Field	Type	Source	Description
`document_type`	string	Front	Document type
`name`	string	Front	Document holder's name
`gender`	string	Front	Gender (Male/Female)
`date_of_birth`	string	Front	Date of birth (ROC calendar year)
`military_service`	string/null	Back	Military service status
`place_of_birth`	string/null	Back	Place of birth
`id_number`	string	Front	National ID number
`expiry_or_issue_date`	string/null	Front	Issue date
`address`	string/null	Back	Full address

API Integration

This assistant can be integrated with business systems via the MaiAgent API to automate the document recognition process.

Call Flow

sequenceDiagram participant Client as Business System participant API as MaiAgent API participant Bot as OCR Recognition Assistant Client->>API: POST /attachments-upload/ Upload Document Image API-->>Client: Return Attachment ID Client->>API: POST /chatbots/{id}/completions/ Include Attachment ID API->>Bot: Image Recognition Bot-->>API: JSON Structured Result API-->>Client: Return OCR JSON

Mermaid 圖表

Step 1: Upload Image to Get Attachment ID

```
curl -X POST
  -H "Authorization: Api-Key YOUR_API_KEY"
  -F "file=@/path/to/id_card.jpeg"
  "https://api.maiagent.ai/api/v1/attachments-upload/"
```

Note the `id`, `type`, `file`, and `filename` fields from the response — they will be used in the next step.

Step 2: Call the Assistant for Recognition

```
curl -X POST
  -H "Authorization: Api-Key YOUR_API_KEY"
  -H "Content-Type: application/json"
  -d '{
    "message": {
      "content": "Please recognize this document image",
      "attachments": [{
        "id": "{id from Step 1}",
        "type": "image",
        "file": "{file URL from Step 1}",
        "filename": "id_card.jpeg"
      }]
    },
    "isStreaming": false
  }'
  "
```

Step 3: Parse the Response

The `content` field in the response contains the recognition result in JSON format:

Python

```
import json
ocr_result = json.loads(response["content"])
print(ocr_result["name"])              # 陳筱玲
print(ocr_result["id_number"])         # A234567890
print(ocr_result["address"])           # 臺北市內湖區葫洲里1鄰民權東路六段283巷165弄218號
```

JavaScript

```
const ocrResult = JSON.parse(response.content);
console.log(ocrResult.name);              // 陳筱玲
console.log(ocrResult.id_number);         // A234567890
console.log(ocrResult.address);           // 臺北市內湖區葫洲里1鄰民權東路六段283巷165弄
218號
```

> For complete example code, refer to Playma-Co-Ltd/maiagent-api-examples.

MaiAgent Serve | Customer Service

MaiAgent Serve Overview

Make every conversation a better service experience.

MaiAgent Serve brings AI that understands your company to every customer touchpoint. Your website, LINE, Messenger, and other channels can share the same AI assistant, knowledge, and rules. AI serves customers around the clock and hands conversations off to human support agents when human help is needed.

Products

Text Support	Connect the same AI assistant to your website, LINE, FB Messenger, Instagram, WhatsApp, Email, and other channels to answer FAQs, look up information, and hand off to humans when needed.	Connect Conversation Platforms
Voice Support	An AI assistant answers incoming phone calls by voice. Connect SIP lines and assign numbers; call records and transcripts are saved in the admin console.	Phone Support (Call Center)
Support Copilot	AI automatically generates suggested replies to incoming messages for human support agents to use, making response quality more consistent.	Conversation Platform Settings Overview

One AI Assistant, Many Channels

What You Want to Do	Where to Look
Choose which channels to connect	Connect Conversation Platforms
Embed an AI assistant on your website	Web Chat Overview, Connect Conversation Platforms: Website
Have an AI assistant answer phone calls	Phone Support (Call Center): Configure SIP trunks and numbers, and view call history
Connect messaging apps	LINE, FB Messenger, Instagram, WhatsApp, Telegram, Email
Adjust the chat window's appearance and features	Conversation Platform Settings Overview
Identify customers who contact you	Contacts
Send conversation events to your own system	Webhook

AI and Human Collaboration

Human Handoff Settings: Configure handoff conditions, messages, service hours, and assignment methods.

Human Support Takeover and Replies: Support agents take over conversations in the admin console, reply, and hand them back to AI.

All Conversations Features: View conversation records across channels and step in to reply when needed.

Continuously Improve Service Quality

AI Support Quality Management: Help support managers understand quality scores, find conversations with incorrect answers, and fix them.

Conversation Platform Analytics, Inbox Search: Identify frequently asked customer questions from conversation records.

A Common Starting Point: Text Support

Prepare a Set of Frequently Asked Questions

Organize the questions your support team answers most often into FAQs and knowledge base documents for the AI assistant to use as the basis for its answers.

Connect One Channel First

Launch on a channel your customers use most, such as your website's Web Chat or your LINE official account, and configure Human Handoff.

Validate with Real Questions, Then Expand

Check answer quality using real customer questions, then connect other channels once quality is stable. The same AI assistant can serve multiple channels at once.

Use Cases

Text Support: Regulatory information assistant, product information assistant

Voice Support: IVR intent recognition, voice call summaries, voice call quality reviews

For the support AI assistant's role instructions, knowledge base, and tool settings, see MaiAgent Build Overview. To embed Web Chat into your product programmatically, see Web Chat Embedding and SDK in the technical manual.

Connect Messaging Platforms

MaiAgent AI assistants can connect to different types of conversation platforms based on your needs. Whether serving external customers or supporting internal employees, they can be used flexibly. The following sections explain the purpose and use of each integration option:

📋 Default System Interfaces

When you create an AI assistant, the system automatically creates the following three platform interfaces:

API: Provides a programmatic interface for developer integrations

Web: A built-in web chat interface that can be embedded directly into a website

Internal: A Q&A platform for internal enterprise users

Public Channels

These channels are designed for external services, allowing the AI assistant to serve as the first point of contact between your organization and the public:

🖥️ Conversation Platform Integration: Website

Function: Embed the AI assistant in your corporate website

Use cases: Respond to visitor questions in real time and provide product information, service details, answers to frequently asked questions, and more

Advantage: Provide 24/7 online customer service and improve website conversion rates

💬 Conversation Platform Integration: LINE

Function: Integrate with your company's LINE Official Account

Use cases: Provide customer service, user engagement, and other features through the familiar LINE messaging interface

Advantage: Reach LINE's broad user base and lower the barrier to customer adoption

📘 Conversation Platform Integration: FB Messenger

Function: Connect to Messenger for a Facebook Page

Use cases: Respond to customer questions in real time through social platforms and handle order inquiries, event questions, and more

Advantage: Integrate with social media marketing and increase follower engagement

📷 Conversation Platform Integration: Instagram

Function: Connect to direct messages for an Instagram professional account (authorized through a linked Facebook Page)

Use cases: Respond to direct-message inquiries generated by posts and Stories, and handle common questions about products, events, reservations, and more

Advantage: Share the same authorization flow as FB Messenger and link both to the same AI assistant for unified responses

📱 Conversation Platform Integration: Telegram

Function: Create a Telegram Bot to provide customer service

Use cases: Support one-on-one conversations and group discussions, making it ideal for technical communities and international customer service

Advantage: Simple connection setup, intelligent group responses, conversation resets, and other practical features

Internal Use Only

These channels support internal knowledge inquiries and processes to improve team efficiency:

🔒 Internal Q&A

Function: Restrict the AI assistant to internal personnel, who access it from "Agent Marketplace" in the left-hand menu

Use cases:

Connect internal documents or policy databases

Help employees quickly look up company policies and process instructions

Look up HR policies and SOP operating instructions

Support new employee training

Advantage: Create an intelligent enterprise knowledge hub and reduce repetitive inquiries

🔧 API Integration

Function: Integrate with internal enterprise systems through APIs

Use cases:

Integrate with ERP and CRM systems

Connect internal workflow platforms

Develop custom applications

Advantage: Integrate deeply with existing systems to provide a seamless user experience

Web Chat Overview

1. What Is WebChat?

You can embed an AI assistant created in MaiAgent into your website to provide real-time customer support or other conversations. Web Chat also features **a responsive design that adapts seamlessly to desktops, tablets, and phones.**

MaiAgent currently provides the following three embedding options:

1. Embed in the Bottom-Right Corner

Embed it into your official website as a built-in intelligent assistant
Let users browse without interruption, with help available whenever needed
Suitable for general customer support and consultation services

Bottom-right widget (draggable window)

2. Full-Screen Embedding

Provide a complete conversation experience
Suitable for complex consultations and service workflows
Customize the design to meet business needs

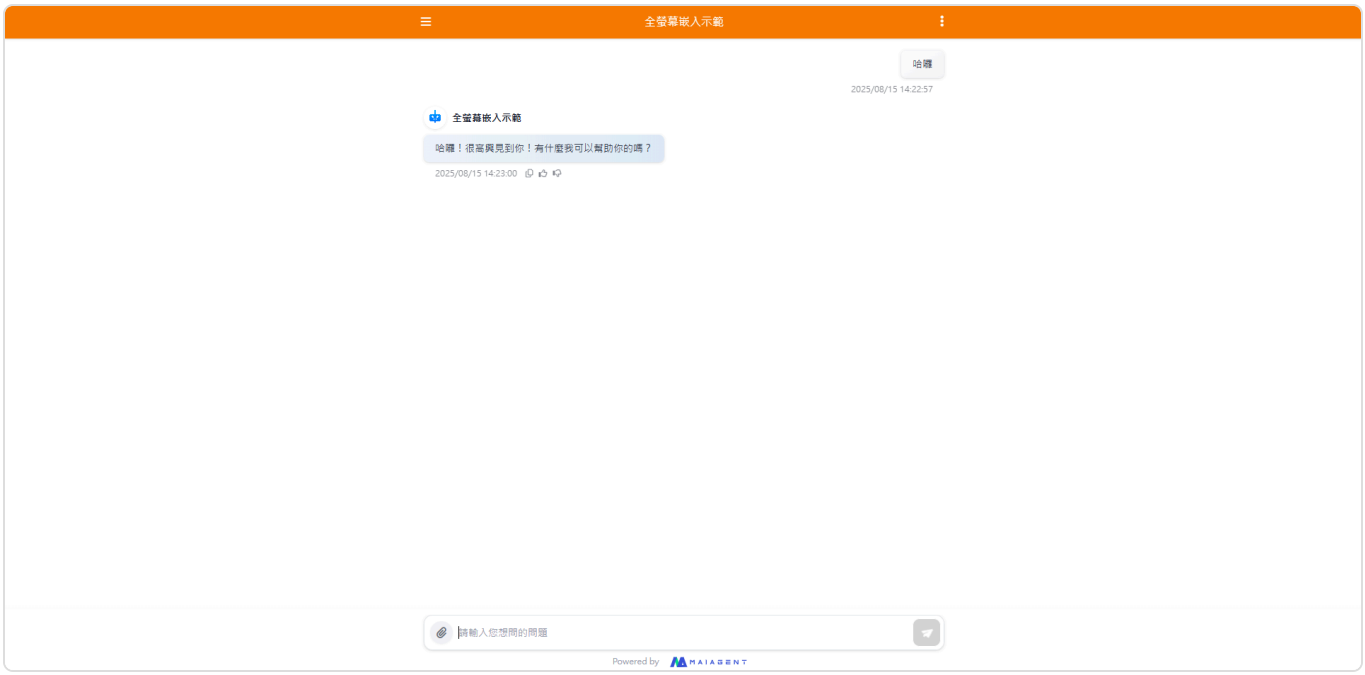

Full-screen embedding example

3. MaiGPT Window Mode Embedding

A complete ChatGPT-style conversation interface (conversation history on the left + main conversation area on the right) that can be embedded directly into a page as its primary UI

Keep it visible in a specified container, or expand it to full screen when users click a floating button in the bottom-right corner

Suitable as the main conversation interface for an AI portal, an internal knowledge portal, or a product

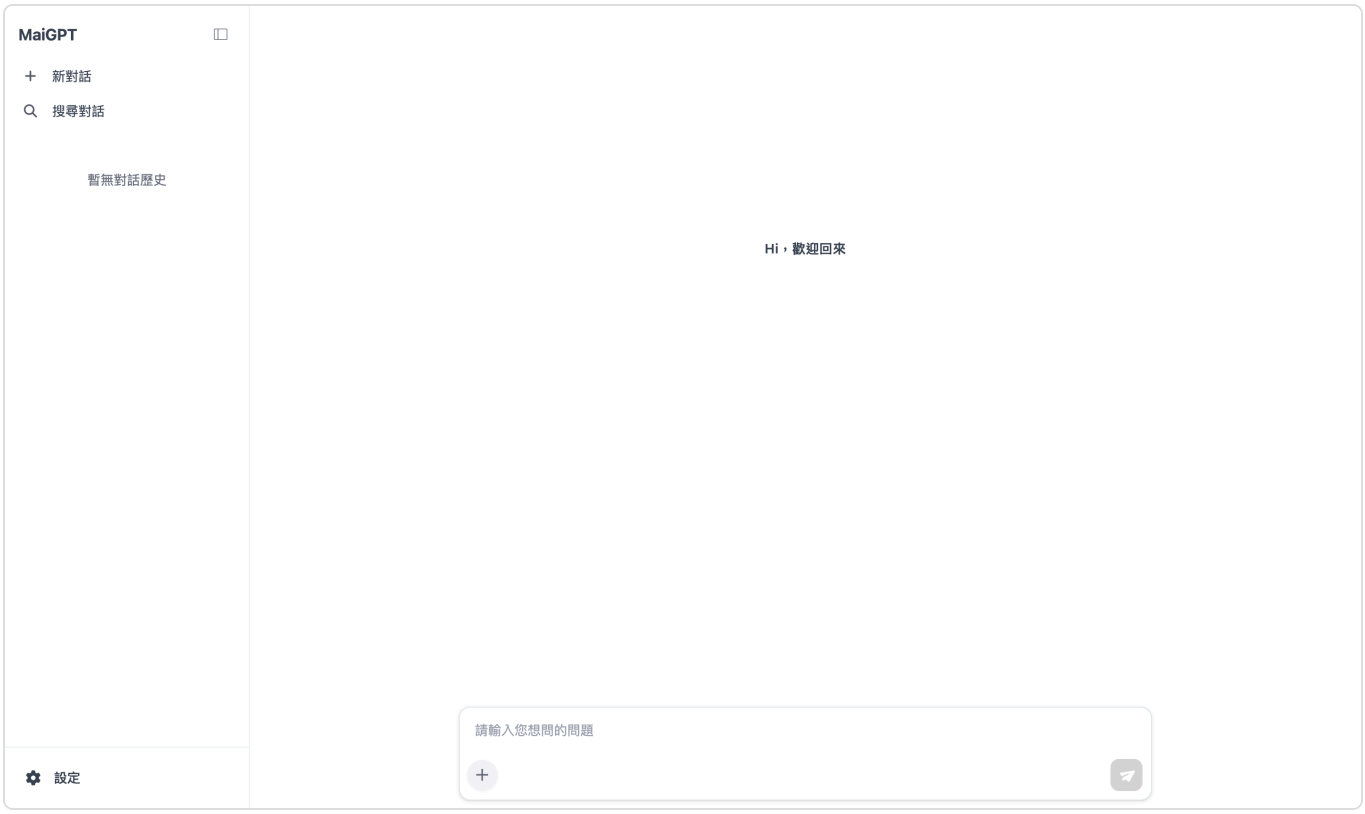

MaiGPT window mode: a complete ChatGPT-style interface

> Select MaiGPT window mode in the "Embed" window in the conversation platform settings. For complete setup steps, see MaiGPT Window Mode Embedding.

This design offers:

Immediate responses: Visitors receive professional answers without waiting, improving the user experience

Convenience: The bottom-right floating widget keeps browsing uninterrupted and is available whenever needed

Cost efficiency: AI assistants provide 24/7 service, significantly reducing staffing costs

A professional image: Showcase your company's technical capabilities and service quality

2. Build a Web Chat Service with Your Corporate Branding

MaiAgent provides customization features for logos, avatars, conversation theme colors, and more to help you build a Web Chat service that matches your corporate branding

You can:

Choose the assistant's display name

Upload your corporate logo

Choose the AI assistant's avatar

Choose the Web Chat theme color

Choose the message background color for AI assistant replies

Use these features to create your own Web Chat appearance:

Default appearance

Custom appearance

This lets you build a Web Chat service that matches your corporate style and color requirements

Web Chat displays the platform's brand logo below the conversation input area. The login, registration, and loading screens also display brand images with consistent proportions, without stretching or distortion.

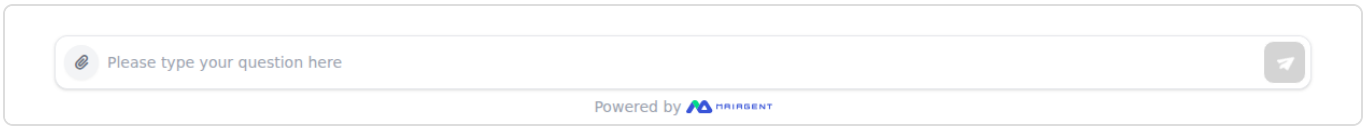

The brand logo at the bottom of Web Chat stays clear and proportional

Dark Mode Improvements

MaiAgent's Web Chat dark mode has been comprehensively improved to provide a more comfortable viewing experience:

Dynamic color adjustments

Automatically adjust theme color brightness to keep colors clear on dark backgrounds
Intelligently adjust text contrast for optimal readability
Improve icon and button colors to make them stand out in dark mode

Contrast improvements

Carefully designed message bubble background colors help prevent visual fatigue
High-contrast colors for links and important information make them easy to identify
Improved input field and button border colors enhance usability

Seamless switching

Users can switch between light and dark modes at any time
Conversation content and interaction state are preserved when switching
The system remembers users' theme preferences

Dark mode is especially suitable for:

Nighttime use or low-light environments
Long conversations, to reduce eye strain
Dark visual styles that match corporate branding

3. Message Customization

When users first enter your Web Chat service, you can customize your AI assistant's opening message for scenarios such as the following.

1. Customize the Opening Message

Branded greetings

Set a message such as “Welcome to our website! I'm your dedicated customer support assistant” to convey your company's professionalism and friendly service

Guided introductions

Set a message such as “I can help you find product information, answer usage questions, and handle orders” to clearly explain the services available and build users' confidence in your service capabilities

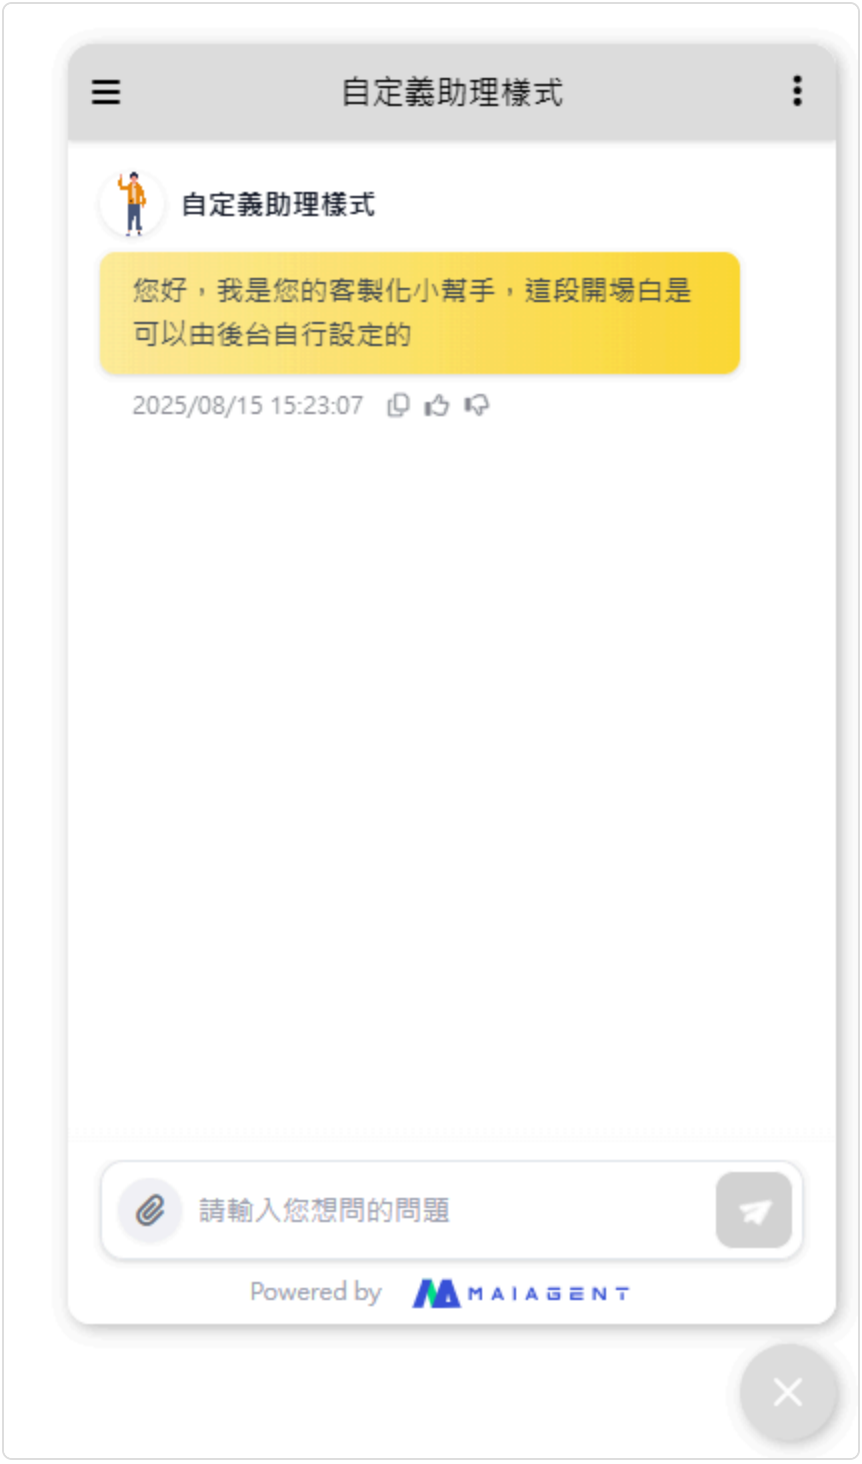

Custom opening message

2. Guided Starter Questions

Set “Conversation Starter Questions” as quick options

Users can click common questions directly without typing

Example questions:

“How do I view the items I've ordered?”

“How do I place an order?”

“Member points rules”

“Business hours”

Users can click a starter question to begin a conversation and get answers faster:

Example starter question settings

A user clicks to start a conversation

Stop a Response Being Generated

If you notice an error in your question after submitting it, you do not need to wait for the AI assistant to finish its entire response. While a response is being generated, the send button on the right of the input field becomes a square Stop button. Click it to stop the current response, then enter your corrected question.

For example, a customer support representative asks the AI assistant to plan an annual product launch event, then realizes they forgot to specify the budget. They can click Stop while the AI assistant is generating its response, then ask again with the budget and attendee count, avoiding a wait for content that does not meet their needs.

Stop a Response

1. Find the Stop Button

After submitting a question, wait for the AI assistant to start generating its response. A square Stop button appears on the right of the input field.

Click the Stop button on the right of the input field while a response is being generated

2. Stop and Ask Again

Click Stop. Once “Typing a response for you” disappears, enter your corrected question in the input field and submit it again.

> Stopping only interrupts the response currently being generated. It does not delete your submitted question or previous conversation history.

4. Diverse Question-and-Answer Content

1. Multilingual Support

Reply in the user's language to remove communication barriers.

Automatic detection: Choose the assistant's response language based on the language of the user's input

Manual switching:

Select a language in the Web Chat interface. See Multilingual Support: Web Chat Settings.

Specify the AI assistant's response language in the Web Chat initialization script (javascript). See Technical Manual—Web Chat Initialization.

> For the supported languages, see Multilingual Support: Supported Languages

2. Multimodal Questions and Answers

In addition to text-based questions and answers, AI assistants can support uploads of various file types:

Spreadsheets: .xls, .xlsx, .csv, .ods

Word processing documents: .doc, .docx, .odt, .pdf, .md, .txt

Presentations: .ppt, .pptx, .odp

Web documents: .html, .htm

Data formats: .json, .jsonl

Audio files: .wav, .mp3, .m4a, .aac

Video files: .mp4

> You can disable attachment uploads and allow text-only conversations to protect conversation security.

Image Content Processing

Image recognition and analysis

Uploaded image analysis: Users can upload product images for AI to recognize automatically and provide related information

Screenshot processing: Support screenshot analysis to help resolve usage issues

Image descriptions: Automatically generate descriptions of image content and related recommendations

Image replies

Product images: Include actual product images in replies

Visual instructions: Provide step-by-step screenshots

Examples:

A user asks where to configure a MaiAgent assistant, and the assistant replies with an image of the settings screen

A user sends a screenshot to confirm whether the path is correct, and the assistant interprets it correctly before replying

The assistant replies with an image

The assistant interprets the user's screenshot

Document Processing Capabilities

Supported document formats

PDF documents: Automatically parse PDF content and answer related questions
Word documents: Process .doc/.docx documents
Excel spreadsheets: Analyze spreadsheet data and provide insights
PowerPoint presentations: Extract presentation content and answer questions

Document interaction features

Document queries: Answer specific questions based on uploaded documents

Data analysis: Automatically analyze tabular data and provide statistics in tables

Document summaries: Generate summaries of key content in long documents

Example 1: Ask the AI assistant to present information in a table and provide links to reference materials

Organize information in a table

Provide external links

Example 2: Ask the assistant to analyze sales content and provide a tabular analysis

Example 3: Upload a docx file and ask the assistant to help analyze issues with CV writing

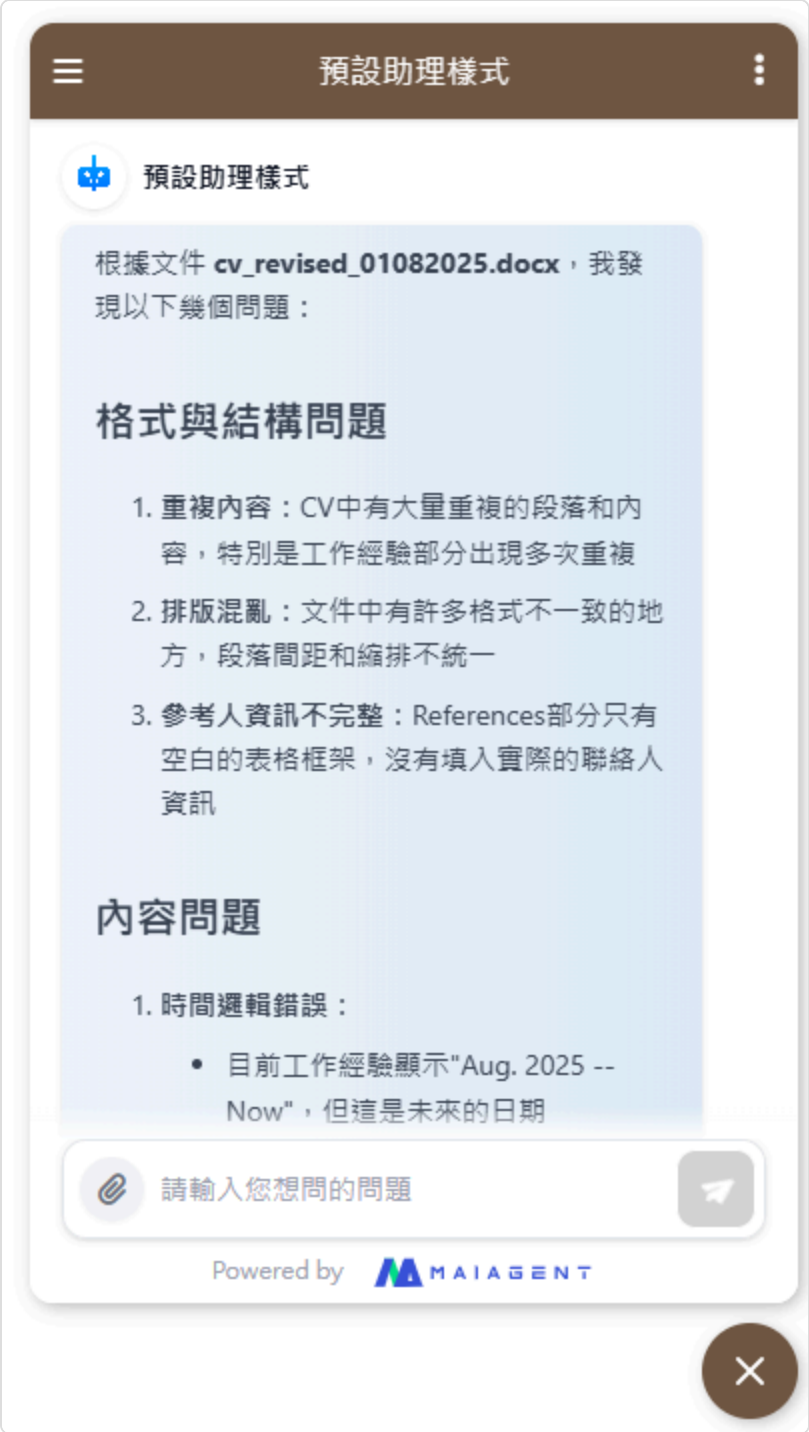

> AI assistants use Markdown for all replies to organize response content clearly.

5. Conversation Memory and Sharing

1. Conversation History Memory

To provide a continuous service experience, WebChat automatically remembers users' conversation history so each conversation can build on previous ones.

What is remembered

Conversation content: Complete question-and-answer records, including text, images, documents, and more

Service status: Unfinished queries and pending issues

Users can view past conversation records in the conversation list on the left of Web Chat, or start a new conversation

Conversation memory can be used for:

Follow-up consultations: A user asks “What was the tent model you recommended last time?” and AI immediately recalls it and replies

Progress tracking: A user says “Continue my previous order inquiry,” and AI automatically retrieves the relevant information

Personalized service: Provide more accurate recommendations and suggestions based on past conversations

2. Conversation Sharing

If you want to share a conversation with friends or colleagues after consulting an AI assistant, or your internal team needs to collaborate on customer issues, use MaiAgent to generate a conversation sharing link:

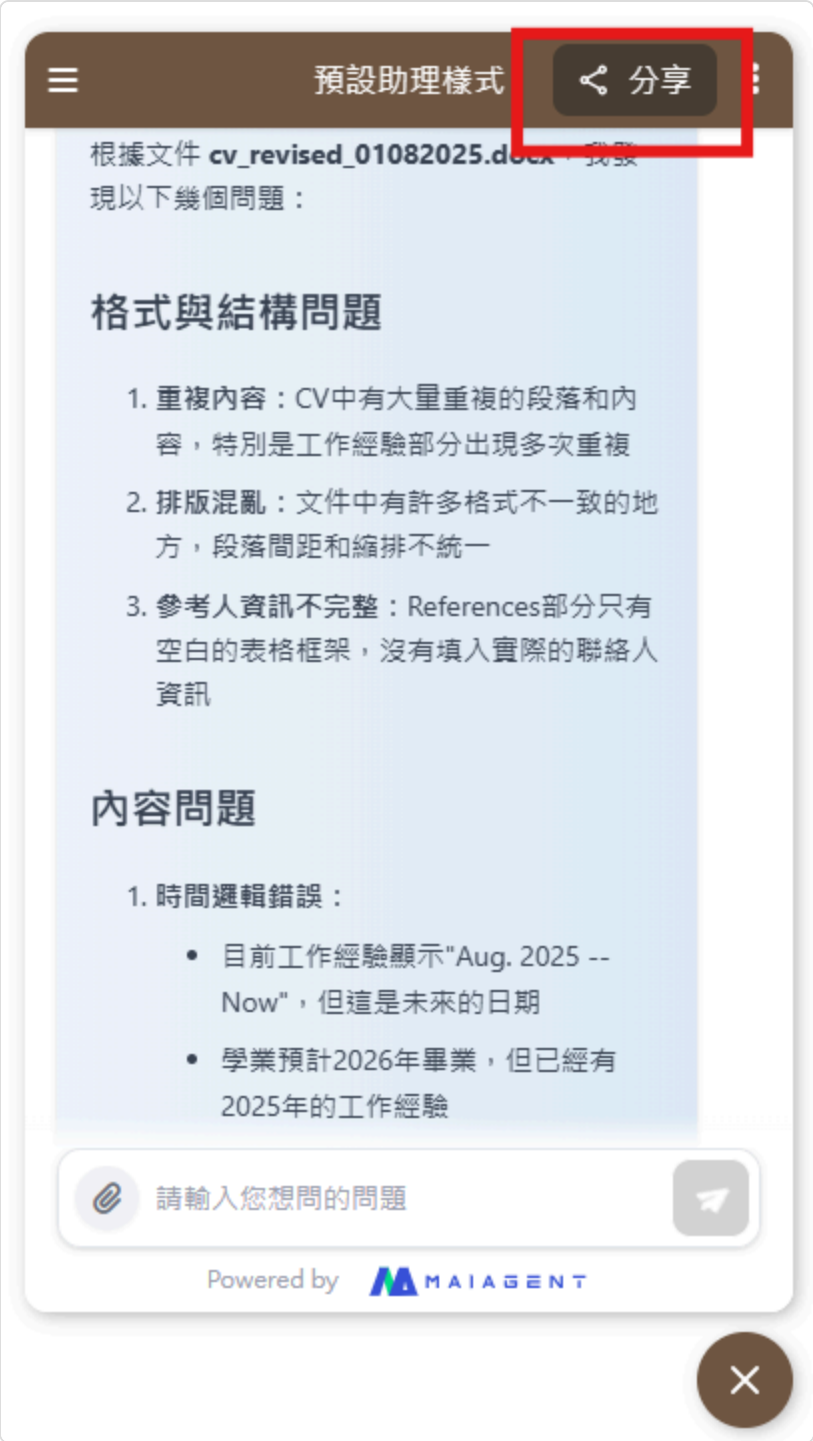

Click Share in the upper-right corner of the chat window, then copy the public link in the pop-up and send it to others.

Shared content

Complete conversation records: Include all text, images, documents, and other content

Conversation summaries: Automatically generate summaries of key points

Related resources: Product links, documents, and other resources mentioned in the conversation

Conversation sharing can be used for:

Customer sharing: “This camping gear recommendation is great. I'll share it with my hiking group”

Internal collaboration: Customer support representatives share complex issues with the technical support team

Training: Share typical conversation examples as employee training materials

3. Conversation Timeout Management

To ensure conversation quality and system performance, Web Chat includes intelligent conversation timeout management:

Intelligent filtering

Automatically filter greetings: The system identifies and filters opening greetings so they are excluded from timeout calculations

Correctly reset message counts: The message counter resets correctly when a conversation restarts

Improved timeout detection: Only substantive conversation content counts toward timeouts, providing more accurate conversation management

Conversation continuity

Ensure users do not encounter unexpected timeouts during normal conversations

Allow conversations to continue after long pauses or time spent thinking

The system intelligently determines whether a conversation is still active

How it works

Start a conversation: A user sends a greeting (such as “Hello”), which the system excludes from timeout calculations

Enter the conversation: The user begins asking substantive questions, and the timeout timer starts

Pause the conversation: If the user temporarily leaves, the system preserves the conversation state

Resume: The user returns to continue the conversation, and the message count updates correctly

This improvement makes conversations smoother and more natural, so users do not need to worry about interruptions caused by spending too long thinking.

6. Choose Whether to Display Tools Used and Cited Documents

1. Display the Tools Used

You can choose whether users can see the list of tools the AI assistant uses during reasoning. If enabled, the following screen appears while the AI assistant generates a response:

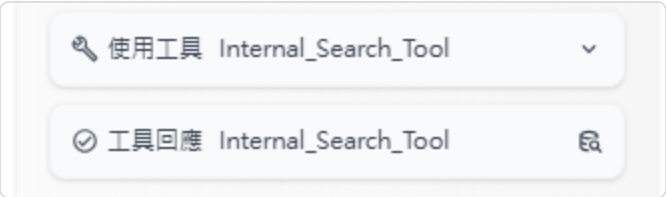

Use this list to check whether the AI assistant is using the tools correctly.

2. Display Cited Documents

Each AI answer automatically generates a list of cited documents, clearly identifying information sources to improve credibility and traceability.

Display settings

Businesses can choose whether users can see the list of cited tools and reference document names. If enabled, the following list appears at the bottom of the AI assistant's reply. Click “Citation Node” to view the referenced document excerpt:

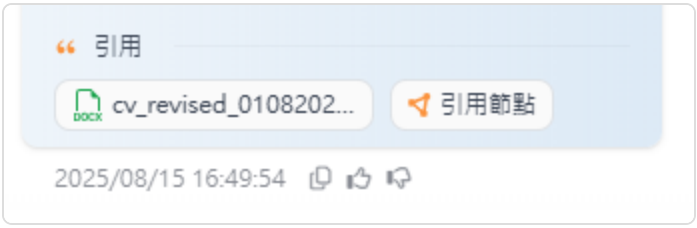

Choose whether to disclose citation information according to the use case:

Public customer support: Display all citations to build trust

Internal support: Hide sensitive documents to protect confidential information

Educational institutions: Display learning resources for students to consult

7. Conversation Access Control — A Secure, Controlled Conversation Environment

1. Login Settings

Require login

To keep conversation content secure and controlled, businesses can enable login and require users to sign in before using AI assistant services. Once configured, visitors who click WebChat are first taken to a

login page. Only users who successfully log in can start a conversation with the AI assistant

Security benefits

Authentication: Ensure each conversation user has an identifiable identity

Access control: Provide services according to the user's identity

Conversation tracking: Fully record each user's conversation history

Data protection: Prevent unauthorized access to sensitive information

2. Three Login Methods

AD (Active Directory) login

Use case: Internal employees, integrating with an existing AD system

Features: Single sign-on (SSO) without additional account management

Setup: Configure AD server connection details and synchronization rules

Practical application: Employees sign in directly with their corporate accounts and automatically receive the corresponding permissions

Keycloak login

Use case: Businesses that need complex identity management, with support for multiple identity providers

Features: Support standard protocols including OAuth2, SAML, and OpenID Connect

Setup: Configure the Keycloak server, user pool, and identity providers

Practical application: Integrate user identities across multiple systems and manage permissions centrally

Extension: Let the AI assistant use the visitor's Keycloak login identity when calling MCP/API tools. See Use the Keycloak Login Identity for Tool Calls

MaiAgent login

Use case: Simple user management requirements and fast deployment

Features: Built-in user management with straightforward, intuitive operation

Setup: Create user accounts and permissions in MaiAgent

Practical application: Create a dedicated login system for a specific group of users

Once enabled, the following screen requires users to log in before starting a conversation:

3. Use the Keycloak Login Identity for Tool Calls

After a visitor logs in to Web Chat through Keycloak, the AI assistant can automatically include **the visitor's own Keycloak login credential** when calling MCP tools or API tools on their behalf. When your backend receives the request, it can use the same Keycloak system to verify who the visitor is and return only the data they are allowed to view. You do not need to configure separate accounts for every tool or ask visitors to log in again.

This feature is off by default. When it is off, all tool calls work exactly as before.

Use Case

A company's internal service portal uses Keycloak to manage employee accounts and Web Chat as an employee assistant. An IT administrator enables this feature on the conversation platform and adds the host of the company's leave management system to the allowlist. After logging in to Web Chat, an employee asks “How many days of annual leave do I have left?” The AI assistant calls the leave system's API tool with that employee's own login credential in the request. The leave system identifies the employee from the credential and returns only their attendance and leave data. When another employee asks, they receive their own data.

Prerequisites

Your organization has completed Keycloak integration. Visitors must be members of your organization. When they first log in through your organization's configured Keycloak, the system automatically creates an account and adds them to the organization (see Third-Party Login: First Login).

In the conversation platform's "Login Settings," "Enable Login" is on, and "Login Source" is set to Keycloak.

The MCP/API tools that will receive credentials have been created and assigned to the AI assistant used by this conversation platform. The receiving system can validate credentials issued by your Keycloak Realm.

Setup Steps

Settings location: In the left menu, select "Customer Conversations > Conversation Platforms," choose a website-type conversation platform, and open the "Login Settings" tab. This section appears after you select Keycloak as the login source.

Turn on "Use the Keycloak Login Identity for Tool Calls."

Read the red warning below the switch, "The visitor's own login credential will be forwarded to the following hosts," and confirm that you understand where credentials will be sent.

In "Tool Hosts Allowed to Receive Credentials," enter the tool hosts that should receive credentials, one per line (or separated by commas):

Enter only the hostname, without `https://`, a port number, or a path, such as `mcp.example.com`.

`*.example.com` matches `example.com` itself and all its subdomains. A hostname without `*.` only matches that exact host.

Hostnames are matched against the tool's "MCP Server URL" or "API URL."

Save the settings.

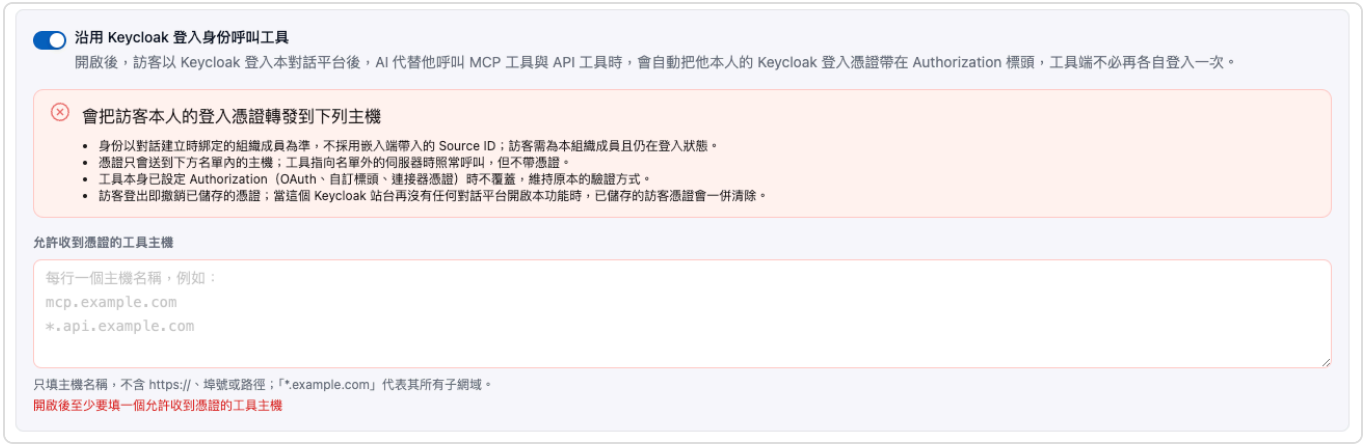

Enabling the feature displays a credential forwarding warning and requires tool hosts allowed to receive credentials

When the switch is on, the list must contain at least one host. If the list is empty or contains an invalid hostname, the screen displays "Enter at least one tool host allowed to receive credentials when enabled" or "The following hostnames are invalid," and the settings are not submitted.

After enabling this feature, visitors must log in again. The platform saves Keycloak credentials only when a visitor logs in, and only if a conversation platform has this feature enabled. Ask visitors who were already logged in before it was enabled to log out and log back in. Subsequent tool calls will then include their credentials.

What the Receiving System Gets

When the AI assistant calls an MCP/API tool on a listed host, the request includes an additional HTTP header:

```
Authorization: Bearer <訪客的 Keycloak access token>
```

This is the standard access token issued to the visitor by your Keycloak Realm. The engineer handling the integration can identify the user through standard Keycloak token validation (verifying the signature, expiration, and Realm). When the credential expires, the platform first renews it with the same Keycloak using the visitor's login session, then sends the request. The tool does not need to handle renewal.

When Credentials Are Included and When They Are Not

Scenario	Result
The visitor has logged in with Keycloak, is still logged in, and the tool host is on the list	Includes the visitor's own credentials
The tool host is not on the list	Calls the tool as usual, without credentials
The tool already has `Authorization` configured (OAuth, API tool headers, contact MCP credentials, etc.)	Does not overwrite it; retains the tool's original authentication method
The visitor has logged out or is not a member of this organization	Does not include credentials
The visitor enters the conversation through Source ID access without login	Does not include credentials; identity is based on the organization member bound when the conversation was created, rather than the Source ID supplied by the embedding site
The toggle is off	Behaves as it did before this feature was enabled

The visitor's complete login credentials are forwarded. A system that receives them can access other services protected by the same Keycloak instance as that visitor. Only list hosts you trust that actually need to identify users, and avoid overly broad wildcards.

> When a visitor logs out, the platform revokes their stored credentials. When no conversation channel on the same Keycloak site has this feature enabled (because the toggle was turned off, a different login provider was selected, or the channel was deleted), the platform also clears all stored visitor credentials.

8. Data Collection (Pre-chat Form)

When data collection is enabled, visitors see a form before starting a conversation and must complete it to enter the chat. Their entries are saved to the visitor's contact, so support agents can see them directly on the conversation page. This is useful when you need to know who the visitor is first, such as collecting a company name and business registration number for dealer support, or a name and Email for event inquiries.

1. Configure the Data Collection Form

Path: Customer Conversations → Conversation Channels → Select the Web Chat channel → Data Collection tab.

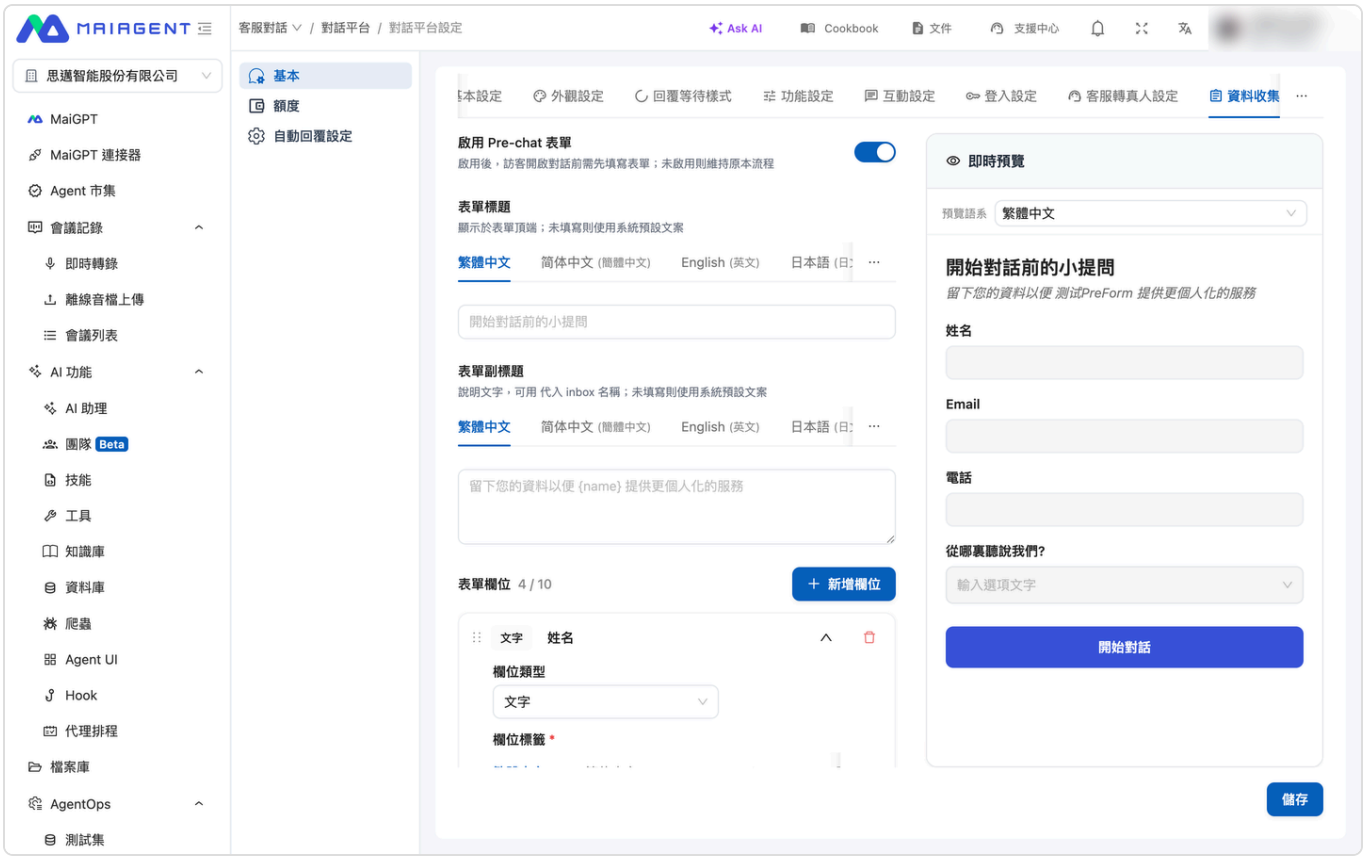

The Data Collection tab: configure the toggle, title, and fields on the left; Live Preview on the right shows the form as you edit it

Enable the Form

Turn on Enable Pre-chat Form. When disabled, visitors enter the conversation directly, following the original flow.

Set the Title and Subtitle (Optional)

Form Title: Appears at the top of the form. If left blank, the default text is “A few questions before we start chatting”.

Form Subtitle: Explanatory text for the form. Use `{name}` to insert the conversation channel name. If left blank, the default text is “Leave your details so {name} can provide more personalized service”.

Both support multiple languages. Switch language tabs to enter each version separately.

Add Fields

Click Add Field. Each field has these settings:

Setting	Description
Field Type	Four types: text, Email, phone, and dropdown
Field Label	The field name visitors see; can be set separately for each language
Placeholder	Gray hint text inside the field
Required	When selected, visitors cannot submit without completing the field
Mapped Contact Field	Determines where the data is saved in the contact: Name (Contact.name), Email (Contact.email), Phone (Contact.phoneNumber), or Custom Field (Metadata)

Dropdown fields also have an Options list. Select Allow Custom Input for “Other” to let visitors enter their own text after selecting “Other”.

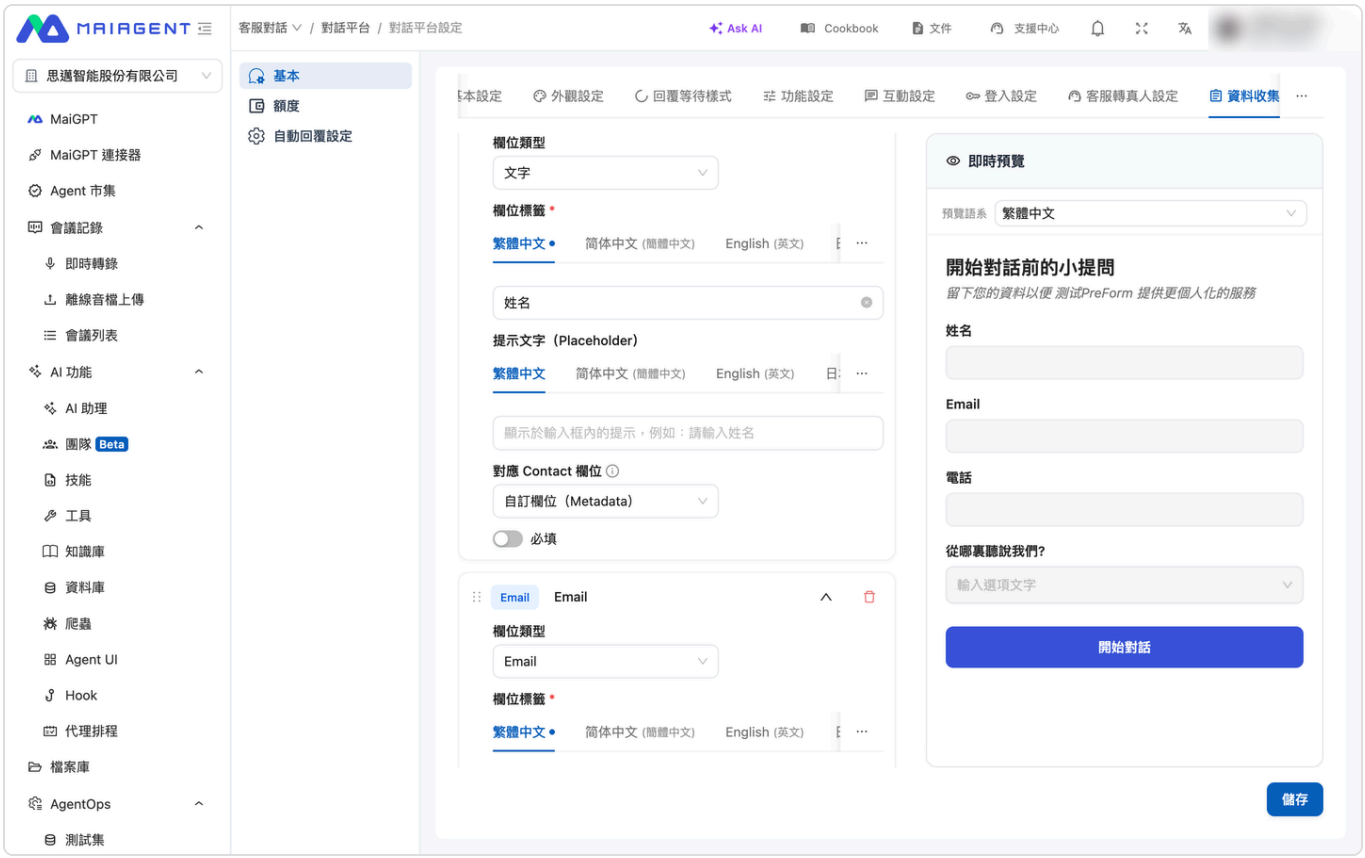

Configure each field's type, label (multilingual), placeholder, mapped Contact field, and required status

Drag fields to reorder them. A form can contain up to 10 fields. The preview on the right updates immediately; switch the Preview Language to check each language version.

Save

After checking the preview, save your settings. They take effect immediately in Web Chat.

> Contact field mapping restrictions Name can only map to a Text field, Email to an Email field, and Phone to a Phone field. Incompatible types display a disabled-state hint. Name, Email, and Phone can each be mapped by only one form field. For other fields (such as business registration number or company name), select Custom Field (Metadata).

2. Visitor Form Flow

When visitors open Web Chat, they first see the form: its title, subtitle, configured fields, and the language switcher in the upper-right corner. After completing it, they click Start Conversation to enter the chat.

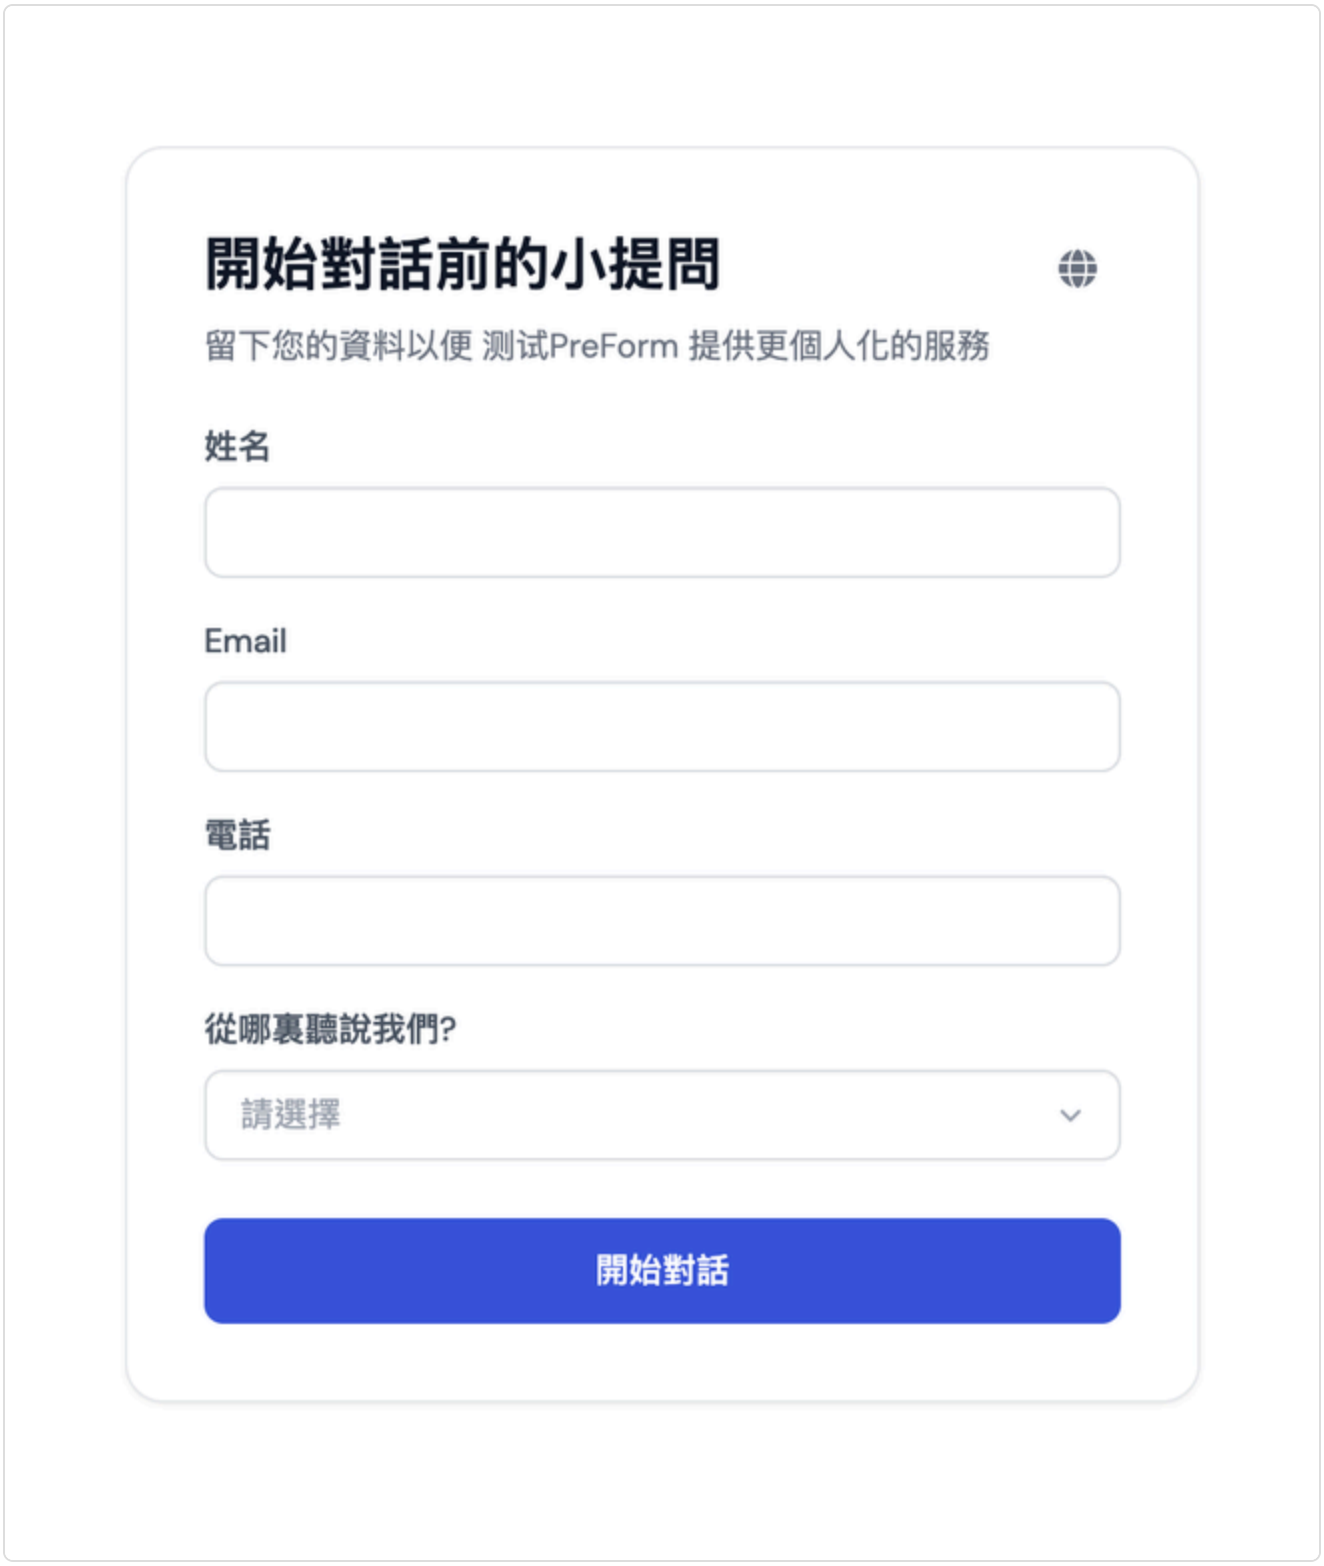

Visitors see the form first when opening Web Chat

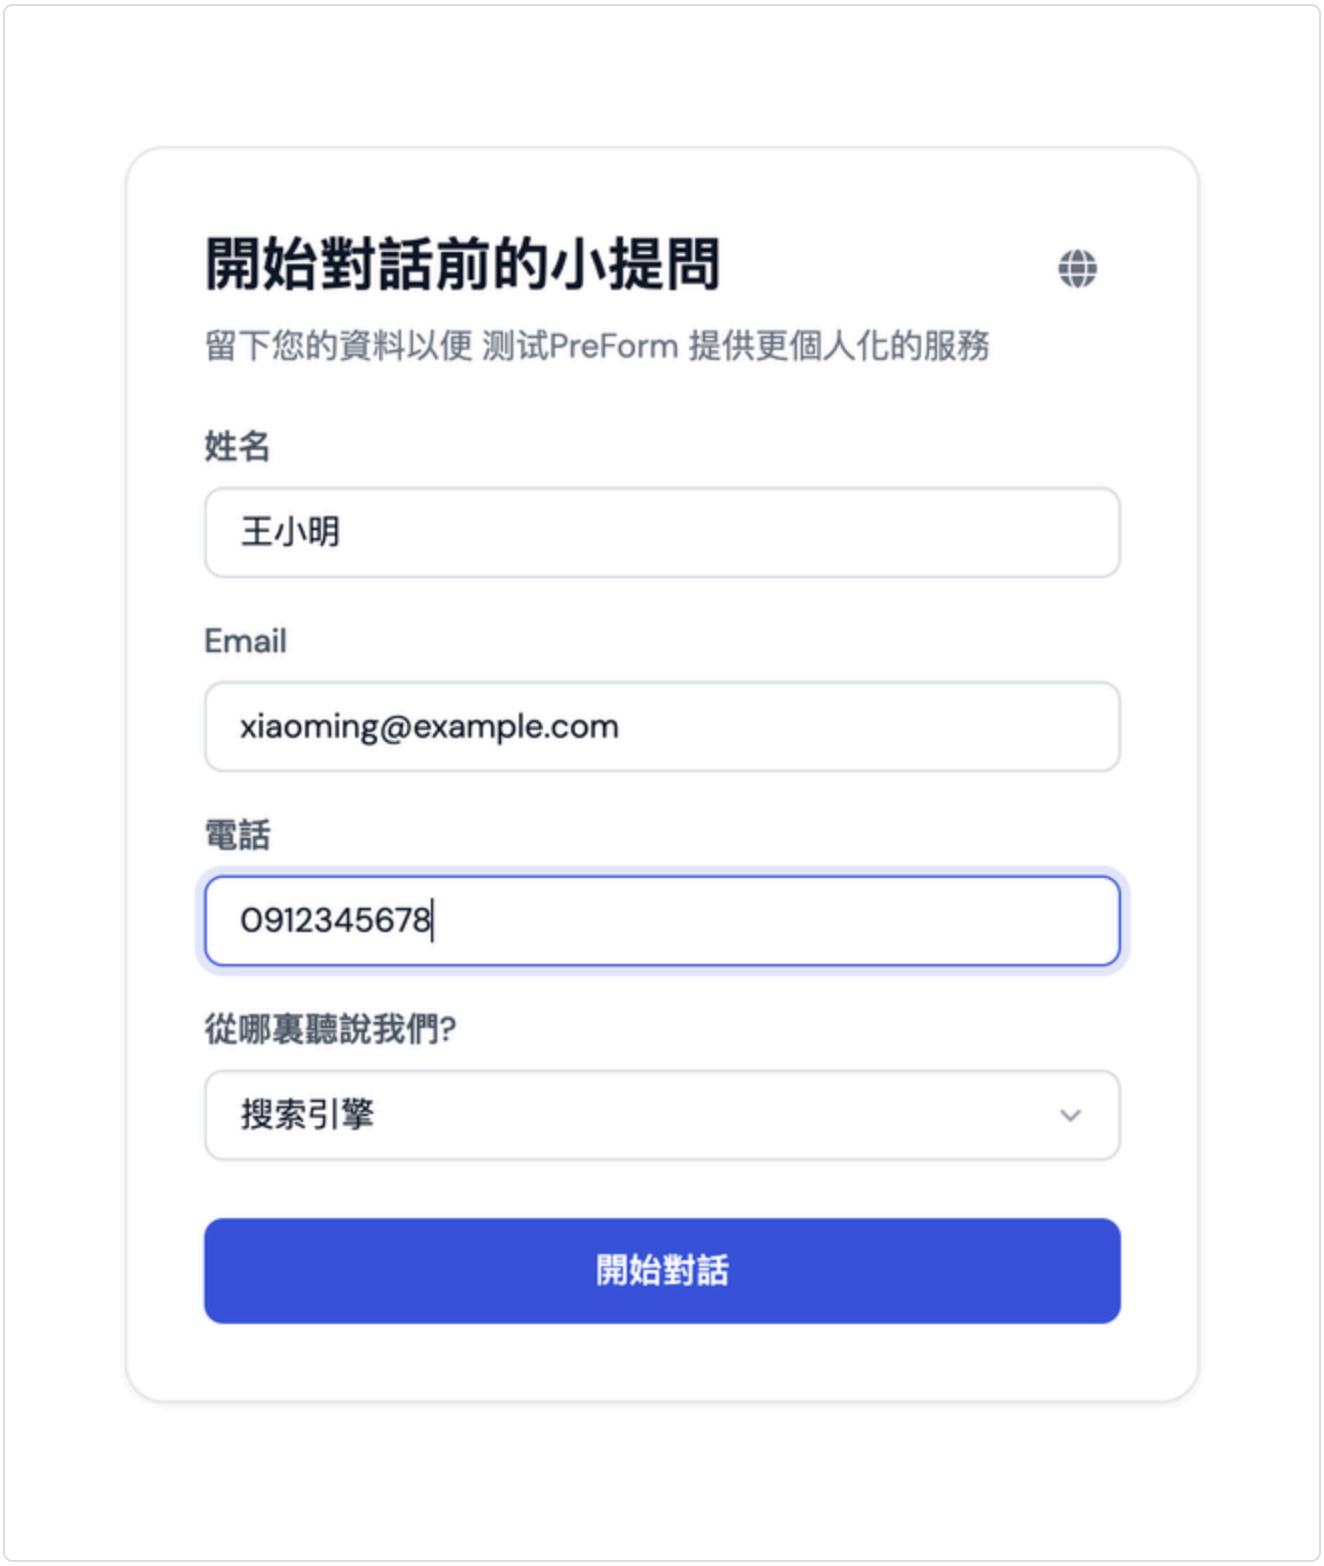

Click Start Conversation after completing the form

If a required field is empty or an Email or phone number has an invalid format, an error appears below the field and submission is blocked.

Once a visitor completes the form in a browser, they will not be asked to complete it again when reopening the same Web Chat in that browser.

To correct their information, visitors can select Fill In Your Details Again from the ⋮ menu in the upper-right corner of Web Chat. The form appears again.

Fill In Your Details Again in the upper-right ⋮ menu

3. Use with Login: Map Data to the Login Account

If the same Web Chat has both login and data collection enabled, the flow is: **Log in → Complete the form → Start the conversation**. The submitted data is linked to the user's login account rather than the browser.

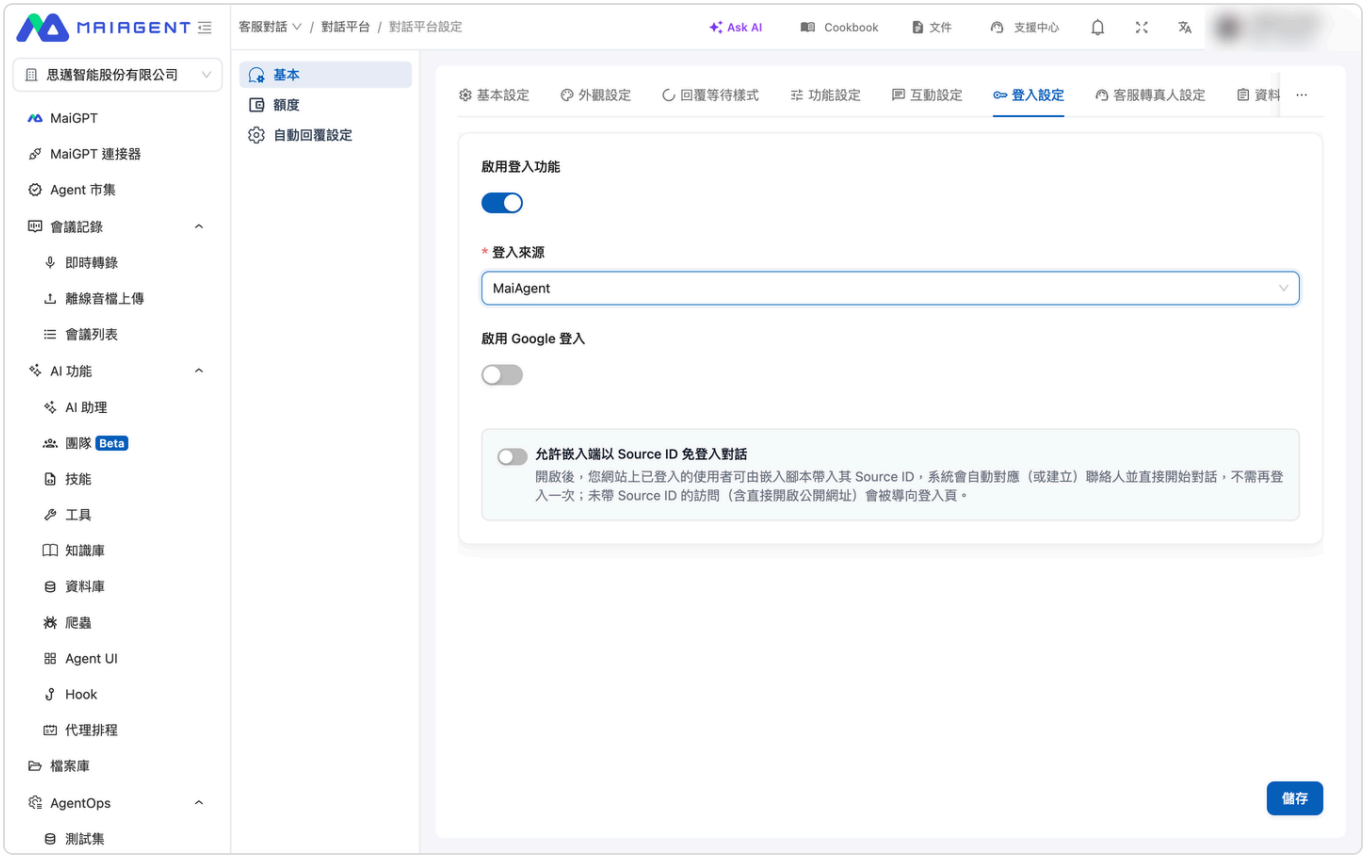

The Login Settings tab: turn on Enable Login and select a login provider; this tab is independent of Data Collection

① Visitors first see the login screen when opening Web Chat

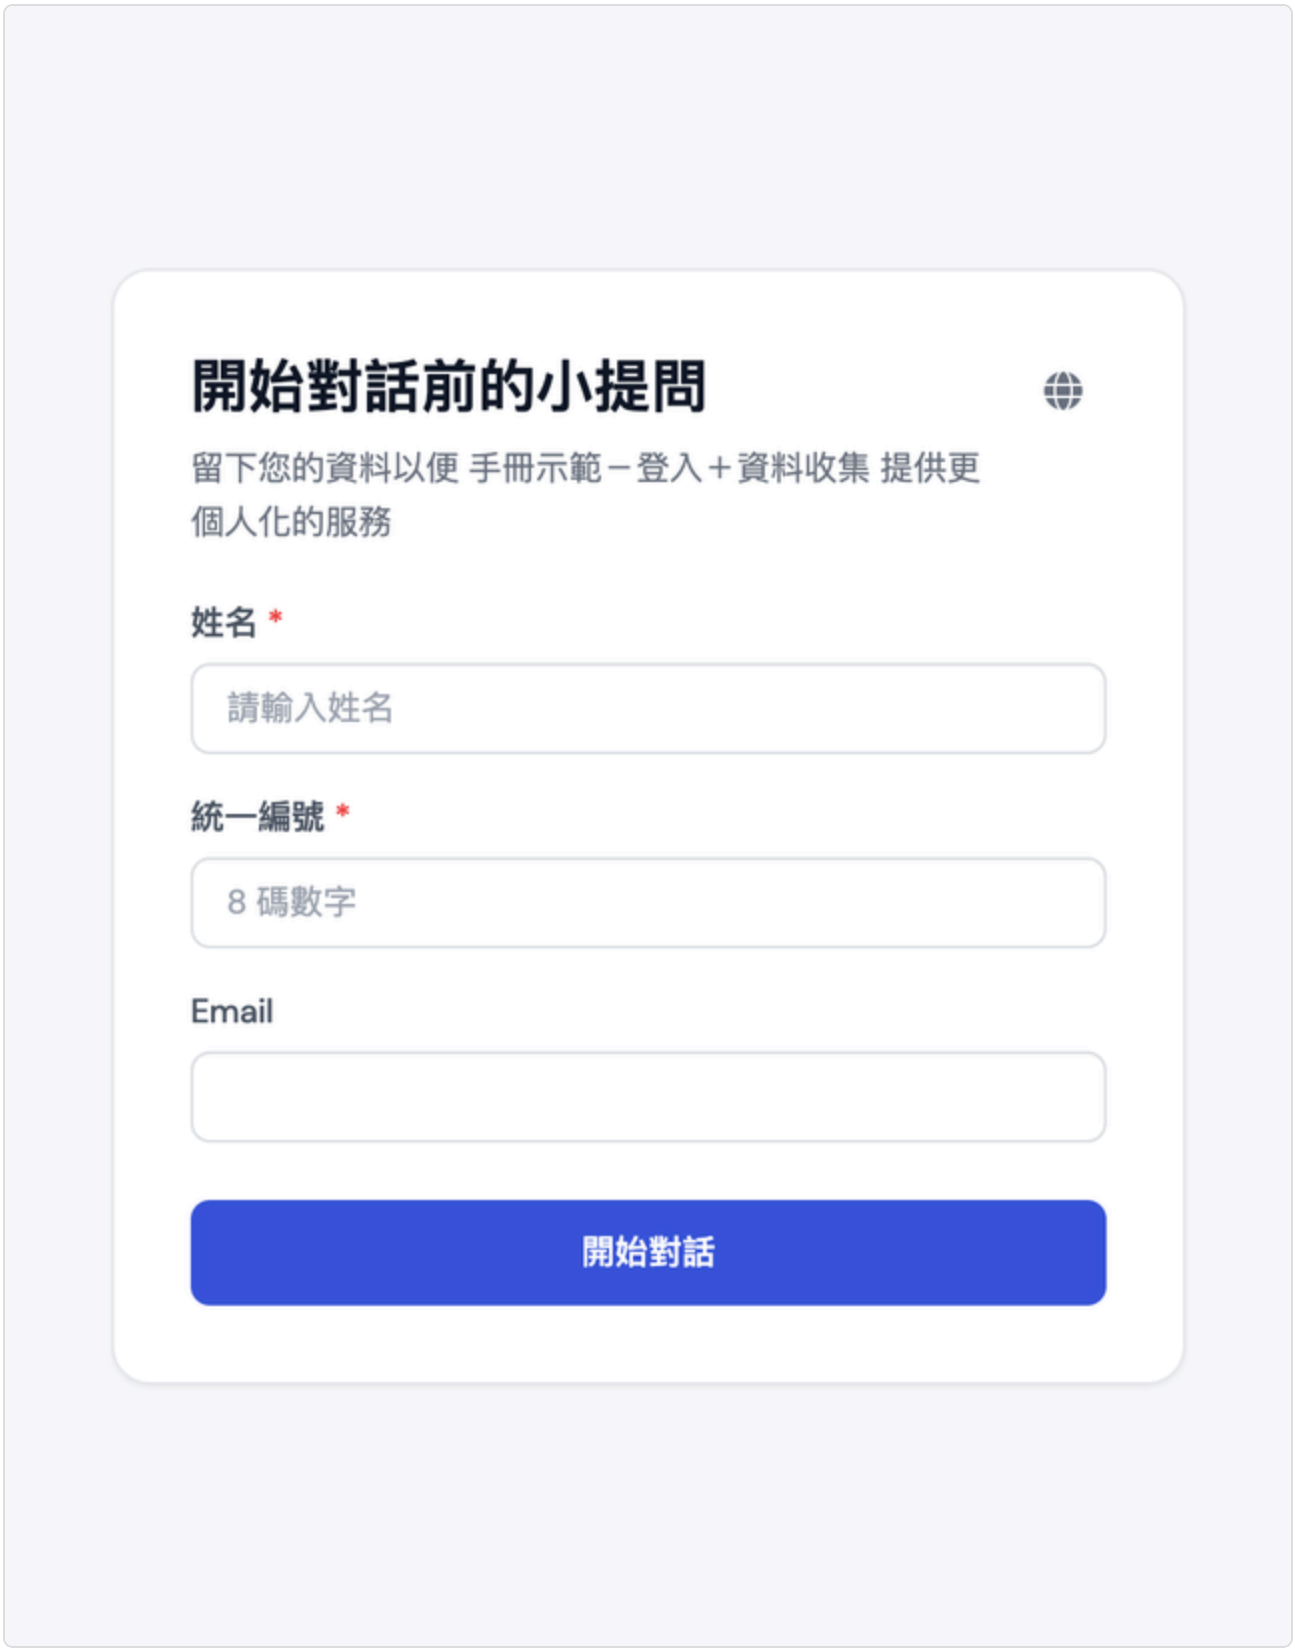

② The data collection form appears only after login

Scenario	Web Chat Behavior
The account is used for the first time and has not completed the form	Shows the form after login; creates a contact linked to the account after submission
The same account logs in on a different device or browser	Skips the form and starts the conversation with the account's existing data
Someone previously completed the form on this device without logging in, and the account itself has no data	Shows an “Is this your information?” confirmation screen (see below)
The account already has data, and this device also has data submitted while logged out	Uses the account's data without asking or merging
A logged-in user clicks Fill In Your Details Again	Updates the account's existing contact without creating another one

Confirmation Screen: Is This Your Information?

If someone previously completed the form on this device without logging in, and the currently logged-in account has no data, the form does not appear immediately after login. Instead, a confirmation screen first shows the previously entered information from this device (partially masked), with two buttons:

Yes, This Is Mine: Links that information to the current account. The same account will not need to complete the form again on any device.

No: Does not merge the data and leaves the original record unchanged, then shows the form for the user to complete once.

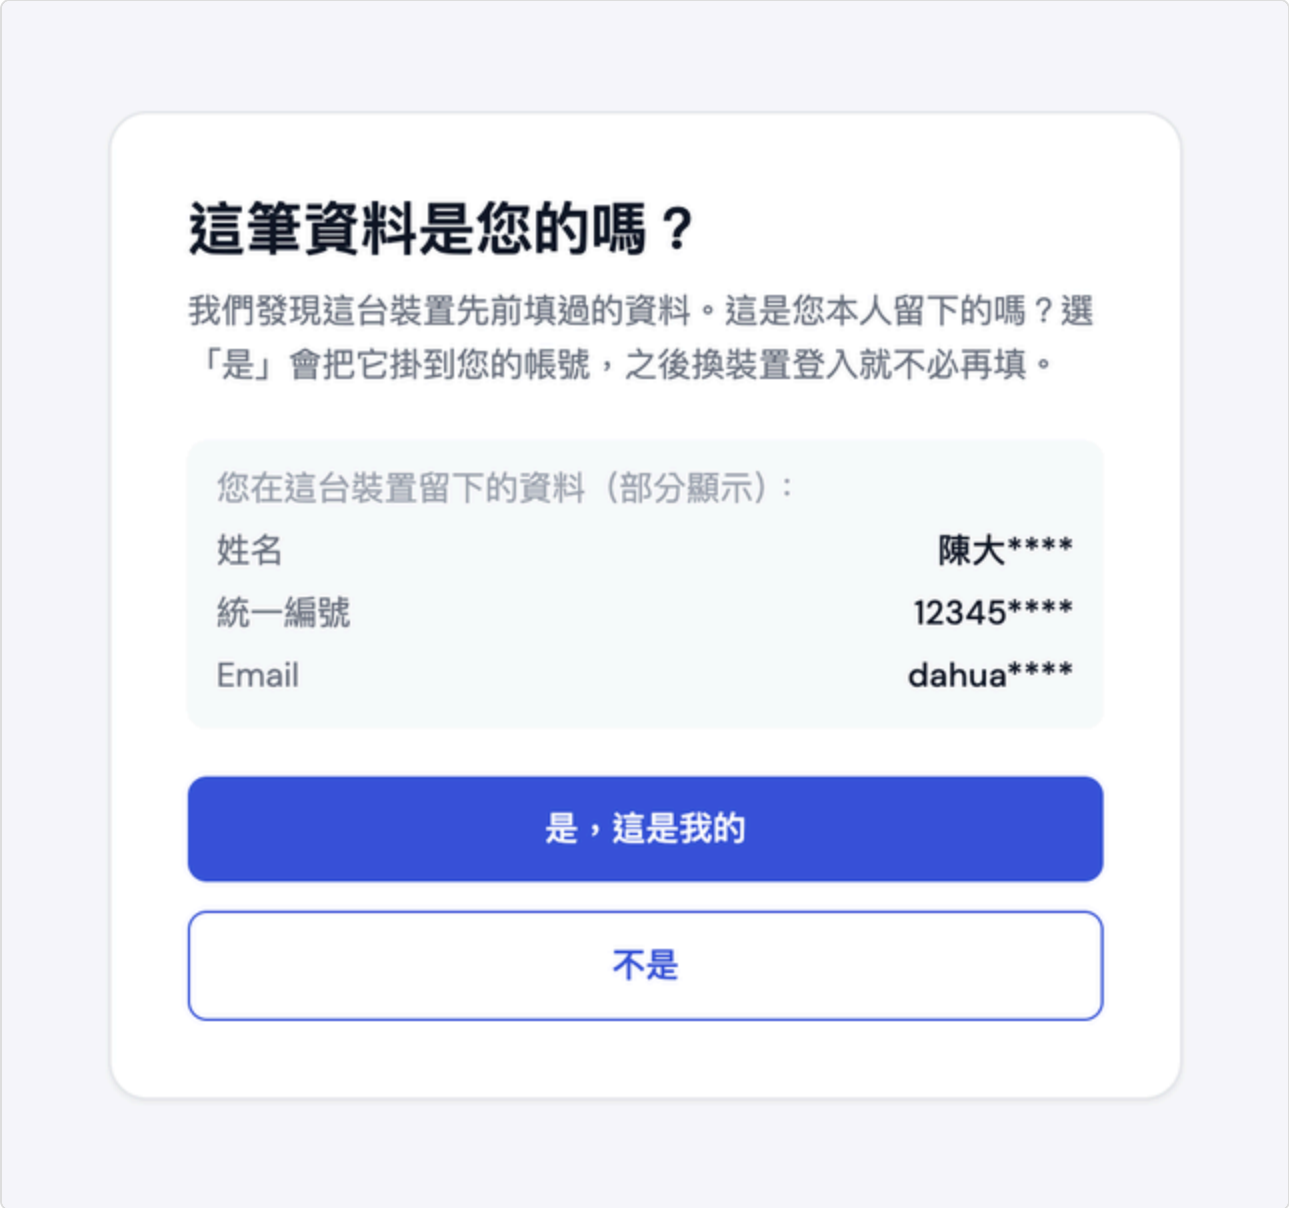

The confirmation screen partially masks the information so users can identify whether they provided it

> Web Chat channels without login enabled are unaffected; they still use the browser to determine whether the form has been completed. The system does not automatically fill the form with account data. The first time an account is used, the user must complete the form even if their data already exists in the admin console.

After entering the conversation, both the AI assistant and the support console use this contact linked to the account:

The AI assistant addresses the user by the name entered in the form

The admin conversation page shows the contact and the business registration number collected by the form

4. View Collected Data in the Admin Console

View in a conversation: Customer Conversations → All Conversations → Open a conversation. The Contact Information section on the right lists each item collected by the form. Fields mapped to Email and Phone appear as “Email” and “Phone Number”; custom fields use the labels you configured.

Contact Information on the right of the conversation page lists each item collected by the form

> The Contact column in the contact list shows the Contact's name. If the form's name field maps to Custom Field (Metadata) instead of Name (Contact.name), contacts for visitors who are not logged in appear as “Anonymous”, and their names only appear in custom attributes. To show visitor names directly in the list, map the name field to Name (Contact.name).

View on the contact page: Customer Conversations → Contacts → Open a contact. The Basic Information tab shows the name and phone number, and the Custom Attributes tab shows other form fields.

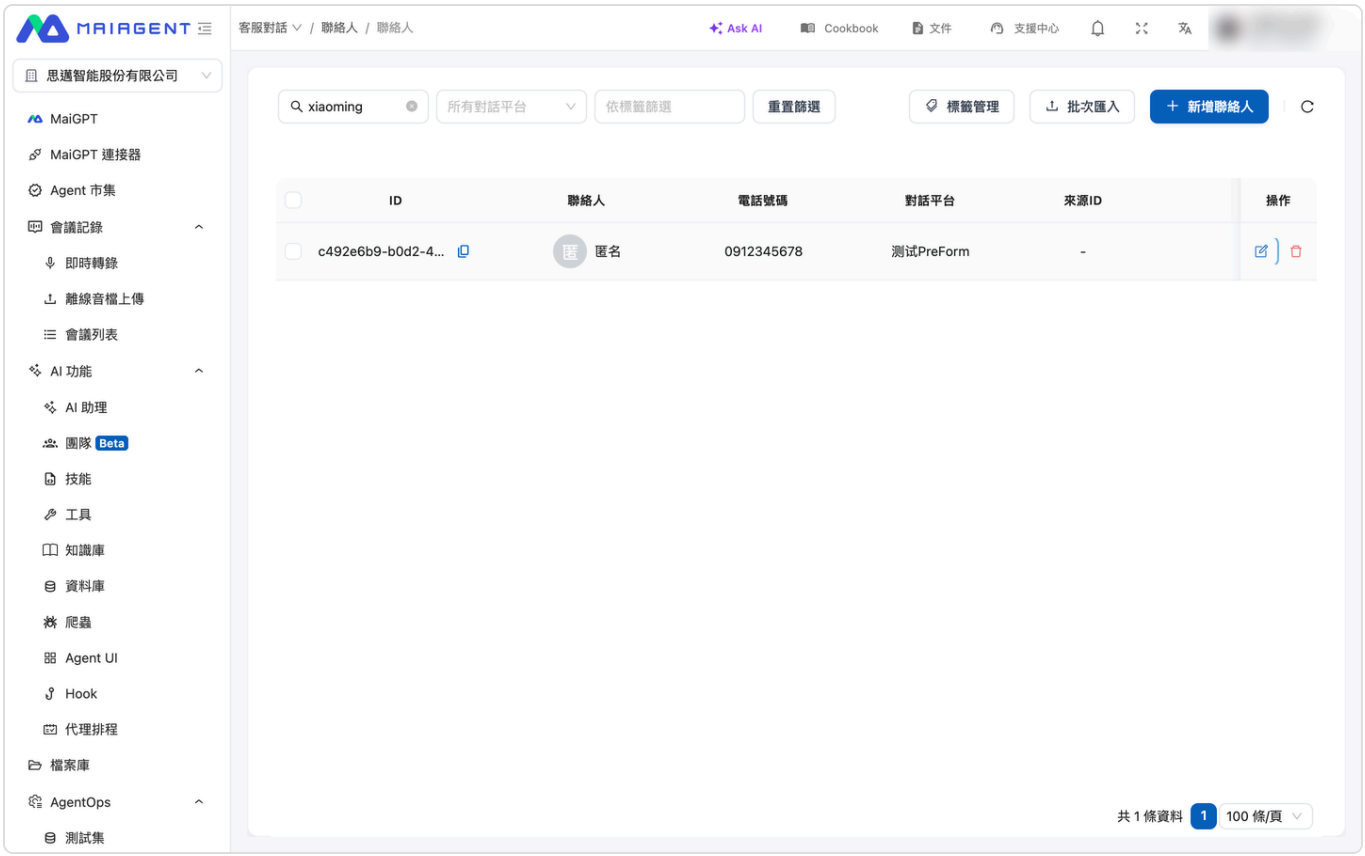

Search the contact list by name, Email, or phone number

The Custom Attributes tab lists the value of each form field

Data submitted by logged-in users is linked to their accounts, so the contact is no longer an unassigned anonymous contact. The contact name depends on the form settings: if a name field maps to Name (Contact.name), the name entered in the form is shown; if there is no name field, the login account's name is shown.

What to Do When a Connection Fails

For example, a visitor reports to a website administrator that Web Chat shows a connection error when opened. The visitor first waits for the system to retry automatically. If the error screen remains, they capture the full screen, including the failed stage, error category, and time, and send it to the administrator to help identify the issue faster.

If a network interruption occurs while Web Chat loads the AI assistant settings, it first retries automatically and keeps displaying the loading screen during that time. If the connection still fails after retrying, the error screen shows:

The loading stage that failed

The error category, along with the HTTP status code if available

The local time when the error occurred

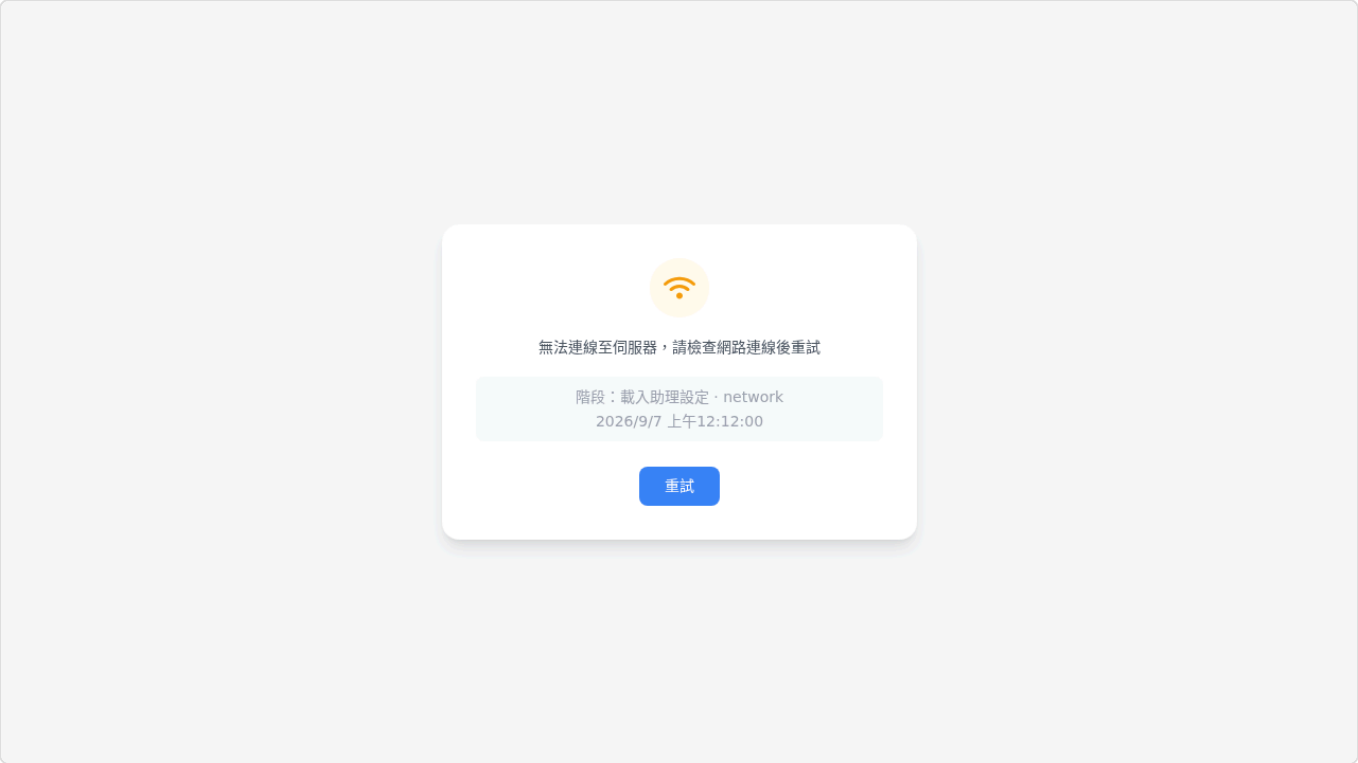

After automatic retries fail, the screen lists the failed stage, error category, and time

If the system has received an HTTP error status code, the server has responded. In this case, it displays the diagnostic error screen immediately without automatically retrying.

> First confirm that the device can access the internet, then click Retry on the error screen. If the problem persists, include all diagnostic information in the screenshot and avoid capturing conversation content or personal data.

Review Earlier Messages While a Reply Is Being Generated

While the AI assistant is still generating a reply, visitors can scroll up to review earlier messages. The reply continues to appear below, but the screen stays where the visitor scrolled instead of being pulled down by new text. Once they move away from the latest content, Back to Latest appears above the input box.

Usage Scenario

A visitor asks the AI assistant to compare several plans at once. Before the answer is finished, she scrolls up to check the first plan's price. The screen stays in place even as more text appears below. After checking, she clicks Back to Latest. The screen returns to her last question so she can continue reading the rest of the comparison.

Steps

1. Scroll Up

While the AI assistant generates a reply, scroll up using the mouse wheel, touch gestures, or scrollbar.

2. Read in Place

The screen stays where you scrolled. Newly generated content does not change your current position.

The reply is still being generated (the square stop button is on the right of the input box); scrolling up shows Back to Latest above the input box

3. Back to Latest

Click Back to Latest to return to the last question and resume automatically following the latest content.

Connect a Messaging Platform: Website

1. Open the Settings Page

1. Option One: Add a Website Support Conversation

In the left menu, click “Customer Conversations > Conversation Channels”, click “＋Connect Conversation Channel” in the upper-right corner, and select “Website”.

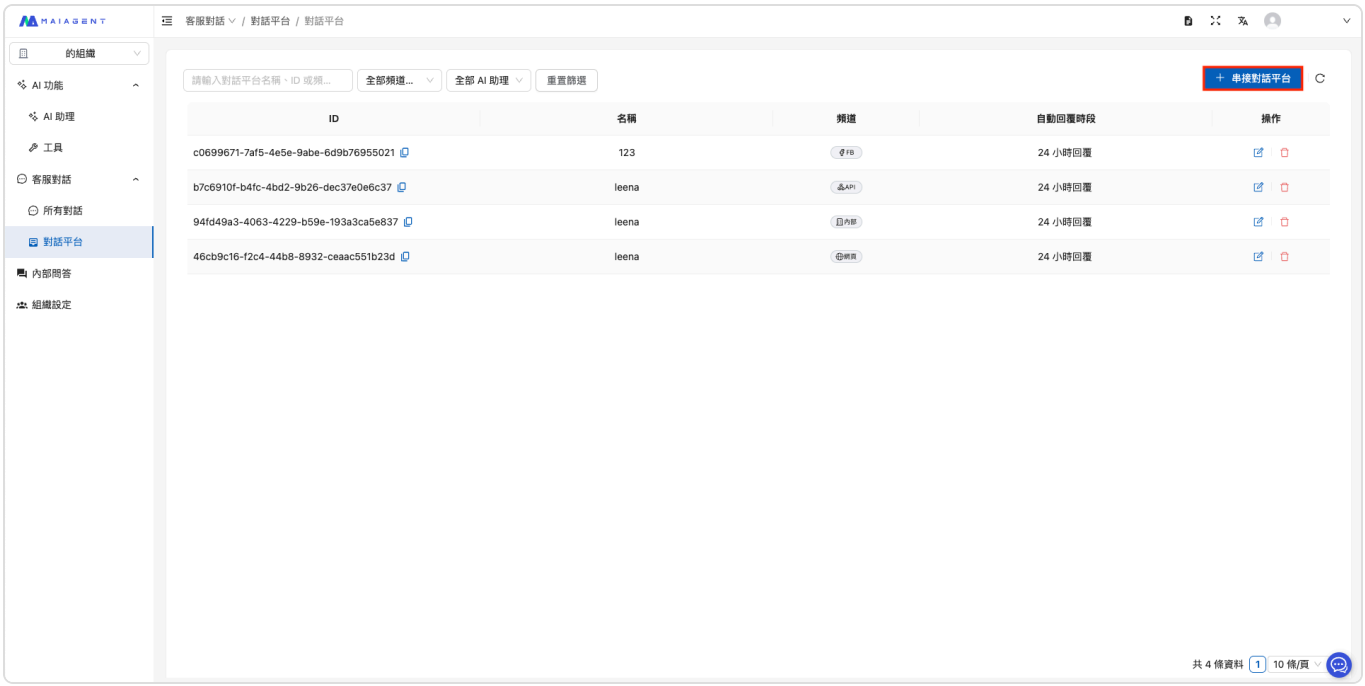

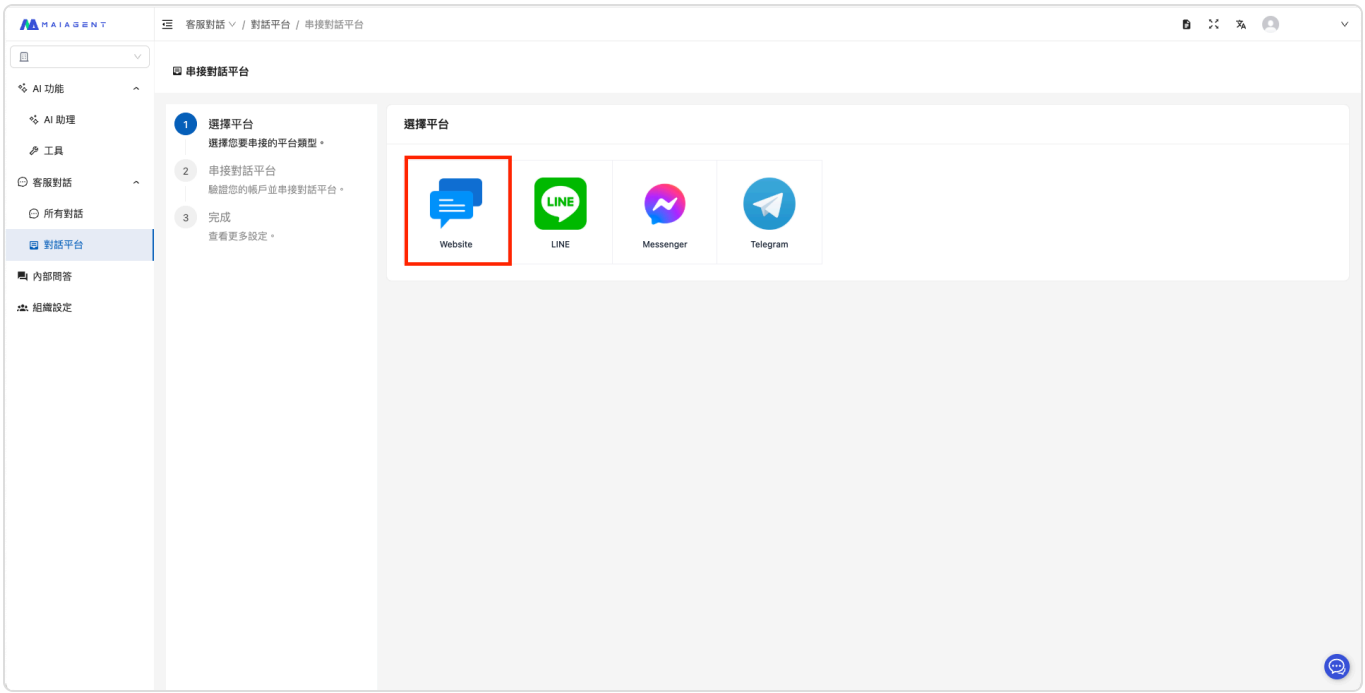

2. Option Two: Link a Conversation Channel to an AI Assistant

In the left sidebar, go to AI Features → AI Assistants, select the assistant you want to configure, and switch to the Conversation Channels tab. Find the website conversation channel you want to adjust and click Channel Settings.

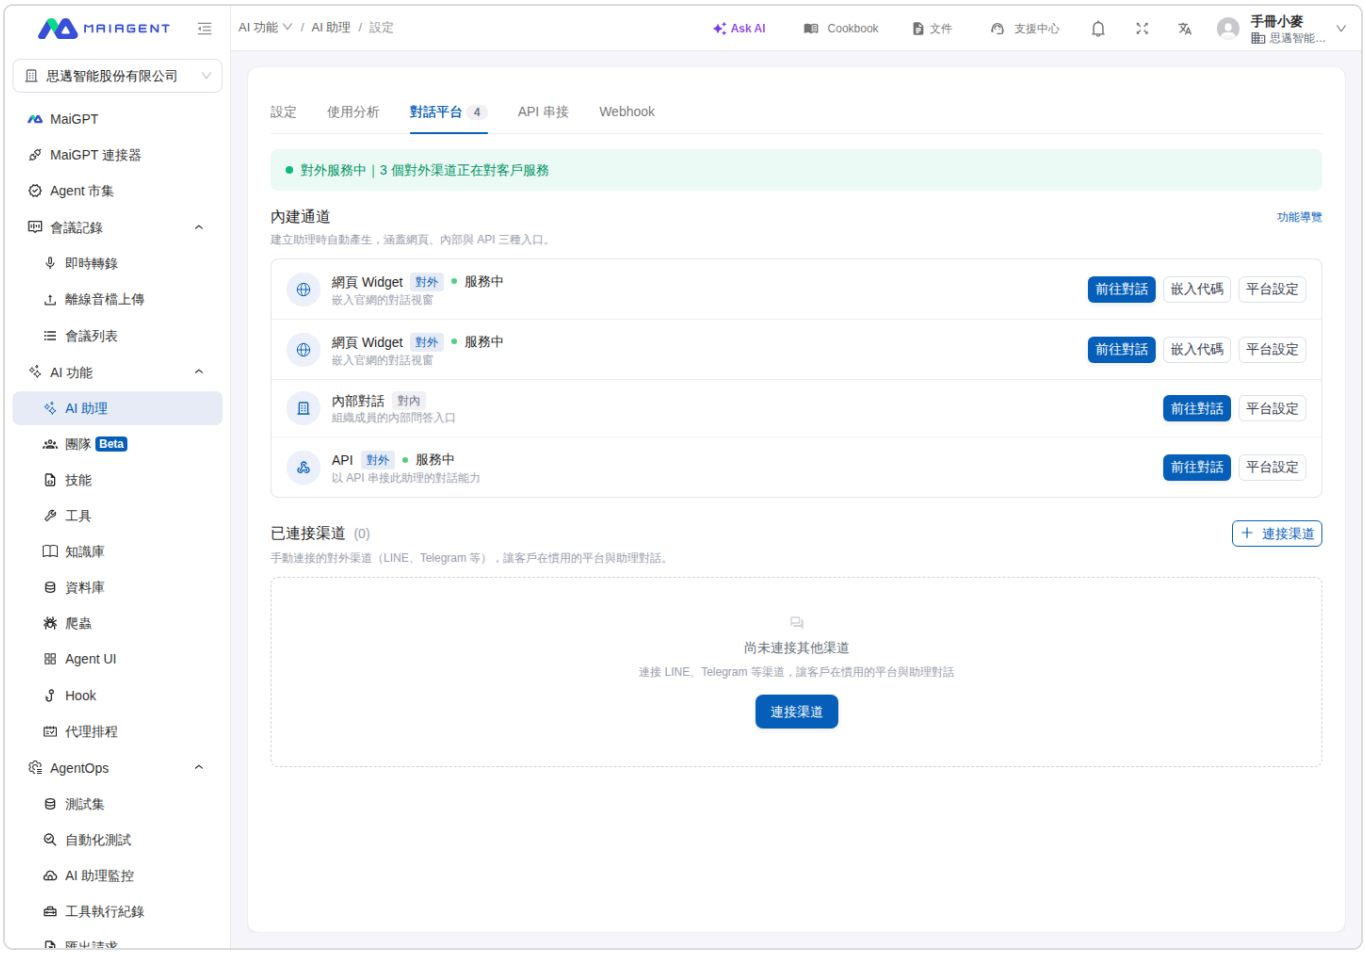

Click Channel Settings on the AI assistant's Conversation Channels tab.

After configuring the channel, click Back to AI Assistant Settings at the top of the page to return to the original assistant's Conversation Channels tab. This return link remains available when switching between Basic, Quota, and Auto-reply Settings.

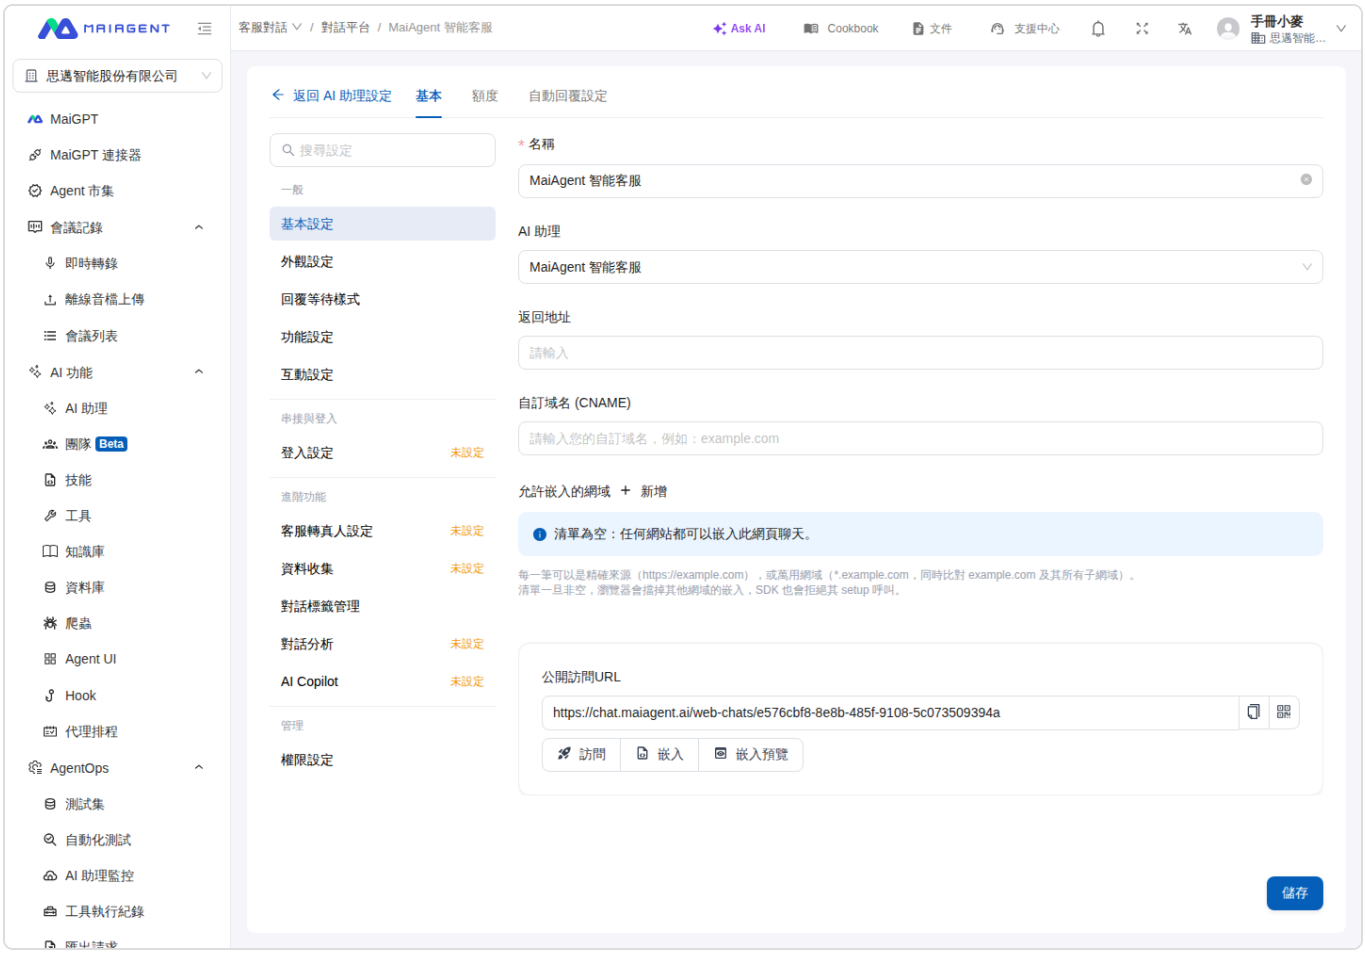

After configuring the channel, return directly to the original AI assistant's settings page.

> Back to AI Assistant Settings appears only when you enter from the AI assistant's Conversation Channels tab. If you enter directly through Conversation Channels in the left menu, the page retains its original navigation.

Usage Scenario: Return to the Assistant After Adjusting Website Support

A support manager needs to update the website support settings for the Product Inquiry Assistant. She clicks Channel Settings on that AI assistant's Conversation Channels tab. After making changes, she clicks Back to AI Assistant Settings to return to the same assistant's channel list and check other support entry points, without finding the assistant again in the assistant list.

3. Option Three: Conversation Channel Actions

In “Customer Conversations > Conversation Channels” in the left sidebar, select the AI assistant you want to configure and click the “Actions” icon to open the conversation channel configuration.

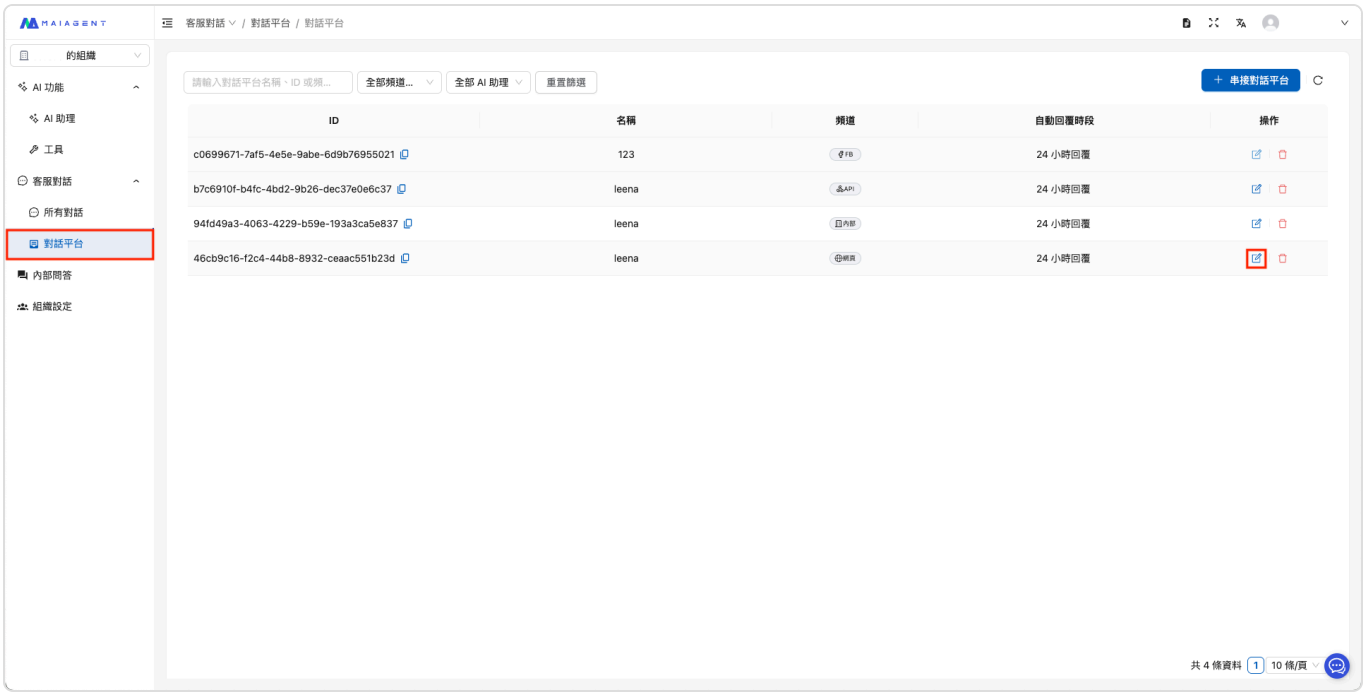

2. Basic Settings

The five categories in the General group on the Basic tab each have a dedicated page with details:

1. General Settings

For the channel name, AI assistant, return URL, custom domain, allowed embedding domains, public access URL, and three embedding methods, see General Settings.

Allowed Embedding Domains

For how the embedding allowlist works, setup steps, and the checklist before enabling it, see General Settings: Allowed Embedding Domains.

2. Appearance Settings

For the logo, avatar, display settings, and light/dark theme colors, see Appearance Settings.

3. Feature Settings

For enabling the channel, allowing file uploads, document citations, OCR, tool display mode, question set option display mode, and AI awareness settings, see Feature Settings.

4. Interaction Settings

For greetings, conversation starter questions, including greeting options in Agent context, conversation display settings, satisfaction surveys, and quota exhaustion messages, see Interaction Settings.

Include Greeting Options in Agent Context

For behavior before and after enabling this option, applicable scenarios, and considerations, see Interaction Settings: Include Greeting Options in Agent Context.

5. Login Settings

Configure the SSO Login Provider

MaiAgent currently supports SSO (Single Sign-On) through Keycloak integration.

Switch providers using the "Login Provider" dropdown.

Select Maiagent or Keycloak as the authentication method.

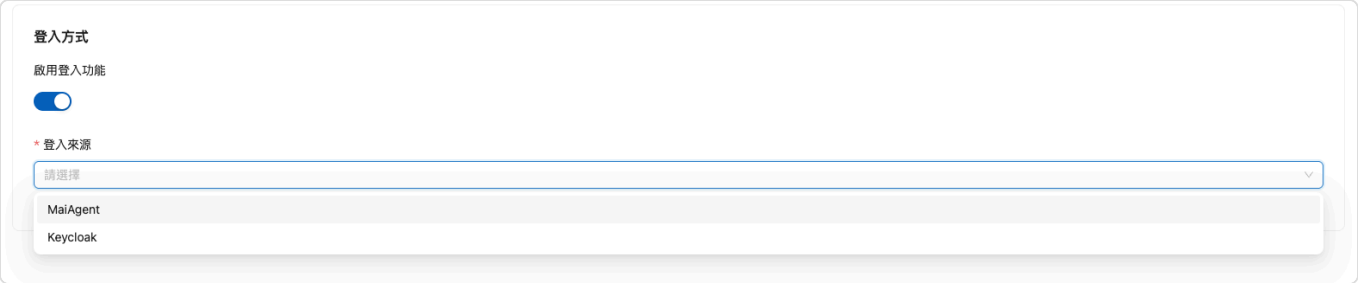

Conversation Access Without Login from the Embedding Site

If you embed Web Chat in your own website and users have already logged in to your site, enable "Allow Conversation Access Without Login via Source ID from the Embedding Site" so they can chat directly. For setup and signature integration instructions, see Conversation Access Without Login via Contact Source ID.

Use the Keycloak Login Identity to Call Tools

When Keycloak is selected as the login provider, enable "Use Keycloak Login Identity to Call Tools" to automatically include the visitor's own Keycloak login credentials when the AI assistant calls MCP/API tools. For setup and considerations, see Use the Keycloak Login Identity to Call Tools.

3. Auto-reply Settings

Configure the time periods when the bot replies here.

Click the toggle next to each day to select which days of the week the bot replies.

Set a time period: choose the start and end times.

Add a time period: if you need multiple periods, click “+Add Time Period” and configure the new period.

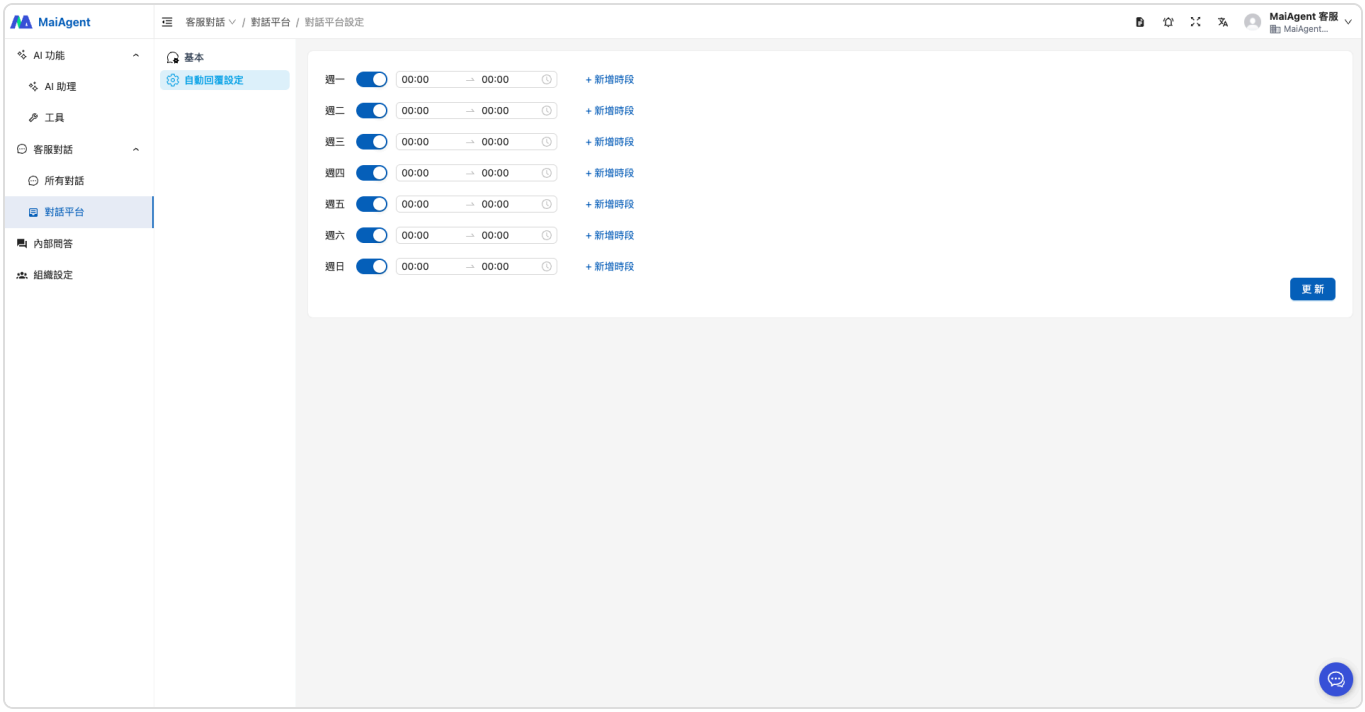

Chat Without Logging In Using a Contact Source ID

When you embed Web Chat in your website, users are often already logged in to your site. This setting lets them start chatting directly with the user identifier from your system, without going through another login screen, while the public chat URL still blocks anonymous access.

MaiAgent calls this identifier the **Source ID**, and uses this name in both the administration interface and the embed script. When the platform receives a Source ID, it maps it to (or creates) a **contact**. That user's conversation history is associated with this contact.

1. What this setting solves

Before this setting, you had two choices for login settings:

Enable login: Users were still asked to log in again, even though the embedding site had already verified their identity.
Disable login: Anyone could use the public chat URL anonymously.

With this setting enabled, you can achieve both goals: users with a Source ID and a valid signature can chat directly, while all visitors without them are redirected to the login page.

> This setting is disabled by default. When disabled, all existing behavior remains unchanged.

2. Enable the setting

Setting location: In the left menu, go to " Customer Service Conversations > Chat Platforms ", select a website-type chat platform, and open the " Login Settings " tab.

1. Select a login source first

This toggle depends on the login settings. You must first select a source (MaiAgent／AD／Keycloak／LINE) under " Login Source ". Until you select one, the toggle is unavailable and the interface displays " Please select a login source first ".

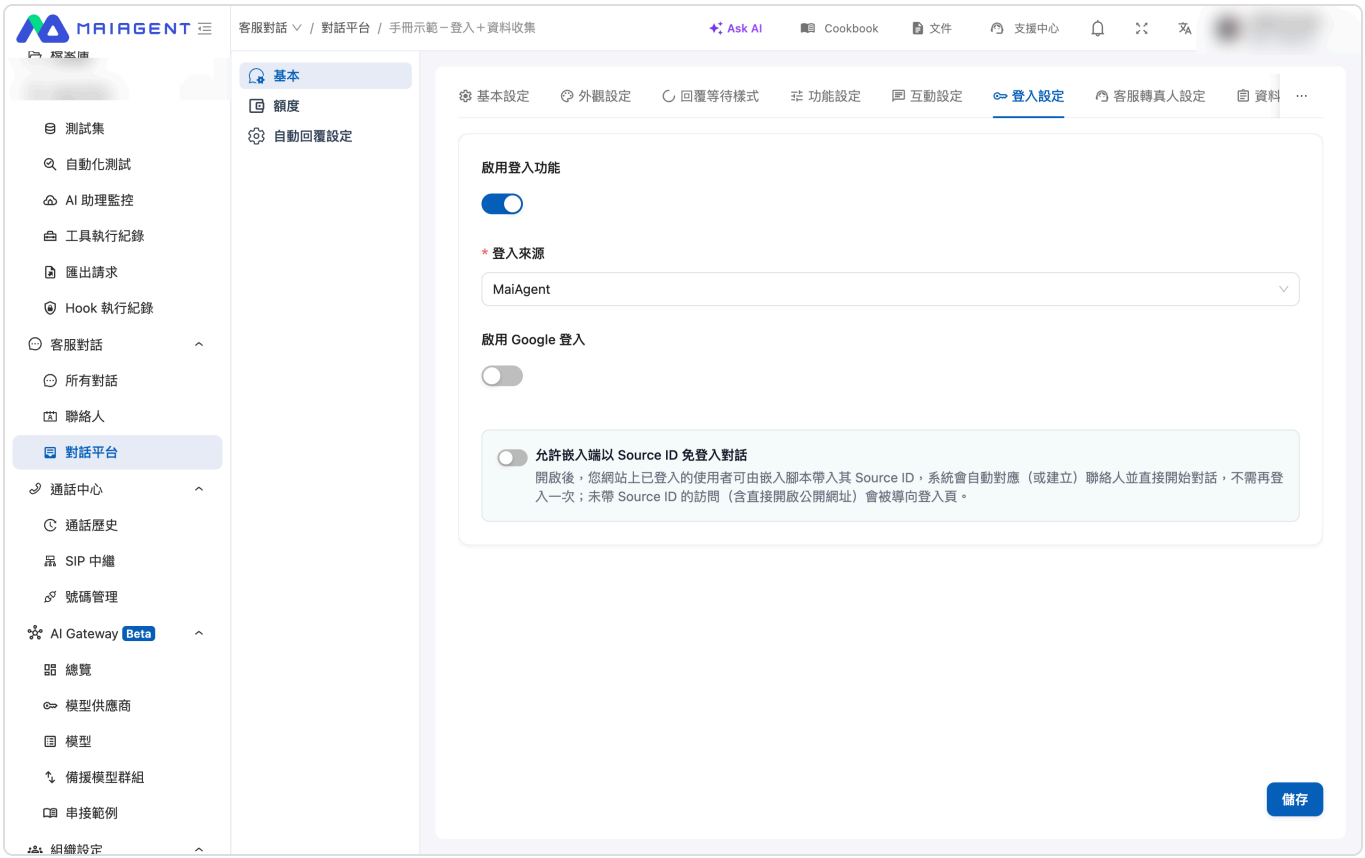

The Login Settings tab, with this toggle disabled by default

2. Enable the toggle

Enable " Allow embedded clients to chat without logging in using Source ID ". The verification method, signing secret, and integration examples appear only after you enable it.

The settings section shown after enabling the toggle; the verification method defaults to "Require signature verification"

3. Select a verification method

Verification method	Description	Use case
Require signature verification (recommended)	The embedded client must include a signature calculated by your backend. Access is granted only after verification succeeds. Only a server holding the secret can start a conversation on behalf of a user.	Production environments
Source ID only	No signature is verified; access is granted when a Source ID is provided.	Testing or closed internal networks

When you select "Source ID only", the Source ID acts as an access pass. If your Source IDs are guessable values such as email addresses, student IDs, or employee IDs, others can impersonate those users to chat and read their conversation history. For scenarios that require verified identities, use signature mode or AD／Keycloak／MaiAgent login instead.

A red risk warning appears when "Source ID only" is selected

4. Obtain the signing secret

After you select signature mode and save the settings, the platform generates a signing secret. In the secret field, you can:

Click " Show " to view the full secret (masked by default)

Click " Copy " to share it with your engineer

Click " Regenerate " to issue a new secret

> Store the secret in an environment variable on your server. Never put it in frontend code. Regenerating the secret immediately invalidates the old one. Update your backend at the same time, or all signature verification attempts will fail.

5. Set the signature validity period

Choose 1, 5, or 15 minutes. The default is 5 minutes. Expired or previously used signatures are rejected.

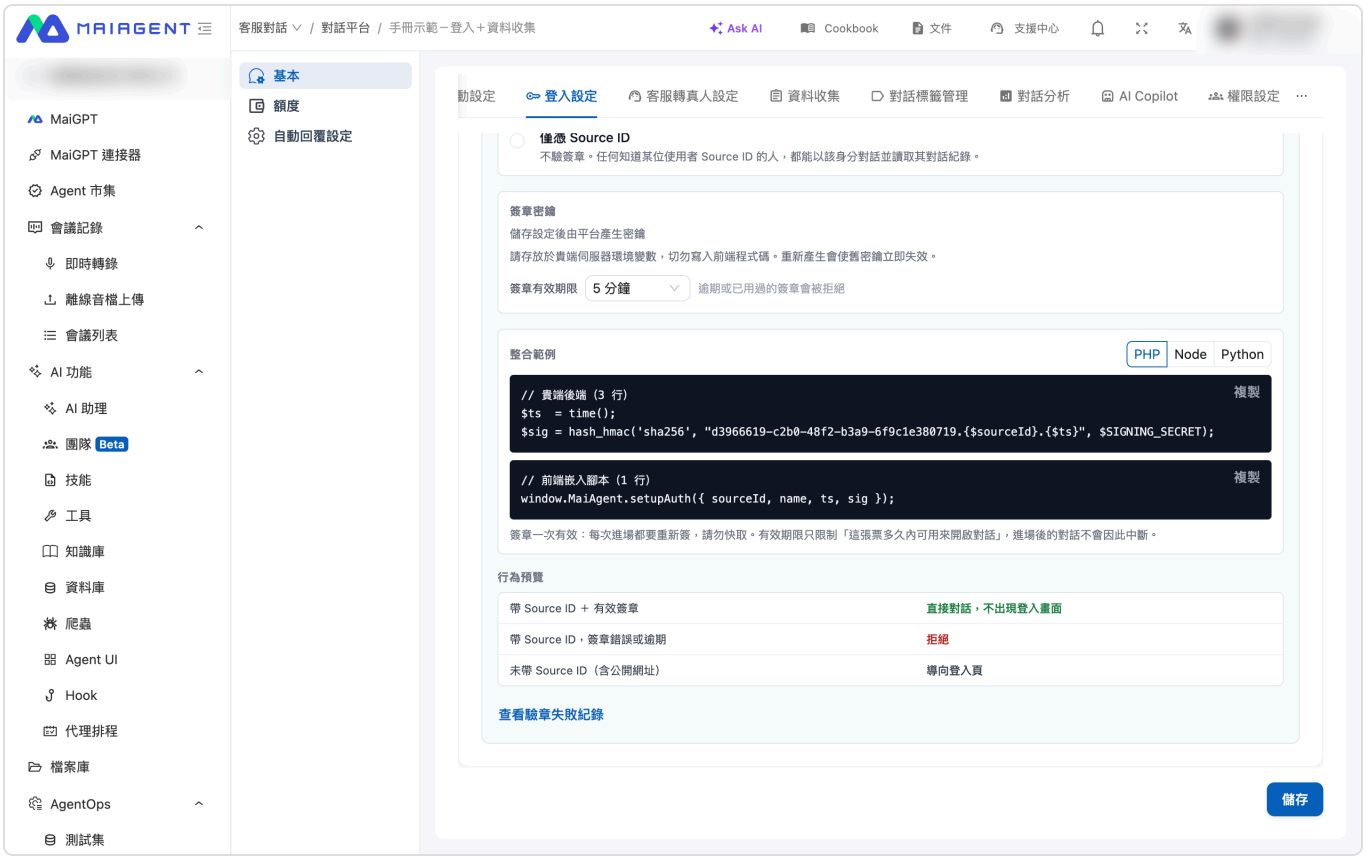

Signing secret, signature validity period, integration examples, and behavior preview

> The validity period limits "how long this signature can be used to start a conversation", rather than how long the conversation lasts. Once a user has authenticated, they can chat normally and will not be logged out after 5 minutes.

3. How your engineer integrates it

The " Integration Examples " section in the administration interface includes this chat platform's identifier, so the examples can be copied and used directly. For complete parameter descriptions, verification rules, and troubleshooting, ask your engineer to refer to Web Chat Embedding and SDK ・ Signature Verification.

1. Calculate the signature on the backend

The signature is `HMAC-SHA256(secret, "{chat platform identifier}.{Source ID}.{timestamp}")`, passed as a hex string. The timestamp is in Unix seconds (UTC), generated by your backend.

PHP

{% code overflow="wrap

```
$ts  = time();
$sig = hash_hmac('sha256', "{$webChatId}.{$sourceId}.{$ts}", $SIGNING_SECRET);
```

Node

{% code overflow="wrap

```
const ts = Math.floor(Date.now()  1000);
const sig = crypto.createHmac('sha256', SIGNING_SECRET)
  .update(`${webChatId}.${sourceId}.${ts}`).digest('hex');
```

Python

{% code overflow="wrap

```
ts = int(time.time())
sig = hmac.new(SIGNING_SECRET.encode(),
  f"{web_chat_id}.{source_id}.{ts}".encode(), hashlib.sha256).hexdigest()
```

2. Pass the values to the embed script on the frontend

```
window.MaiAgent.setupAuth({ sourceId, name, ts, sig });
```

3. Integration notes

Generate a new signature for every visit: Each signature can be used successfully only once. The embedded client must not cache signatures.

Calculate signatures on the backend: Once the secret appears in frontend code, it is effectively public.

Keep clocks accurate: The system allows 60 seconds of clock skew. Timestamps outside this range are rejected.

Use the correct chat platform identifier: The signature includes the chat platform identifier. A signature calculated with another chat platform's identifier will not pass verification.

4. Behavior comparison

Incoming request	Result
Includes Source ID + valid signature	Starts chatting directly without a login screen
Includes Source ID with an invalid or expired signature	Rejected
Reuses the same signature	First use succeeds; second use is rejected
No Source ID (including directly opening the public URL)	Redirected to the login page
Verification method is "Source ID only" + includes Source ID	Starts chatting directly without signature verification
This toggle is disabled	Retains the original login settings behavior

> All rejection scenarios return the same external response, without revealing the reason for failure or whether a particular Source ID exists. The reason appears only in the signature verification failure log in the administration interface.

5. Signature verification failure log

If you encounter integration issues, click " View Signature Verification Failure Log " in the settings section to see recent failures, so your engineer can troubleshoot independently. The log does not contain secrets or full signatures.

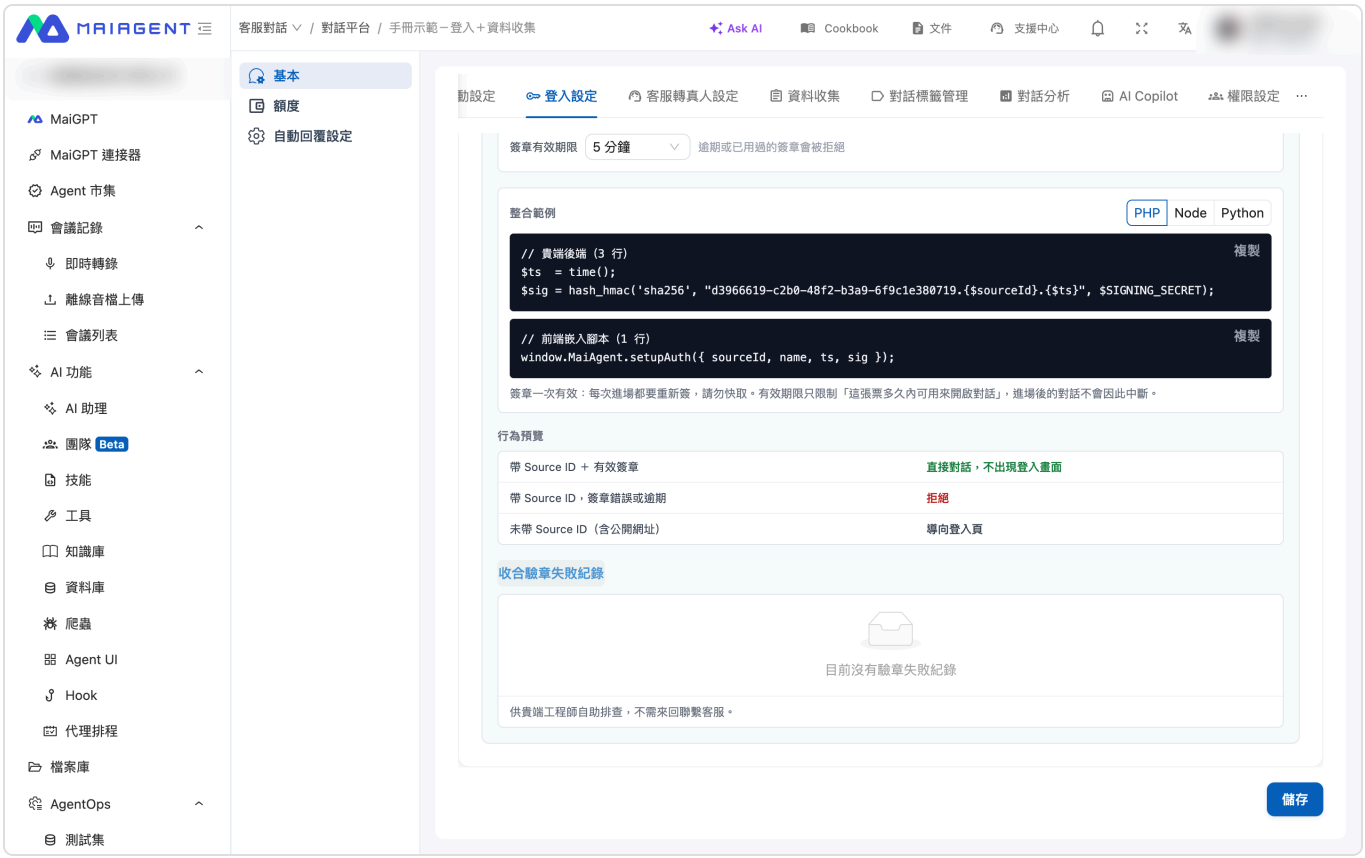

The expanded signature verification failure log

Reason	Common causes
Signature expired	Server clock skew, or the embedded client cached a signature
Timestamp exceeds the allowed future range	Server clock is more than 60 seconds ahead
Signature mismatch (secrets may differ)	The backend secret differs from the one in the administration interface, or the signature string was constructed incorrectly
Missing timestamp or signature	The embedded client did not provide `ts` or `sig`
Signature already used	The same signature was sent a second time
Replay prevention cache unavailable	A temporary platform issue; the system rejects access rather than granting it
Source ID maps to multiple contacts	Duplicate contact records with the same Source ID exist on this chat platform

6. Frequently asked questions

- Will the user's conversation end when the signature expires?
- Can one Source ID map to multiple contacts?
- Will enabling this setting affect existing embeds?
- Can users continue chatting after logging out of my website?
- Will AI tool calls include the login credentials of a user who enters with a Source ID?

Citation Display

This article introduces the citation display optimization features in the WebChat chat window, including multilingual support and visual presentation improvements.

What Is Citation Display?

When the AI Assistant answers a question using information from specific documents or data in the Knowledge Base, the system displays "citations" below the response, letting end users know where the information comes from.

This feature is especially suitable for scenarios that require source verification, such as:

Regulatory inquiry assistant: Displays the source of legal provisions

Policy consultation assistant: Displays policy document sources

Product customer service assistant: Displays product manuals or official documents

Medical information assistant: Displays medical literature sources

Citation Display Location

Citations appear below the AI Assistant's response content in a clear format:

```
AI Assistant response content...

📚 Citations:
• Product User Manual - Chapter 3
• FAQ - Return and Exchange Policy
```

Multilingual Support

Automatic Language Adaptation

WebChat citation display now supports automatic multilingual adaptation. When end users switch languages, the citation interface text automatically switches to the corresponding language.

Supported Languages

Traditional Chinese: 引用來源、參考資料

Simplified Chinese: 引用来源、参考资料

English: Citation、References

Language Switch Timing

WebChat automatically switches languages based on the following scenarios:

Browser language settings: On first load, the display language is based on the user's browser language

Manual switch: When users switch languages in WebChat, citation text updates accordingly

Website language environment: If your website supports multiple languages, WebChat automatically matches the website's language settings

> For WebChat language configuration, refer to: Web Chat Overview

Visual Presentation Optimization

In addition to multilingual support, the visual presentation of citations has been optimized for clearer readability.

Optimization Items

Spacing adjustments: More appropriate spacing between citations and response content for clearer visual hierarchy

Font styling: Citations use a different font style to distinguish them from the main content

Color treatment: Citations use lighter colors to avoid interfering with the readability of main content

Icon markers: Unified icons mark citations for better recognition

Dark Mode Support

Citation display also supports dark mode. When end users' devices use a dark theme, the text and background colors of citations automatically adjust for clear visibility against dark backgrounds.

Dark Mode Adaptation Items

Text color automatically adjusts to lighter shades

Background color adapts to the dark theme

Maintains contrast with main content

> Dark mode activates automatically based on the user's device settings; no manual configuration is needed.

Practical Use Cases

🏛️ Government Policy Consultation Platform

Usage scenario:
When citizens inquire about the eligibility requirements for a subsidy policy, the AI Assistant not only provides the answer but also displays the source link to the policy announcement.

Benefits:
* Increases information credibility
* Allows citizens to review the full announcement
* Reduces concerns about information accuracy

⚖️ Legal Consultation Assistant

Usage scenario:
When users inquire about specific legal provisions, the AI Assistant displays the source and article number of the relevant law.

Benefits:
* Provides complete legal basis
* Users can verify the original legal provisions
* Enhances the professionalism of the consultation service

🏥 Medical Information Platform

Usage scenario:
When patients inquire about disease-related information, the AI Assistant indicates which medical literature or health education article the information comes from.

Benefits:
* Provides information source transparency
* Helps patients understand the credibility of the information
* Complies with medical information provision standards

Admin Settings

Enable Citation Display

The citation display feature is enabled by default. To adjust the settings:

Go to the "AI Assistant" settings page
Select "Response Settings"
Find the "Show Citations" option
Toggle the switch to enable or disable

Customize Citation Format

You can add metadata to Knowledge Base documents to customize the citation display format:

Document title: The document name displayed in citations
Chapter number: The specific chapter or page number displayed
Publication date: The document's publication or update date
External link: A link to the original document

> For Knowledge Base metadata configuration, refer to: Document Management: Tags and Metadata

FAQ

Q1: Will citations show all referenced documents?

Not necessarily. The system displays the most relevant citations, typically no more than 3-5, to avoid information overload.

Q2: Can citations be hidden?

Yes. Disable the "Show Citations" option in the AI Assistant settings.

Q3: Can the citation language be fixed to a specific language?

Citations automatically match the WebChat language setting. If you want to fix it to a specific language, you need to lock the language option in the WebChat settings.

Q4: Do citations support clickable links?

Yes. If you have configured external links in your Knowledge Base documents, citations will display as clickable links.

Q5: What if citations are hard to read in dark mode?

The system has been optimized for dark mode. If issues persist, it may be a browser cache problem. Try clearing the cache or refreshing the page.

Markdown Rendering

This article introduces the message content rendering optimization features in the WebChat window, making AI Assistant responses clearer and easier to read.

What Is Message Content Rendering?

Message content rendering refers to how WebChat converts the AI Assistant's text responses into the visual presentation that users see. AI Assistant responses may include:

Plain text: General text descriptions

Markdown formatting: Headers, lists, bold, italics, and other formatting

Code blocks: Code examples or commands

Tables: Structured data presentation

Links: External links or reference materials

Good rendering makes this content easier to read and understand.

Optimization Items

1. Fully Upgraded Markdown Rendering Engine

WebChat now uses an enhanced Markdown rendering engine that supports richer formatting and more precise rendering.

Supported Markdown Formats

Format Type	Syntax Example	Rendering Description
Headers	`# Heading 1`, `## Heading 2`	Clear heading hierarchy for content segmentation
Bold	`bold text`	Emphasizes key content
Italic	`*italic text*`	Secondary emphasis or supplementary notes
Lists	`- Item 1`, `1. Item 1`	Ordered or unordered lists
Inline code	`` `code` ``	Inline code or commands

Format Type	Syntax Example	Rendering Description
Code blocks	``` code ```	Multi-line code with syntax highlighting support

Tables	`| Column 1`
	`| Column 2`
	`|`
		Structured data display

Links	`link text`	Clickable hyperlinks
Blockquotes	`> quote content`	Quotations or special notes

2. Dark Mode Visual Optimization

When end users' devices use a dark theme, WebChat automatically adjusts the visual presentation of message content to ensure clarity and readability on dark backgrounds.

Dark Mode Adjustments

Text and background contrast: Automatically adjusts text color for clear display on dark backgrounds

Message feedback buttons: Adjusts button colors and borders to remain visible in dark mode

Table visuals: Adjusts table borders and background colors for improved readability

Blockquote backgrounds: Adjusts blockquote background tones to differentiate from the main content

3. Streaming Rendering Performance Improvements

WebChat uses real-time streaming to display AI Assistant responses, so users don't need to wait for the full response to be generated. This update optimizes streaming rendering performance:

Faster rendering: Markdown formatting renders more quickly, reducing perceived latency

Improved smoothness: Text appears character by character more smoothly, with no noticeable lag

Format accuracy: Markdown formatting renders correctly even during the streaming process

> Users can see the already-generated content in real time while the AI Assistant is still composing its response, without waiting for the complete answer.

Practical Use Cases

💻 Technical Documentation Assistant

Scenario:
When developers ask for code examples, the AI Assistant provides responses with code blocks.

Optimization Effects:
* Code blocks support syntax highlighting for easier reading
* Code displays clearly in dark mode
* Developers can directly copy and use the code

Example Response:

Here's an example of reading a CSV file in Python:

```
python
import pandas as pd

df = pd.read_csv('data.csv')
print(df.head())
```

This code reads a CSV file and displays the first 5 rows.

```
### 📊 Data Analysis Assistant

Scenario:
Users ask for sales data analysis results, and the AI Assistant presents
them in table format.

Optimization Effects:
* Tables are clearly formatted with aligned columns
* Table borders and background colors adapt to dark mode
* Large datasets remain readable

Example Response:
```

Here are the top three selling products this quarter:

Product Name	Units Sold	Revenue
Product A	1,250	$125,000
Product B	980	$98,000
Product C	856	$85,600

Sales data as of the 20th of this month.

📚 Teaching Assistant

Scenario:
Students ask about learning steps, and the AI Assistant explains using an ordered list.

Optimization Effects:
* List items are clearly numbered
* Heading hierarchy is distinct, making step order easy to follow
* Key content is emphasized with bold text

Example Response:

Recommended Steps for Learning Python

Basic Syntax

Variables and data types

Conditionals and loops

Advanced Concepts

Functions and modules

Object-oriented programming

Hands-on Practice

Complete small projects

```
Contribute to open source


Each step is recommended to take 2-3 weeks of study.
```

Administrator Settings

Markdown Rendering Settings

WebChat's Markdown rendering is enabled by default. If your use case doesn't require Markdown formatting, you can adjust it in the settings:

Go to "Publish" → "WebChat Settings"

Find the "Message Rendering" settings section

Select the rendering mode:

Full Markdown: Supports all Markdown formats (default)

Basic formatting: Supports only bold, italic, and links

Plain text: No formatting rendered

Dark Mode Settings

Dark mode activates automatically based on user device settings. If you want WebChat to use a fixed light or dark theme:

Go to "Publish" → "WebChat Settings"

Find the "Appearance" settings section

Select the theme mode:

Auto: Based on user device settings (default)

Light: Always use the light theme

Dark: Always use the dark theme

> For detailed WebChat appearance settings, refer to: Connect Chat Platform: Website

Performance Impact

Rendering Performance

Performance impact of the optimized Markdown rendering engine:

Load time: First load increases by approximately 0.1-0.2 seconds (virtually unnoticeable)
Memory usage: Increases by approximately 1-2 MB (minimal impact on modern browsers)
Rendering speed: Approximately 30% faster compared to the previous version

Compatibility

The optimized rendering engine supports the following browsers:

Chrome
Edge (version 90+)
Firefox (version 88+)
Safari (version 14+)
Mobile browsers (iOS Safari, Android Chrome)

> Older browser versions may not correctly display some Markdown formats. It is recommended to prompt users to update their browsers.

FAQ

Q1: Why doesn't my AI Assistant's response have formatting?

Possible causes:
* The AI Assistant is not using Markdown syntax
* Markdown rendering is disabled
* The browser version is too old

Q2: What if some content is hard to read in dark mode?

Try the following:
1. Clear the browser cache
2. Refresh the page
3. Ensure you are using the latest browser version

Q3: Can I customize the syntax highlighting colors for code blocks?

Currently, syntax highlighting colors are fixed. Future versions may support customization.

Q4: How is long table content displayed?

Tables automatically support horizontal scrolling to ensure complete content display without affecting the layout of other content.

Q5: Will rendering optimization affect display on mobile devices?

No. Rendering optimization has been adapted for mobile devices to ensure clear display on small screens as well.

Embed MaiGPT in Window Mode

In addition to the common "floating button in the bottom-right corner," MaiAgent's Web Chat SDK introduces **MaiGPT window mode**, which lets you embed a **full ChatGPT-style conversation interface** (conversation history list on the left + main conversation area on the right) directly into any position on your page as the primary UI, rather than a small widget tucked in a corner.

> This page explains "how to embed MaiGPT window mode on your website." If you have not yet created a Web Chat platform, please refer to Integrate with Platform: Website to complete the basic setup first.

Where to Start? Get the Embed Code

The embed code on this page is obtained from Basic Settings. Follow the steps below:

Access the Chat Platform Settings

Go to "Customer Service > Chat Platforms", select your Web Chat platform, and click "Operations" to enter the settings page.

Click "Embed"

In Basic Settings, find the Public Access URL section and click the "Embed" button to open the embed window.

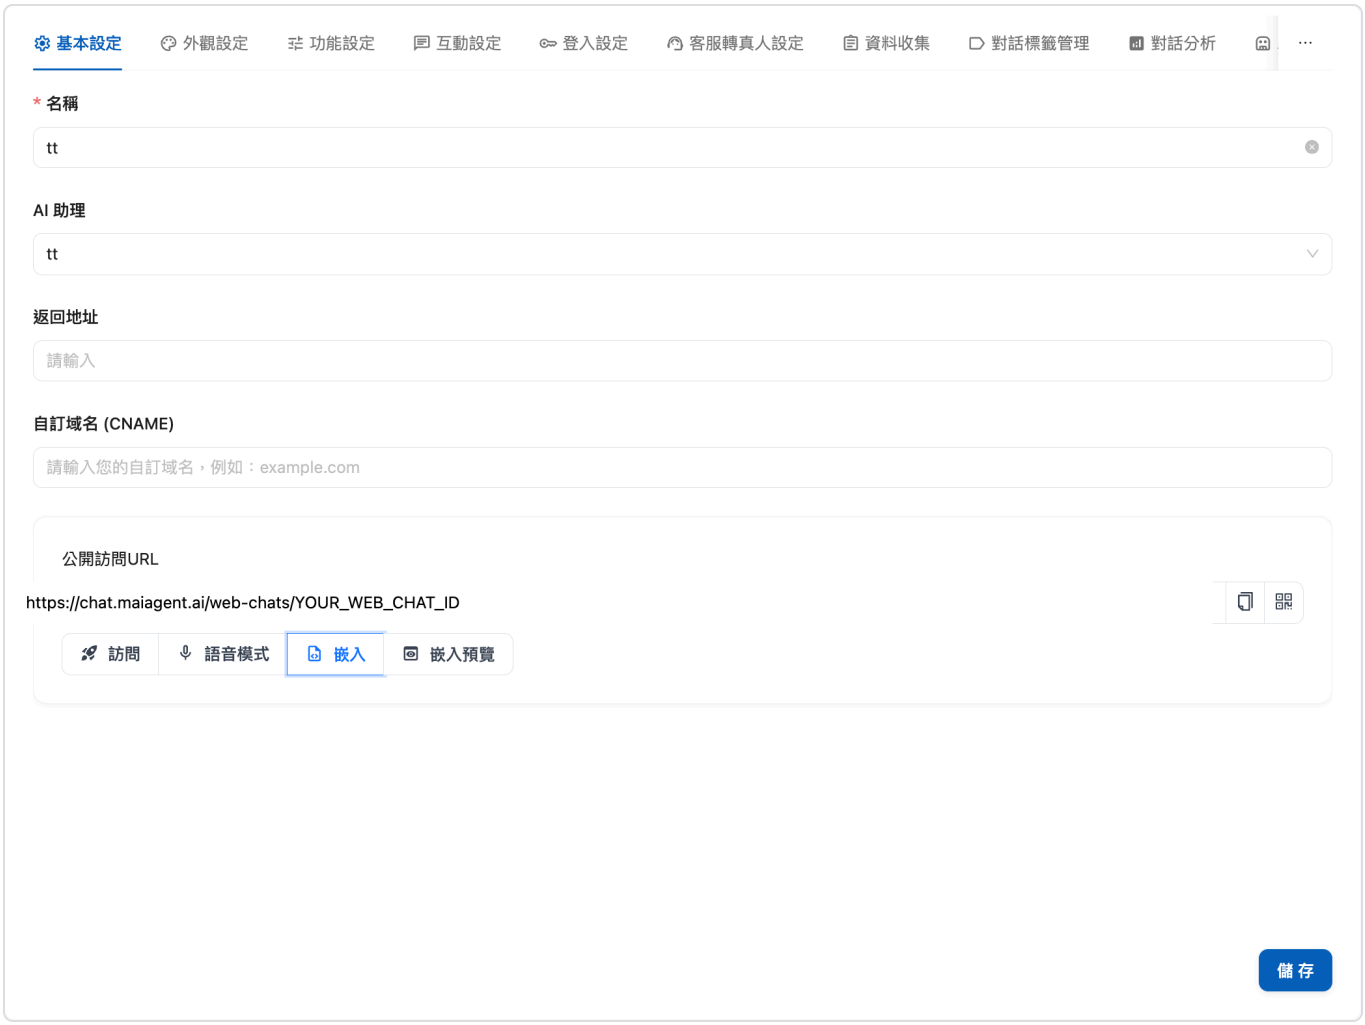

Basic Settings > Embed

Select the MaiGPT Embed Method

Switch between three embed methods at the top of the window and select "MaiGPT":

Embed Method	Description
JavaScript	Loaded via